

# ADOPT-01: Observability Maturity Model

---

**Series:** ADOPT | **Notebook:** 1 of 5 | **Created:** March 2026 | **Last Updated:** 03/12/2026

## Overview

---

Observability maturity is not a binary state — organizations progress through distinct levels as their practices, tooling, and culture evolve. This notebook introduces a five-level maturity model for Dynatrace-powered observability, provides assessment criteria for each level, and demonstrates how to use DQL queries to gauge your current position. Understanding where you stand today is the first step toward building a roadmap for improvement.

---

## Table of Contents

---

- [1. The Five Maturity Levels](#)
  - [2. Level 1 — Reactive Monitoring](#)
  - [3. Level 2 — Proactive Monitoring](#)
  - [4. Level 3 — Data-Driven Observability](#)
  - [5. Level 4 — Predictive Observability](#)
  - [6. Level 5 — Autonomous Operations](#)
  - [7. Assessing Your Current Level](#)
  - [8. Mapping Dynatrace Features to Maturity Levels](#)
  - [9. Summary and Next Steps](#)
-

# Prerequisites

| Requirement           | Details                                                                                |
|-----------------------|----------------------------------------------------------------------------------------|
| Dynatrace Environment | SaaS or Managed with Grail enabled                                                     |
| Permissions           | storage:logs:read , storage:metrics:read , storage:entities:read , storage:events:read |
| Data                  | At least 24 hours of ingested monitoring data                                          |
| Audience              | Platform engineers, SREs, engineering managers, and leadership stakeholders            |

## 1. The Five Maturity Levels

The observability maturity model defines five progressive levels. Each level builds on the previous one, adding capabilities and shifting organizational behavior.

| Level | Name        | Characteristics                                               | Key Indicator                                               |
|-------|-------------|---------------------------------------------------------------|-------------------------------------------------------------|
| 1     | Reactive    | Alert-driven, manual troubleshooting, siloed tools            | "We find out about problems from users"                     |
| 2     | Proactive   | Threshold-based alerts, basic dashboards, some automation     | "We detect problems before users report them"               |
| 3     | Data-Driven | Centralized observability, SLOs defined, DQL-powered analysis | "We make decisions based on telemetry data"                 |
| 4     | Predictive  | AI-driven anomaly detection, forecasting, capacity planning   | "We predict problems before they occur"                     |
| 5     | Autonomous  | Self-healing, automated remediation, closed-loop operations   | "The platform resolves problems without human intervention" |

Most organizations operate between Level 1 and Level 3. The goal is not necessarily to reach Level 5 everywhere, but to deliberately choose the appropriate level for each service based on business criticality.

## 2. Level 1 — Reactive Monitoring

---

### Characteristics

- Monitoring is an afterthought, added post-deployment
- Alerts are noisy and frequently ignored
- Troubleshooting relies on SSH-ing into hosts and reading log files
- No standardized observability tooling across teams
- Mean Time to Detect (MTTD) is measured in hours, often reported by end users

### Assessment Criteria

| Criterion        | Level 1 Indicator              |
|------------------|--------------------------------|
| Agent coverage   | < 50% of hosts instrumented    |
| Data sources     | Only infrastructure metrics    |
| Alerting         | Static thresholds or none      |
| Dashboards       | Ad-hoc, per-team, no standards |
| Incident process | Manual, tribal knowledge       |

### Dynatrace Features at This Level

- OneAgent deployment on critical hosts
- Basic infrastructure monitoring
- Default Davis AI alerting (often not yet configured)

## 3. Level 2 — Proactive Monitoring

---

### Characteristics

- Monitoring is part of the deployment checklist
- Threshold-based alerts detect common failure modes
- Centralized dashboards exist but may not be maintained

- Teams respond to problems before user impact escalates
- MTTD is measured in minutes

## Assessment Criteria

| Criterion        | Level 2 Indicator                       |
|------------------|-----------------------------------------|
| Agent coverage   | 50-80% of hosts instrumented            |
| Data sources     | Infrastructure + application metrics    |
| Alerting         | Threshold-based, some Davis AI          |
| Dashboards       | Standardized templates for key services |
| Incident process | Defined runbooks for top 10 issues      |

## Dynatrace Features at This Level

- Full-stack OneAgent deployment
- Davis AI problem detection enabled
- Custom metric alerting profiles
- Basic synthetic monitoring for key URLs

## 4. Level 3 — Data-Driven Observability

---

### Characteristics

- Observability data drives architectural and operational decisions
- SLOs are defined and tracked for critical services
- Teams use DQL and Grail for ad-hoc investigation
- Log, trace, and metric data are correlated
- Observability is treated as a platform service

## Assessment Criteria

| Criterion        | Level 3 Indicator                            |
|------------------|----------------------------------------------|
| Agent coverage   | > 80% of hosts and services instrumented     |
| Data sources     | Infrastructure + APM + logs + traces + RUM   |
| Alerting         | SLO-based alerting, reduced noise            |
| Dashboards       | Business-context dashboards, executive views |
| Incident process | Data-driven RCA, blameless postmortems       |

### Dynatrace Features at This Level

- Grail data lakehouse with DQL queries
- OpenPipeline for log routing and enrichment
- SLO definitions and burn-rate alerting
- Distributed tracing with span analytics
- Real User Monitoring (RUM) and Session Replay

## 5. Level 4 — Predictive Observability

---

### Characteristics

- AI-driven anomaly detection replaces static thresholds
- Capacity planning uses forecasting models
- Problems are identified before they cause user impact
- Teams focus on trends and patterns rather than individual alerts

## Assessment Criteria

| Criterion        | Level 4 Indicator                              |
|------------------|------------------------------------------------|
| Agent coverage   | > 95% with auto-discovery                      |
| Data sources     | All telemetry types + business events          |
| Alerting         | AI-driven, predictive, low false-positive rate |
| Dashboards       | Predictive views, forecasting charts           |
| Incident process | Automated triage and enrichment                |

### Dynatrace Features at This Level

- Davis AI with custom anomaly detection
- Metric forecasting and trend analysis
- Business event correlation
- Ownership and team assignment for entities

## 6. Level 5 — Autonomous Operations

---

### Characteristics

- Self-healing systems automatically remediate known issues
- Workflows trigger automated responses to problems
- Human intervention is the exception, not the norm
- Continuous optimization is automated

## Assessment Criteria

| Criterion        | Level 5 Indicator                                   |
|------------------|-----------------------------------------------------|
| Agent coverage   | 100% with automated rollout                         |
| Data sources     | Full telemetry + external integrations              |
| Alerting         | Event-driven automation, minimal manual alerts      |
| Dashboards       | Real-time executive cockpits                        |
| Incident process | Closed-loop: detect → diagnose → remediate → verify |

## Dynatrace Features at This Level

- Dynatrace Workflows for automated remediation
- AutomationEngine with event-driven triggers
- Configuration-as-code (Monaco) for drift detection
- Full API-driven operations

## 7. Assessing Your Current Level

The following DQL queries help you assess where your organization stands today. Each query targets a key indicator of maturity.

### 7.1 Monitored Entity Coverage

Understanding how many entities Dynatrace is monitoring gives a baseline for agent deployment coverage.

```
// Count monitored entities by type to assess coverage breadth
fetch dt.entity.host
| summarize host_count = count()
| append [fetch dt.entity.service | summarize service_count = count()]
| append [fetch dt.entity.process_group | summarize process_group_count = count()]
| append [fetch dt.entity.application | summarize application_count = count()]
```

## 7.2 Data Ingestion Diversity

Mature organizations ingest multiple telemetry types. This query checks which data sources are actively flowing into Grail.

```
// Check log ingestion volume over the last 24 hours
fetch logs, from:-24h
| summarize log_count = count(), distinct_sources = countDistinct(log.source)
```

```
// Check span ingestion to verify distributed tracing is active
fetch spans, from:-24h
| summarize span_count = count(), distinct_services =
countDistinct(dt.entity.service)
```

## 7.3 Davis AI Problem Detection Activity

Active Davis AI problem detection is a key indicator of Level 2+ maturity. If few or no problems are being detected, it may indicate insufficient instrumentation or misconfigured alerting.

```
// Count Davis problems in the last 7 days by status
fetch dt.davis.problems, from:-7d
| summarize problem_count = count(), by:{event.status}
| sort problem_count desc
```

## 7.4 Alerting Quality Check

A high ratio of duplicate or frequent events suggests noisy alerting — a characteristic of lower maturity levels.

```
// Assess alerting noise: frequent and duplicate problems vs unique problems
fetch dt.davis.problems, from:-7d
| summarize
    total = count(),
    frequent = countIf(dt.davis.is_frequent_event == true),
    duplicate = countIf(dt.davis.is_duplicate == true)
| fieldsAdd unique = total - frequent - duplicate
| fieldsAdd noise_ratio = round((toDouble(frequent + duplicate) /
toDouble(total)) * 100, decimals: 1)
```



#### Interpreting the noise ratio:

- < 20% — Healthy alerting (Level 3+)
- 20-50% — Moderate noise, review alerting profiles (Level 2)
- > 50% — Significant noise, alerting needs overhaul (Level 1)

## 8. Mapping Dynatrace Features to Maturity Levels

Use this table to identify which Dynatrace capabilities to adopt next based on your target maturity level.

| Feature                        | L1 | L2 | L3 | L4 | L5 |
|--------------------------------|----|----|----|----|----|
| OneAgent (Infrastructure)      | X  | X  | X  | X  | X  |
| OneAgent (Full-Stack APM)      |    | X  | X  | X  | X  |
| Davis AI Problem Detection     |    | X  | X  | X  | X  |
| Synthetic Monitoring           |    | X  | X  | X  | X  |
| Log Management (Grail)         |    |    | X  | X  | X  |
| Distributed Tracing (Spans)    |    |    | X  | X  | X  |
| OpenPipeline                   |    |    | X  | X  | X  |
| SLO Definitions                |    |    | X  | X  | X  |
| DQL / Notebooks                |    |    | X  | X  | X  |
| Real User Monitoring           |    |    | X  | X  | X  |
| Business Events                |    |    |    | X  | X  |
| Davis AI Forecasting           |    |    |    | X  | X  |
| Ownership & Teams              |    |    |    | X  | X  |
| Workflows (Automation)         |    |    |    |    | X  |
| Configuration-as-Code (Monaco) |    |    |    |    | X  |
| AutomationEngine               |    |    |    |    | X  |

**Tip:** You do not need to adopt every feature at every level. Focus on the features that address your most pressing gaps.

## 9. Summary and Next Steps

---

### Key Takeaways

- Observability maturity is a spectrum from reactive (Level 1) to autonomous (Level 5)
- Assessment should be data-driven — use DQL queries to measure coverage, ingestion diversity, and alerting quality
- Not every service needs Level 5 maturity — align investment with business criticality
- Dynatrace features map directly to maturity levels, providing a clear adoption path

### Next Steps

- Proceed to **ADOPT-02: Platform Health Assessment** to build a detailed scorecard of your current Dynatrace deployment
- Use the maturity assessment results to prioritize features for adoption
- Share the maturity model with leadership to align on observability investment priorities

---

*This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.*