

AIOPS-05: AI Models — Causal, Predictive, and Generative

Series: AIOPS — Dynatrace Intelligence | **Notebook:** 5 of 8 | **Created:** May 2026 |
Last Updated: 05/05/2026

Overview

Behind every Dynatrace Intelligence surface there's a specific model doing specific work. This notebook names the models, explains what each one is responsible for, and answers the questions admins ask about them: where do they live, what do they cost, and what is (or is not) bring-your-own.

Audience: Platform architect, AI / governance lead, anyone evaluating Dynatrace AI for regulatory or procurement reasons.

Outcome: A clear separation between what the platform's AI is, where each piece runs, and what is configurable vs. fixed.




Table of Contents

1. [The Three Model Categories Recap](#)
 2. [Causal Correlation Analysis](#)
 3. [Predictive AI — Baselines and Forecasts](#)
 4. [Seasonal Baseline](#)
 5. [Generative Models — Where They Live](#)
 6. [Bring Your Own Model — What's Possible](#)
 7. [Cross-Series Pointers](#)
-

Prerequisites

Requirement	Details
Dynatrace Environment	SaaS Gen3
Apps	Notebooks app for the analyzer demos in this notebook
Permissions	<code>davis:analyzers:execute</code> for analyzer queries; standard read permissions otherwise
Background	AIOPS-01 for the umbrella; AIOPS-02 for anomaly detection mechanisms (this notebook is the <i>model</i> perspective on the same material)

1. The Three Model Categories Recap

Category	What it does	Where it runs	Configurable?
Causal	Topology-aware root cause	Davis backend	No (Smartscape itself is the configuration)
Predictive	Baselines, forecasts, novelty	Davis analyzer engine	Yes — sensitivity, lookback, scope

Category	What it does	Where it runs	Configurable?
Generative	LLM-backed Assist & summaries	Hosted / managed by Dynatrace	Limited — you can't swap the LLM

Three different model families. Three different architectural approaches. Treat them differently.

2. Causal Correlation Analysis

The single biggest differentiator of Dynatrace AI vs. correlation-based AIOps tools.

Causal AI does not run a statistical model over symptom signals to guess at co-occurrence. It walks the Smartscape topology graph and uses the dependency edges to attribute symptoms to upstream causes. Two events on dependent entities go into one problem; two events on independent entities stay separate.

Inputs: the Davis event stream (`dt.davis.events`) plus the live Smartscape graph.

Output: records in `dt.davis.problems` , each with a `root_cause_entity_id` and an array of `affected_entity_ids` .

Operator's job: make Smartscape complete. Tag entities, instrument every service, declare custom dependencies where auto-discovery doesn't reach.

3. Predictive AI — Baselines and Forecasts

Predictive AI is a **family of statistical models** wrapped as Davis analyzers. Each model is purpose-built; each is callable directly from MCP and from the Anomaly Detection app.

Model	Purpose	MCP tool
Adaptive baseline	Detect anomalies vs. learned recent baseline	<code>mcp__dynatrace__adaptive-anomaly-detector</code>
Seasonal baseline	Detect anomalies vs. recurring weekly/daily pattern	<code>mcp__dynatrace__seasonal-baseline-anomaly-detector</code>

Model	Purpose	MCP tool
Static threshold	Test against a fixed limit	<code>mcp__dynatrace__static-threshold-analyzer</code>
Novelty detection	Flag never-seen patterns	<code>mcp__dynatrace__timeseries-novelty-detection</code>
Forecast	Project a series forward	<code>mcp__dynatrace__timeseries-forecast</code>

All five accept a `timeseries` query as input. Their output is a structured analyzer result — anomaly intervals, forecast bands with confidence intervals, novelty flags. Embeddable in workflows for scheduled analysis.

Configuration knobs:

- *Lookback window* (how far back is "learned")
- *Sensitivity* (how aggressive to be about flagging anomalies)
- *Scope* (per-entity, per-dimension, global)

Tuning these is the difference between alert fatigue and silent failure.

4. Seasonal Baseline

Worth a closer look because it's where most teams underuse the platform.

Seasonal baseline learns the metric's recurring patterns — daily, weekly, and where enough data exists, yearly. Detection considers the *expected value at this time of week*, not just "is the metric high."

The use case: any metric that varies with business activity. Login traffic peaks at 8 AM and noon. Database write volume peaks at end-of-month. ETL job durations peak on Sunday night. A static threshold or even an auto-adaptive baseline misses these — they see the seasonal peak as either noise or as the new normal.

Common mistake: applying seasonal baseline to a metric that doesn't have seasonality. The model still produces output, but the bands are essentially the global baseline plus statistical noise.

Below is an example that exercises the analyzer engine — pick a metric you know has weekly seasonality (host CPU on a business-hours app, for instance) and feed it in.

```
// Hourly host CPU for the last 14 days, per host
// (Feed this query to mcp__dynatrace__seasonal-baseline-anomaly-detector for
evaluation.)
timeseries cpu = avg(dt.host.cpu.usage),
  by:{dt.entity.host},
  from:-14d,
  interval:1h
| limit 100
```

5. Generative Models — Where They Live

Generative AI in Dynatrace is hosted and managed — you do not see or swap the underlying LLM. Three things live in this category:

1. **DQL2NL skill** — translates DQL ↔ natural language
2. **Problem summaries** — narrative explanations in the Problems app
3. **Dynatrace Assist** — chat surface across the product

What you don't get: model selection, fine-tuning on customer data, prompt-template customization. Dynatrace owns the model surface so the data privacy guarantees are stable across customers.

What you do get: vector-based grounding over Dynatrace docs and (via problem context) your own observability data when the assistant is invoked from a problem. Refer to the Davis CoPilot data privacy and security policy for current details.

6. Bring Your Own Model — What's Possible

Practical answer in 2026: **the platform LLM is fixed; your AI integrations are not.**

What that means in practice:

Surface	BYO model?	Notes
Dynatrace Assist	No	Hosted; not configurable
Problem summaries	No	Hosted; not configurable

Surface	BYO model?	Notes
Workflow AI tasks	Yes	Workflow tasks can call external LLMs (OpenAI, Anthropic, Bedrock, Vertex) — see AIOPS-06
MCP-based agentic flows	Yes	Your agent (Claude Code, Cursor, GitHub Copilot) chooses its own LLM
Dynatrace AI app (community)	Varies	Community apps can wrap external models — see AIOPS-06

Architectural takeaway: keep tenant-grounded, in-product AI on the platform's models. Use BYO models for orchestration, automation, and any operation that benefits from a model your team can fine-tune or audit independently.

7. Cross-Series Pointers

- **AIOPS-02** — the operator's perspective on the predictive models
 - **AIOPS-06** — agentic / BYO model integrations via workflows and MCP
 - **AUTOM-05/06** — anomaly detection settings as code
-

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.