

# Dynatrace SDKs

Series: AUTOM | Notebook: 6 of 8 | Created: January 2026

## Table of Contents

- 1. [Introduction](#)
- 2. [TypeScript SDK](#)
- 3. [Python SDK](#)
- 4. [Common Patterns](#)
- 5. [Building Applications](#)
- 6. [Next Steps](#)

## Prerequisites

Before starting this notebook, ensure you have:

| Requirement             | Description                |
|-------------------------|----------------------------|
| Node.js 18+             | For TypeScript SDK         |
| Python 3.9+             | For Python SDK             |
| API Token               | Token with required scopes |
| Development Environment | VS Code or similar IDE     |

# Learning Objectives

---

By the end of this notebook, you will:

- Understand the Dynatrace SDK architecture
  - Know how to use TypeScript and Python clients
  - Be able to query data and manage configurations programmatically
  - Build custom automation applications
- 

## 1. Introduction

---

Dynatrace provides official SDKs for programmatic access to the platform. These SDKs are auto-generated from OpenAPI specifications, ensuring they stay current with API changes.

### Available SDKs

| SDK            | Language              | Use Case                |
|----------------|-----------------------|-------------------------|
| @dynatrace-sdk | TypeScript/JavaScript | Web apps, Node.js tools |
| dt-sdk         | Python                | Scripts, data science   |

### SDK vs Direct API

| Aspect          | SDK        | Direct API |
|-----------------|------------|------------|
| Type safety     | Yes        | No         |
| Auto-completion | Yes        | No         |
| Authentication  | Built-in   | Manual     |
| Pagination      | Handled    | Manual     |
| Error handling  | Structured | Raw HTTP   |

## Client Libraries

| Client                | Purpose                     |
|-----------------------|-----------------------------|
| <b>QueryClient</b>    | Execute DQL queries         |
| <b>SettingsClient</b> | Manage Settings 2.0 objects |
| <b>EntitiesClient</b> | Query entity topology       |
| <b>MetricsClient</b>  | Query and ingest metrics    |
| <b>EventsClient</b>   | Query and ingest events     |
| <b>LogsClient</b>     | Query and ingest logs       |

## 2. TypeScript SDK

### Installation

```
npm install @dynatrace-sdk/client-query
npm install @dynatrace-sdk/client-classic-environment-v2
npm install @dynatrace-sdk/client-settings-v2
```

### Configuration

```
import { setEnvConfig } from '@dynatrace-sdk/client-core';

// Configure environment
setEnvConfig({
  environmentUrl: process.env.DT_TENANT_URL,
  apiToken: process.env.DT_API_TOKEN
});
```

## Query Example

```
import { queryClient } from '@dynatrace-sdk/client-query';

async function queryHosts() {
  const result = await queryClient.query({
    body: {
      query: `
        fetch dt.entity.host
        | fieldsAdd name = entity.name, osType
        | sort name asc
        | limit 10
      `,
      defaultTimeframeStart: 'now-24h',
      defaultTimeframeEnd: 'now'
    }
  });

  console.log('Hosts:', result.result?.records);
  return result.result?.records;
}
```

---

## Settings Management

```
import { settingsObjectsClient } from '@dynatrace-sdk/client-classic-environment-v2';

// List management zones
async function listManagementZones() {
  const response = await settingsObjectsClient.getSettingsObjects({
    schemaIds: 'builtin:management-zones',
    pageSize: 100
  });

  return response.items;
}

// Create management zone
async function createManagementZone(name: string) {
  const response = await settingsObjectsClient.postSettingsObjects({
    body: [{
      schemaId: 'builtin:management-zones',
      scope: 'environment',
      value: {
        name: name,
        rules: []
      }
    }]
  });

  return response;
}
```

## Entity Queries

```
import { entitiesClient } from '@dynatrace-sdk/client-classic-environment-v2';

async function getServices() {
  const response = await entitiesClient.getEntities({
    entitySelector: 'type(SERVICE)',
    fields: '+properties,+tags',
    pageSize: 100
  });

  return response.entities;
}
```

---

## 3. Python SDK

### Installation

```
pip install dt-sdk
```

### Configuration

```
import os
from dynatrace_sdk import DynatraceClient

# Initialize client
client = DynatraceClient(
    environment_url=os.environ['DT_TENANT_URL'],
    api_token=os.environ['DT_API_TOKEN']
)
```

## Query Example

```
from dynatrace_sdk.query import QueryService

def query_hosts():
    query_service = QueryService(client)

    result = query_service.query(
        query="""
            fetch dt.entity.host
            | fieldsAdd name = entity.name, osType
            | sort name asc
            | limit 10
        """,
        default_timeframe_start='now-24h',
        default_timeframe_end='now'
    )

    return result.records
```

---

## Settings Management

```
from dynatrace_sdk.settings import SettingsService

def list_management_zones():
    settings = SettingsService(client)

    zones = settings.list_objects(
        schema_ids='builtin:management-zones',
        page_size=100
    )

    return zones

def create_management_zone(name: str):
    settings = SettingsService(client)

    result = settings.create_object(
        schema_id='builtin:management-zones',
        scope='environment',
        value={
            'name': name,
            'rules': []
        }
    )

    return result
```



## Pagination Handling

```
def get_all_hosts():
    """Iterate through all pages of hosts."""
    all_hosts = []
    next_page_key = None

    while True:
        response = client.entities.get_entities(
            entity_selector='type(HOST)',
            page_size=500,
            next_page_key=next_page_key
        )

        all_hosts.extend(response.entities)

        if not response.next_page_key:
            break
        next_page_key = response.next_page_key

    return all_hosts
```

---

## 4. Common Patterns

---

### Bulk Configuration Export

```
// TypeScript: Export all settings for backup
import { settingsObjectsClient, settingsSchemasClient } from '@dynatrace-
sdk/client-classic-environment-v2';
import * as fs from 'fs';

async function exportAllSettings() {
  // Get all available schemas
  const schemas = await
settingsSchemasClient.getAvailableSchemaDefinitions();

  const export_data: Record<string, any[]> = {};

  for (const schema of schemas.items || []) {
    const objects = await settingsObjectsClient.getSettingsObjects({
      schemaIds: schema.schemaId,
      pageSize: 500
    });

    if (objects.items && objects.items.length > 0) {
      export_data[schema.schemaId] = objects.items;
    }
  }

  fs.writeFileSync('settings-export.json', JSON.stringify(export_data, null,
2));
  return export_data;
}
```

## Error Handling

```
import { ApiError } from '@dynatrace-sdk/client-core';

async function safeQuery(query: string) {
  try {
    const result = await queryClient.query({
      body: { query }
    });
    return { success: true, data: result };
  } catch (error) {
    if (error instanceof ApiError) {
      console.error(`API Error: ${error.status} - ${error.message}`);
      return { success: false, error: error.message };
    }
    throw error;
  }
}
```

---

## Rate Limiting

```
import time
from functools import wraps

def rate_limited(max_per_second: float):
    """Decorator to rate limit API calls."""
    min_interval = 1.0 / max_per_second
    last_called = [0.0]

    def decorator(func):
        @wraps(func)
        def wrapper(*args, **kwargs):
            elapsed = time.time() - last_called[0]
            if elapsed < min_interval:
                time.sleep(min_interval - elapsed)
            result = func(*args, **kwargs)
            last_called[0] = time.time()
            return result
        return wrapper
    return decorator

@rate_limited(10) # Max 10 requests per second
def api_call(entity_id: str):
    return client.entities.get_entity(entity_id=entity_id)
```

## Async Operations

```
// TypeScript: Concurrent queries with rate limiting
import pLimit from 'p-limit';

const limit = pLimit(5); // Max 5 concurrent requests

async function queryMultipleHosts(hostIds: string[]) {
  const queries = hostIds.map(hostId =>
    limit(() => queryClient.query({
      body: {
        query: `
          fetch dt.entity.host
          | filter id == "${hostId}"
        `
      }
    }
  ));
  return Promise.all(queries);
}
```

---

## 5. Building Applications

---

### CLI Tool Example

```
#!/usr/bin/env ts-node
import { setEnvConfig } from '@dynatrace-sdk/client-core';
import { queryClient } from '@dynatrace-sdk/client-query';
import { Command } from 'commander';

const program = new Command();

program
  .name('dt-cli')
  .description('Dynatrace CLI tool')
  .version('1.0.0');

program
  .command('hosts')
  .description('List all hosts')
  .action(async () => {
    setEnvConfig({
      environmentUrl: process.env.DT_TENANT_URL!,
      apiToken: process.env.DT_API_TOKEN!
    });

    const result = await queryClient.query({
      body: {
        query: 'fetch dt.entity.host | fields entity.name'
      }
    });

    console.table(result.result?.records);
  });

program.parse();
```

## Reporting Script

```
import pandas as pd
from dynatrace_sdk import DynatraceClient
from dynatrace_sdk.query import QueryService

def generate_host_report():
    """Generate a host inventory report."""
    client = DynatraceClient(
        environment_url=os.environ['DT_TENANT_URL'],
        api_token=os.environ['DT_API_TOKEN']
    )

    query_service = QueryService(client)

    result = query_service.query(
        query="""
            fetch dt.entity.host
            | fieldsAdd
              name = entity.name,
              osType,
              cpuCores = entity.detected.properties[`host.cpuCores`],
              memoryTotal = entity.detected.properties[`host.memoryTotal`]
            | sort name asc
        """
    )

    df = pd.DataFrame(result.records)
    df.to_csv('host_inventory.csv', index=False)

    return df
```

---

## Best Practices

| Practice              | Description                               |
|-----------------------|-------------------------------------------|
| Environment variables | Never hardcode credentials                |
| Type safety           | Use TypeScript types or Python type hints |
| Error handling        | Catch and handle API errors               |
| Pagination            | Always handle paginated responses         |
| Rate limiting         | Respect API limits                        |
| Logging               | Add structured logging                    |

## 6. Next Steps

### When to Use SDKs

| Scenario              | Tool Choice                |
|-----------------------|----------------------------|
| Custom reporting tool | SDK                        |
| One-off query         | Direct API or Notebooks    |
| Config deployment     | Monaco or Terraform        |
| Integration app       | SDK                        |
| Complex logic         | SDK with JavaScript/Python |

### Continue the Series

| Next Notebook               | Focus                         |
|-----------------------------|-------------------------------|
| AUTOM-07: CI/CD Integration | GitOps patterns and pipelines |



## Additional Resources

- [Dynatrace Developer Portal](#)
  - [TypeScript SDK Documentation](#)
  - [Python SDK on PyPI](#)
  - [API Reference](#)
- 

## Summary

---

In this notebook, you learned:

- How to use the TypeScript and Python SDKs
- Common patterns for queries, settings, and entities
- Building custom CLI tools and reporting scripts
- Best practices for SDK usage

**Key Takeaway:** SDKs provide the most flexibility for custom applications. Use them when you need programmatic access, complex logic, or integration with other systems. For simple config management, Monaco or Terraform are often simpler.

---

*Continue to AUTOM-07: CI/CD Integration to learn GitOps patterns.*

---

*This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.*