

CLOUD-99: Best Practice Summary

Series: CLOUD | Notebook: 99 | Created: March 2026 | Last Updated: 03/26/2026

Overview

This notebook distills every actionable best practice from the CLOUD series (CLOUD-01 through CLOUD-08) into definitive, priority-ranked tables. Each entry specifies the exact setting, value, or action to take. No hedging — these are the recommended configurations.

Table of Contents

- 1. [Connection & Authentication](#)
- 2. [AWS Integration](#)
- 3. [Azure Integration](#)
- 4. [GCP Integration](#)
- 5. [Kubernetes \(EKS / AKS / GKE\)](#)
- 6. [Serverless Monitoring](#)
- 7. [Log Ingestion](#)
- 8. [Cost Optimization](#)
- 9. [Multi-Cloud Governance](#)
- 10. [Alerting & Observability](#)

1. Connection & Authentication

#	Best Practice	Recommended Setting/Value	Priority	Source
1	Use Clouds app direct connections for SaaS	Clouds app > Add connection — no ActiveGate required	Critical	CLOUD-01
2	Reserve classic ActiveGate polling for Managed only	ActiveGate-based polling for Dynatrace Managed or strict network isolation only	Critical	CLOUD-01

#	Best Practice	Recommended Setting/Value	Priority	Source
3	AWS: Use IAM role-based auth (STS AssumeRole)	IAM Role with trust policy to Dynatrace AWS account; never use long-lived access keys in production	Critical	CLOUD-02
4	AWS: Attach ReadOnlyAccess or custom least-privilege policy	Never grant AdministratorAccess to the monitoring role	Critical	CLOUD-02
5	Azure: Use Azure Native Dynatrace Service	Deploy from Azure Marketplace for zero-infrastructure setup with unified billing	Recommended	CLOUD-05
6	Azure: Use managed identity when ActiveGate runs on Azure VM	Eliminates credential rotation entirely	Critical	CLOUD-05
7	Azure: Assign Reader role at subscription scope	Scope: subscription or management group; role: Reader or Monitoring Reader	Critical	CLOUD-05
8	Azure: Use certificate-based auth for production (non-native)	Client secret for dev only; certificates for production Entra ID app registrations	Recommended	CLOUD-05
9	GCP: Create dedicated service account with minimal roles	Assign roles/monitoring.viewer , roles/compute.viewer , roles/container.viewer , roles/cloudasset.viewer	Critical	CLOUD-06
10	GCP: Use dedicated project for monitoring resources	Isolate Dynatrace service accounts, Pub/Sub topics, and billing in a separate GCP project	Recommended	CLOUD-06
11	GCP: Use Helm on GKE for push-based integration	Deploy via Pub/Sub for metrics and logs; no ActiveGate needed for SaaS	Recommended	CLOUD-06

2. AWS Integration

#	Best Practice	Recommended Setting/Value	Priority	Source
12	Enable services in tiers, not all at once	Tier 1 (always): EC2, ELB, RDS, Lambda. Tier 2: S3, SQS, SNS, DynamoDB, ECS/EKS. Tier 3: CloudFront, Route 53, Kinesis, Step Functions. Tier 4: SageMaker, Bedrock, IoT	Critical	CLOUD-02
13	Use tag-based resource filtering	Tag key: <code>dynatrace-monitored</code> , value: <code>true</code> ; only monitored resources are discovered	Critical	CLOUD-02
14	Scope monitoring to active regions only	Disable regions with no workloads in the Clouds app connection settings	Recommended	CLOUD-02
15	Account for CloudWatch metric delays in alert windows	Standard metrics: 5–10 min delay. Set alert evaluation window $\geq$ 15 min for cloud metrics	Critical	CLOUD-02
16	Enable CloudWatch detailed monitoring for critical EC2	Detailed monitoring: 1-min granularity (additional cost); standard: 5-min	Recommended	CLOUD-02
17	Combine cloud integration with OneAgent on compute	Cloud integration = cloud context (tags, AZ, instance type); OneAgent = deep observability (traces, code-level). Use both	Critical	CLOUD-01
18	Review CloudWatch billing dashboard monthly	Correlate API cost spikes with Dynatrace service additions	Recommended	CLOUD-02

3. Azure Integration

#	Best Practice	Recommended Setting/Value	Priority	Source
19	Map Azure resource groups to Dynatrace Segments	One Segment per resource group or application for access control and scoped views	Critical	CLOUD-05
20	Propagate Azure tags to Dynatrace	Tags must flow through automatically; verify <code>environment</code> , <code>team</code> , <code>cost-center</code> appear in Dynatrace	Recommended	CLOUD-05

#	Best Practice	Recommended Setting/Value	Priority	Source
21	Use naming convention for resource groups	Pattern: <code>rg-<application>-<environment>-<region></code> (e.g., <code>rg-myapp-prod-eastus</code>)	Recommended	CLOUD-05
22	Forward Azure logs via Event Hub or Diagnostic Settings	Use Azure Diagnostic Settings to stream to Dynatrace; Event Hub for high-volume	Recommended	CLOUD-07

4. GCP Integration

#	Best Practice	Recommended Setting/Value	Priority	Source
23	Use Workload Identity for GKE pods	Replace service account JSON keys with Workload Identity for DynaKube and integration pods	Critical	CLOUD-06
24	Tag all GCP resources with standard labels	Required labels: <code>env</code> , <code>team</code> , <code>service</code> , <code>cost-center</code> ; propagate as Dynatrace tags	Recommended	CLOUD-06
25	Propagate GCP project ID as a Dynatrace tag	Tag key: <code>gcp.project</code> ; enables cross-project filtering and Segment creation	Recommended	CLOUD-06
26	Create Dynatrace Segments per GCP project	One Segment per project for access control, cost attribution, and scoped dashboards	Recommended	CLOUD-06
27	Monitor Pub/Sub message volume for cost control	Filter logs at Cloud Logging export to reduce Pub/Sub throughput and costs	Recommended	CLOUD-06
28	Use Cloud Monitoring metrics scope for multi-project	Aggregate metrics from multiple projects into a single metrics scope rather than separate integrations	Optional	CLOUD-06

5. Kubernetes (EKS / AKS / GKE)

#	Best Practice	Recommended Setting/Value	Priority	Source
29	Deploy DynaKube in CloudNativeFullStack mode for standard node pools	<code>spec.oneAgent.cloudNativeFullStack</code> — full agent injection + infrastructure monitoring	Critical	CLOUD-03

#	Best Practice	Recommended Setting/Value	Priority	Source
30	Use ApplicationMonitoring mode for Fargate and GKE Autopilot	<code>spec.oneAgent.applicationMonitoring</code> with <code>useCSIDriver: false</code> for Fargate	Critical	CLOUD-03, CLOUD-06
31	Add tolerations for tainted node groups	<code>tolerations: [{effect: NoSchedule, key: node-role, operator: Exists}]</code> to ensure OneAgent runs on all nodes	Critical	CLOUD-03
32	Enable kubernetes-monitoring capability on ActiveGate	<code>spec.activeGate.capabilities: [kubernetes-monitoring, routing]</code>	Critical	CLOUD-03
33	Use nodeSelector to exclude Windows nodes	<code>nodeSelector: {kubernetes.io/os: linux}</code> in DynaKube spec	Recommended	CLOUD-03
34	Enforce namespace labels for cost allocation	Require <code>team</code> , <code>cost-center</code> , <code>environment</code> labels on all namespaces	Recommended	CLOUD-03
35	Set resource requests and limits on all pods	Accurate requests enable fair cost attribution and prevent resource contention	Recommended	CLOUD-03
36	Use Dynatrace for workload monitoring, Container Insights for control plane	Keep CloudWatch Container Insights / Azure Monitor for control plane logs; Dynatrace for distributed tracing and AI alerting	Recommended	CLOUD-03
37	Use IRSA (AWS) or Workload Identity (GCP) for pod IAM	Never mount service account JSON keys into pods; use cloud-native identity federation	Critical	CLOUD-03, CLOUD-06

6. Serverless Monitoring

#	Best Practice	Recommended Setting/Value	Priority	Source
38	Install Dynatrace Lambda Layer on all Lambda functions	Enables distributed tracing, code-level visibility, and direct log collection	Critical	CLOUD-04

#	Best Practice	Recommended Setting/Value	Priority	Source
39	Use meaningful Lambda function names	Avoid auto-generated names; use <code><service>-<function>-<env></code> pattern for identification	Recommended	CLOUD-04
40	Monitor concurrency headroom	Alert when <code>cloud.aws.lambda.concurrentExecutions</code> exceeds 80% of account limit (default: 1,000/region)	Critical	CLOUD-04
41	Track cold start ratio via spans	Query <code>faas.coldstart == true</code> from spans; high ratio indicates need for provisioned concurrency or SnapStart	Recommended	CLOUD-04
42	Use error-to-invocation ratio, not raw error count	<code>error_pct = errors / invocations * 100</code> ; alert threshold: > 1% for critical functions	Critical	CLOUD-04
43	Correlate API Gateway latency with Lambda duration	End-to-end latency = API GW overhead + Lambda duration; monitor both	Recommended	CLOUD-04
44	Monitor DynamoDB call duration from Lambda spans	Filter: <code>span.kind == "client"</code> , <code>db.system == "dynamodb"</code> ; alert on p95 > baseline	Recommended	CLOUD-04
45	Set alert evaluation window longer than metric delay	Cloud metric delay: 5–10 min. Minimum evaluation window for Lambda alerts: 15 min	Critical	CLOUD-02

7. Log Ingestion

#	Best Practice	Recommended Setting/Value	Priority	Source
46	Use Amazon Data Firehose for CloudWatch log forwarding	Fully managed, auto-scaling, no custom code. Buffer: 1 MB or 60 seconds. Enable GZIP compression	Critical	CLOUD-07

#	Best Practice	Recommended Setting/Value	Priority	Source
47	Migrate legacy Lambda log forwarder to Firehose	<code>dynatrace-aws-log-forwarder</code> is deprecated — do not use for new deployments	Critical	CLOUD-07
48	Use Lambda Layer log collection for Lambda functions	When Lambda Layer is deployed for tracing, use its built-in log collection instead of Firehose for Lambda logs	Recommended	CLOUD-07
49	Apply subscription filter patterns at CloudWatch level	Filter pattern examples: <code>?"ERROR"</code> <code>?"WARN"</code> <code>?"CRITICAL"</code> for Lambda; <code>""</code> (all) for application services	Critical	CLOUD-07
50	Drop debug/trace logs at source	Pre-filter: exclude DEBUG and TRACE in CloudWatch subscription filters. Estimated savings: 40–60%	Critical	CLOUD-07
51	Filter health check log lines at source	Remove <code>/health</code> , <code>/ready</code> , <code>/live</code> endpoint noise before forwarding. Savings: 10–30%	Recommended	CLOUD-07
52	Use OpenPipeline for enrichment, routing, and fine-grained filtering	Add <code>cloud.provider</code> , route by log group to specific Grail buckets, drop remaining noise	Recommended	CLOUD-07
53	Route logs to purpose-specific Grail buckets	Route <code>/aws/lambda/*</code> → <code>lambda_logs</code> , <code>/ecs/*</code> → <code>application_logs</code> , <code>/aws/rds/*</code> → <code>database_logs</code>	Recommended	CLOUD-07
54	Enable S3 backup on Firehose delivery stream	Backup failed or all records to S3 for disaster recovery and compliance	Recommended	CLOUD-07
55	Azure: Use Event Hub or Diagnostic Settings for log forwarding	Stream activity and resource logs directly to Dynatrace	Recommended	CLOUD-07
56	GCP: Use Pub/Sub via GKE integration for log forwarding	Push-based delivery; filter at Cloud Logging export level to reduce volume	Recommended	CLOUD-07

8. Cost Optimization

#	Best Practice	Recommended Setting/Value	Priority	Source
57	Disable unused cloud service monitoring	Only enable AWS/Azure/GCP services your applications depend on; disable the rest	Critical	CLOUD-02
58	Tag-based filtering is the #1 cost lever	Monitor only resources tagged <code>dynatrace-monitored: true</code> ; reduces both DPS and cloud API costs	Critical	CLOUD-02
59	Disable individual metrics you do not need	Per-service metric selection in Clouds app or Settings UI	Recommended	CLOUD-02
60	Consider CloudWatch Metric Streams at scale	Streaming can be cheaper than polling for >100 resources with frequent metric retrieval	Optional	CLOUD-02
61	Set bucket-level retention policies in Grail	Low-value logs: 7–14 days. Standard logs: 35 days. Compliance logs: 365+ days	Critical	CLOUD-07
62	Monitor log volume weekly	Query: <code>fetch logs, from:-7d \ makeTimeseries log_count = count(), interval:1h</code> — investigate spikes immediately	Recommended	CLOUD-07
63	Rightsize overprovisioned instances	Alert threshold: CPU < 10% sustained over 24h. Action: downsize or terminate	Recommended	CLOUD-08
64	Identify zero-traffic services for decommission	Query services with no spans/requests over 7 days; flag as decommission candidates	Optional	CLOUD-08
65	Enforce K8s resource requests/limits	Pods without requests/limits cause uncontrolled resource usage and inaccurate cost allocation	Recommended	CLOUD-03, CLOUD-08

9. Multi-Cloud Governance

#	Best Practice	Recommended Setting/Value	Priority	Source
66	Apply consistent tag schema across all providers	Required tags: <code>cloud-provider</code> , <code>environment</code> , <code>team</code> , <code>application</code> , <code>cost-center</code> , <code>region</code>	Critical	CLOUD-08
67	Normalize region names across providers	Use normalized values: <code>us-east</code> , <code>eu-west</code> , <code>ap-south</code> — not provider-specific names	Recommended	CLOUD-08
68	Create Dynatrace Segments per cloud provider + environment	One Segment per combination (e.g., <code>AWS-Production</code> , <code>Azure-Staging</code>) for access control	Critical	CLOUD-08
69	Enforce tagging compliance with automation	Reject or alert on untagged cloud resources; query: <code>fetch dt.entity.host \\\ filter isNull(tags)</code>	Critical	CLOUD-08
70	Define SLOs at the service level, not infrastructure level	SLOs are cloud-agnostic: measure user experience (latency, error rate, availability)	Recommended	CLOUD-08
71	Route alerts by team, not by cloud provider	Dynatrace Workflows route to the application owner regardless of which cloud has the issue	Critical	CLOUD-08
72	Use unified runbooks referencing DQL queries	Runbooks should use Dynatrace DQL, not provider-specific CLI tools	Recommended	CLOUD-08
73	Conduct monthly utilization reviews across all providers	Review underutilized hosts (CPU < 10%), untagged resources, and zero-traffic services	Recommended	CLOUD-08

10. Alerting & Observability

#	Best Practice	Recommended Setting/Value	Priority	Source
74	Use Dynatrace Intelligence for automatic cross-cloud anomaly detection	Enable Anomaly Detection on all monitored entities; do not rely solely on static thresholds	Critical	CLOUD-08

#	Best Practice	Recommended Setting/Value	Priority	Source
75	Set alert evaluation windows ≥ 15 min for cloud metrics	Cloud metric delay: 5–15 min. A 1-min evaluation window on 5-min metrics produces false negatives	Critical	CLOUD-02
76	Use OneAgent metrics for sub-minute alerting	For alerts requiring < 5 min evaluation, use <code>dt.host.*</code> or <code>dt.process.*</code> metrics from OneAgent, not cloud metrics	Critical	CLOUD-02
77	Combine cloud integration with OneAgent on all compute	Cloud integration: infrastructure context. OneAgent: processes, traces, code-level diagnostics. Both required for full-stack visibility	Critical	CLOUD-01
78	Build unified health dashboards showing all providers	Components: host CPU (all providers), active detected problems, service error rates, log error trends, K8s container metrics	Recommended	CLOUD-08
79	Deduplicate alerts using Dynatrace Intelligence correlation	Dynatrace Intelligence automatically correlates related issues across providers; do not create redundant static alerts	Recommended	CLOUD-08
80	Forward control plane logs to Dynatrace for unified analysis	EKS: Container Insights logs. AKS: Azure Monitor logs. GKE: Cloud Logging. Forward all to Grail for cross-platform analysis	Recommended	CLOUD-03, CLOUD-06

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.