

DBMON-05: Query Analysis

Series: DBMON | **Notebook:** 5 of 6 | **Created:** March 2026 | **Last Updated:** 03/12/2026

Overview

This notebook provides deep query analysis techniques for identifying database performance problems. You will learn how to detect N+1 query anti-patterns, analyze query frequency and duration distributions, identify queries that may lack proper indexes, detect connection pool exhaustion, and use query normalization for effective grouping. These techniques apply across all database technologies.

Table of Contents

- [1. N+1 Query Detection](#)
 - [2. Query Frequency Analysis](#)
 - [3. Duration Distribution Analysis](#)
 - [4. Missing Index Detection](#)
 - [5. Connection Pool Analysis](#)
 - [6. Query Normalization and Grouping](#)
 - [7. Optimization Prioritization](#)
 - [8. Summary and Next Steps](#)
-

Prerequisites

Requirement	Details
Dynatrace Environment	SaaS or Managed with Grail enabled
OneAgent	Deployed on application hosts with database clients
Permissions	<code>storage:spans:read</code>
Data	Significant application traffic with diverse database query patterns
Prior Reading	DBMON-01 through DBMON-03 for database monitoring context

1. N+1 Query Detection

The N+1 query problem occurs when an application executes one query to fetch a list of items, then executes N additional queries to fetch related data for each item. This is the most common database anti-pattern and can cause severe performance degradation.

Detection Strategy

N+1 patterns have these characteristics:

- A single normalized query pattern is executed many times within a short window
- The queries share the same parent span (same HTTP request or transaction)
- Individual query execution times are fast, but total time is high due to volume

We detect this by looking for query patterns with unusually high call counts per trace.

```
// Detect N+1 patterns – query patterns with high repetition per trace
fetch spans, from:-1h
| filter isNotNull(db.system) and isNotNull(db.statement)
| summarize calls_per_trace = count(),
             total_ms = sum(duration) / 1000000.0,
             avg_ms = avg(duration) / 1000000.0,
             by:{trace.id, db.statement, db.system}
| filter calls_per_trace >= 10
| sort calls_per_trace desc
| limit 25
```

The query above shows query patterns that are repeated 10 or more times within a single trace. A high `calls_per_trace` value (especially 50+) strongly suggests an N+1 pattern.

Tip: Focus on patterns where `avg_ms` is low but `total_ms` is high. This indicates many fast individual queries that accumulate to a significant total. The fix is typically to replace N individual queries with a single batch query using `IN` clauses or joins.

```
// Aggregate N+1 candidates – which query patterns are most frequently
repeated?
fetch spans, from:-1h
| filter isNotNull(db.system) and isNotNull(db.statement)
| summarize calls_per_trace = count(), by:{trace.id, db.statement, db.system}
| filter calls_per_trace >= 10
| summarize affected_traces = count(),
             avg_calls_per_trace = avg(calls_per_trace),
             max_calls_per_trace = max(calls_per_trace),
             by:{db.statement, db.system}
| sort affected_traces desc
| limit 15
```

2. Query Frequency Analysis

Understanding which queries run most often helps prioritize optimization. A query that runs millions of times per hour has far more impact than a slow query that runs once a day.

```
// Top 20 most frequent queries – highest call volume
fetch spans, from:-1h
| filter isNotNull(db.system) and isNotNull(db.statement)
| summarize call_count = count(),
              total_time_ms = sum(duration) / 1000000.0,
              avg_ms = avg(duration) / 1000000.0,
              p95_ms = percentile(duration, 95) / 1000000.0,
              by:{db.statement, db.system}
| sort call_count desc
| limit 20
```

```
// Query frequency trend – detect changing patterns over 6 hours
fetch spans, from:-6h
| filter isNotNull(db.system)
| makeTimeseries total_queries = count(),
                  unique_patterns = countDistinct(db.statement),
                  interval:15m
```

```
// Queries called by the most services – shared queries are high-impact
optimization targets
fetch spans, from:-1h
| filter isNotNull(db.system) and isNotNull(db.statement)
| summarize call_count = count(),
              calling_services = countDistinct(dt.entity.service),
              avg_ms = avg(duration) / 1000000.0,
              by:{db.statement, db.system}
| filter calling_services >= 2
| sort calling_services desc
| limit 15
```

3. Duration Distribution Analysis

Analyzing how query durations are distributed reveals bimodal patterns (e.g., cached vs uncached responses), tail latency problems, and outlier behavior.

```
// Duration percentile breakdown by database system
fetch spans, from:-1h
| filter isNotNull(db.system)
| summarize p50_ms = percentile(duration, 50) / 1000000.0,
             p90_ms = percentile(duration, 90) / 1000000.0,
             p95_ms = percentile(duration, 95) / 1000000.0,
             p99_ms = percentile(duration, 99) / 1000000.0,
             max_ms = max(duration) / 1000000.0,
             total = count(),
             by:{db.system}
| fieldsAdd tail_ratio = round(p99_ms / p50_ms, decimals:1)
| sort total desc
```

The `tail_ratio` column shows how much worse the 99th percentile is compared to the median. A ratio above 10 suggests significant tail latency — a small percentage of queries take disproportionately long.

```
// Latency bucket distribution – visualize where queries cluster
fetch spans, from:-1h
| filter isNotNull(db.system)
| fieldsAdd duration_ms = duration / 1000000.0
| fieldsAdd bucket = if(duration_ms < 0.1, then:"<0.1ms",
                        else:if(duration_ms < 1, then:"0.1-1ms",
                                else:if(duration_ms < 10, then:"1-10ms",
                                        else:if(duration_ms < 100, then:"10-100ms",
                                                else:if(duration_ms < 1000, then:"100ms-1s",
                                                        else:">1s")))))
| summarize query_count = count(), by:{db.system, bucket}
| sort db.system asc, bucket asc
```

4. Missing Index Detection

Queries that consistently have high latency and high call count may indicate missing database indexes. While Dynatrace cannot directly inspect index usage, we can identify candidates by correlating query patterns with performance characteristics.

Heuristics for Missing Indexes

Signal	What It Suggests
SELECT with high avg latency and high volume	Table scan on frequently-read table
P95 >> P50 for the same query pattern	Some executions hit different data sizes
Query latency increasing over time	Table growing without proper indexing
High variance in execution time	Inconsistent query plans

```
// Index candidate detection – high-volume queries with high latency and variance
fetch spans, from:-1h
| filter isNotNull(db.system) and isNotNull(db.statement)
| filter isNotNull(db.operation) and db.operation == "SELECT"
| summarize call_count = count(),
              avg_ms = avg(duration) / 1000000.0,
              p50_ms = percentile(duration, 50) / 1000000.0,
              p95_ms = percentile(duration, 95) / 1000000.0,
              stddev_ms = stddev(duration) / 1000000.0,
              by:{db.statement, db.system}
| filter call_count >= 50 and avg_ms > 10
| fieldsAdd variance_ratio = round(stddev_ms / avg_ms, decimals:2)
| sort avg_ms desc
| limit 15
```

```
// Detect queries getting slower over time – potential index degradation
fetch spans, from:-24h
| filter isNotNull(db.system) and isNotNull(db.statement)
| filter isNotNull(db.operation) and db.operation == "SELECT"
| makeTimeseries avg_ms = avg(duration) / 1000000.0,
                  call_count = count(),
                  by:{db.system},
                  interval:1h
```

5. Connection Pool Analysis

Connection pool exhaustion causes application threads to wait for available connections, leading to increased response times and timeouts. We detect connection pressure through concurrency analysis and error pattern monitoring.

```
// Concurrent database calls per service – identify connection pressure
fetch spans, from:-1h
| filter isNotNull(db.system)
| summarize concurrent_calls = count(),
    unique_db_targets = countDistinct(server.address),
    avg_ms = avg(duration) / 1000000.0,
    by:{dt.entity.service, db.system}
| fieldsAdd service_name = entityName(dt.entity.service,
type:"dt.entity.service")
| sort concurrent_calls desc
| limit 20
```

```
// Database connection errors – timeout and connection refused patterns
fetch spans, from:-6h
| filter isNotNull(db.system) and otel.status_code == "ERROR"
| summarize error_count = count(),
    by:{db.system, server.address, otel.status_message}
| sort error_count desc
| limit 20
```

```
// Database error rate trend – detect connection pool exhaustion episodes
fetch spans, from:-6h
| filter isNotNull(db.system)
| makeTimeseries total = count(),
    errors = countIf(otel.status_code == "ERROR"),
    interval:5m
```

6. Query Normalization and Grouping

Dynatrace automatically normalizes SQL queries by replacing literal values with `?` placeholders. This allows grouping identical query patterns regardless of parameter values. Understanding normalization helps you interpret `db.statement` values correctly.

Normalization Examples

Original Query	Normalized (db.statement)
SELECT * FROM users WHERE id = 42	SELECT * FROM users WHERE id = ?
INSERT INTO orders (id, amount) VALUES (1, 99.50)	INSERT INTO orders (id, amount) VALUES (?, ?)
UPDATE products SET stock = 5 WHERE sku = 'ABC123'	UPDATE products SET stock = ? WHERE sku = ?

```
// Query pattern diversity – how many unique normalized patterns per database?
fetch spans, from:-1h
| filter isNotNull(db.system) and isNotNull(db.statement)
| summarize unique_patterns = countDistinct(db.statement),
              total_calls = count(),
              by:{db.system, db.namespace}
| fieldsAdd avg_calls_per_pattern = round(toDouble(total_calls) /
toDouble(unique_patterns), decimals:1)
| sort total_calls desc
```

```
// Identify one-off queries – patterns called only once (may indicate dynamic SQL)
fetch spans, from:-1h
| filter isNotNull(db.system) and isNotNull(db.statement)
| summarize call_count = count(), by:{db.statement, db.system}
| filter call_count == 1
| summarize one_off_count = count(), by:{db.system}
| sort one_off_count desc
```

Warning: A high number of one-off query patterns suggests dynamic SQL generation or poor query parameterization. This prevents the database from caching execution plans and can degrade performance significantly.

7. Optimization Prioritization

Not all slow queries are worth optimizing. Use the **impact score** to prioritize: combine call frequency, average duration, and error rate into a single ranking.


```
// Query optimization priority – ranked by total time impact
fetch spans, from:-1h
| filter isNotNull(db.system) and isNotNull(db.statement)
| summarize call_count = count(),
             total_time_ms = sum(duration) / 1000000.0,
             avg_ms = avg(duration) / 1000000.0,
             p95_ms = percentile(duration, 95) / 1000000.0,
             error_count = countIf(otel.status_code == "ERROR"),
             by:{db.statement, db.system}
| fieldsAdd error_rate_pct = round((toDouble(error_count) /
toDouble(call_count)) * 100, decimals:2)
| sort total_time_ms desc
| limit 20
```

The `total_time_ms` column represents the cumulative time spent on each query pattern. Optimizing the top entries yields the greatest overall performance improvement.

Priority	Criteria	Action
Critical	High total time + high error rate	Fix errors first, then optimize
High	High total time + high call count + low avg	Likely N+1 — batch the queries
Medium	High total time + low call count + high avg	Likely missing index — add index
Low	Low total time regardless of avg	Minimal impact — defer optimization

8. Summary and Next Steps

In this notebook you learned:

- How to detect N+1 query anti-patterns using per-trace call count analysis
- Query frequency analysis to identify the most common database operations
- Duration distribution analysis to reveal tail latency and bimodal patterns
- Heuristic-based missing index detection using latency and variance metrics
- Connection pool pressure analysis through concurrency and error monitoring
- Query normalization concepts and one-off query detection

- Impact-based optimization prioritization

Next Steps

- **DBMON-06: Dashboards and Alerting** — Build database monitoring dashboards and configure alerting rules
-

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.