

DBMON-06: Dashboards and Alerting

Series: DBMON | **Notebook:** 6 of 6 | **Created:** March 2026 | **Last Updated:** 03/12/2026

Overview

This notebook covers building database monitoring dashboards and configuring alerting for database health. You will learn dashboard design patterns for database KPIs, how to create alert-ready queries for slow query detection, connection pool thresholds, error rate monitoring, and database-specific SLO definitions. These queries are designed to be directly usable in Dynatrace dashboards and metric events.

Table of Contents

1. [Database KPIs](#)
 2. [Health Overview Dashboard](#)
 3. [Response Time Monitoring](#)
 4. [Error Rate Alerting](#)
 5. [Throughput and Capacity](#)
 6. [Slow Query Alerting](#)
 7. [Database SLO Definitions](#)
 8. [Summary](#)
-

Prerequisites

Requirement	Details
Dynatrace Environment	SaaS or Managed with Grail enabled
OneAgent	Deployed on application hosts with database clients
Permissions	storage:spans:read , storage:metrics:read , storage:entities:read
Data	Active database traffic across monitored services
Prior Reading	DBMON-01 through DBMON-05 for context on database monitoring techniques

1. Database KPIs

Effective database monitoring revolves around a small set of key performance indicators. Each KPI maps to a specific dashboard tile and potential alert condition.

KPI	What It Measures	Dashboard Tile	Alert Threshold
Response Time (P95)	Query latency at the 95th percentile	Timeseries chart	> 500ms sustained
Error Rate	Percentage of failed database calls	Single value + trend	> 1% of total calls
Throughput	Queries per second	Timeseries chart	Sudden drop > 50%
Slow Query Count	Queries exceeding threshold	Counter + list	> 10 per 5min window
Connection Errors	Timeout and refused connections	Counter	Any occurrence
Top Queries by Time	Highest-impact query patterns	Table	Total time > SLO budget

2. Health Overview Dashboard

The health overview provides a single-pane summary of all database systems. Use these queries as dashboard tiles for an executive-level view.

```
// Dashboard tile: Database health summary – one row per database system
fetch spans, from:-1h
| filter isNotNull(db.system)
| summarize total_calls = count(),
             avg_ms = avg(duration) / 1000000.0,
             p95_ms = percentile(duration, 95) / 1000000.0,
             error_count = countIf(otel.status_code == "ERROR"),
             slow_count = countIf(duration > 500000000),
             unique_databases = countDistinct(db.namespace),
             by:{db.system}
| fieldsAdd error_rate_pct = round((toDouble(error_count) /
toDouble(total_calls)) * 100, decimals:2)
| fieldsAdd slow_rate_pct = round((toDouble(slow_count) /
toDouble(total_calls)) * 100, decimals:2)
| sort total_calls desc
```

```
// Dashboard tile: Total database calls – single value
fetch spans, from:-1h
| filter isNotNull(db.system)
| summarize total_db_calls = count()
```

```
// Dashboard tile: Database calls per service – identify heaviest consumers
fetch spans, from:-1h
| filter isNotNull(db.system)
| summarize call_count = count(),
             avg_ms = avg(duration) / 1000000.0,
             error_count = countIf(otel.status_code == "ERROR"),
             by:{dt.entity.service}
| fieldsAdd service_name = entityName(dt.entity.service,
type:"dt.entity.service")
| sort call_count desc
| limit 10
```

3. Response Time Monitoring

Response time monitoring is the most critical aspect of database dashboards. These queries provide both real-time and trend views.

```
// Dashboard tile: Database response time trend by system (6-hour view)
fetch spans, from:-6h
| filter isNotNull(db.system)
| makeTimeseries p95_ms = percentile(duration, 95) / 1000000.0,
                  by:{db.system},
                  interval:5m
```

```
// Dashboard tile: Response time by database instance (server.address)
fetch spans, from:-1h
| filter isNotNull(db.system) and isNotNull(server.address)
| summarize call_count = count(),
              avg_ms = avg(duration) / 1000000.0,
              p95_ms = percentile(duration, 95) / 1000000.0,
              p99_ms = percentile(duration, 99) / 1000000.0,
              by:{db.system, server.address, db.namespace}
| sort p95_ms desc
| limit 15
```

```
// Dashboard tile: P50 vs P95 vs P99 comparison – all databases combined
fetch spans, from:-6h
| filter isNotNull(db.system)
| makeTimeseries p50_ms = percentile(duration, 50) / 1000000.0,
                  p95_ms = percentile(duration, 95) / 1000000.0,
                  p99_ms = percentile(duration, 99) / 1000000.0,
                  interval:5m
```

4. Error Rate Alerting

Database errors (connection timeouts, deadlocks, constraint violations) should trigger alerts when they exceed normal baseline levels.

```
// Alert query: Database error rate per 5-minute window
fetch spans, from:-1h
| filter isNotNull(db.system)
| makeTimeseries total = count(),
                  errors = countIf(otel.status_code == "ERROR"),
                  by:{db.system},
                  interval:5m
```

```
// Alert query: Error breakdown by type – classify error categories
fetch spans, from:-1h
| filter isNotNull(db.system) and otel.status_code == "ERROR"
| summarize error_count = count(),
             by:{db.system, server.address, otel.status_message}
| sort error_count desc
| limit 20
```

```
// Alert query: Error rate trend over 24 hours – detect escalating problems
fetch spans, from:-24h
| filter isNotNull(db.system)
| makeTimeseries total = count(),
                  errors = countIf(otel.status_code == "ERROR"),
                  interval:30m
```

Recommended Alert Thresholds

Condition	Severity	Alert When
Error rate > 5% for 5 minutes	Critical	Immediate page
Error rate > 1% for 15 minutes	Warning	Notification
Any connection refused error	Critical	Immediate page
Deadlock detected	Warning	Notification

5. Throughput and Capacity

Monitoring query throughput helps detect traffic anomalies and plan for capacity. A sudden drop in throughput may indicate an upstream failure, while a spike may indicate a runaway process.

```
// Dashboard tile: Queries per minute by database system
fetch spans, from:-6h
| filter isNotNull(db.system)
| makeTimeseries qpm = count(), by:{db.system}, interval:1m
```

```
// Dashboard tile: Operation mix over time – read vs write trend
fetch spans, from:-6h
| filter isNotNull(db.system) and isNotNull(db.operation)
| fieldsAdd op_type = if(
  in(db.operation, {"SELECT", "find", "Query", "GetItem", "ReadItem",
"GET", "HGET", "get", "search"}),
  then:"READ",
  else:"WRITE")
| makeTimeseries op_count = count(), by:{op_type}, interval:5m
```

```
// Dashboard tile: Hourly query volume comparison – today vs yesterday
fetch spans, from:-24h
| filter isNotNull(db.system)
| summarize today_count = count()
| append [
  fetch spans, from:-48h, to:-24h
  | filter isNotNull(db.system)
  | summarize yesterday_count = count()
]
```

6. Slow Query Alerting

Slow query alerts detect when query performance degrades beyond acceptable thresholds. Configure these as metric events in Dynatrace.

```
// Alert query: Slow query count per 5-minute window (> 500ms threshold)
fetch spans, from:-1h
| filter isNotNull(db.system)
| filter duration > 500000000
| makeTimeseries slow_count = count(), by:{db.system}, interval:5m
```

```
// Alert query: Current slow query detail – for investigation
fetch spans, from:-15m
| filter isNotNull(db.system)
| filter duration > 500000000
| fields timestamp, db.system, db.namespace, db.operation,
    db.statement, server.address,
    duration_ms = duration / 1000000.0,
    dt.entity.service
| fieldsAdd service_name = entityName(dt.entity.service,
type:"dt.entity.service")
| sort duration_ms desc
| limit 20
```

```
// Alert query: P95 response time exceeding SLO threshold
fetch spans, from:-6h
| filter isNotNull(db.system)
| makeTimeseries p95_ms = percentile(duration, 95) / 1000000.0,
    by:{db.system, server.address},
    interval:5m
```

Slow Query Alert Thresholds

Database Type	Warning Threshold	Critical Threshold	Window
SQL (PostgreSQL, MySQL, etc.)	P95 > 200ms	P95 > 500ms	5 minutes
NoSQL (MongoDB, DynamoDB)	P95 > 100ms	P95 > 300ms	5 minutes
Cache (Redis, Memcached)	P95 > 5ms	P95 > 20ms	5 minutes
Search (Elasticsearch)	P95 > 500ms	P95 > 2000ms	5 minutes

7. Database SLO Definitions

Service Level Objectives (SLOs) for databases define the expected performance contract. Use these queries to measure SLO compliance.

Recommended Database SLOs

SLO	Target	Measurement
Availability	99.9% of calls succeed	Error rate < 0.1%
Latency	95% of calls under threshold	P95 < vendor-specific threshold
Throughput	No more than 50% drop from baseline	Queries/min vs 7-day average

```
// SLO measurement: Database availability – success rate over 24 hours
fetch spans, from:-24h
| filter isNotNull(db.system)
| summarize total = count(),
              successful = countIf(otel.status_code != "ERROR"),
              by:{db.system}
| fieldsAdd availability_pct = round((toDouble(successful) / toDouble(total))
* 100, decimals:3)
| sort availability_pct asc
```

```
// SLO measurement: Latency compliance – percentage of calls under threshold
fetch spans, from:-24h
| filter isNotNull(db.system)
| summarize total = count(),
              under_threshold = countIf(duration < 500000000),
              by:{db.system}
| fieldsAdd latency_slo_pct = round((toDouble(under_threshold) /
toDouble(total)) * 100, decimals:2)
| sort latency_slo_pct asc
```

```
// SLO trend: Hourly availability over 24 hours
fetch spans, from:-24h
| filter isNotNull(db.system)
| makeTimeseries total = count(),
                  errors = countIf(otel.status_code == "ERROR"),
                  by:{db.system},
                  interval:1h
```

8. Summary

In this notebook you learned:

- The key database KPIs to monitor: response time, error rate, throughput, slow queries, and connection health
- Dashboard-ready queries for a unified database health overview
- Response time monitoring with P50/P95/P99 trend analysis
- Error rate alerting queries and recommended thresholds
- Throughput monitoring including read/write ratio trends and day-over-day comparison
- Slow query alerting with vendor-specific thresholds
- Database SLO definitions and compliance measurement queries

Series Complete

This is the final notebook in the DBMON series. For a complete database monitoring implementation:

1. **DBMON-01** — Understand fundamentals and establish baselines
2. **DBMON-02** — Deep dive into SQL database monitoring
3. **DBMON-03** — NoSQL database monitoring patterns
4. **DBMON-04** — Cache and messaging system monitoring
5. **DBMON-05** — Advanced query analysis and optimization
6. **DBMON-06** — (This notebook) Dashboards, alerting, and SLOs

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.