

FAQ-09: When Should I Query a Metric Instead of Raw Logs?

Series: FAQ — Frequently Asked Questions | **Reference:** 09 — When to Query a Metric Instead of Raw Logs | **Created:** June 2026 | **Last Updated:** 06/16/2026

Overview

Logs are where the detail lives — the exact line, the stack trace, the request that failed. But most of what teams build *on top of* logs isn't detail; it's a number over time: *how many errors per minute, what's the p95 of this endpoint, is the 5xx rate climbing.*

Those questions get answered by running a `fetch logs` aggregation, and when that aggregation sits behind a dashboard tile or an alert, it runs again and again — every refresh, every evaluation interval, forever.

That is the optimization opportunity. In Dynatrace on Grail, **scanning logs is a billed activity** and **querying a metric effectively is not**. The same "errors per minute" answer costs real money when it re-scans logs on every dashboard render, and almost nothing when it reads a pre-aggregated metric. Out-of-the-box (OOTB) metrics already answer many of these questions; for the rest, you can extract a metric from the log stream once, at ingest, and point every recurring query at it instead.

This entry explains the economics behind that trade, gives you a one-question rule for deciding which to use, shows what's already available OOTB before you extract anything, and covers the two things that quietly undo the savings — cardinality and the fact that extraction doesn't backfill.

Table of Contents

1. [Short Answer](#)
2. [Why This Matters — The DPS Query Economics](#)

3. [The Decision — Recurring vs. One-Shot](#)
 4. [Use OOTB Metrics Before You Extract Anything](#)
 5. [Extracting a Metric From Logs](#)
 6. [Cardinality — The One Thing That Undoes the Savings](#)
 7. [Extraction Is Forward-Only — and the Retention Split](#)
 8. [Recommended Approach](#)
 9. [Common Gotchas](#)
-

Prerequisites

Requirement	Details
Audience	Platform engineers, SREs, and observability owners building dashboards, alerts, and SLOs on log data — and anyone watching DPS query consumption climb
Format	Decision-support document — explains the metric-vs-log-query trade and how to act on it; not a hands-on lab
Deployment	Dynatrace SaaS with Grail; Log Monitoring via OneAgent or OpenPipeline; DPS (Dynatrace Platform Subscription) licensing
Related topic series	OPLOGS-03 (OpenPipeline metric extraction from logs), OPLOGS-07 (analytics & dashboard queries), OPMIG-07 (metric & event extraction), OPIPE-03 (sampling-aware metrics from spans), OPIPE-04 (cardinality management), FINOPS-01 (querying DPS consumption), FINOPS-03 (Cut/Tune/Filter optimization), ORGNZ (bucket & retention strategy)
Related FAQ	FAQ-08 (how OneAgent decides which logs to collect — what's in the stream in the first place)

1. Short Answer

Ask one question of every log-based query: **how often does it run?**

If the query runs...	Use	Because
Repeatedly — dashboard tiles, alerts, anomaly detection, SLOs	A metric (OOTB if one exists, otherwise extracted from the log at ingest)	The aggregate is computed once at ingest and read cheaply thereafter
Once — root-cause analysis, forensic lookup, ad-hoc exploration	Raw logs (<code>fetch logs</code> with a tight time range and early filters)	You need the actual lines, and you only pay for the single scan

The decision is *not* "logs are bad, metrics are good." Logs carry the detail you cannot get any other way. The decision is about **where a query lives**: a `fetch logs` aggregation behind a dashboard re-scans the same data on every render, while a metric behind that same tile reads a tiny pre-aggregated series. Move the *recurring aggregate* to a metric; keep the *raw stream* for investigation.

Two moves implement this:

1. **Reach for an OOTB metric first** (§4) — Dynatrace already produces service, host, process, and consumption metrics that teams frequently re-derive from logs by hand.
2. **Extract a metric from the log at ingest** (§5) — when no OOTB metric fits, an OpenPipeline metric-extraction rule turns "count of matching log records, by a few dimensions" into a first-class metric, computed once as the data arrives.

Sources: [Extract metrics from log data \(DT docs\)](#), [DPS Log Management & Analytics \(DT docs\)](#).

Derived: the "how often does it run" rule is an authoring synthesis of the DPS billing model (§2) — the docs describe the capabilities billed, not this decision heuristic.

2. Why This Matters — The DPS Query Economics

Under DPS, each Grail data type is billed across distinct **capabilities**. The asymmetry that drives this whole entry is in the last column:

Signal (powered by Grail)	Ingest & Process	Retain	Query
Logs	billed (bytes)	billed (bytes)	billed (bytes scanned)

Signal (powered by Grail)	Ingest & Process	Retain	Query
Events	billed	billed	billed
Traces	billed	billed	billed
Metrics	billed (data points)	billed	— no Query capability —

Logs, events, and traces each carry a **Query** capability: when you scan them with DQL, you consume billable bytes (over and above the query volume bundled into *Retain with Included Queries*). **Metrics have no Query capability at all** — they bill on ingest and retention, and reading them back with `timeseries` does not draw a separate query charge.


Now put that against how dashboards and alerts behave. A dashboard tile re-runs its query on every load and every auto-refresh. An alert or anomaly detector re-evaluates on a fixed interval, continuously, for as long as it exists. A `fetch logs` aggregation in either of those places pays the log **Query** cost *every single time*. The identical answer sourced from a metric pays it *never*.

This is why the move is almost pure upside for recurring workloads: you are not losing the log data (it still flows, still retained per your bucket policy) — you are removing the repeated scan. In the FINOPS-03 **Cut / Tune / Filter** framework this is a *Tune* lever: you keep the signal and change only the surface you query against.

The flip side keeps you honest: for a **one-shot investigation**, a single `fetch logs` scan is cheap and a metric can't show you the offending line. Don't pre-build metrics for questions you ask once.

Sources: [DPS Log Management & Analytics \(DT docs\)](#), [DPS Metrics \(DT docs\)](#), [DPS Traces \(DT docs\)](#). **Derived:** the "Metrics have no Query capability" contrast is read from comparing the billed-capability lists on the DPS Logs vs. DPS Metrics pages — neither page states the comparison; FINOPS-01 documents the same capability split from a live tenant.

3. The Decision — Recurring vs. One-Shot



The same answer takes two very different cost paths depending on where the query lives:

```
// RECURRING — dashboard tile or alert. Reads a pre-aggregated metric.  
// No log Query capability consumed.  
timeseries sum(log.http.requests), from:-24h, by:{status}
```

```
// ONE-SHOT — root-cause investigation. Scans the raw log stream once.  
// Billed on bytes scanned, so keep the window tight and filter early.  
fetch logs, from:-1h  
| filter k8s.namespace.name == "prod" and loglevel == "ERROR"  
| summarize errors = count(), by:{dt.entity.host}
```

The first query is what belongs behind a tile that thousands of people load and a detector that fires every minute. The second is what you run when something broke and you need to see it. Both are correct — for their job. The anti-pattern is putting the *second* shape behind a dashboard, where it quietly re-scans logs forever.

A quick test for any existing tile or alert: **if its query has a `fetch logs ... | summarize` shape and it lives somewhere that re-runs on a schedule, it's a candidate to move to a metric.**

Sources: [makeTimeseries \(DT docs\)](#), [Aggregation functions \(DT docs\)](#), [Extract metrics from log data \(DT docs\)](#). OPLGS-07 covers the dashboard-query patterns in depth.

4. Use OOTB Metrics Before You Extract Anything

Before building an extraction rule, check whether Dynatrace already produces the metric. Teams routinely re-derive from logs numbers the platform emits for free:

You want...	Often re-derived from logs as...	OOTB metric already exists
Request throughput / error rate / latency	Counting request/error log lines	Service RED metrics — the <code>builtin:service.*</code> family (request count, failure rate, response time) from OneAgent-instrumented services
Host CPU / memory / disk	Parsing agent or OS logs	<code>dt.host.cpu.usage</code> , <code>dt.host.memory.usage</code> , and the <code>dt.host.*</code> family
Process resource use	Parsing process logs	The <code>dt.process.*</code> family
Log volume / ingest cost	<code>fetch logs \ summarize count() over a long window</code>	DPS consumption metrics (the <code>dt.billing.*</code> / <code>dt.system.events</code> surfaces — see FINOPS-01)
Did this log source stop sending?	Counting recent records	<code>log.source.file_status</code> / <code>log.source.ingest_status</code> (see FAQ-08)

How to check: open the metric browser / **Data explorer** and search the entity or domain before writing a pipeline rule. If an OOTB metric answers the question, you're done — zero extraction config, zero added cardinality, and it already has historical depth.

Exact metric keys vary by Dynatrace version and by what's deployed in your tenant — confirm the key in your own metric browser rather than assuming it. The point stands

regardless of the precise key: **the cheapest extraction is the one you don't have to do because the metric already exists.**

Sources: [Built-in metrics / metric browser \(DT docs\)](#), [DPS Metrics \(DT docs\)](#). FINOPS-01 documents the DPS consumption metric surfaces; FAQ-08 documents the `log.source.*` fields. **Derived:** the "re-derived from logs" column is a synthesis of common practice — the docs list the metrics, not the log-based work they replace.

5. Extracting a Metric From Logs

When no OOTB metric fits — a business count, an application-specific status, a value embedded in a log line — extract one at ingest with an **OpenPipeline metric-extraction processor**. The aggregate is computed once, as records arrive, and stored as a metric you then query cheaply forever.

```
# OpenPipeline log pipeline – metric-extraction processor
- name: extract-request-metric
  type: metric
  enabled: true
  condition: contains(content, "duration=")
  metricKey: log.request.duration      # becomes a first-class metric
  dimensions:
    - service: k8s.namespace.name      # bounded, low-cardinality
    - method: extracted_method
    - status: extracted_status
  value: extracted_duration_ms          # omit to count matching records
```

Common shapes:

Question	Metric	Dimensions
How many requests, by outcome?	<code>log.request.count</code> (count of matches)	service, method, status
What's the latency distribution?	<code>log.request.duration</code> (extract numeric field)	service, endpoint
How many of business event X?	<code>log.business.<event>.count</code>	a few business dimensions

A few framing points:

- **OpenPipeline is the modern path.** Classic Log Monitoring also supports log-metric definitions; new work should use OpenPipeline metric extraction. See **OPLOGS-03 §3** for the full processor walkthrough and **OPMIG-07** for the metric-&-event-extraction deep dive.
- **Log-derived metrics are exact, not sampled.** Span-derived metrics must be made *sampling-aware* because only a fraction of spans are kept (**OPIPE-03**); logs aren't head/tail-sampled, so a count extracted from logs reflects every matching record.
- **Count vs. value.** Omit `value` to count matching records; set it to a parsed numeric field to track a measurement. Parse the field first (DPL) so the value and dimensions exist on the record when the metric processor runs.

Sources: [Extract metrics from log data \(DT docs\)](#), [OpenPipeline processors \(DT docs\)](#), [Dynatrace Pattern Language \(DT docs\)](#). OPLOGS-03, OPMIG-07, and OPIPE-03 carry the implementation depth.

6. Cardinality — The One Thing That Undoes the Savings

A metric's cost scales with its **cardinality**: the number of unique dimension-value combinations it produces. Cardinality is the **product** of its dimensions' distinct values:

```
data points ≈ (distinct services) × (distinct methods) × (distinct statuses)
× ...
```

Three bounded dimensions (say 40 services × 5 methods × 6 statuses = 1,200 series) is cheap and useful. Add one **unbounded** dimension — `request.id`, `user.id`, `trace_id`, a raw URL with query strings — and the product explodes into the millions. At that point you've traded an expensive log query for an expensive metric, which is worse: you pay continuously and the metric is unusable.

Guardrails:

- **Only bounded dimensions on extracted metrics:** status code, service, method, region, environment, log level. These have a small, stable set of values.

- **Never** put per-request or per-user identifiers on a metric. If you need to pivot by them, that's a *log* question (one-shot, §3) — or a trace.
- **Reduce before you extract:** normalize URLs to route templates, bucket numeric values, drop the noisy dimension at the parse step so it never reaches the metric processor.

This is the single most common way a log-to-metric optimization backfires, and it has the same root cause every time: a high-cardinality label that someone added "just in case." **OPIPE-04 (Cardinality Management)** covers the controls in depth; **FINOPS-03 §9** walks a worked high-cardinality-metric remediation.

Sources: [DPS Metrics \(DT docs\)](#), [Extract metrics from log data \(DT docs\)](#). OPIPE-04 and FINOPS-03 carry the cardinality-control patterns. **Derived:** the cross-product cost model is standard dimensional-metric behavior applied to log-extracted metrics.

7. Extraction Is Forward-Only — and the Retention Split

Two operational facts shape how you roll this out.

Extraction is forward-only. An OpenPipeline metric-extraction rule acts on records *as they arrive*. A rule you create today begins producing its metric from today — it does **not** retroactively manufacture metric history from logs already ingested. Practical consequences:

- Stand up the extraction rule **before** you need the trend, not the day you build the dashboard.
- For the gap, query the raw logs directly (a one-shot scan over the historical window) until the metric accumulates enough depth.
- When migrating an existing dashboard from logs to a metric, run both in parallel for a bit and cut over once the metric's history covers your default dashboard window.

Split the retention. Once the recurring questions are answered by a metric, the raw logs no longer need to live long *for those questions*:

- Keep the **metric** retained long — it's small and cheap, so multi-month trend tiles stay inexpensive.

- Keep the **raw logs** in a shorter-retention bucket sized for investigation and compliance, not for powering dashboards.

That split is where the saving compounds: long-horizon dashboards read the cheap long-lived metric, while the expensive raw logs age out on a short clock. Bucket and retention design lives in the **ORGNZ** series; retention as a cost lever is **FINOPS-03 §4 (Tune)**.

Sources: [Extract metrics from log data \(DT docs\)](#), [DPS Log Management & Analytics \(DT docs\)](#).

ORGNZ covers bucket/retention design; FINOPS-03 covers retention tuning. **Derived:** the "forward-only" property follows from OpenPipeline processing at ingest time; the retention-split recommendation synthesizes the bucket model with the per-capability billing in §2.

8. Recommended Approach

1. **Inventory the recurring log queries.** List the dashboard tiles, alerts, anomaly detectors, and SLOs whose queries have a `fetch logs ... | summarize` shape. These are your candidates — they pay the log Query cost on every run.
2. **Check for an OOTB metric first (§4).** Search the metric browser for the service, host, process, or consumption metric that already answers the question. If it exists, repoint the tile/alert and you're done.
3. **Extract a metric for the rest (§5).** Build an OpenPipeline metric-extraction rule with a clear `metricKey` and a small set of **bounded** dimensions. Parse any needed value/dimension fields earlier in the pipeline.
4. **Hold the line on cardinality (§6).** No per-request, per-user, or per-trace dimensions. Sanity-check the dimension cross-product before enabling.
5. **Cut over forward, not backward (§7).** Enable the rule, let the metric accumulate, then migrate the tile/alert. Keep raw-log access for the historical gap and for investigation.
6. **Split retention (§7).** Retain the metric long; shorten the raw-log bucket to what investigation and compliance actually need.
7. **Keep raw logs for what they're for.** Root-cause analysis, forensics, and ad-hoc exploration stay on `fetch logs` with tight time ranges and early filters — that's the one-shot path, and it's cheap.

8. **Measure the result.** Use FINOPS-01's query-side consumption queries to confirm log Query consumption dropped after the cutover.

9. Common Gotchas

Symptom	Likely cause	Where
Log Query consumption keeps climbing	A dashboard/alert still runs <code>fetch logs \</code> <code>summarize</code> on a schedule	§2, §3
New extracted metric has no history	Extraction is forward-only — it doesn't backfill old logs	§7
Extracted metric is expensive / unusable	An unbounded dimension (request/user/trace ID) blew up cardinality	§6
"I lost the detail when I switched to a metric"	Metrics answer <i>how many/how fast</i> , not <i>which line</i> — keep raw logs for the detail	§1, §3
Built a metric that already existed OOTB	Skipped the metric-browser check before extracting	§4
Span-extracted metric undercounts vs. logs	Spans are sampled; make span metrics sampling-aware (log metrics aren't sampled)	§5, OPIPE-03
Trend tile still slow/costly over long windows	Raw logs retained long for dashboards instead of the cheap metric	§7
Dashboard cutover broke during the gap	Migrated before the metric had enough history	§7

Sources: [Extract metrics from log data \(DT docs\)](#), [DPS Log Management & Analytics \(DT docs\)](#), [DPS Metrics \(DT docs\)](#) — each row maps to the capability, rule, or constraint cited in the referenced section.

 **Disclaimer:** This content is AI-generated, community-driven, and **not supported by Dynatrace**. DPS capabilities, billed units, and OpenPipeline behavior evolve across releases —

always verify the current billing model and metric-extraction mechanics against the [Dynatrace documentation](#) and your own tenant before relying on a specific cost claim.