

IAM-04: Policy Authoring and Management

Series: IAM — IAM Administration | **Notebook:** 4 of 12 | **Created:** January 2026 |
Last Updated: 04/26/2026

Mastering Dynatrace Policy Syntax

Policies are the heart of Dynatrace Gen3 IAM. They define what actions users can perform. This notebook provides a comprehensive guide to policy authoring, from basic syntax to advanced patterns.

Table of Contents

1. [Policy Fundamentals](#)
 2. [Policy Statement Syntax](#)
 3. [Services and Actions Reference](#)
 4. [Condition Expressions](#)
 5. [Common Policy Patterns](#)
 6. [Default vs Custom Policies](#)
 7. [Policy Testing and Validation](#)
 8. [GitOps for Policies](#)
 9. [Policy Assessment Queries](#)
-

Prerequisites

Requirement	Details
Dynatrace Environment	SaaS with Gen3 IAM enabled
Permissions	account-iam-admin to create/modify policies
Prior Knowledge	IAM-01 through IAM-03

Sprint 1.337 (April 2026): Permission Updates

New `frontend.name` permission field for Grail RUM data. Sprint 1.337 introduced `frontend.name` as a permission-relevant field on `metric` and `dt.entity.*` Smartscape data. You can now scope IAM policies to specific frontend applications at the **record-level** and **field-level**, instead of relying on bucket-only isolation:

```
ALLOW storage:metric:read
  WHERE storage:metric.frontend.name == "checkout-web"
  OR    storage:metric.frontend.name == "account-web";

ALLOW storage:user.sessions:read
  WHERE storage:user.sessions.frontend.name STARTS WITH "public-";
```

Use this when multiple teams own different frontends but share a single RUM tenant — fine-grained ABAC without splitting buckets.

`dt.security_context` is now a primary field at source. OneAgent now enriches all telemetry with `dt.security_context` (and `dt.cost.*`) as **top-level primary fields** at ingestion on Latest Dynatrace tenants. The `MATCH()` ABAC patterns covered later in this notebook (and across IAM-04, IAM-05, IAM-07) work directly against the primary-field value with no need for OpenPipeline parse processors on OneAgent-instrumented data. The structured-context spread completed in PR #50 already aligns with this — the sprint-1.337 change makes the field guarantee tenant-wide rather than pipeline-dependent.

1. Policy Fundamentals

Policies define **what actions** users can perform. They're separate from boundaries (which control **what data** users can see).

Policy Scope Levels

Scope	Applies To	Managed By
Account	All environments in account	Account Admin
Environment	Single environment	Environment Admin

Policy Assignment

Policies are assigned to **groups**, not individual users:


```
Group: dt-checkout-editors
├─ Policy: environment-editor (default)
├─ Policy: log-management (custom)
└─ Boundary: checkout-services-only
```

Policy Types

Type	Description	Examples
Default Policies	Pre-built by Dynatrace	<code>environment-viewer</code> , <code>environment-editor</code>
Custom Policies	Created by you	<code>log-read-only</code> , <code>dashboard-manager</code>

2. Policy Statement Syntax

Every policy contains one or more **statements** that define allowed actions.



Basic Statement Structure

```
ALLOW <service>:<resource>:<action>;
```

Examples:

Statement	Meaning
ALLOW storage:logs:read	Read all logs
ALLOW storage:spans:read	Read all spans
ALLOW settings:objects:write	Modify settings
ALLOW document:documents:read	Read dashboards/notebooks

Wildcards

Use * to match multiple values:

Statement	Matches
ALLOW storage:*:read	Read any storage type (logs, spans, metrics)
ALLOW storage:logs:*	Any action on logs (read, write)
ALLOW *:*:read	Read access to everything

Multiple Statements

Combine statements to build complete policies:

```
ALLOW storage:logs:read;
ALLOW storage:spans:read;
ALLOW storage:metrics:read;
ALLOW document:documents:read;
```

Note: Statements are additive. Multiple ALLOW statements combine to grant broader access.

3. Services and Actions Reference

Core Services

Service	Purpose	Common Resources
storage	Grail data access	logs , spans , metrics , events , bizevents
settings	Configuration objects	objects , schemas
document	Notebooks, dashboards	documents , directShares
environment	Environment management	extensions , activeGates , oneAgents
state	Entity management	entities , problems
automation	Workflows, AutomationEngine	workflows , calendars

Common Actions

Action	Description
read	View/query data
write	Create/modify data

Action	Description
delete	Remove data
share	Share with others
execute	Run (for workflows)

Storage Service Details

Resource	Description	Actions
logs	Log records	read, write
spans	Distributed traces	read, write
metrics	Metric data	read, write
events	Platform events	read, write
bizevents	Business events	read, write
buckets	Storage buckets	read, write, delete

Settings Service Details

Resource	Description	Actions
objects	Configuration objects	read, write, delete
schemas	Schema definitions	read

Document Service Details

Resource	Description	Actions
documents	Dashboards, notebooks	read, write, delete, share
directShares	Direct share links	write

4. Condition Expressions

Conditions add fine-grained control to policy statements.

Condition Syntax

```
ALLOW <service>;<resource>;<action> WHERE <condition>;
```

Condition Operators

Operator	Description	Example
<code>=</code>	Equals	<code>storage:dt.security_context = "team-a"</code>
<code>!=</code>	Not equals	<code>storage:dt.security_context != "restricted"</code>
<code>IN</code>	In list	<code>storage:dt.security_context IN ("team-a", "team-b")</code>
<code>startsWith</code>	Prefix match	<code>settings:schemaId startsWith "builtin:alerting"</code>
<code>contains</code>	Contains substring	<code>settings:schemaId contains "custom"</code>
<code>MATCH</code>	Wildcard pattern match	<code>storage:dt.security_context MATCH('*/app:easytrade')</code>

MATCH() VS startsWith : Use `MATCH()` for flexible wildcard patterns anywhere in the value (e.g. middle segment). Use `startsWith` when matching a fixed leading prefix.

Common Condition Fields

Domain	Field	Description
Storage	<code>storage:dt.security_context</code>	Security context on data
Storage	<code>storage:bucket-name</code>	Bucket name (for team isolation)

Domain	Field	Description
Settings	<code>settings:schemaId</code>	Settings schema
Settings	<code>settings:dt.security_context</code>	Security context on settings

Security Context vs Bucket Conditions

Condition Type	Use Case	Example
Security Context	Flexible, changeable access	<code>storage:dt.security_context = "team-a"</code>
Bucket	Hard data isolation	<code>storage:bucket-name = "team-a-logs"</code>
Both	Defense in depth	Both conditions in policy + boundary

Condition Examples

Scope to security context (exact):

```
ALLOW storage:logs:read WHERE storage:dt.security_context = "checkout"
```

Scope to specific bucket (team isolation):

```
ALLOW storage:logs:read WHERE storage:bucket-name = "checkout_logs"
```

Scope to settings schema:

```
ALLOW settings:objects:write WHERE settings:schemaId startsWith  
"builtin:alerting"
```

Multiple conditions (AND):

```
ALLOW storage:logs:read WHERE storage:dt.security_context = "team-a" AND  
storage:bucket-name = "team-a-logs"
```

Multiple values (IN):

```
ALLOW storage:logs:read WHERE storage:dt.security_context IN ("team-a",  
"team-b", "shared")
```

Multiple buckets (team access):

```
ALLOW storage:logs:read WHERE storage:bucket-name IN ("checkout_logs",
"shared_logs")
ALLOW storage:spans:read WHERE storage:bucket-name IN ("checkout_spans",
"shared_spans")
```

Structured Security Context Design

For organisations with multiple teams and cross-cutting access requirements, a flat single-value context (e.g. "checkout") scales poorly. Instead, encode multiple dimensions into a structured key-value string:

```
comp:<component>/bu:<business-unit>/app:<application>;
```

Examples:

Entity	Security Context
Easytrade application	comp:app/bu:digital/app:easytrade
Easytrade database	comp:db/bu:digital/app:easytrade
Easytrade load balancer	comp:lb/bu:digital/app:easytrade

 **Dimension order matters for Classic entities.** Dynatrace Classic IAM uses prefix-based matching (`startsWith`). Place the dimension you need transversal access on **first** — typically `comp` for infrastructure teams.

Using `MATCH()` for flexible policy conditions:

App team — access all components within their application:


```
// 3rd Gen signals
ALLOW storage:metrics:read WHERE storage:dt.security_context
MATCH('*/app:easytrade');
ALLOW storage:logs:read WHERE storage:dt.security_context
MATCH('*/app:easytrade');
ALLOW storage:spans:read WHERE storage:dt.security_context
MATCH('*/app:easytrade');
```

Transversal database team — access all database components across all applications:

```
// 3rd Gen signals
ALLOW storage:metrics:read WHERE storage:dt.security_context
MATCH('comp:db*');
ALLOW storage:logs:read WHERE storage:dt.security_context
MATCH('comp:db*');
ALLOW storage:spans:read WHERE storage:dt.security_context
MATCH('comp:db*');
// Classic entities (prefix-based)
ALLOW storage:entities:read WHERE storage:dt.security_context
startsWith('comp:db');
```

Note: `MATCH()` applies to 3rd Gen (Grail) data. Classic entity access (`storage:entities:read`) must use `startsWith` with explicit component prefixes. See **IAM-05** for the corresponding boundary patterns.

5. Common Policy Patterns



Pattern 1: Read-Only Access

```
// Read all data, no write access
ALLOW storage:*:read;
ALLOW document:documents:read;
ALLOW state:*:read;
ALLOW settings:objects:read;
```

Pattern 2: Team-Scoped Access (Security Context)

```
// Full access to team's data only (flexible, query-time filtering)
ALLOW storage:*:read WHERE storage:dt.security_context = "checkout-team";
ALLOW storage:*:write WHERE storage:dt.security_context = "checkout-team";
ALLOW settings:objects:* WHERE settings:dt.security_context = "checkout-team";
```

Pattern 3: Bucket-Based Team Isolation (Recommended for Strict Isolation)

Use this pattern when teams must have physically separated data with no cross-team visibility.

```
// Grant access to team's specific buckets (hard data isolation)
ALLOW storage:logs:read WHERE storage:bucket-name = "checkout_logs";
ALLOW storage:logs:write WHERE storage:bucket-name = "checkout_logs";
ALLOW storage:spans:read WHERE storage:bucket-name = "checkout_spans";
ALLOW storage:spans:write WHERE storage:bucket-name = "checkout_spans";
ALLOW storage:metrics:read WHERE storage:bucket-name = "checkout_metrics";
ALLOW storage:metrics:write WHERE storage:bucket-name = "checkout_metrics";

// Also grant access to shared buckets
ALLOW storage:logs:read WHERE storage:bucket-name = "shared_logs";
ALLOW storage:spans:read WHERE storage:bucket-name = "shared_spans";
```

Bucket + Security Context (Defense in Depth):

```
// Combine bucket restriction with security context for maximum isolation
ALLOW storage:logs:read WHERE storage:bucket-name = "checkout_logs"
    AND storage:dt.security_context = "checkout";
```

Pattern 4: Log Management

```
// Read all logs, manage log pipelines
ALLOW storage:logs:read;
ALLOW settings:objects:* WHERE settings:schemaId startsWith
"builtin:logmonitoring";
```

Pattern 5: Dashboard Creator

```
// Create and share dashboards
ALLOW document:documents:*;
ALLOW storage:*:read;
```

Pattern 6: Alerting Manager

```
// Manage alerting configuration
ALLOW settings:objects:* WHERE settings:schemaId startsWith
"builtin:alerting";
ALLOW settings:objects:* WHERE settings:schemaId startsWith
"builtin:problem";
ALLOW automation:workflows:*;
```

Pattern 7: SRE On-Call Access

```
// Broad read access + problem management
ALLOW storage:*:read;
ALLOW state:problems:*;
ALLOW document:documents:read;
ALLOW settings:objects:read;
```

Pattern 8: Security Auditor

```
// Read everything, write nothing
ALLOW storage:*:read;
ALLOW settings:*:read;
ALLOW document:*:read;
ALLOW state:*:read;
```

Pattern 9: Compliance-Scoped Access (Bucket-Based)

```
// Access only to PCI-compliant data bucket
ALLOW storage:logs:read WHERE storage:bucket-name = "pci_logs";
ALLOW storage:spans:read WHERE storage:bucket-name = "pci_spans";
// No access to non-PCI buckets
```

Choosing Between Patterns

Requirement	Recommended Pattern
Teams can share data flexibly	Pattern 2 (Security Context)
Teams must never see each other's data	Pattern 3 (Buckets)
Compliance/regulatory data separation	Pattern 9 (Bucket-Based)
Temporary project access	Pattern 2 (Security Context)
Cost allocation by team	Pattern 3 (Buckets)

6. Default vs Custom Policies

Built-in Default Policies

Dynatrace provides default policies for common use cases:

Policy	Description	Use Case
environment-viewer	Read-only access	Stakeholders, auditors
environment-editor	Read + write (most things)	Team members
environment-admin	Full environment control	Administrators
account-viewer	View account settings	Cross-team visibility
account-iam-admin	Manage IAM	IAM administrators

When to Use Default Policies

Scenario	Recommendation
Simple permission levels	Use defaults
Quick onboarding	Use defaults
Standard roles	Use defaults

When to Create Custom Policies

Scenario	Recommendation
Team-scoped data access	Custom policy
Specific capability (logs only)	Custom policy
Compliance requirements	Custom policy
Least privilege needs	Custom policy

Combining Default and Custom

Groups can have both:

```
Group: dt-checkout-editors
├─ Default Policy: environment-viewer (base access)
└─ Custom Policy: checkout-write-access (team-specific writes)
```

This provides base read access plus targeted write permissions.

7. Policy Testing and Validation

Testing Methodology

1. **Create policy in test environment** (or account)
2. **Assign to test group** with a test user
3. **Verify intended access** works
4. **Verify unintended access** is blocked

5. **Document results** and adjust

Common Validation Tests

Test	Expected Result
Query logs in scope	Success
Query logs out of scope	Empty or denied
Modify settings in scope	Success
Modify settings out of scope	Denied
Create dashboard	Success (if allowed)

Debugging Policy Issues

If access isn't working as expected:

1. **Check group membership** - Is user in the correct group?
2. **Check policy assignment** - Is policy assigned to group?
3. **Check boundary** - Does boundary allow access to the entity?
4. **Check conditions** - Do condition values match?
5. **Check security context** - Is data tagged correctly?

See **IAM-09: Troubleshooting Access Issues** for detailed debugging.

8. GitOps for Policies

Version control your policies for audit trail and peer review.

Policy as Code Structure

```
iam-config/  
├── policies/  
│   ├── team-checkout-access.yaml  
│   ├── team-payments-access.yaml  
│   ├── sre-oncall.yaml  
│   └── security-auditor.yaml  
├── groups/  
│   └── group-assignments.yaml  
└── README.md
```

Policy YAML Format (Monaco)

```
# policies/team-checkout-access.yaml  
configs:  
  - id: checkout-team-policy  
    type:  
      settings:  
        schema: builtin:iam.policy  
    config:  
      name: checkout-team-access  
      description: "Access policy for Checkout team"  
      statements:  
        - effect: ALLOW  
          service: storage  
          resource: "*"   
          action: read  
          condition: "storage:dt.security_context = 'checkout'"   
        - effect: ALLOW  
          service: storage  
          resource: "*"   
          action: write  
          condition: "storage:dt.security_context = 'checkout'"
```

GitOps Workflow

1. **Branch:** Create feature branch for policy change
2. **Edit:** Modify policy YAML
3. **Review:** Submit PR, get peer review
4. **Test:** Deploy to test environment
5. **Approve:** Merge after validation

6. **Deploy:** CI/CD applies to production

Benefits

Benefit	Description
Audit Trail	Git history shows who changed what, when
Peer Review	Changes reviewed before apply
Rollback	Easy revert to previous state
Documentation	Policies are self-documenting
Consistency	Same process for all changes

See **AUTOM-03: Monaco** for full implementation details.

9. Policy Assessment Queries

Use these queries to audit policy-related activity.

```
// Track policy changes in audit log
fetch logs, from: now() - 30d
| filter matchesPhrase(log.source, "audit")
| filter matchesPhrase(content, "policy")
| fields timestamp, content
| sort timestamp desc
| limit 50
```

```
// Find access denied events
fetch logs, from: now() - 7d
| filter matchesPhrase(log.source, "audit")
| filter matchesPhrase(content, "denied") or matchesPhrase(content,
"unauthorized")
| fields timestamp, content
| sort timestamp desc
| limit 100
```

```
// Summarize permission-related audit events
fetch logs, from: now() - 7d
| filter matchesPhrase(log.source, "audit")
| filter matchesPhrase(content, "permission") or matchesPhrase(content,
"policy")
| summarize eventCount = count(), by:{content}
| sort eventCount desc
| limit 20
```

Next Steps

With policies defined, complete your access control setup:

Recommended Path

1. **IAM-05: Boundary Design Patterns** - Control what entities users can see
2. **IAM-06: User Lifecycle and Provisioning** - Automate user management
3. **IAM-07: Audit Logging and Compliance** - Monitor policy effectiveness
4. **IAM-10: Templated Policy-Group Assignments** - Reusable parameterized policies for team/environment scoping at scale
5. **IAM-11: Policy Persona Workshop** - Design persona-based policies from stakeholder requirements using a structured 6-goal methodology

Policy Checklist

Before moving on, ensure you have:

- ☐ Understood policy statement syntax
 - ☐ Reviewed services and actions reference
 - ☐ Designed policies for your team structure
 - ☐ Decided on default vs custom policy usage
 - ☐ Established testing methodology
 - ☐ Considered GitOps for policy management
-

Summary

In this notebook, you learned:

- Policy fundamentals and scope levels
 - Statement syntax: `ALLOW <service>:<resource>:<action>`
 - Services and actions reference
 - Condition expressions for fine-grained control
 - Seven common policy patterns
 - When to use default vs custom policies
 - Policy testing and validation approaches
 - GitOps for policy version control
-

References

- [Policy Statements Reference](#)
 - [Permission Management](#)
 - [Account Policies](#)
-

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.