

# K8S-10: Metadata Telemetry Enrichment

---

**Series:** K8S — Kubernetes Monitoring | **Notebook:** 10 of 13 | **Created:** January 2026 | **Last Updated:** 04/25/2026

## Enriching All Telemetry with Kubernetes Metadata

---

Kubernetes metadata enrichment automatically adds labels and annotations from your Kubernetes resources to all telemetry signals. This is the **recommended approach** for adding context to your observability data because it enriches everything: metrics, logs, traces, events, and entities.

---

## Table of Contents

---

1. [Why Metadata Enrichment?](#)
  2. [Enrichment Methods Comparison](#)
  3. [DynaKube Configuration](#)
  4. [Configuring Settings-Based Enrichment](#)
  5. [Namespace Selectors](#)
  6. [Understanding Enrichment Files](#)
  7. [Querying Enriched Data](#)
  8. [Use Cases](#)
  9. [Settings API](#)
  10. [Troubleshooting](#)
-

## OneAgent Attribute Enrichment (OneAgent 1.331+)

**Requires:** OneAgent version **1.331** or later

OneAgent can enrich **all telemetry signals** (metrics, spans, logs, events, entities) with custom metadata at the source — before data reaches the Dynatrace platform. This is more efficient than server-side tagging (auto-tags) because enrichment happens on the host and propagates to all Smartscape nodes.

**Primary Fields** (standardized from Semantic Dictionary):

- `dt.security_context` — data governance and access control
- `dt.cost.costcenter` — cost allocation
- `dt.cost.product` — product attribution

**Primary Tags** (custom key-value pairs):

- `primary_tags.environment` — environment identification (production, staging, etc.)
- `primary_tags.team` — team ownership
- `primary_tags.business_unit` — organizational unit

**Configuration:**

```
# Set during OneAgent installation
Dynatrace-OneAgent-Linux.sh --set-host-
tag="primary_tags.environment=production" --set-host-
tag="dt.security_context=confidential"

# Set on existing agents via oneagentctl
oneagentctl --set-host-tag="primary_tags.environment=production"
oneagentctl --set-host-tag="dt.cost.costcenter=12345"

# Per-process via environment variable (overrides host-level)
DT_TAGS="primary_tags.team=platform primary_tags.environment=production"
```

**Benefits over auto-tagging:**

|                   |                                |                                                                   |
|-------------------|--------------------------------|-------------------------------------------------------------------|
| Aspect            | Auto-Tags (Server-Side)        | Attribute Enrichment (Agent-Side)                                 |
| -----             | -----                          | -----                                                             |
| When applied      | After data arrives at platform | At the source, before transmission                                |
| Scope             | Entity tags only               | All signals: metrics, spans, logs, events, entities               |
| Grail integration | Limited                        | Full — feeds OpenPipeline routing, bucket assignment, permissions |

| Cost allocation | Not supported | `dt.cost.costcenter` , `dt.cost.product` fields |  
| Security context | Not supported | `dt.security_context` for data governance |

See: [Primary Grail fields and tags enrichment through OneAgent](#)

## Prerequisites

| Requirement           | Details                                                  |
|-----------------------|----------------------------------------------------------|
| Dynatrace Environment | SaaS with Grail                                          |
| Permissions           | <code>settings.read</code> , <code>settings.write</code> |
| DynaKube              | Deployed with <code>metadataEnrichment</code> enabled    |
| Kubernetes            | Namespaces with labels/annotations to enrich             |

## 1. Why Metadata Enrichment?

Kubernetes metadata enrichment solves a critical observability challenge: **connecting telemetry to business context**.

### Common Use Cases

| Use Case                   | Metadata Example                  | Benefit                      |
|----------------------------|-----------------------------------|------------------------------|
| Cost allocation            | <code>cost-center: finance</code> | Attribute costs to teams     |
| Environment identification | <code>env: production</code>      | Filter by deployment stage   |
| Team ownership             | <code>team: checkout</code>       | Route alerts to owners       |
| Compliance tagging         | <code>compliance: pci-dss</code>  | Identify regulated workloads |
| Application grouping       | <code>app: ecommerce</code>       | Group related services       |

### What Gets Enriched

With settings-based enrichment, **all signals** receive metadata:

| Signal Type  | Enriched? | Notes                                 |
|--------------|-----------|---------------------------------------|
| Logs         | Yes       | All container and pod logs            |
| Metrics      | Yes       | Including Kubernetes platform metrics |
| Spans/Traces | Yes       | Distributed tracing data              |
| Events       | Yes       | Kubernetes events                     |
| Entities     | Yes       | Smartscape topology                   |

## Primary Grail Fields and Tags

In addition to settings-based enrichment, Dynatrace automatically propagates certain **Primary Grail Fields** across all signal types. These fields are indexed, require no configuration, and are the foundation for cross-signal data organization:

| Primary Grail Field  | DQL Field                       | Propagation |
|----------------------|---------------------------------|-------------|
| Kubernetes Cluster   | <code>k8s.cluster.name</code>   | All signals |
| Kubernetes Namespace | <code>k8s.namespace.name</code> | All signals |
| Host Group           | <code>dt.host_group.id</code>   | All signals |
| AWS Account          | <code>aws.account.id</code>     | All signals |


**Primary Grail Tags** (prefix `primary_tags.` ) extend this concept for custom dimensions like `primary_tags.app` or `primary_tags.team` . Use them when Primary Grail Fields don't align with your organizational structure (e.g., shared infrastructure hosting multiple applications).

**Cost allocation fields** like `dt.cost.costcenter` and `dt.cost.product` also propagate to service metrics, enabling chargeback reporting across teams.

**See also: ORGNZ-10: Advanced Segment Definitions** covers Primary Grail Fields in depth, including enrichment approaches for dedicated vs. shared infrastructure scenarios.

## 2. Enrichment Methods Comparison

Dynatrace offers multiple ways to add Kubernetes metadata. Choose based on your needs:



### Why Settings-Based is Recommended

Settings-based enrichment:

- Enriches **all signals**: logs, metrics, traces, events, entities
- Single configuration point (Dynatrace Settings)
- Leverages existing namespace labels
- No application changes required
- Works with Kubernetes platform metrics

Manual pod annotations:

- Does NOT enrich: Kubernetes metrics, events, Smartscape entities, Prometheus metrics
- Requires changes to every pod spec

**Warning:** Do not mix settings-based enrichment with manual pod annotations. Using both simultaneously may cause conflicts and unexpected behavior.

### 3. DynaKube Configuration

---

Metadata enrichment must be enabled in your DynaKube custom resource.

#### Enable Metadata Enrichment

```
apiVersion: dynatrace.com/v1beta5
kind: DynaKube
metadata:
  name: dynakube
  namespace: dynatrace
spec:
  apiUrl: https://ENVIRONMENT_ID.live.dynatrace.com/api

  # Enable metadata enrichment (enabled by default)
  metadataEnrichment:
    enabled: true

  oneAgent:
    cloudNativeFullStack: {}

  activeGate:
    capabilities:
      - kubernetes-monitoring
```

#### Disable Metadata Enrichment (if needed)

```
# Using annotation
kubectl annotate dynakube -n dynatrace dynakube \
  feature.dynatrace.com/disable-metadata-enrichment="true"
```

Or in the DynaKube spec:

```
spec:
  metadataEnrichment:
    enabled: false
```

## 4. Configuring Settings-Based Enrichment

### Access Settings

1. Navigate to **Settings > Cloud and virtualization > Kubernetes telemetry enrichment**
2. Select **Add rule**

### Rule Configuration Options

| Field                | Description                                | Example |
|----------------------|--------------------------------------------|---------|
| <b>Metadata type</b> | Source type (annotation or label)          | Label   |
| <b>Key</b>           | The label/annotation key to capture        | team    |
| <b>Prefix</b>        | Optional prefix for the enriched attribute | k8s.    |

### Example: Capture Team Label

If your namespaces have:

```
apiVersion: v1
kind: Namespace
metadata:
  name: checkout
  labels:
    team: checkout-team
    cost-center: engineering
    env: production
```

Create enrichment rules:

| Rule | Metadata Type | Key         | Prefix |
|------|---------------|-------------|--------|
| 1    | Label         | team        | k8s.   |
| 2    | Label         | cost-center | k8s.   |
| 3    | Label         | env         | k8s.   |

Result: All telemetry from the `checkout` namespace will have:

- `k8s.team: checkout-team`
- `k8s.cost-center: engineering`
- `k8s.env: production`

## 5. Namespace Selectors

---

By default, enrichment rules apply to **all namespaces**. Use namespace selectors to limit scope.

### DynaKube Namespace Selector

If your DynaKube uses a `namespaceSelector`, ensure it matches the namespaces you want to enrich:

```
spec:
  metadataEnrichment:
    enabled: true
    namespaceSelector:
      matchLabels:
        dynatrace-enrich: enabled
```

### Labeling Namespaces for Enrichment

```
# Enable enrichment for a namespace
kubectl label namespace checkout dynatrace-enrich=enabled

# Verify labels
kubectl get namespace checkout --show-labels
```

### Match Expressions

For more complex selection:



```
namespaceSelector:
  matchExpressions:
    - key: env
      operator: In
      values:
        - production
        - staging
```

## 6. Understanding Enrichment Files

The Dynatrace Operator creates enrichment files that are used by OneAgent and applications.

### File Locations

| File                   | Location                       | Purpose              |
|------------------------|--------------------------------|----------------------|
| dt_metadata.json       | /var/lib/dynatrace/enrichment/ | JSON format metadata |
| dt_metadata.properties | /var/lib/dynatrace/enrichment/ | Properties format    |
| dt_host_metadata.json  | /var/lib/dynatrace/enrichment/ | Host-level metadata  |

### Example dt\_metadata.json Content

```
{
  "k8s.namespace.name": "checkout",
  "k8s.pod.name": "checkout-api-7d9f8c6b4d-x2k9m",
  "k8s.team": "checkout-team",
  "k8s.cost-center": "engineering",
  "k8s.env": "production"
}
```

### Accessing Enrichment Files in Applications

For manual instrumentation or custom applications:

```

import json
import os

# Load Dynatrace metadata
enrichment_paths = [
    '/var/lib/dynatrace/enrichment/dt_metadata.json',
    '/var/lib/dynatrace/enrichment/dt_host_metadata.json'
]

metadata = {}
for path in enrichment_paths:
    if os.path.exists(path):
        with open(path) as f:
            metadata.update(json.load(f))

print(metadata)

```

## 7. Querying Enriched Data

---

Once enrichment is configured, you can filter and group by the enriched attributes.

```

// Query logs by enriched team label
fetch logs, from:-1h
| filter isNotNull(k8s.namespace.name)
| summarize logCount = count(), by:{k8s.namespace.name}
| sort logCount desc
| limit 20

```

```

// Group metrics by cost center (enriched label)
timeseries avgCpuMillicores = avg(dt.kubernetes.container.cpu_usage),
from:-1h, by:{k8s.namespace.name}
| sort avgCpuMillicores desc

```

```

// Find all spans from production environment
fetch spans, from:-1h
| filter isNotNull(k8s.namespace.name)
| summarize
    spanCount = count(),
    avgDuration = avg(duration),
    by:{k8s.namespace.name, service.name}
| sort spanCount desc
| limit 20

```

## 8. Use Cases

---

### Cost Allocation by Team

Label namespaces with cost centers:

```
kubectl label namespace checkout cost-center=checkout-team
kubectl label namespace catalog cost-center=catalog-team
kubectl label namespace shared cost-center=platform
```

Create enrichment rule for `cost-center` label, then query:

```
timeseries totalCpu = sum(dt.kubernetes.container.cpu_usage), from:-1h, by:
{k8s.cost-center}
```

### Pipeline Routing

Route logs to different buckets based on enriched metadata. In OpenPipeline:

```
processors:
  - type: route
    rules:
      - condition: "k8s.env == 'production'"
        destination: "prod_logs_365d"
      - condition: "k8s.env == 'staging'"
        destination: "staging_logs_35d"
```

### Security Context Assignment

Use enriched metadata in security context for fine-grained access:

```
processors:
  - type: security-context
    rules:
      - condition: "k8s.team == 'checkout-team'"
        context: "team:checkout"
```

### Grail Permissions

Create IAM policies based on Kubernetes attributes:

```
ALLOW storage:buckets:read WHERE storage:bucket-name STARTSWITH "default_";  
ALLOW storage:logs:read WHERE storage:k8s.namespace.name = "checkout";
```

## 9. Settings API

---

Manage enrichment rules programmatically via the Settings API.

### Schema

The schema for Kubernetes telemetry enrichment is:

```
builtin:kubernetes.generic.metadata.enrichment
```

### List Current Rules

```
curl -X GET  
"https://ENVIRONMENT_ID.live.dynatrace.com/api/v2/settings/objects" \  
-H "Authorization: Api-Token YOUR_TOKEN" \  
-H "Content-Type: application/json" \  
-d '{  
  "schemaIds": ["builtin:kubernetes.generic.metadata.enrichment"],  
  "scopes": ["environment"]  
}
```

## Create Rule via API

```
curl -X POST
"https://ENVIRONMENT_ID.live.dynatrace.com/api/v2/settings/objects" \
-H "Authorization: Api-Token YOUR_TOKEN" \
-H "Content-Type: application/json" \
-d '[
  {
    "schemaId": "builtin:kubernetes.generic.metadata.enrichment",
    "scope": "environment",
    "value": {
      "enabled": true,
      "metadataType": "LABEL",
      "key": "team",
      "prefix": "k8s."
    }
  }
]
```

## 10. Troubleshooting

### Enrichment Not Appearing

| Symptom                 | Cause                       | Solution                            |
|-------------------------|-----------------------------|-------------------------------------|
| No enriched attributes  | Rules not propagated        | Wait 45 minutes after rule creation |
| Some namespaces missing | namespaceSelector mismatch  | Verify DynaKube namespaceSelector   |
| Pods not enriched       | metadataEnrichment disabled | Check DynaKube spec                 |

### Verify DynaKube Status

```
# Check DynaKube configuration
kubectl -n dynatrace get dynakube -o yaml | grep -A5 metadataEnrichment

# Check for disable annotation
kubectl -n dynatrace get dynakube -o yaml | grep disable-metadata-enrichment
```

## Verify Enrichment Files

```
# Check enrichment directory exists
kubectl exec -it <pod-name> -- ls -la /var/lib/dynatrace/enrichment/

# View enrichment content
kubectl exec -it <pod-name> -- cat
/var/lib/dynatrace/enrichment/dt_metadata.json
```

## Timing Considerations

| Action                  | Propagation Time           |
|-------------------------|----------------------------|
| New rule created        | Up to 45 minutes           |
| Rule modified           | Up to 45 minutes           |
| DynaKube deployed first | Immediate (if rules exist) |

**Tip:** If you configure enrichment rules **before** deploying DynaKube, the rules take effect immediately when DynaKube is applied.

## Summary

In this notebook, you learned:

- Why settings-based metadata enrichment is the recommended approach
- How to enable enrichment in DynaKube
- Creating enrichment rules in Settings UI
- Using namespace selectors to scope enrichment
- Understanding enrichment file locations and formats
- Querying enriched data with DQL
- Practical use cases: cost allocation, routing, security
- Troubleshooting common issues

## References

---

- [Metadata enrichment of all telemetry originating from Kubernetes](#)
  - [Configure enrichment directory](#)
  - [Settings API - Kubernetes Telemetry Enrichment schema](#)
  - [Set up Grail permissions for telemetry from Kubernetes](#)
- 

*This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.*