

K8S-10: Metadata Telemetry Enrichment

Series: K8S | **Notebook:** 10 of 12 | **Created:** January 2026 | **Last Updated:** 02/05/2026

Enriching All Telemetry with Kubernetes Metadata

Kubernetes metadata enrichment automatically adds labels and annotations from your Kubernetes resources to all telemetry signals. This is the **recommended approach** for adding context to your observability data because it enriches everything: metrics, logs, traces, events, and entities.

Table of Contents

1. [Why Metadata Enrichment?](#)
 2. [Enrichment Methods Comparison](#)
 3. [DynaKube Configuration](#)
 4. [Configuring Settings-Based Enrichment](#)
 5. [Namespace Selectors](#)
 6. [Understanding Enrichment Files](#)
 7. [Querying Enriched Data](#)
 8. [Use Cases](#)
 9. [Settings API](#)
 10. [Troubleshooting](#)
-

Prerequisites

Requirement	Details
Dynatrace Environment	SaaS with Grail
Permissions	<code>settings.read</code> , <code>settings.write</code>
DynaKube	Deployed with <code>metadataEnrichment</code> enabled
Kubernetes	Namespaces with labels/annotations to enrich

1. Why Metadata Enrichment?

Kubernetes metadata enrichment solves a critical observability challenge: **connecting telemetry to business context**.

Common Use Cases

Use Case	Metadata Example	Benefit
Cost allocation	<code>cost-center: finance</code>	Attribute costs to teams
Environment identification	<code>env: production</code>	Filter by deployment stage
Team ownership	<code>team: checkout</code>	Route alerts to owners
Compliance tagging	<code>compliance: pci-dss</code>	Identify regulated workloads
Application grouping	<code>app: ecommerce</code>	Group related services

What Gets Enriched

With settings-based enrichment, **all signals** receive metadata:

Signal Type	Enriched?	Notes
Logs	Yes	All container and pod logs
Metrics	Yes	Including Kubernetes platform metrics
Spans/Traces	Yes	Distributed tracing data
Events	Yes	Kubernetes events
Entities	Yes	Smartscape topology

Primary Grail Fields and Tags

In addition to settings-based enrichment, Dynatrace automatically propagates certain **Primary Grail Fields** across all signal types. These fields are indexed, require no configuration, and are the foundation for cross-signal data organization:

Primary Grail Field	DQL Field	Propagation
Kubernetes Cluster	<code>k8s.cluster.name</code>	All signals
Kubernetes Namespace	<code>k8s.namespace.name</code>	All signals
Host Group	<code>dt.host_group.id</code>	All signals
AWS Account	<code>aws.account.id</code>	All signals

Primary Grail Tags (prefix `primary_tags.`) extend this concept for custom dimensions like `primary_tags.app` or `primary_tags.team`. Use them when Primary Grail Fields don't align with your organizational structure (e.g., shared infrastructure hosting multiple applications).

Cost allocation fields like `dt.cost.costcenter` and `dt.cost.product` also propagate to service metrics, enabling chargeback reporting across teams.

See also: ORGNZ-10: Advanced Segment Definitions covers Primary Grail Fields in depth, including enrichment approaches for dedicated vs. shared infrastructure scenarios.

2. Enrichment Methods Comparison

Dynatrace offers multiple ways to add Kubernetes metadata. Choose based on your needs:



Why Settings-Based is Recommended

Settings-based enrichment:

- Enriches **all signals**: logs, metrics, traces, events, entities
- Single configuration point (Dynatrace Settings)
- Leverages existing namespace labels
- No application changes required
- Works with Kubernetes platform metrics

Manual pod annotations:

- Does NOT enrich: Kubernetes metrics, events, Smartscape entities, Prometheus metrics
- Requires changes to every pod spec

Warning: Do not mix settings-based enrichment with manual pod annotations. Using both simultaneously may cause conflicts and unexpected behavior.

3. DynaKube Configuration

Metadata enrichment must be enabled in your DynaKube custom resource.

Enable Metadata Enrichment

```
apiVersion: dynatrace.com/v1beta3
kind: DynaKube
metadata:
  name: dynakube
  namespace: dynatrace
spec:
  apiUrl: https://ENVIRONMENT_ID.live.dynatrace.com/api

  # Enable metadata enrichment (enabled by default)
  metadataEnrichment:
    enabled: true

  oneAgent:
    cloudNativeFullStack: {}

  activeGate:
    capabilities:
      - kubernetes-monitoring
```

Disable Metadata Enrichment (if needed)

```
# Using annotation
kubectl annotate dynakube -n dynatrace dynakube \
  feature.dynatrace.com/disable-metadata-enrichment="true"
```

Or in the DynaKube spec:

```
spec:
  metadataEnrichment:
    enabled: false
```

4. Configuring Settings-Based Enrichment

Access Settings

1. Navigate to **Settings > Cloud and virtualization > Kubernetes telemetry enrichment**
2. Select **Add rule**

Rule Configuration Options

Field	Description	Example
Metadata type	Source type (annotation or label)	Label
Key	The label/annotation key to capture	team
Prefix	Optional prefix for the enriched attribute	k8s.

Example: Capture Team Label

If your namespaces have:

```
apiVersion: v1
kind: Namespace
metadata:
  name: checkout
  labels:
    team: checkout-team
    cost-center: engineering
    env: production
```

Create enrichment rules:

Rule	Metadata Type	Key	Prefix
1	Label	team	k8s.
2	Label	cost-center	k8s.
3	Label	env	k8s.

Result: All telemetry from the `checkout` namespace will have:

- `k8s.team: checkout-team`
- `k8s.cost-center: engineering`
- `k8s.env: production`

5. Namespace Selectors

By default, enrichment rules apply to **all namespaces**. Use namespace selectors to limit scope.

DynaKube Namespace Selector

If your DynaKube uses a `namespaceSelector`, ensure it matches the namespaces you want to enrich:

```
spec:
  metadataEnrichment:
    enabled: true
    namespaceSelector:
      matchLabels:
        dynatrace-enrich: enabled
```

Labeling Namespaces for Enrichment

```
# Enable enrichment for a namespace
kubectl label namespace checkout dynatrace-enrich=enabled

# Verify labels
kubectl get namespace checkout --show-labels
```

Match Expressions

For more complex selection:

```
namespaceSelector:
  matchExpressions:
    - key: env
      operator: In
      values:
        - production
        - staging
```

6. Understanding Enrichment Files

The Dynatrace Operator creates enrichment files that are used by OneAgent and applications.

File Locations

File	Location	Purpose
dt_metadata.json	/var/lib/dynatrace/enrichment/	JSON format metadata
dt_metadata.properties	/var/lib/dynatrace/enrichment/	Properties format
dt_host_metadata.json	/var/lib/dynatrace/enrichment/	Host-level metadata

Example dt_metadata.json Content

```
{
  "k8s.namespace.name": "checkout",
  "k8s.pod.name": "checkout-api-7d9f8c6b4d-x2k9m",
  "k8s.team": "checkout-team",
  "k8s.cost-center": "engineering",
  "k8s.env": "production"
}
```

Accessing Enrichment Files in Applications

For manual instrumentation or custom applications:


```

import json
import os

# Load Dynatrace metadata
enrichment_paths = [
    '/var/lib/dynatrace/enrichment/dt_metadata.json',
    '/var/lib/dynatrace/enrichment/dt_host_metadata.json'
]

metadata = {}
for path in enrichment_paths:
    if os.path.exists(path):
        with open(path) as f:
            metadata.update(json.load(f))

print(metadata)

```

7. Querying Enriched Data

Once enrichment is configured, you can filter and group by the enriched attributes.

```

// Query logs by enriched team label
fetch logs, from:-1h
| filter isNotNull(k8s.namespace.name)
| summarize logCount = count(), by:{k8s.namespace.name}
| sort logCount desc
| limit 20

```

```

// Group metrics by cost center (enriched label)
timeseries avgCpuMillicores = avg(dt.kubernetes.container.cpu_usage),
from:-1h, by:{k8s.namespace.name}
| sort avgCpuMillicores desc

```

```

// Find all spans from production environment
fetch spans, from:-1h
| filter isNotNull(k8s.namespace.name)
| summarize
    spanCount = count(),
    avgDuration = avg(duration),
    by:{k8s.namespace.name, service.name}
| sort spanCount desc
| limit 20

```

8. Use Cases

Cost Allocation by Team

Label namespaces with cost centers:

```
kubectl label namespace checkout cost-center=checkout-team
kubectl label namespace catalog cost-center=catalog-team
kubectl label namespace shared cost-center=platform
```

Create enrichment rule for `cost-center` label, then query:

```
fetch dt.metrics
| filter metric.key == "dt.kubernetes.container.cpu_usage"
| summarize totalCpu = sum(value), by:{k8s.cost-center}
```

Pipeline Routing

Route logs to different buckets based on enriched metadata. In OpenPipeline:

```
processors:
  - type: route
    rules:
      - condition: "k8s.env == 'production'"
        destination: "prod_logs_365d"
      - condition: "k8s.env == 'staging'"
        destination: "staging_logs_35d"
```

Security Context Assignment

Use enriched metadata in security context for fine-grained access:

```
processors:
  - type: security-context
    rules:
      - condition: "k8s.team == 'checkout-team'"
        context: "team:checkout"
```

Grail Permissions

Create IAM policies based on Kubernetes attributes:

```
ALLOW storage:buckets:read WHERE storage:bucket-name STARTSWITH "default_";  
ALLOW storage:logs:read WHERE storage:k8s.namespace.name = "checkout";
```

9. Settings API

Manage enrichment rules programmatically via the Settings API.

Schema

The schema for Kubernetes telemetry enrichment is:

```
builtin:kubernetes.generic.metadata.enrichment
```

List Current Rules

```
curl -X GET  
"https://ENVIRONMENT_ID.live.dynatrace.com/api/v2/settings/objects" \  
-H "Authorization: Api-Token YOUR_TOKEN" \  
-H "Content-Type: application/json" \  
-d '{  
  "schemaIds": ["builtin:kubernetes.generic.metadata.enrichment"],  
  "scopes": ["environment"]  
'
```

Create Rule via API

```
curl -X POST
"https://ENVIRONMENT_ID.live.dynatrace.com/api/v2/settings/objects" \
-H "Authorization: Api-Token YOUR_TOKEN" \
-H "Content-Type: application/json" \
-d '[
{
  "schemaId": "builtin:kubernetes.generic.metadata.enrichment",
  "scope": "environment",
  "value": {
    "enabled": true,
    "metadataType": "LABEL",
    "key": "team",
    "prefix": "k8s."
  }
}
]'
```

10. Troubleshooting

Enrichment Not Appearing

Symptom	Cause	Solution
No enriched attributes	Rules not propagated	Wait 45 minutes after rule creation
Some namespaces missing	namespaceSelector mismatch	Verify DynaKube namespaceSelector
Pods not enriched	metadataEnrichment disabled	Check DynaKube spec

Verify DynaKube Status

```
# Check DynaKube configuration
kubectl -n dynatrace get dynakube -o yaml | grep -A5 metadataEnrichment

# Check for disable annotation
kubectl -n dynatrace get dynakube -o yaml | grep disable-metadata-enrichment
```

Verify Enrichment Files

```
# Check enrichment directory exists
kubectl exec -it <pod-name> -- ls -la /var/lib/dynatrace/enrichment/

# View enrichment content
kubectl exec -it <pod-name> -- cat
/var/lib/dynatrace/enrichment/dt_metadata.json
```

Timing Considerations

Action	Propagation Time
New rule created	Up to 45 minutes
Rule modified	Up to 45 minutes
DynaKube deployed first	Immediate (if rules exist)

Tip: If you configure enrichment rules **before** deploying DynaKube, the rules take effect immediately when DynaKube is applied.

Summary

In this notebook, you learned:

- Why settings-based metadata enrichment is the recommended approach
- How to enable enrichment in DynaKube
- Creating enrichment rules in Settings UI
- Using namespace selectors to scope enrichment
- Understanding enrichment file locations and formats
- Querying enriched data with DQL
- Practical use cases: cost allocation, routing, security
- Troubleshooting common issues

References

- [Metadata enrichment of all telemetry originating from Kubernetes](#)
 - [Configure enrichment directory](#)
 - [Settings API - Kubernetes Telemetry Enrichment schema](#)
 - [Set up Grail permissions for telemetry from Kubernetes](#)
-

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.