

# K8S-14: Kubernetes Deployment Guide

---

**Series:** K8S | **Notebook:** 14 of 14 | **Type:** LAB | **Created:** April 2026 | **Last Updated:** 04/04/2026

## Overview

---

Step-by-step guide to deploy Dynatrace monitoring on Kubernetes with all recommended production settings. Covers operator installation, DynaKube configuration, metadata enrichment, build propagation, segments, and verification.

This is a hands-on deployment notebook. Each section includes commands to run and queries to verify. Follow the sections in order for a complete production deployment.

---

## Table of Contents

---

1. [Prerequisites](#)
  2. [Create Access Tokens](#)
  3. [Install Dynatrace Operator](#)
  4. [Configure DynaKube CR \(Production Recommended\)](#)
  5. [Label Namespaces for Monitoring](#)
  6. [Configure Telemetry Enrichment Rules](#)
  7. [Configure Build Propagation & Release Detection](#)
  8. [Create Segments](#)
  9. [Verify Deployment](#)
  10. [Monitoring-the-Monitoring](#)
  11. [Recommended Next Steps](#)
  12. [Summary](#)
-

# 1. Prerequisites

| Requirement               | Details                                                                                 |
|---------------------------|-----------------------------------------------------------------------------------------|
| <b>Kubernetes Cluster</b> | Version 1.24+ with <code>kubectl</code> access and cluster-admin permissions            |
| <b>Helm</b>               | Version 3.x installed locally                                                           |
| <b>Dynatrace Tenant</b>   | SaaS environment with admin access                                                      |
| <b>Network</b>            | Outbound 443 to <code>*.live.dynatrace.com</code> and <code>*.apps.dynatrace.com</code> |
| <b>Knowledge</b>          | Completed K8S-01 (Fundamentals) and K8S-02 (DynaKube Deployment)                        |

## Verify Your Environment

Run these commands to confirm your cluster is ready:

```
# Verify Kubernetes version (must be 1.24+)
kubectl version --short

# Verify Helm version (must be 3.x)
helm version --short

# Verify cluster access
kubectl get nodes
```

## 2. Create Access Tokens

The Dynatrace Operator requires two tokens: an **API token** for cluster communication and a **data ingest token** for sending telemetry.

### API Token

1. Navigate to **Access Tokens > Generate new token**
2. Select the template: **Kubernetes: Dynatrace Operator**
3. This template includes the required scopes:

| Scope                            | Purpose                                   |
|----------------------------------|-------------------------------------------|
| InstallerDownload                | Download OneAgent and ActiveGate binaries |
| activeGateTokenManagement.create | Create ActiveGate auth tokens             |
| entities.read                    | Read entity topology                      |
| settings.read                    | Read environment settings                 |
| settings.write                   | Write environment settings                |

## Data Ingest Token

1. Navigate to **Access Tokens > Generate new token**
2. Select the template: **Kubernetes: Data Ingest**
3. This template includes the required scopes:

| Scope                     | Purpose                                |
|---------------------------|----------------------------------------|
| metrics.ingest            | Send Kubernetes and Prometheus metrics |
| logs.ingest               | Send container and pod logs            |
| openTelemetryTrace.ingest | Send distributed traces                |

**Important:** Always use the predefined templates. Manually selecting scopes risks missing required permissions, which causes silent failures that are difficult to diagnose.

## Save Your Tokens

Copy both tokens immediately after generation. You cannot retrieve them later.

```
# Store tokens as environment variables for the next steps
export API_TOKEN="dt0c01.XXXXXXXX.YYYYYYYYYYYYYYYY"
export DATA_INGEST_TOKEN="dt0c01.XXXXXXXX.ZZZZZZZZZZZZZZZZ"
```

### 3. Install Dynatrace Operator

---

#### Step 1: Install the Operator via Helm

```
helm install dynatrace-operator oci://public.ecr.aws/dynatrace/dynatrace-operator \
  --create-namespace \
  --namespace dynatrace \
  --atomic
```

The `--atomic` flag ensures the installation is rolled back automatically if any component fails to start.

#### Step 2: Create the Token Secret

```
kubectl -n dynatrace create secret generic dynakube \
  --from-literal=apiToken=$API_TOKEN \
  --from-literal=dataIngestToken=$DATA_INGEST_TOKEN
```

#### Step 3: Verify Operator Pods

```
kubectl get pods -n dynatrace
```

Expected output:

| NAME                                | READY | STATUS  | RESTARTS | AGE |
|-------------------------------------|-------|---------|----------|-----|
| dynatrace-operator-xxxxxxxxxx-xxxxx | 1/1   | Running | 0        | 30s |
| dynatrace-webhook-xxxxxxxxxx-xxxxx  | 1/1   | Running | 0        | 30s |

**Note:** The operator and webhook pods must both be `Running` before applying the DynaKube CR. If either pod is in `CrashLoopBackOff`, check the logs with `kubectl logs -n dynatrace <pod-name>`.

### 4. Configure DynaKube CR (Production

## Recommended)

---

The DynaKube custom resource defines how Dynatrace monitors your cluster. The configuration below represents production-recommended settings with all key features enabled.

### Complete DynaKube YAML

Save this as `dynakube.yaml` and customize the placeholder values:

```

apiVersion: dynatrace.com/v1beta6
kind: DynaKube
metadata:
  name: dynakube
  namespace: dynatrace
  annotations:
    # Fail pod startup if injection fails – surfaces problems immediately
    feature.dynatrace.com/injection-failure-policy: "fail"
    # Opt-in injection – only labeled namespaces get instrumented
    feature.dynatrace.com/automatic-injection: "false"
    # Enable build label propagation for Release Inventory
    feature.dynatrace.com/label-version-detection: "true"
    # Extend CSI mount timeout for cloud providers with slow volume
    attachment
    feature.dynatrace.com/max-csi-mount-timeout: "15m"
spec:
  apiUrl: https://&lt;ENVIRONMENT_ID&gt;.live.dynatrace.com/api

  # Network zone – isolates traffic per cluster
  networkZone: &lt;cluster-name&gt;

  # Metadata enrichment – flows K8s labels into all telemetry
  metadataEnrichment:
    enabled: true

  oneAgent:
    # Host group – scopes dashboards and alerts per cluster
    hostGroup: &lt;cluster-name&gt;
    cloudNativeFullStack:
      # Opt-in namespace selector
      namespaceSelector:
        matchLabels:
          dt-monitoring: "true"

    # Host properties for entity-level context
    args:
      - --set-host-property=environment=production
      - --set-host-property=team=platform
      - --set-host-property=cost-center=infrastructure

    # Tolerate control plane nodes
    tolerations:
      - effect: NoSchedule
        key: node-role.kubernetes.io/control-plane
        operator: Exists

  # ActiveGate – K8s API monitoring + routing

```

```

activeGate:
  capabilities:
    - routing
    - kubernetes-monitoring
    - metrics-ingest
  replicas: 2
  resources:
    requests:
      cpu: 500m
      memory: 1Gi
    limits:
      cpu: 2000m
      memory: 2Gi
  topologySpreadConstraints:
    - maxSkew: 1
      topologyKey: topology.kubernetes.io/zone
      whenUnsatisfiable: ScheduleAnyway
      labelSelector:
        matchLabels:
          dynatrace.com/component: activegate

# Log monitoring
logMonitoring: {}

```

## Apply the DynaKube CR

```

# Replace placeholders, then apply
kubectl apply -f dynakube.yaml

```

## Setting Explanations

| Setting                                                       | Value                                 | Why It Matters                                                          |
|---------------------------------------------------------------|---------------------------------------|-------------------------------------------------------------------------|
| <code>injection-failure-policy:</code><br><code>"fail"</code> | Fail pod startup on injection error   | Surfaces misconfiguration immediately instead of running uninstrumented |
| <code>automatic-injection:</code><br><code>"false"</code>     | Opt-in injection via namespace labels | Prevents accidental injection into system namespaces                    |

| Setting                                      | Value                                            | Why It Matters                                                         |
|----------------------------------------------|--------------------------------------------------|------------------------------------------------------------------------|
| <code>label-version-detection: "true"</code> | Propagate <code>app.kubernetes.io/version</code> | Populates Release Inventory for deployment tracking                    |
| <code>max-csi-mount-timeout: "15m"</code>    | Extended CSI driver timeout                      | Prevents pod failures on cloud providers with slow EBS/disk attachment |
| <code>networkZone</code>                     | Per-cluster zone                                 | Isolates ActiveGate traffic; required for multi-cluster                |
| <code>metadataEnrichment.enabled</code>      | <code>true</code>                                | Flows K8s labels/annotations into all telemetry signals                |
| <code>cloudNativeFullStack</code>            | Full-stack mode                                  | Full code-level visibility with automatic injection                    |
| <code>namespaceSelector</code>               | <code>dt-monitoring: "true"</code>               | Only labeled namespaces get OneAgent injection                         |
| <code>hostGroup</code>                       | Per-cluster name                                 | Groups hosts for scoped alerting and dashboards                        |
| <code>args: --set-host-property</code>       | Custom properties                                | Adds environment, team, cost-center to every host entity               |
| <code>tolerations</code>                     | Control plane toleration                         | Ensures DaemonSet runs on all nodes including control plane            |
| <code>activeGate.replicas: 2</code>          | High availability                                | Survives single-pod failure; required for production                   |

| Setting                                          | Value                  | Why It Matters                                              |
|--------------------------------------------------|------------------------|-------------------------------------------------------------|
| <code>topologySpreadConstraints</code>           | Zone-aware scheduling  | Distributes ActiveGate pods across availability zones       |
| <code>capabilities: routing</code>               | ActiveGate routing     | Routes OneAgent traffic through ActiveGate                  |
| <code>capabilities: kubernetes-monitoring</code> | K8s API integration    | Enables cluster-level monitoring, events, and workload data |
| <code>capabilities: metrics-ingest</code>        | Prometheus scraping    | Enables annotation-based Prometheus metric collection       |
| <code>logMonitoring: {}</code>                   | Log collection enabled | Collects container stdout/stderr logs                       |

## 5. Label Namespaces for Monitoring

Because the DynaKube uses `automatic-injection: "false"` with a `namespaceSelector`, only namespaces labeled with `dt-monitoring=true` will receive OneAgent injection.

### Label Application Namespaces

```
# Label each application namespace
kubectl label namespace <app-namespace-1> dt-monitoring=true
kubectl label namespace <app-namespace-2> dt-monitoring=true

# Verify labeled namespaces
kubectl get namespaces -l dt-monitoring=true
```

### Namespaces to NOT Label

These system namespaces should **never** receive the `dt-monitoring=true` label:

| Namespace       | Reason                                                               |
|-----------------|----------------------------------------------------------------------|
| kube-system     | Core Kubernetes components — injection can cause instability         |
| kube-public     | Cluster info namespace — no application workloads                    |
| kube-node-lease | Node heartbeat namespace — no application workloads                  |
| cert-manager    | Certificate management — injection interferes with TLS operations    |
| dynatrace       | Dynatrace's own namespace — self-monitoring is handled internally    |
| istio-system    | Service mesh control plane — injection conflicts with Envoy sidecars |

**Tip:** After labeling, trigger a rolling restart of deployments in those namespaces to activate injection: `kubectl rollout restart deployment -n <namespace> .`

## 6. Configure Telemetry Enrichment Rules

Telemetry enrichment flows Kubernetes labels and annotations into all telemetry signals (logs, metrics, traces, events). This enables filtering, cost allocation, and IAM scoping.

### Settings Path

**Settings > Cloud and Virtualization > Kubernetes Telemetry Enrichment**

Schema ID: `builtin:kubernetes.generic.metadata.enrichment`

### Recommended Enrichment Rules

| # | Metadata Type   | Key  | Mapped To           | Purpose                                                  |
|---|-----------------|------|---------------------|----------------------------------------------------------|
| 1 | Namespace label | team | dt.security_context | Enables record-level IAM — teams see only their own data |

| # | Metadata Type   | Key                       | Mapped To          | Purpose                                 |
|---|-----------------|---------------------------|--------------------|-----------------------------------------|
| 2 | Namespace label | cost-center               | dt.cost.costcenter | Enables FinOps cost allocation per team |
| 3 | Namespace label | app.kubernetes.io/part-of | dt.cost.product    | Enables product-level cost attribution  |

## Label Your Namespaces

The enrichment rules map **existing** Kubernetes labels to Dynatrace fields. Ensure your namespaces carry these labels:

```
# Add organizational labels to namespaces
kubectl label namespace <app-namespace> team=backend
kubectl label namespace <app-namespace> cost-center=ecommerce
kubectl label namespace <app-namespace> app.kubernetes.io/part-of=checkout-platform

# Verify labels
kubectl get namespace <app-namespace> --show-labels
```

## Propagation Timing

| Action                       | Propagation Time                 |
|------------------------------|----------------------------------|
| New enrichment rule created  | Up to 45 minutes                 |
| Namespace label added        | Immediate (next telemetry cycle) |
| Pod restart after enrichment | Immediate                        |

**Important:** If you create enrichment rules after DynaKube is already deployed, existing pods must be restarted to pick up the enriched metadata. Run `kubectl rollout restart deployment -n <namespace>` for affected namespaces.

**See also:** K8S-10 (Metadata Telemetry Enrichment) for a deep dive into enrichment methods, file locations, and advanced use cases.

# 7. Configure Build Propagation & Release Detection

Build propagation connects your Kubernetes deployment metadata to the Dynatrace Release Inventory. This gives you deployment tracking, version comparison, and release health monitoring.

## What Is Already Enabled

The DynaKube in Section 4 includes:

```
annotations:  
  feature.dynatrace.com/label-version-detection: "true"
```

This tells the Dynatrace Operator to read version information from standard Kubernetes labels.

## Required Pod Labels

Your pods (or the Deployments/StatefulSets that create them) must carry these standard labels:

| Label                     | Purpose           | Example                   |
|---------------------------|-------------------|---------------------------|
| app.kubernetes.io/version | Version string    | 1.4.2 , 2025.03.15-hotfix |
| app.kubernetes.io/part-of | Application group | checkout-platform         |
| app.kubernetes.io/name    | Component name    | payment-service           |

Example Deployment snippet:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: payment-service
  labels:
    app.kubernetes.io/name: payment-service
    app.kubernetes.io/part-of: checkout-platform
    app.kubernetes.io/version: "1.4.2"
spec:
  template:
    metadata:
      labels:
        app.kubernetes.io/name: payment-service
        app.kubernetes.io/part-of: checkout-platform
        app.kubernetes.io/version: "1.4.2"
```

## Optional: Custom Release Mapping

If your deployment model uses non-standard labels, add namespace annotations to override the defaults:

```
apiVersion: v1
kind: Namespace
metadata:
  name: checkout
  annotations:
    mapping.release.dynatrace.com/version: "metadata.labels['app-version']"
    mapping.release.dynatrace.com/product: "metadata.labels['product-name']"
    mapping.release.dynatrace.com/stage: "metadata.namespace"
```

## Verify in Dynatrace

After deploying workloads with the required labels:

1. Navigate to **Releases** in the Dynatrace UI
2. Confirm your services appear with version information
3. Deploy a new version and verify the Release Inventory updates

## 8. Create Segments

---

Segments are the Gen3 replacement for Management Zones. They provide scoped views of your environment using enriched metadata fields.

### Settings Path

**Settings > Ownership & Organization > Segments**

### Recommended Segments

Create segments based on the enrichment rules configured in Section 6:

| Segment Name                             | Filter                                               | Use Case                          |
|------------------------------------------|------------------------------------------------------|-----------------------------------|
| Per-team: Team:<br>Backend               | <code>dt.security_context == "backend"</code>        | Team-scoped dashboards and alerts |
| Per-team: Team:<br>Frontend              | <code>dt.security_context == "frontend"</code>       | Team-scoped dashboards and alerts |
| Per-cluster: Cluster:<br>Production      | <code>k8s.cluster.name == "prod-us-east-1"</code>    | Cluster-level views               |
| Per-namespace:<br>Namespace: Checkout    | <code>k8s.namespace.name == "checkout"</code>        | Application-scoped monitoring     |
| Per-environment:<br>Environment: Staging | <code>k8s.cluster.name == "staging-us-east-1"</code> | Environment isolation             |

## Creating a Segment via API

```
curl -X POST
"https://<ENVIRONMENT_ID>.apps.dynatrace.com/platform/classic/environment-api/v2/settings/objects" \
-H "Authorization: Api-Token <TOKEN>" \
-H "Content-Type: application/json" \
-d '{
  "schemaId": "builtin:ownership.segments",
  "scope": "environment",
  "value": {
    "name": "Team: Backend",
    "description": "All telemetry owned by the backend team",
    "variables": {"type": "query", "value": "fetch logs | filter
dt.security_context == \"backend\""}
  }
}'
```

**Note:** Segments replace Management Zones in Dynatrace Gen3. Unlike Management Zones, segments are dynamic and work across all Grail data — not just entity-based data.

**See also:** ORGNZ-06 (Segments & Data Access) for comprehensive segment design patterns.

## 9. Verify Deployment

Run these DQL queries to confirm every component is working correctly.

### 9.1 Verify Dynatrace Components Are Running

```
// Verify Dynatrace components are running
fetch logs, from:-1h
| filter k8s.namespace.name == "dynatrace"
| summarize count = count(), by:{k8s.pod.name}
| sort count desc
| limit 20
```

Expected: You should see pods for `dynatrace-operator` , `dynatrace-webhook` , `dynakube-oneagent-*` , and `dynakube-activegate-*` .

## 9.2 Verify Host Monitoring

```
// Check all monitored hosts
fetch dt.entity.host
| fieldsKeep entity.name, host_group, state
| sort entity.name asc

// Alternative: Smartscape on Grail (entity.name → name)
// smartscapeNodes HOST
// | fieldsKeep name, host_group, state
// | sort name asc
```

Expected: Every node in your cluster should appear with the `host_group` matching your `<cluster-name>` value from the DynaKube.

## 9.3 Verify Container Metrics Are Flowing

```
// Top containers by CPU
timeseries avgCpu = avg(dt.kubernetes.container.cpu_usage), from:-1h, by:
{k8s.container.name, k8s.namespace.name}
| fieldsAdd avgCpuValue = arrayAvg(avgCpu)
| sort avgCpuValue desc
| limit 10
```

Expected: Container CPU metrics from your labeled namespaces. If this returns no data, verify namespace labels and wait for data to propagate (up to 5 minutes).

## 9.4 Verify Metadata Enrichment

```
// Check enrichment is working – dt.security_context should be populated
fetch logs, from:-1h
| filter isNotNull(dt.security_context)
| summarize count = count(), by:{dt.security_context}
| sort count desc
```

Expected: Logs grouped by the `team` label you assigned to namespaces. If empty, verify enrichment rules are configured (Section 6) and wait up to 45 minutes for

propagation.

## 9.5 Verify Build Propagation (Release Inventory)

```
// Check Release Inventory population
fetch dt.entity.process_group_instance
| filter isNotNull(softwareVersion)
| fieldsKeep entity.name, softwareVersion
| limit 10

// Alternative: Smartscape on Grail (entity.name → name)
// smartscapeNodes PROCESS
// | filter isNotNull(softwareVersion)
// | fieldsKeep name, softwareVersion
// | limit 10
```

Expected: Process groups with version information from `app.kubernetes.io/version` labels. If empty, verify your pods carry the required labels (Section 7).

## 9.6 Verify ActiveGate Health

```
// ActiveGate pod CPU – should show 2 pods across zones
timeseries agCpu = avg(dt.kubernetes.container.cpu_usage), from:-1h,
  filter:{k8s.namespace.name == "dynatrace" and
  matchesValue(k8s.container.name, "*activegate*")},
  by:{k8s.pod.name}
| fieldsAdd avgCpu = arrayAvg(agCpu)
| fields k8s.pod.name, avgCpu
```

Expected: Two ActiveGate pods with CPU metrics. If only one pod appears, check the topology spread constraints and verify your cluster has multiple availability zones.

# 10. Monitoring-the-Monitoring

---

After deployment, set up ongoing health checks for the Dynatrace components themselves.

## 10.1 OneAgent Memory Usage per Node

```
// OneAgent memory consumption per node
timeseries oaMem = avg(dt.kubernetes.container.memory_working_set), from:-1h,
  filter:{k8s.namespace.name == "dynatrace" and
  matchesValue(k8s.container.name, "*oneagent*")},
  by:{k8s.pod.name}
| fieldsAdd avgMemMB = arrayAvg(oaMem) / 1048576
| sort avgMemMB desc
```

Healthy range: 300-800 MB per OneAgent pod. If consistently above 1 GB, check for excessive process groups or deep monitoring of high-cardinality applications.

## 10.2 Dynatrace Component Failures

```
// K8s events for OOMKilled, CrashLoopBackOff in dynatrace namespace
fetch events, from:-24h
| filter k8s.namespace.name == "dynatrace"
| filter event.kind == "K8S_EVENT"
| fields timestamp, event.name, event.kind, k8s.pod.name
| sort timestamp desc
| limit 30
```

Expected: No `OOMKilled` or `CrashLoopBackOff` events. If present, increase resource limits in the DynaKube CR.

## 10.3 Metric Data Gaps

```
// Check for gaps in host CPU metrics (indicates OneAgent interruptions)
timeseries hostCpu = avg(dt.host.cpu.usage), from:-6h, by:{dt.entity.host}
| fieldsAdd dataPoints = arraySize(hostCpu)
| fieldsAdd expectedPoints = 72
| fieldsAdd completeness = toDouble(dataPoints) / toDouble(expectedPoints) *
100.0
| filter completeness < 95.0
| fields dt.entity.host, completeness, dataPoints
| sort completeness asc
```

Expected: All hosts at 100% completeness. Hosts below 95% indicate OneAgent restarts, pod evictions, or network interruptions to the Dynatrace backend.

## 11. Recommended Next Steps

With the base deployment complete, these additional configurations maximize the value of your Dynatrace Kubernetes monitoring:

| Priority | Configuration              | What It Enables                                                 | Reference     |
|----------|----------------------------|-----------------------------------------------------------------|---------------|
| High     | OpenPipeline log routing   | Route logs to buckets by namespace, set retention tiers         | OPLOGS series |
| High     | Custom Grail buckets       | Separate retention for prod vs. non-prod, compliance data       | ORGNZ-02      |
| High     | anomaly detectors          | Automated baseline alerting for CPU, memory, pod restarts       | K8S-09        |
| Medium   | Kubernetes dashboards      | Cluster health overview, workload summary, namespace drill-down | DASH series   |
| Medium   | Alerting workflows         | detected problem → Slack, Teams, PagerDuty, Jira                | WFLOW series  |
| Medium   | Prometheus scraping        | Collect custom app metrics via pod annotations                  | K8S-13        |
| Low      | Multi-cluster federation   | Unified view across dev/staging/prod clusters                   | K8S-04        |
| Low      | GitOps DynaKube management | Version-control DynaKube CR with ArgoCD or Flux                 | K8S-03        |

## 12. Summary

### Deployment Checklist

| Step                      | Status | Notes                                       |
|---------------------------|--------|---------------------------------------------|
| 1. Prerequisites verified | [ ]    | K8s 1.24+, Helm 3.x, network access         |
| 2. Access tokens created  | [ ]    | API token + Data Ingest token via templates |
| 3. Operator installed     | [ ]    | Helm chart + token secret                   |

| Step                           | Status | Notes                                             |
|--------------------------------|--------|---------------------------------------------------|
| 4. DynaKube applied            | [ ]    | v1beta6 with all production settings              |
| 5. Namespaces labeled          | [ ]    | <code>dt-monitoring=true</code> on app namespaces |
| 6. Enrichment rules configured | [ ]    | team, cost-center, part-of                        |
| 7. Build labels added          | [ ]    | <code>app.kubernetes.io/version</code> on pods    |
| 8. Segments created            | [ ]    | Per-team, per-cluster, per-namespace              |
| 9. Deployment verified         | [ ]    | All DQL verification queries passing              |
| 10. Monitoring-the-monitoring  | [ ]    | OneAgent health, AG health, data gaps             |

## Key Takeaways

| Concept             | Recommendation                                                                      |
|---------------------|-------------------------------------------------------------------------------------|
| Injection model     | Opt-in via <code>namespaceSelector</code> — never inject into system namespaces     |
| Failure policy      | <code>fail</code> — surface problems immediately, do not run uninstrumented         |
| ActiveGate          | 2 replicas with zone-aware topology spread                                          |
| Metadata enrichment | Enable from day one — retrofitting is expensive                                     |
| Build propagation   | Standard K8s labels ( <code>app.kubernetes.io/*</code> ) — no custom tooling needed |
| Segments            | Replace Management Zones for Gen3 — use enriched fields as filters                  |
| Verification        | Run DQL checks after every deployment change                                        |

## Additional Resources

- [Dynatrace Operator Helm Chart](#)
- [DynaKube Custom Resource Reference](#)

- [Kubernetes Monitoring Overview](#)
  - [Metadata Enrichment Guide](#)
  - [Release Inventory](#)
  - [Segments Documentation](#)
- 

*This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.*