

Namespace Organization and Boundaries

Series: K8S | **Notebook:** 6 of 9 | **Created:** January 2026

Organizing Kubernetes Monitoring with Namespaces

Namespaces provide logical boundaries in Kubernetes for resource isolation, access control, and organizational structure. This notebook covers namespace strategies and how to leverage them in Dynatrace for filtered views, access control, and cost allocation.

Table of Contents

1. Namespace Strategies
 2. Namespace Monitoring in Dynatrace
 3. Resource Quotas and Limits
 4. Namespace-Based Access Control
 5. Multi-Tenant Clusters
 6. Cost Allocation by Namespace
 7. Namespace Best Practices
 8. Next Steps
-

Prerequisites

Requirement	Details
Dynatrace Environment	SaaS with Kubernetes monitoring
DynaKube	Deployed with namespace selector configured
Permissions	entities.read , metrics.read
Knowledge	K8S-01 Fundamentals, K8S-05 Workload Monitoring

1. Namespace Strategies

Common Namespace Patterns

Strategy	Structure	Use Case
By Environment	dev , staging , prod	Single cluster, multiple envs
By Team	team-checkout , team-catalog	Team ownership
By Application	app-frontend , app-backend	Application isolation
By Function	monitoring , logging , ingress	Infrastructure services
Hybrid	team-checkout-prod	Combined approach

Namespace Naming Conventions

<prefix>-<owner>-<env>

Examples:

app-checkout-prod

team-platform-dev

infra-monitoring

System Namespaces

Namespace	Purpose	Monitor?
kube-system	Core K8s components	Yes (critical)
kube-public	Public resources	Minimal
kube-node-lease	Node heartbeats	No
default	Catch-all	Discourage use

```
// List all namespaces with workload counts
fetch dt.entity.cloud_application_namespace
| fields entity.name, kubernetesClusterName, tags
| sort kubernetesClusterName asc, entity.name asc
| limit 50
```

2. Namespace Monitoring in Dynatrace

DynaKube Namespace Selector

Control which namespaces receive OneAgent injection:

```
spec:
  namespaceSelector:
    matchLabels:
      monitoring: dynatrace
```

Or exclude specific namespaces:

```
spec:
  namespaceSelector:
    matchExpressions:
      - key: monitoring
        operator: NotIn
        values:
          - disabled
```

Namespace Labels for Filtering

Label	Purpose	Example
team	Team ownership	team: checkout
env	Environment	env: production
cost-center	Billing allocation	cost-center: eng-123
monitoring	Injection control	monitoring: enabled

```
// Namespace resource usage summary
timeseries cpuUsage = avg(dt.containers.cpu.usage_percent), by:
{k8s.namespace.name}
| sort avg(cpuUsage) desc
| limit 15
```

```
// Memory usage by namespace
timeseries memUsage = avg(dt.containers.memory.usage_percent), by:
{k8s.namespace.name}
| sort avg(memUsage) desc
| limit 15
```

3. Resource Quotas and Limits

ResourceQuota Monitoring

ResourceQuotas limit total resource consumption per namespace:

```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: compute-quota
  namespace: team-checkout
spec:
  hard:
    requests.cpu: "10"
    requests.memory: 20Gi
    limits.cpu: "20"
    limits.memory: 40Gi
    pods: "50"
```

LimitRange for Defaults

```
apiVersion: v1
kind: LimitRange
metadata:
  name: default-limits
  namespace: team-checkout
spec:
  limits:
    - default:
        cpu: "500m"
        memory: "512Mi"
      defaultRequest:
        cpu: "100m"
        memory: "128Mi"
      type: Container
```

Quota Utilization Metrics

Metric	Alert When
CPU quota usage	>80% of quota
Memory quota usage	>80% of quota
Pod count	Approaching limit

```
// CPU requests by namespace (quota tracking)
timeseries cpuRequests = sum(dt.kubernetes.workload.requests_cpu), by:
{k8s.namespace.name}
| sort avg(cpuRequests) desc
| limit 15
```

```
// Memory requests by namespace
timeseries memRequests = sum(dt.kubernetes.workload.requests_memory), by:
{k8s.namespace.name}
| sort avg(memRequests) desc
| limit 15
```

4. Namespace-Based Access Control

Dynatrace IAM with Boundaries

Use Dynatrace boundaries to restrict visibility by namespace:

Policy Example (Grant access to checkout namespace):

```
{
  "name": "pol-namespace-checkout-viewer",
  "statement": {
    "effect": "ALLOW",
    "permissions": ["environment:roles:viewer"],
    "conditions": [
      {
        "kubernetes.namespace": "checkout"
      }
    ]
  }
}
```

Segment by Namespace

Create segments for namespace-based filtering:

Segment	Filter	Use Case
seg-checkout	k8s.namespace.name == "checkout"	Team view
seg-production	k8s.namespace.name matches "*-prod"	Prod only
seg-platform	k8s.namespace.name in ("monitoring","logging")	Platform team

Kubernetes RBAC Alignment

Align Dynatrace access with Kubernetes RBAC:

K8s Role	Dynatrace Access
namespace-admin	Full namespace visibility
namespace-viewer	Read-only namespace data
cluster-admin	All namespaces

5. Multi-Tenant Clusters

Tenant Isolation Strategies

Level	Mechanism	Dynatrace Support
Soft	Namespace + NetworkPolicy	Boundaries, segments
Hard	Separate clusters	Separate DynaKubes
Virtual	vCluster	Namespace labels

DynaKube per Tenant (Hard Isolation)

For strong isolation, use separate DynaKube instances:

```
# Tenant A
apiVersion: dynatrace.com/v1beta3
kind: DynaKube
metadata:
  name: dynakube-tenant-a
  namespace: dynatrace
spec:
  apiUrl: https://tenant-a.live.dynatrace.com/api
  namespaceSelector:
    matchLabels:
      tenant: a
```

Shared DynaKube (Soft Isolation)

Single DynaKube with boundary-based access control:

```
spec:
  oneAgent:
    cloudNativeFullStack: {}
  # All namespaces monitored, access controlled via IAM
```

```
// Workload distribution across namespaces
fetch dt.entity.cloud_application
| summarize workloadCount = count(), by:{cloudApplicationNamespaces}
| sort workloadCount desc
| limit 20
```

6. Cost Allocation by Namespace

Resource-Based Cost Allocation

Calculate costs based on resource consumption:

Resource	Cost Metric	Calculation
CPU	Core-hours	CPU usage × hours × rate
Memory	GB-hours	Memory usage × hours × rate
Storage	GB-hours	PVC size × hours × rate

Labels for Cost Centers

```
apiVersion: v1
kind: Namespace
metadata:
  name: team-checkout
  labels:
    cost-center: "eng-checkout-123"
    team: "checkout"
    budget-owner: "alice@example.com"
```

```
// Resource consumption by namespace (for cost allocation)
timeseries
  avgCpu = avg(dt.containers.cpu.usage_percent),
  avgMem = avg(dt.containers.memory.usage_percent),
  by:{k8s.namespace.name}
| sort avg(avgCpu) desc
| limit 15
```

7. Namespace Best Practices

Naming and Organization

Practice	Reason
Use consistent naming	Easier filtering and automation
Apply standard labels	Cost allocation, access control
Avoid <code>default</code> namespace	Encourages explicit organization
Document ownership	Clear responsibility

Resource Management

Practice	Reason
Set ResourceQuotas	Prevent resource hogging
Configure LimitRanges	Ensure defaults
Monitor quota usage	Proactive capacity planning

Security

Practice	Reason
Apply NetworkPolicies	Namespace isolation
Use RBAC per namespace	Least privilege
Separate sensitive workloads	Compliance requirements

Monitoring

Practice	Reason
Label namespaces for filtering	Easy DQL queries
Create namespace-based dashboards	Team-specific views
Set up namespace alerts	Targeted notifications

Next Steps

With namespace organization in place, proceed to:

Next Notebook	Topic
K8S-07: Events and Logs	Kubernetes event analysis
K8S-08: DQL for Kubernetes	Advanced query patterns
K8S-09: Troubleshooting	Debugging K8s monitoring

Summary

In this notebook, you learned:

- Namespace strategies and naming conventions
- DynaKube namespace selector configuration
- Resource quota and limit monitoring
- Namespace-based access control with boundaries
- Multi-tenant cluster patterns
- Cost allocation by namespace
- Namespace best practices

References

- [Kubernetes Namespaces](#)
 - [Resource Quotas](#)
 - [Dynatrace Boundaries](#)
-

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.