

M2S-09: Step 9 — Optimize: Validate, Optimize, and Decommission

Series: M2S | **Notebook:** 9 of 9 | **Phase:** Run | **Step:** Optimize | **Created:** March 2026 | **Last Updated:** 03/30/2026

The migration is functionally complete. Agents are reporting, configurations are applied, integrations are reconnected, SaaS-exclusive features are adopted, and users are trained. This final step closes the loop: validate that every success criterion is met, optimize the SaaS environment for long-term performance, obtain stakeholder sign-off, and decommission the Managed cluster.

M2S Migration Journey — 3 Phases / 9 Steps

Plan: 1. Discover | 2. Strategize | 3. Design

Upgrade: 4. Prepare | 5. Execute | 6. Integrate

Run: 7. Expand | 8. Enable | **9. Optimize**

Table of Contents

1. [Migration Validation](#)
2. [Performance Optimization](#)
3. [Davis AI Baseline Establishment](#)
4. [Data Retention Optimization](#)
5. [Migration Success Metrics](#)
6. [Stakeholder Sign-Off](#)
7. [Managed Decommissioning](#)
8. [Series Conclusion](#)
9. [Step Completion Checklist](#)

Optimization Use Cases

With the migration complete, focus optimization efforts on these SaaS-enabled use cases:

Use Case	SaaS Capability	Priority
Automated incident management	Workflows + problem triggers	High
Automated remediation / self-healing	Workflows + HTTP request actions	High
Automated release validation	Site Reliability Guardian	Medium
Digital experience management	RUM, Session Replay, Synthetic	High
Application security management	Runtime vulnerability detection, code-level protection	High
Log management and analytics	Grail + OpenPipeline	High
Business analytics	Business events + Grail	Medium

Operational Procedure Review

Review and update operational procedures now that you're on SaaS:

Procedure	Update Required
Release schedules	Dynatrace SaaS updates bi-weekly automatically — no action needed
Notifications and alerting	Review thresholds now that Davis AI baselines are established
High availability	SaaS provides 99.5%+ SLA — document this in your HA plans
Capacity planning	No more cluster capacity planning — SaaS scales with licensing

Procedure	Update Required
Decommission old infrastructure	Remove Managed cluster servers, old ActiveGates, old network rules

Prerequisites

Requirement	Details
Steps 1-8 Complete	All migration phases finished — Discovery through Enable
SaaS Tenant Active	All agents reporting, configurations applied, integrations connected
Entity Inventory from Step 1	Original Managed entity counts for comparison
Managed Environment Running	Still accessible for comparison (not yet decommissioned)
Stakeholder Access	Ability to obtain sign-off from platform, security, application, and executive stakeholders

Order of Operations — You Are Here

This notebook covers operations **9-11** of the 11-step Order of Operations (see **M2S-02: Step 2 — Strategize**) defined in Step 2:

#	Operation	Status
1	Assess — Inventory current environment	Step 4: Prepare 
2	Provision — SaaS tenant and access (SSO)	Step 4: Prepare 
3	Install — New ActiveGates in parallel	Step 4: Prepare 
4	Migrate — Configuration and integrations	Step 5: Execute 
5	Rebuild — Dashboards, zones, alerts	Step 5: Execute 
6	Redirect — OneAgents to SaaS	Step 5: Execute 
7	Reconnect — Integrations and extensions	Step 6: Integrate 

#	Operation	Status
8	Migrate — Remaining config and integrations	Step 6: Integrate 
9	Validate — Data flow and performance	This notebook
10	Cutover — Full switch to SaaS	This notebook
11	Decommission — Managed environment	This notebook

1. Migration Validation

Validation is the most critical activity in this step. A migration is not complete until every success criterion is verified with data. Compare current SaaS entity counts and data flows against the inventory captured during Step 1 (Discover).

1.1 Entity Coverage

Run each query below and compare the results against your Managed inventory from Step 1. Any discrepancy must be investigated before proceeding.

```
// Host count – compare to Step 1 inventory
fetch dt.entity.host
| summarize hostCount = count()
```

```
// Service count – compare to Step 1 inventory
fetch dt.entity.service
| summarize serviceCount = count()
```

```
// Application count – compare to Step 1 inventory
fetch dt.entity.application
| summarize appCount = count()
```

```
// Process group count – compare to Step 1 inventory
fetch dt.entity.process_group
| summarize processGroupCount = count()
```

```
// Synthetic monitor count – compare to Step 1 inventory
fetch dt.entity.synthetic_test
| summarize syntheticCount = count()
```

Entity Coverage Comparison

Fill in the table below with the results from the queries above and the inventory from Step 1:

Entity Type	Expected (Step 1 Inventory)	Actual (SaaS)	Status
Hosts	—	—	[] Match
Services	—	—	[] Match
Applications	—	—	[] Match
Process Groups	—	—	[] Match
Synthetic Monitors	—	—	[] Match
ActiveGates	—	—	[] Match

Note: Small discrepancies in service or process group counts are expected — SaaS may discover additional services through improved auto-detection. Investigate any entity type where the SaaS count is *lower* than the Managed inventory.

1.2 Data Flow Validation

Verify that all three pillars of observability — metrics, logs, and traces — are flowing continuously into SaaS with no gaps.

```
// Metric continuity – count hosts with active CPU data
timeseries avgCpu = avg(dt.host.cpu.usage), from:-1h, by:{dt.entity.host}
| fieldsAdd hasData = arrayAvg(avgCpu)
| filter isNotNull(hasData)
| summarize hostsWithData = count()
```

```
// Log ingestion continuity – volume per 5-minute bucket
fetch logs, from:-1h
| summarize logCount = count(), by:{time_bucket = bin(timestamp, 5m)}
| sort time_bucket desc
| limit 24
```

```
// Trace/span continuity – volume per 5-minute bucket
fetch spans, from:-1h
| summarize spanCount = count(), by:{time_bucket = bin(timestamp, 5m)}
| sort time_bucket desc
| limit 12
```

```
// Log sources – verify all expected sources are present
fetch logs, from:-1h
| summarize count = count(), by:{log.source}
| sort count desc
| limit 10
```

1.3 Data Gaps Detection

Gaps in data flow indicate agents that lost connectivity during migration or misconfigured network routes. A healthy migration shows consistent counts across all time buckets.

```
// Count reporting hosts per 5-minute bucket – drops indicate gaps
timeseries avgCpu = avg(dt.host.cpu.usage), from:-6h, by:{dt.entity.host}
| fieldsAdd hasData = arrayAvg(avgCpu)
| filter isNotNull(hasData)
| summarize hostsReporting = count()
```

```
// Service request volume over time – look for unexpected drops
fetch spans, from:-6h
| filter span.kind == "server"
| makeTimeseries requestCount = count(), interval: 5m
```

Tip: Run these gap-detection queries across a 6-hour or 24-hour window. A single 5-minute bucket with zero data is acceptable during a maintenance window, but multiple consecutive empty buckets indicate a connectivity problem that must be resolved.

2. Performance Optimization

With data flowing and entities validated, optimize the SaaS environment for long-term operational efficiency. Focus on query performance, dashboard responsiveness, and alert signal quality.

2.1 Query Performance

Queries migrated from Managed may not follow SaaS best practices. Review and optimize all frequently-used queries:

Optimization	Technique	Impact
Add time filters	Always use <code>from:</code> on every <code>fetch</code> and <code>timeseries</code> query	Reduces scan scope dramatically
Filter early	Apply <code>filter</code> immediately after <code>fetch</code> , before any transforms	Reduces records processed in downstream stages
Select fields early	Use <code>fieldsKeep</code> or <code>fieldsRemove</code> to drop unneeded columns	Reduces memory usage per record
Use aggregations	Replace raw-data queries with <code>summarize</code> or <code>makeTimeseries</code>	Returns summary instead of millions of rows
Target buckets	Use <code>bucket:</code> parameter to limit scan to specific Grail buckets	Avoids scanning irrelevant data
Limit results	Add <code>limit</code> as the final command	Prevents oversized result sets

2.2 Dashboard Optimization

Slow dashboards are the most visible symptom of unoptimized queries. Prioritize dashboards used by operations and executives:

Issue	Root Cause	Solution
Slow loading	Too many tiles, unoptimized queries	Reduce tile count to < 20; optimize each tile's DQL

Issue	Root Cause	Solution
Tile timeouts	Missing time filters, broad scan scope	Add <code>from:</code> and <code>bucket:</code> to every query
Excessive data	Raw data tiles instead of aggregations	Replace with <code>summarize</code> or <code>makeTimeseries</code>
Redundant queries	Multiple tiles querying the same data	Use shared variables or combine into fewer queries

2.3 Alert Tuning

Alert quality is critical in the first weeks after migration. Expect initial noise as Davis AI establishes baselines:

Symptom	Likely Cause	Action
Too many alerts	Thresholds too sensitive for SaaS baseline	Temporarily raise thresholds; revisit after 2 weeks
Missing alerts	Notification channels not fully reconnected	Re-verify every channel from Step 6
Noisy alerts	Transient issues during baseline period	Add minimum duration conditions (e.g., 5 consecutive minutes)
Duplicate alerts	Overlapping alerting rules migrated from Managed	Consolidate into fewer, broader rules

Important: Do NOT disable alerting during the baseline period. Instead, route non-critical alerts to a staging channel (e.g., a dedicated Slack channel) so the team can review without alert fatigue.

3. Davis AI Baseline Establishment

Davis AI uses machine learning to establish behavioral baselines for every entity in the environment. In a new SaaS environment, Davis starts with no historical data and must learn normal behavior from scratch.

Baseline Timeline

Baseline Type	Time Required	What Davis Learns
Response time	2-7 days	Normal service response patterns
Error rates	2-7 days	Expected error levels per service
Resource usage	7-14 days	CPU, memory, disk patterns per host
Traffic patterns	7-14 days	Daily and weekly traffic cycles
Seasonal patterns	4-6 weeks	Monthly or periodic business cycles

What to Expect During Baseline Period

Week	Behavior
Week 1	Higher alert volume — Davis flags many deviations as it lacks history
Week 2	Alert volume decreases — short-term baselines established
Weeks 3-4	Near-normal alert volume — weekly patterns learned
Weeks 5-6	Stable — Davis has sufficient history for accurate anomaly detection

Warning: Expect more alerts in the first 1-2 weeks. This is normal and expected. Do NOT disable Davis AI or suppress problem detection. Instead, review each alert to confirm Davis is detecting real issues versus baseline noise. The learning period is essential for long-term accuracy.

Monitor Davis Problem Trends

Track the problem trend over the first weeks post-migration. A healthy baseline establishment shows declining problem counts over time:

```
// Problem trend over last 7 days – expect a declining curve as baselines establish
fetch dt.davis.problems, from:-7d
| makeTimeseries problemCount = count(default: 0), interval: 1d
```

```
// Problem breakdown by category – identify which types are noisiest
fetch dt.davis.problems, from:-7d
| summarize count = count(), by:{event.category}
| sort count desc
```

```
// Mean time to resolve – track improvement over baseline period
fetch dt.davis.problems, from:-7d
| filter event.status == "CLOSED"
| filter dt.davis.is_frequent_event == false and dt.davis.is_duplicate ==
false
| makeTimeseries avgDurationHours = avg(toLong(resolved_problem_duration) /
3600000000000.0), time: event.end
```

4. Data Retention Optimization

Grail provides flexible data retention through bucket configuration. Optimize retention to balance compliance requirements, operational needs, and cost.

Retention Strategy by Data Type

Data Type	Recommended Retention	Rationale
Logs (operational)	35 days	Standard troubleshooting window
Logs (compliance/audit)	90-365 days	Regulatory requirements
Metrics	5 years (aggregated)	Built-in aggregated retention; no action needed
Spans/traces	35 days	Sufficient for root cause analysis
Business events	90-365 days	Business intelligence and trend analysis
Davis problems	365 days	Year-over-year reliability comparison

Bucket Targeting for Query Performance

Use the `bucket:` parameter in queries to target specific Grail buckets and avoid scanning irrelevant data:

```
// Target specific log buckets instead of scanning all logs
fetch logs, bucket:{"app_logs_*"}, from:-1h
| filter loglevel == "ERROR"
| summarize count = count(), by:{log.source}
```

Cost Optimization Tips

Technique	Impact
Set shorter retention for high-volume, low-value data	Directly reduces storage costs
Route verbose debug logs to a separate bucket with 7-day retention	Keeps debug data available without long-term cost
Use <code>samplingRatio:</code> for exploratory queries on large datasets	Reduces query cost for ad-hoc analysis
Archive compliance data to external storage if Grail retention exceeds needs	Avoids paying Grail rates for cold storage

5. Migration Success Metrics

Define and measure success quantitatively. Each metric has a target — the migration is complete when all targets are met.

Metric	Target	Actual	Status
Host coverage	100% of Managed inventory	—	[]
Service coverage	100% of Managed inventory	—	[]
Application coverage	100% of Managed inventory	—	[]
Data gaps	< 15 minutes total	—	[]
Alert delivery	100% of channels verified	—	[]

Metric	Target	Actual	Status
Dashboard availability	100% of dashboards rendering	—	[]
Integration success	100% of integrations reconnected	—	[]
User access	100% of users can log in via SSO	—	[]

Note: Do not proceed to stakeholder sign-off until every metric shows a passing status. Any failing metric must be resolved first.

6. Stakeholder Sign-Off

Formal sign-off marks the official end of the migration project. Each stakeholder group validates the aspects of the migration relevant to their responsibility.

Sign-Off Matrix

Stakeholder	What They Validate	Approval	Date
Platform Team Lead	Entity coverage, data flow, alert delivery, dashboard availability	[]	
Security Team	SSO configuration, IAM policies, data masking, compliance certifications	[]	
Application Teams	Service visibility, trace data, application-specific dashboards	[]	
Operations Team	Alert channels, runbook updates, incident workflow connectivity	[]	
Executive Sponsor	Overall migration success, budget, timeline adherence	[]	

Sign-Off Preparation

Before scheduling sign-off meetings, prepare the following:

1. **Migration summary report** — Entity counts, data flow validation results, success metrics
 2. **Issue log** — Any open issues with resolution timeline
 3. **Risk register** — Remaining risks (e.g., Davis baselines still establishing)
 4. **Decommission plan** — Timeline and steps for Managed shutdown
 5. **Training completion report** — From Step 8 (Enable)
-

7. Managed Decommissioning

This is the point of no return. Decommissioning the Managed environment is the final act of the migration. Follow the order of operations carefully — rushing this step risks data loss and operational disruption.

Order of Operations

These correspond to items 9-11 in the overall migration order of operations: Validate, Cutover, Decommission.

Step	Action	Timing	Responsible
1	Confirm no agents pointing to Managed — Verify every OneAgent and ActiveGate is reporting to SaaS	Immediate	Platform Team
2	Keep Managed running for comparison — Run both environments in parallel for a validation period	2-4 weeks post-migration	Platform Team
3	Export final Managed data and reports — Download any data, dashboards, or reports not migrated	Before shutdown	Platform Team
4	Archive SaaS Upgrade Assistant deploy result CSVs — These document what was migrated and any failures	Before shutdown	Platform Team

Step	Action	Timing	Responsible
5	Archive entity inventory and config mapping docs — Preserve the Step 1 inventory and configuration mapping for audit	Before shutdown	Platform Team
6	Decommission Managed cluster — Shut down Managed servers, release infrastructure	After validation period + sign-off	Infrastructure Team

Pre-Decommission Verification

```
// Final host count in SaaS – this is the definitive count before Managed shutdown
fetch dt.entity.host
| summarize finalHostCount = count()
```

```
// Final service count in SaaS
fetch dt.entity.service
| summarize finalServiceCount = count()
```

```
// Verify no data gaps in the last 24 hours before decommission
fetch logs, from:-24h
| summarize logCount = count(), by:{time_bucket = bin(timestamp, 1h)}
| sort time_bucket asc
```

What to Archive Before Shutdown

Artifact	Format	Storage Location
SaaS Upgrade Assistant deploy result CSVs	CSV	Team SharePoint / Wiki
Entity inventory from Step 1	Spreadsheet	Team SharePoint / Wiki
Configuration mapping document	Spreadsheet or Markdown	Git repository
Managed dashboard screenshots	PNG/PDF	Team SharePoint

Artifact	Format	Storage Location
Migration timeline and decision log	Document	Project management tool
Stakeholder sign-off records	PDF or email	Project management tool

Post-Decommission Cleanup

After Managed is shut down:

Task	Details
Revoke Managed API tokens	Ensure no scripts still reference Managed endpoints
Update DNS records	Remove or redirect any DNS entries pointing to Managed cluster
Release infrastructure	Return servers, release cloud resources, cancel hosting contracts
Update documentation	Ensure all runbooks, architecture diagrams, and onboarding guides reference SaaS only
Notify stakeholders	Send formal notification that Managed is decommissioned

8. Series Conclusion

Congratulations on completing the M2S: Managed-to-SaaS Migration series. Here is a summary of what was accomplished across all nine steps:

Step	Phase	Achievement
1. Discover	Plan	Understood SaaS differences, inventoried Managed environment, assessed readiness
2. Strategize	Plan	Defined migration approach, timeline, and risk mitigation strategy

Step	Phase	Achievement
3. Design	Plan	Created target SaaS architecture, network topology, and IAM structure
4. Prepare	Upgrade	Provisioned SaaS tenant, deployed ActiveGates, configured SSO
5. Execute	Upgrade	Migrated configurations via SaaS Upgrade Assistant, migrated OneAgents
6. Integrate	Upgrade	Reconnected dashboards, alerting, CI/CD, ITSM, extensions, and API scripts
7. Expand	Run	Adopted SaaS-exclusive capabilities — Grail, Notebooks, Workflows, Davis Copilot
8. Enable	Run	Trained users, updated documentation, established support processes
9. Optimize	Run	Validated migration, optimized performance, obtained sign-off, decommissioned Managed

Ongoing Success

The migration is complete, but the journey continues. Here are the key activities for the first 90 days post-migration:

Timeframe	Focus
Weeks 1-2	Monitor Davis AI baseline establishment; review and tune alerts
Weeks 3-4	Optimize dashboards and queries; resolve any remaining integration issues
Weeks 5-8	Deepen adoption of SaaS-exclusive features; expand automation workflows
Weeks 9-12	Conduct first quarterly review; measure ROI against pre-migration baseline

Additional Resources

- [SaaS Upgrade Assistant Documentation](#)
- [Upgrading from Dynatrace Managed to SaaS](#)

- [Dynatrace Documentation](#)
 - [Dynatrace Community](#)
 - [Dynatrace University](#)
 - [Dynatrace Community: Upgrade to SaaS](#)
-

9. Step Completion Checklist

The migration is complete when every item below is confirmed.

Checkpoint	Status
Entity counts match inventory from Step 1	[]
No data gaps > 15 minutes in metrics, logs, or traces	[]
All alerts delivering correctly to all channels	[]
All dashboards showing data with acceptable performance	[]
Davis AI baselines establishing (2+ weeks of data)	[]
Query and dashboard performance optimized	[]
Data retention configured per compliance requirements	[]
Stakeholder sign-off obtained from all parties	[]
Managed environment running for comparison period (2-4 weeks)	[]
Migration artifacts archived (deploy results, inventories, mapping docs)	[]
Managed environment decommissioned	[]

Summary

In Step 9, you:

- Validated migration completeness by comparing SaaS entity counts against the Step 1 inventory

- Verified continuous data flow across metrics, logs, and traces with no gaps
- Optimized query performance, dashboards, and alert configuration for SaaS
- Understood the Davis AI baseline establishment timeline and what to expect
- Configured data retention to balance compliance, operational needs, and cost
- Measured migration success against quantitative targets
- Obtained formal stakeholder sign-off
- Executed the Managed decommissioning order of operations
- Archived all migration artifacts for future reference

Key Takeaway: A migration is only complete when every success criterion is validated with data, every stakeholder has signed off, and the Managed environment is safely decommissioned. The DQL queries in this notebook provide the evidence needed to close the project with confidence.

Thank you for completing the M2S: Managed-to-SaaS Migration Best Practices series!

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.