

MZ2POL-04: Policies and Boundaries

Series: MZ2POL | **Notebook:** 5 of 8 | **Created:** December 2025 | **Last Updated:** 01/28/2026

Overview

This notebook provides a comprehensive guide to **IAM Policies** and **Policy Boundaries** - the two components that together replace Management Zone access control. Policies define **WHAT** users can do, while boundaries define **WHERE** those permissions apply.

Table of Contents

1. [Understanding the Relationship](#)
 2. [Policy Fundamentals](#)
 3. [Default Policies](#)
 4. [Condition and Boundary Syntax](#)
 5. [Creating Boundaries](#)
 6. [Best Practice: Complete Boundary for MZ Restriction](#)
 7. [Recommended Group-Policy-Boundary Structure](#)
 8. [Mapping MZ Access to Policies + Boundaries](#)
 9. [Security Context and Bucket Strategy](#)
 10. [Custom Policies](#)
 11. [Best Practices](#)
 12. [Troubleshooting](#)
 13. [Additional Resources](#)
-

Prerequisites

- Completed MZ2POL-01 through MZ2POL-03
- Access to Dynatrace Account Management
- Policy Management permissions


Learning Objectives

By the end of this notebook, you will:

1. Understand policy statement syntax and structure
 2. Know the default policies and when to use them
 3. Understand boundary syntax and how boundaries restrict policies
 4. Be able to map MZ access patterns to policy + boundary combinations
 5. Know the recommended group-policy-boundary structure for SAML/AD integration
-


1. Understanding the Relationship

Policies vs Boundaries



Component	Purpose	Analogy
Policy	Defines WHAT actions are allowed	"You can read logs"
Boundary	Defines WHERE policies apply	"...but only for team-frontend"

How They Work Together



Key Characteristics

Policies:

- Define permissions (ALLOW statements)
- Can include conditions (WHERE clauses)
- Default policies cover most use cases

Boundaries:

- Separate from policies (decoupled)
- Reusable across multiple policies
- Can only narrow permissions, never expand
- Optional - policies work without them (full scope)

2. Policy Fundamentals

Policy Statement Structure

```
<effect> <service>:<resource>:<action> [WHERE  
<condition>]
```

Component	Description	Examples
Effect	Permission type	ALLOW (only option currently)
Service	Dynatrace service	storage , settings , app
Resource	Resource type	logs , buckets , objects
Action	Operation	read , write , delete
Condition	Optional filter	WHERE field = "value"

Basic Policy Statements

```
// Grant read access to all logs  
ALLOW storage:logs:read  
  
// Grant read/write to buckets  
ALLOW storage:buckets:read  
ALLOW storage:buckets:write  
  
// Grant all actions on a resource  
ALLOW storage:logs:*
```

Common Services and Resources

Service	Resources	Common Actions
storage	logs , spans , events , metrics , buckets	read , write , delete
settings	objects , schemas	read , write , delete
app	apps , functions	run , install
automation	workflows	read , write , run
document	documents	read , write , delete , share

3. Default Policies

Dynatrace provides **default policies** that cover common access patterns. These are **read-only** but cover most use cases.

Dynatrace Access Policies

Policy	Use Case	Key Permissions
Dynatrace Viewer	Read-only access	View dashboards, data, settings
Dynatrace Operator	Basic operations	Viewer + create dashboards, notebooks
Dynatrace Standard User	Regular users	Operator + modify some settings
Dynatrace Professional User	Power users	Standard + advanced features
Dynatrace Admin User	Administrators	Full platform access

Data Access Policies

Policy	Use Case	Key Permissions
Data Viewer	Read monitored data	Query logs, metrics, traces
Data Editor	Modify data configs	Viewer + data configuration

Choosing the Right Default Policy

User Type	→ Dynatrace Policy	+ Data Policy
Executives/Viewers	→ Dynatrace Viewer	+ Data Viewer
Developers	→ Dynatrace Standard	+ Data Viewer
SRE/Operations	→ Dynatrace Professional	+ Data Editor
Platform Admins	→ Dynatrace Admin	+ Data Editor

4. Condition and Boundary Syntax

Policies and boundaries share the same condition syntax.

Supported Operators

Operator	Description	Example
=	Exact match	field = "value"
!=	Not equal	field != "value"
startsWith	Prefix match	field startsWith "prefix"
in	Value in list	field in ("a", "b", "c")
IN	Value in list (alternative)	field IN ("a", "b")

Supported Fields

Field	Description	Used In
<code>environment</code>	Environment restrictions	Boundaries
<code>environment:name</code>	Environment by name	Boundaries
<code>environment:management-zone</code>	MZ-based (transitional)	Boundaries
<code>storage:dt.security_context</code>	Security context	Both
<code>storage:bucket</code>	Grail bucket	Both
<code>settings:schemaId</code>	Settings schema ID	Policies
<code>settings:dt.security_context</code>	Settings security context	Boundaries

Combining Conditions

AND logic (all conditions must match):

```
ALLOW storage:logs:read
  WHERE storage:dt.security_context = "team-a"
  AND storage:bucket = "production_logs"
```

OR logic (multiple statements or lines):

```
// In policies: Multiple statements
ALLOW storage:logs:read WHERE storage:dt.security_context = "team-a"
ALLOW storage:logs:read WHERE storage:dt.security_context = "team-b"

// In boundaries: Each line is OR-combined
storage:dt.security_context = "team-a"
storage:dt.security_context = "team-b"
```

5. Creating Boundaries

Via UI

1. Navigate to **Account Management** → **Identity & Access Management**
2. Select **Policy Boundaries** → **Boundaries** tab
3. Click **Create boundary**
4. Enter:
 - **Boundary name:** Descriptive name
 - **Boundary query:** Restriction conditions
5. Click **Save**

Boundary Examples

Team-Based Boundary:

```
Name: Frontend Team Scope  
Query: storage:dt.security_context = "team-frontend"
```

Environment Boundary:

```
Name: Production Only  
Query: storage:dt.security_context startsWith "prod-"
```

Multi-Team Boundary (OR logic):

```
Name: Platform Teams  
Query:  
storage:dt.security_context = "team-infra"  
storage:dt.security_context = "team-sre"  
storage:dt.security_context = "team-platform"
```

Boundary Limitations

Limitation	Description	Workaround
Max 10 lines	Only 10 conditions per boundary	Create multiple boundaries
No AND in lines	Each line is one condition	Use multiple boundaries for AND
No complex expressions	Limited to basic operators	Simplify conditions

6. Best Practice: Complete Boundary for MZ Restriction

When restricting access to replicate a Management Zone, use **all three domains** in your boundary query:

```
environment:management-zone IN ("LOB5");
storage:dt.security_context IN ("LOB5");
settings:dt.security_context IN ("LOB5");
```

Why All Three Domains?

Domain	What It Restricts
environment:management-zone	Classic entity access (hosts, services, processes)
storage:dt.security_context	Grail data (logs, spans, metrics, events)
settings:dt.security_context	Settings objects

Using only one domain leaves gaps in access control. For complete MZ replacement, include all three.

7. Recommended Group-Policy-Boundary Structure

For organizations using SAML/SSO with Active Directory, this structure provides clear separation of duties:

```
Group: "LOB5-Team" (SAML from Active Directory)
├─ Policy: Dynatrace Viewer (standard read access)
└─ Policy: Dynatrace Professional
    └─ Boundary:
        environment:management-zone IN ("LOB5");
        storage:dt.security_context IN ("LOB5");
        settings:dt.security_context IN ("LOB5");
```

Key Benefits

- **SAML Group:** Ties to existing AD group membership - no separate user management
- **Standard User Policy:** Provides baseline access without boundary (if needed)
- **Pro User + Boundary:** Advanced permissions restricted to team's scope

Alternative: All Policies with Boundary

For stricter isolation, apply the boundary to all policies:

```
Group: "LOB5-Team" (SAML from AD)
├─ Policy: Dynatrace Viewer
│   └─ Boundary: LOB5-Scope
└─ Policy: Dynatrace Professional
    └─ Boundary: LOB5-Scope
```

Binding via UI


1. Navigate to **Group Management**
2. Select or create a group
3. Go to **Permissions** tab
4. Click **Add permission**
5. Select:
 - **Policy:** Choose policy to bind

- **Boundary** (optional): Select restriction
- **Environment**: Select target environment

6. Click **Save**

8. Mapping MZ Access to Policies + Boundaries

Common MZ Permission Patterns



MZ Permission	Equivalent Policy + Boundary
View only	Dynatrace Viewer + Boundary
View + Edit	Dynatrace Standard User + Boundary
Full access to MZ	Dynatrace Professional User + Boundary
Admin in MZ	Custom policy + Boundary

Migration Example: Team-Based MZ

Before (MZ-based):

- Management Zone: "Frontend-Team"
- Rules: Services with tag `team: frontend`
- Users assigned via RBAC

After (Policy + Boundary):

1. Set security context on entities:
 - Services get `dt.security_context = "team-frontend"`
2. Create boundary:
 - Name: Frontend Team Scope
 - Query:
 - `environment:management-zone IN ("Frontend-Team");`
 - `storage:dt.security_context IN ("team-frontend");`
 - `settings:dt.security_context IN ("team-frontend");`
3. Bind to group:
 - Group: Frontend Developers (SAML)
 - Policy: Dynatrace Standard User
 - Boundary: Frontend Team Scope

Migration Example: Environment MZ

Before (MZ-based):

- Management Zone: "Production"
- Rules: Hosts with tag `env:production`
- Users get MZ-filtered view

After (Policy + Boundary):

1. Set security context on entities:
 - Hosts get `dt.security_context = "prod-{region}"`
2. Create boundary:
 - Name: Production Environment
 - Query:
 - `environment:management-zone IN ("Production");`
 - `storage:dt.security_context startsWith "prod-";`
 - `settings:dt.security_context startsWith "prod-";`
3. Bind to group:
 - Group: Production Operators (SAML)
 - Policy: Dynatrace Professional User
 - Boundary: Production Environment

9. Security Context and Bucket Strategy

What Is Security Context?

Security Context (`dt.security_context`) is an entity attribute used for access control. It's the recommended way to scope boundaries for Grail data.

Primary Grail Fields

When migrating from Management Zones, leverage **Primary Grail Fields** for consistent data partitioning:

Field	Source	MZ Migration Use
<code>k8s.cluster.name</code>	Kubernetes	Replace cluster-based MZs
<code>k8s.namespace.name</code>	Kubernetes	Replace namespace-based MZs
<code>aws.account.id</code>	AWS metadata	Replace AWS account MZs
<code>dt.host_group.id</code>	OneAgent config	Replace host group-based MZs

Bucket Strategy for Migration

Buckets partition Grail data and integrate with policies and boundaries for team-level access control.

Bucket Limitations

 **Critical:** Plan buckets carefully during migration - these constraints cannot be changed later.

Limitation	Details
One data type per bucket	Logs, metrics, events, OR spans - not mixed
Names are immutable	Bucket names CANNOT be changed after creation
No data migration	Data CANNOT be moved between buckets
Naming rules	3-100 chars, lowercase alphanumeric, underscores, hyphens only
Maximum buckets	80 per environment (default limit)
Optimal ingest	~1 TB/day per bucket for best query performance
Acceptable ingest	1-3 TB/day per bucket (limited query window)
Maximum ingest	3 TB/day per bucket hard limit

Query Constraints

Limit	Value	Impact
Maximum data scanned	500 GB	Limits queryable time window
Maximum records returned	1,000	Use aggregations for larger datasets
Maximum response payload	1 MB	Large result sets may be truncated

Default Bucket Retentions

Bucket	Retention
default_logs	35 days
default_metrics	15 months
default_spans	10 days
default_events	35 days

Using Buckets for Team Isolation (MZ Replacement)

When replacing team-based MZs, use buckets in **both policies AND boundaries**:

```
// Policy: Grant team access to their specific buckets
ALLOW storage:logs:read WHERE storage:bucket.name = "frontend_logs"
ALLOW storage:spans:read WHERE storage:bucket.name = "frontend_spans"
ALLOW storage:metrics:read WHERE storage:bucket.name = "frontend_metrics"

// Boundary: Restrict to team's buckets and security context
storage:bucket.name IN ("frontend_logs", "frontend_spans",
"frontend_metrics");
storage:dt.security_context IN ("team-frontend");
environment:management-zone IN ("Frontend-Team");
```

Migration Mapping: MZ to Buckets

MZ Pattern	Bucket Strategy
Team MZ (e.g., "Frontend-Team")	Team-specific buckets: frontend_logs , frontend_spans
Environment MZ (e.g., "Production")	Environment buckets: prod_logs , prod_metrics
Regional MZ (e.g., "US-East")	Region buckets: useast_logs , useast_spans
Compliance MZ (e.g., "PCI")	Compliance buckets: pci_logs , pci_events

Complete Team Migration Example

Before (MZ-based):

- Management Zone: "Checkout-Team"
- Users see all data within the MZ

After (Policy + Bucket + Boundary):

```
Group: Checkout-Team (SAML from AD)
├─ Policy: Checkout Data Access
│   └─ ALLOW storage:logs:read WHERE storage:bucket.name = "checkout_logs"
│   └─ ALLOW storage:spans:read WHERE storage:bucket.name = "checkout_spans"
│   └─ ALLOW storage:metrics:read WHERE storage:bucket.name =
"checkout_metrics"
└─ Boundary:
    storage:bucket.name IN ("checkout_logs", "checkout_spans",
"checkout_metrics");
    storage:dt.security_context IN ("checkout");
    environment:management-zone IN ("Checkout-Team");
```

For detailed bucket design guidance, see **ORGNZ-03: Bucket Strategy and Design** which covers naming conventions, retention planning, and cost attribution patterns in depth.

Setting Security Context

Security context can be set via:

- **Entity Enrichment rules** (recommended)
- **Host properties** via OneAgent
- **OpenPipeline processing**
- **Auto-tagging rules** (legacy)

Security Context Naming Strategy

Pattern	Example	Use Case
team-{name}	team-frontend	Team ownership
env-{name}	env-production	Environment separation
region-{code}	region-us-east	Geographic isolation
app-{name}	app-checkout	Application scoping
{env}-{team}	prod-frontend	Combined scoping

Query Entities by Security Context

```
// List services with their security context
// Identify services that need security context assignment
fetch dt.entity.service
| fields entity.name,
        dt.security_context,
        tags,
        managementZones
| sort dt.security_context asc
| limit 50
```

```
// Count entities by security context
// Helps verify security context distribution
fetch dt.entity.service
| summarize count = count(), by:{dt.security_context}
| sort count desc
```

10. Custom Policies

When to Create Custom Policies

- Default policies are too broad or narrow
- Need specific permission combinations
- Require conditional access
- Compliance requirements

Custom Policy Examples

Log Viewer for Specific Team:

```
// Policy Name: Frontend Team Log Access
// Description: Read-only access to frontend service logs

ALLOW storage:logs:read
  WHERE storage:dt.security_context = "team-frontend"

ALLOW storage:spans:read
  WHERE storage:dt.security_context = "team-frontend"
```

Environment-Restricted Admin:

```
// Policy Name: Development Admin
// Description: Full access to development environment only

ALLOW storage:*:*
  WHERE storage:dt.security_context startsWith "dev-"

ALLOW settings:objects:*
  WHERE settings:scope startsWith "environment:dev"
```

Dashboard and Notebook Creator:

```
// Policy Name: Dashboard Creator
// Description: Create and manage dashboards and notebooks

ALLOW document:documents:read
ALLOW document:documents:write
ALLOW document:documents:delete
ALLOW document:documents:share

// Read data for dashboards
ALLOW storage:logs:read
ALLOW storage:metrics:read
ALLOW storage:spans:read
```

11. Best Practices

Do's

-  **Start with default policies** - customize only when needed

- ☒ **Use least privilege** - grant minimum required permissions
- ☒ **Document policies and boundaries** - clear names and descriptions
- ☒ **Test before production** - verify with test users
- ☒ **Use boundaries for scope** - keep policies reusable
- ☒ **Include all three domains** in boundaries for complete MZ replacement
- ☒ **Use SAML groups** tied to AD for team management

Don'ts

- ☒ **Don't create policy per user** - use groups
- ☒ **Don't duplicate default policies** - extend instead
- ☒ **Don't use overly broad wildcards** - be specific
- ☒ **Don't embed conditions in policies** when boundaries are better
- ☒ **Don't forget the settings domain** in boundaries

Naming Conventions

```
// Good names
"Frontend Team - Standard Access"    (policy)
"Production Environment"             (boundary)
"L0B5-Scope"                        (boundary)

// Bad names
"Policy 1"
"Test"
"John's boundary"
```

Audit Checklist

- ☐ Each boundary has a clear purpose
 - ☐ Boundary names follow naming convention
 - ☐ All three domains included (environment, storage, settings)
 - ☐ SAML groups aligned with AD structure
 - ☐ Dependencies tracked (which policies use which boundaries)
 - ☐ Regular review scheduled
-

12. Troubleshooting

Common Issues

Issue	Likely Cause	Solution
User can't access anything	No policy bound	Bind policy to user's group
User sees too much	Boundary too broad or missing	Tighten boundary conditions
User can't see expected data	Missing data policy	Add Data Viewer/Editor
Permissions inconsistent	Multiple conflicting policies	Review all bound policies
Works in classic, not Grail	Only environment domain used	Add storage domain to boundary
Settings not restricted	Missing settings domain	Add settings:dt.security_context

Debugging Steps

1. **Verify group membership:** Is user in correct group?
2. **Check policy bindings:** What policies are bound to group?
3. **Review boundaries:** Are all three domains included?
4. **Test with admin:** Does admin see the data?
5. **Check security context:** Is it set on the entities?

Testing Boundaries

1. **Create test user** in target group
2. **Log in as test user**
3. **Verify access:**
 - Can access expected data?
 - Cannot access restricted data?

4. Test edge cases:

- Entities at boundary edges
 - New entities without security context
-

Summary

In this notebook, you learned:

1. **Policies define WHAT**, boundaries define **WHERE**
2. **Default policies** cover most use cases - customize only when needed
3. **Shared syntax** for conditions in both policies and boundaries
4. **Three domains** needed for complete MZ replacement (environment, storage, settings)
5. **SAML + Policy + Boundary** structure for AD integration
6. **Migration patterns** for team-based and environment-based MZs

Next Steps

Continue to **MZ2POL-05: Segments Implementation** to:

- Create DQL-based data filters
- Replace MZ data filtering with Segments
- Configure cross-app filtering

Additional Resources

- [Working with Policies](#)
 - [IAM Policy Reference](#)
 - [Policy Boundaries Documentation](#)
 - [Grant Access with Security Context](#)
-

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.