

OPIPE-06: Cross-Scope Design Patterns

Series: OPIPE — OpenPipeline Beyond Logs | **Notebook:** 6 of 7 | **Created:** March 2026 | **Last Updated:** 05/06/2026

Correlating Logs, Spans, Metrics, and Events Across Scopes

OpenPipeline's six scopes process data independently — but observability value comes from **correlating** across them. A slow API response (span) may be caused by a database timeout (log) that triggers a detected problem (event) and degrades an SLO metric. This notebook covers design patterns that connect the dots across scopes, cascade processing for derived data, and a production readiness checklist.

For individual scope deep-dives, see **OPIPE-02** (spans), **OPIPE-03** (metrics), **OPIPE-05** (events). For log processing, see **OPLOGS-01**. For trace analysis, see **SPANS-04: Service Dependencies & Flow Analysis**.

Table of Contents

1. [Shared Dimensions: The Correlation Key](#)
2. [Pattern: Same Metric from Multiple Scopes](#)
3. [Pattern: Cascade Processing](#)
4. [Pattern: Unified Bucket Families](#)
5. [When NOT to Use OpenPipeline](#)
6. [Production Readiness Checklist](#)
7. [Summary](#)
8. [References](#)

Prerequisites

Requirement	Details
Dynatrace Environment	SaaS with Grail
Permissions	<code>storage:logs:read</code> , <code>storage:spans:read</code> , <code>storage:metrics:read</code> , <code>storage:events:read</code> , <code>storage:bizevents:read</code>
Recommended	Complete OPIPE-01 through OPIPE-05

1. Shared Dimensions: The Correlation Key

Cross-scope correlation works when different data types share **common dimension values**. If your logs and spans both carry `service.name` and `k8s.namespace.name` , you can join them in DQL.

Correlation Fields Across Scopes

Field	Logs	Spans	Metrics	Events	Bizevents
<code>dt.entity.service</code>	Yes	Yes	Yes	Yes	—
<code>dt.entity.host</code>	Yes	—	Yes	Yes	—
<code>service.name</code>	Sometimes	Yes	—	—	—
<code>k8s.namespace.name</code>	Yes	Yes	Yes	—	—
<code>k8s.deployment.name</code>	Yes	Yes	Sometimes	—	—
<code>trace.id</code>	Sometimes	Yes	—	—	—
<code>dt.security_context</code>	If enriched	If enriched	If enriched	If enriched	If enriched

The Enrichment Strategy

If a critical dimension is missing from one scope, **add it via OpenPipeline enrichment** so that cross-scope queries work:

Scope	Missing Field	How to Add
Logs	<code>service.name</code>	Enrich from <code>dt.entity.service</code> via entity lookup
Spans	<code>k8s.deployment.name</code>	Enrich from <code>k8s.pod.name</code> by stripping replica suffix
Metrics	<code>k8s.namespace.name</code>	Enrich from entity tags or extension configuration

```
// Cross-scope: Correlate error logs with error spans by service entity
fetch logs, from:-1h
| filter loglevel == "ERROR" and isNotNull(dt.entity.service)
| summarize error_logs = count(), by:{dt.entity.service}
| lookup [
  fetch spans, from:-1h
  | filter http.response.status_code >= 500 and
isNotNull(dt.entity.service)
  | summarize error_spans = count(), by:{dt.entity.service}
], sourceField:dt.entity.service, lookupField:dt.entity.service, fields:
{error_spans}
| sort error_logs desc
| limit 10
```

2. Pattern: Same Metric from Multiple Scopes

A powerful validation pattern: extract the **same metric** from two different data sources and compare them. If they diverge, one pipeline has a configuration problem.

Example: Error Count from Logs vs. Spans

Source	Metric Key	Extraction Rule
Logs	<code>log.error_count</code>	Count of <code>loglevel == "ERROR"</code> , by <code>dt.entity.service</code>
Spans	<code>span.error_count</code>	Count of <code>http.response.status_code >= 500</code> , by <code>dt.entity.service</code>

These metrics should track each other. If logs show 10x more errors than spans, it could mean:

- Application-level errors are logged but not reflected in HTTP status codes
- Span sampling is dropping error spans (see **OPIPE-03** on sampling-aware metrics)
- The log pipeline is catching errors from non-HTTP sources (background jobs, message consumers)

The comparison itself is diagnostic — the divergence tells you something about your data.

```
// Compare: Error counts from logs vs. spans by service
fetch logs, from:-1h
| filter loglevel == "ERROR" and isNotNull(dt.entity.service)
| summarize log_errors = count(), by:{dt.entity.service}
| lookup [
    fetch spans, from:-1h
    | filter http.response.status_code >= 500 and
isNotNull(dt.entity.service)
    | summarize span_errors = count(), by:{dt.entity.service}
], sourceField:dt.entity.service, lookupField:dt.entity.service, fields:
{span_errors}
| fieldsAdd ratio = if(isNotNull(span_errors) and span_errors > 0,
    then: round(toDouble(log_errors) / toDouble(span_errors), decimals: 1),
    else: -1.0)
| sort log_errors desc
| limit 10
```

3. Pattern: Cascade Processing

Cascade processing creates a chain of derived data across scopes:

```
Span → generates Event → Event triggers Metric → Metric drives SLO
```

Example: Slow Transaction Cascade

Stage	Scope	Configuration	Output
1. Detect	Spans	Filter: <code>span.kind == "server"</code> AND <code>duration > 5s</code>	Matching spans

Stage	Scope	Configuration	Output
2. Generate	Spans → Events	Event generation rule: create event with <code>event.type = "slow_transaction"</code>	Event record per slow span
3. Extract	Events → Metrics	Metric extraction: <code>slow_transaction.count</code> by <code>service.name</code>	Metric time series
4. Alert	Metrics → SLO	SLO: <code>slow_transaction.count < 10</code> per 5m per service	SLO breach triggers alert

Design Considerations

- **Each stage adds latency** — The full cascade may take seconds from span ingestion to SLO evaluation
- **Failures propagate** — If the span pipeline drops the slow span, the event is never generated, and the metric never counts it
- **Test end-to-end** — After configuring a cascade, verify data flows through every stage

```
// Identify slow transactions that could trigger a cascade
fetch spans, from:-1h
| filter span.kind == "server" and duration > 5000000000
| summarize slow_count = count(), avg_duration_ms = avg(duration / 1ms),
  by:{service.name}
| sort slow_count desc
| limit 10
```

4. Pattern: Unified Bucket Families

When related data across scopes should share the same lifecycle and access controls, use a **bucket family** naming convention:

Bucket Name	Scope	Contents	Retention
<code>checkout_logs</code>	Logs	Checkout service application logs	35 days
<code>checkout_spans</code>	Spans	Checkout service trace data	14 days
<code>checkout_events</code>	Events	Checkout-related detected events	35 days

Bucket Name	Scope	Contents	Retention
checkout_bizevents	Bizevents	Purchase and cart business events	90 days

Benefits

- **Consistent naming** — `checkout_*` makes it obvious which buckets belong together
- **Unified access control** — IAM policies can use `dt.system.bucket` patterns with wildcards
- **Coordinated retention** — Related data expires in a logical sequence (spans first, then logs, business events last)

Querying Across a Bucket Family

```
// Query all buckets in a family to understand data distribution
fetch logs, from:-1h
| filter matchesValue(dt.system.bucket, "checkout_*")
| summarize log_count = count()
| append [
    fetch spans, from:-1h
    | filter matchesValue(dt.system.bucket, "checkout_*")
    | summarize span_count = count()
]
```

5. When NOT to Use OpenPipeline

OpenPipeline is powerful but not always the right tool. Prefer DQL query-time processing when:

Scenario	Why Not OpenPipeline	Better Approach
Exploratory analysis	You do not know what patterns you are looking for yet	DQL ad-hoc queries
Frequently changing logic	Pipeline changes require deployment; DQL changes are instant	DQL with dashboard variables
Complex joins	OpenPipeline cannot join across scopes at ingestion	DQL <code>lookup</code> and <code>join</code> at query time

Scenario	Why Not OpenPipeline	Better Approach
One-time investigations	Configuring a pipeline for a single investigation is overhead	DQL <code>parse</code> with DPL patterns
Historical data	OpenPipeline only processes new data — it cannot reprocess historical records	DQL for retroactive analysis
Correlation dashboards	Cross-scope correlation requires data from multiple tables	DQL with <code>append</code> and <code>lookup</code>

The Rule of Thumb

OpenPipeline for permanent, irreversible actions (drop, mask, route, extract). **DQL for everything else.**

If you are unsure whether a transformation belongs in OpenPipeline or DQL, ask: *"Will I regret not having the raw data?"* If yes, keep it raw and process at query time.

6. Production Readiness Checklist

Before deploying a new or modified OpenPipeline configuration to production:

Pipeline Design

- [] Each pipeline has a single, clear purpose (one source type per pipeline)
- [] Routing rules are ordered from most specific to least specific
- [] Default pipeline catches only unmatched data with short retention
- [] No overlapping routing conditions between pipelines

Processing Rules

- [] Masking rules applied BEFORE any other processing
- [] Drop rules filter noise before extraction (cost efficiency)
- [] Parsing rules tested against representative samples
- [] Enrichment fields have defined values for all conditions (no nulls)

Metric Extraction

- [] Cardinality estimated for each extracted metric (see **OPIPE-04**)
- [] No unique-per-record dimensions (`trace.id` , `span.id` , `content`)
- [] Sampling-aware extraction for span-derived count metrics (see **OPIPE-03**)
- [] Metric keys follow naming convention: `scope.metric_name` (e.g., `span.request_count`)

Security & Governance

- [] `dt.security_context` set on all scopes that need access control
- [] Security events routed to dedicated long-retention buckets
- [] Compliance-relevant data never dropped or sampled
- [] PII masking verified on all scopes that handle user data

Testing

- [] End-to-end data flow verified: ingestion → pipeline → bucket → query
- [] Cascade chains tested: span → event → metric → SLO (if applicable)
- [] Volume monitoring in place to detect unexpected spikes or drops
- [] Rollback plan documented (revert to previous pipeline configuration)

Summary

In this notebook you learned:

- **Shared dimensions** — Cross-scope correlation requires common fields like `dt.entity.service` and `k8s.namespace.name`
- **Same metric from multiple scopes** — Extract matching metrics from logs and spans to validate pipeline health
- **Cascade processing** — Chain span → event → metric → SLO for automated detection and alerting
- **Unified bucket families** — Name related buckets with a common prefix for consistent lifecycle management

- **When NOT to use OpenPipeline** — Keep data raw for exploratory analysis, changing logic, and one-time investigations
 - **Production readiness** — Checklist covering pipeline design, processing rules, metric extraction, security, and testing
-

What's Next?

You have completed the OPIPE series. From here:

- **OPLOGS** — Deep-dive into log processing: **OPLOGS-01: OpenPipeline Fundamentals**
 - **SPANS** — Master trace analysis: **SPANS-01: Fundamentals**
 - **OPMIG** — Migrating from classic logs: **OPMIG-01: Why Migrate**
 - **ORGNZ** — Data organization and governance: **ORGNZ-01: Introduction**
-

References

- [OpenPipeline Documentation](#)
 - [Grail Data Lakehouse](#)
 - [DQL Cross-Data Queries](#)
 - [SLO Configuration](#)
-

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.