



OpenPipeline Fundamentals

Series: OPLOGS | **Notebook:** 1 of 8 | **Created:** December 2025 | **Last Updated:** 03/25/2026

Understanding the Unified Data Ingestion Framework

This notebook introduces OpenPipeline, Dynatrace's unified data processing framework for logs, traces, metrics, and events.

 **New Addition (March 2026)**

Companion Series

- **OPMIG** — Migrating from classic logs to OpenPipeline? Start with **OPMIG-01: Why Migrate**
- **OPIPE** — Ready to process spans, metrics, and events? Continue to **OPIPE-01: OpenPipeline as a Multi-Scope Platform**
- **SPANS** — Need to query and analyze distributed traces? See **SPANS-01: Fundamentals**

Table of Contents

1. [What is OpenPipeline?](#)
2. [OpenPipeline Architecture](#)
3. [Exploring Your OpenPipeline Data](#)
4. [Key OpenPipeline Fields](#)
5. [Data Sources Explained](#)
6. [Environment Summary](#)
7.  [Summary](#)
8.  [Next Steps](#)
9.  [References](#)

Prerequisites

-  Access to a Dynatrace environment with log data
-  DQL query permissions (viewer role minimum)
-  Basic understanding of log management concepts

1. What is OpenPipeline?


OpenPipeline is Dynatrace's unified data ingestion and processing framework that replaces classic log ingestion. It provides:

- **Unified Processing:** Single framework for logs, metrics, traces, and business events
- **Real-time Transformation:** Parse, enrich, mask, and route data at ingestion
- **Grail Storage:** Direct integration with Dynatrace's data lakehouse
- **Flexible Routing:** Send data to different buckets with custom retention
- **Cost Control:** Drop unnecessary data before storage

OpenPipeline vs Classic Log Ingestion

Feature	Classic Logs	OpenPipeline v2.0
Data Processing	Post-ingestion	At ingestion time
Storage	Log Storage v1	Grail Data Lakehouse
Query Language	Limited	Full DQL Support
Retention	Global	Per-bucket configurable
Data Masking	Limited	Full regex support
Parsing	Basic	DPL (Dynatrace Pattern Language)
Custom Routing	No	Yes, by content/source

2. OpenPipeline Architecture



3. Exploring Your OpenPipeline Data

Let's start by discovering what data sources and pipelines are active in your environment.

```
// Discover data sources feeding OpenPipeline
fetch logs, from: now() - 1h
| summarize {log_count = count()}, by: {dt.openpipeline.source}
| sort log_count desc
```

```
// See which pipelines are processing your logs
fetch logs, from: now() - 1h
| summarize {log_count = count()}, by: {dt.openpipeline.pipelines}
| sort log_count desc
```

```
// Check storage bucket distribution
fetch logs, from: now() - 1h
| summarize {log_count = count()}, by: {dt.system.bucket}
| sort log_count desc
```

4. Key OpenPipeline Fields

OpenPipeline adds metadata fields to every log record:

Pipeline Metadata

Field	Description	Example
<code>dt.openpipeline.source</code>	How the log was ingested	<code>oneagent , /api/v2/logs/ingest , /api/v2/otlp/v1/logs</code>
<code>dt.openpipeline.pipelines</code>	Pipeline(s) that processed the log	<code>["logs:pipeline_Default_Pipeline_2798"]</code>
<code>dt.system.bucket</code>	Storage bucket name	<code>default_logs , custom_logs</code>

Core Log Fields

Field	Description
<code>timestamp</code>	When the log was generated
<code>content</code>	The log message body
<code>loglevel</code>	Log severity (ERROR, WARN, INFO, DEBUG, NONE)
<code>status</code>	Status string (alternative to loglevel)
<code>log.source</code>	Source identifier (e.g., "Container Output")
<code>log.iostream</code>	Stream type (stdout, stderr)

```
// View a sample log record with all key fields
fetch logs, from: now() - 1h
| fields timestamp, content, loglevel, status, log.source, log.iostream,
      dt.openpipeline.source, dt.openpipeline.pipelines, dt.system.bucket
| limit 5
```

```
// Analyze log levels in your environment
fetch logs, from: now() - 1h
| summarize {count = count()}, by: {loglevel}
| sort count desc
```

5. Data Sources Explained

OneAgent (`oneagent`)

Logs collected automatically by Dynatrace OneAgent from:

- Container stdout/stderr
- Process log files
- System logs

Log Ingest API (`/api/v2/logs/ingest`)

Logs sent directly via the Dynatrace API:

- Custom application logs
- Third-party integrations
- Cloud provider logs (AWS, Azure, GCP)

OTLP (`/api/v2/otlp/v1/logs`)

OpenTelemetry Protocol logs:

- OpenTelemetry Collector
- Fluent Bit with OTLP output
- Custom OTLP exporters

```
// Compare volume by data source
fetch logs, from: now() - 24h
| summarize {
    log_count = count(),
    unique_hosts = countDistinct(dt.entity.host)
}, by: {dt.openpipeline.source}
| sort log_count desc
```

```
// Logs per hour by source (trend analysis)
fetch logs, from: now() - 24h
| makeTimeseries {log_count = count()}, by: {dt.openpipeline.source}, interval:
1h
```

6. Pipeline Stages Overview

OpenPipeline processes data through ordered stages:

Stage 1: Routing

- Matches incoming data to the appropriate pipeline
- Based on source, content, or metadata

Stage 2: Masking (Security)

- Redacts sensitive data BEFORE processing
- Protects PII, credentials, secrets
- Applied early for security compliance

Stage 3: Filtering

- Drops unwanted records
- Reduces storage costs
- Removes noise (debug logs, health checks)

Stage 4: Processing

- Parses structured data from content
- Adds enrichment fields
- Transforms and normalizes data

Stage 5: Extraction

- Creates metrics from log data
- Generates events and business events

Stage 6: Storage

- Routes to appropriate Grail bucket
- Applies retention policies

7. Environment Summary

Let's get a complete picture of your OpenPipeline environment:

```
// Complete environment summary
fetch logs, from: now() - 1h
| summarize {
    total_logs = count(),
    unique_sources = countDistinct(dt.openpipeline.source),
    unique_buckets = countDistinct(dt.system.bucket),
    unique_hosts = countDistinct(dt.entity.host),
    error_count = countIf(loglevel == "ERROR" OR status == "ERROR"),
    warn_count = countIf(loglevel == "WARN" OR status == "WARN")
}
```

```
// Top log sources by entity
fetch logs, from: now() - 1h
| filter isNotNull(dt.entity.host)
| summarize {log_count = count()}, by: {host.name, log.source}
| sort log_count desc
| limit 15
```



Summary

In this notebook, you learned:

- ✓ **What OpenPipeline is** - Dynatrace's unified data processing framework
- ✓ **Architecture** - Data flow from sources through processing to Grail
- ✓ **Key fields** - `dt.openpipeline.source`, `dt.openpipeline.pipelines`, `dt.system.bucket`
- ✓ **Data sources** - OneAgent, Log Ingest API, OTLP
- ✓ **Pipeline stages** - Routing, Masking, Filtering, Processing, Extraction, Storage



Next Steps

Continue to **OPLOGS-02: Migration Guide** to learn how to migrate from classic log ingestion to OpenPipeline v2.0.

References

- [OpenPipeline Documentation](#)
 - [Grail Data Lakehouse](#)
 - [DQL Reference](#)
-

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.