

OPLOGS-02: Migration to OpenPipeline

Series: OPLOGS — OpenPipeline Logs | **Notebook:** 2 of 8 | **Created:** December 2025 | **Last Updated:** 01/28/2026

Planning and Executing Your Log Migration

This notebook guides you through assessing your current log environment and planning migration to OpenPipeline v2.0.

Table of Contents

1. [Migration Assessment](#)
 2. [Migration Paths](#)
 3. [Planning Your OpenPipeline Configuration](#)
 4. [Migration Validation](#)
 5. [Common Migration Patterns](#)
 6. [Migration Checklist](#)
 7.  [Summary](#)
 8.  [Next Steps](#)
 9.  [References](#)
-

Prerequisites

-  Access to a Dynatrace environment with log data

-  Completed OPLOGS-01 fundamentals
-  Knowledge of your current log sources and volumes

1. Migration Assessment

Before migrating, you need to understand your current log landscape.

Key Questions to Answer:

1. **What log sources do I have?** (OneAgent, API, OTLP)
2. **What is my log volume?** (records/hour, GB/day)
3. **Which logs need parsing?** (unstructured content)
4. **What sensitive data exists?** (PII, credentials)
5. **What can be dropped?** (debug, health checks)

```
// Assessment Query 1: Log sources inventory
fetch logs, from: now() - 24h
| summarize {
    log_count = count(),
    unique_hosts = countDistinct(dt.entity.host)
}, by: {dt.openpipeline.source, log.source}
| sort log_count desc
| limit 20
```

```
// Assessment Query 2: Log volume by hour (capacity planning)
fetch logs, from: now() - 24h
| makeTimeseries {log_count = count()}, interval: 1h
```

```
// Assessment Query 3: Log level distribution
fetch logs, from: now() - 24h
| summarize {
    total = count(),
    errors = countIf(loglevel == "ERROR" OR status == "ERROR"),
    warnings = countIf(loglevel == "WARN" OR status == "WARN"),
    info = countIf(loglevel == "INFO" OR status == "INFO"),
    debug = countIf(loglevel == "DEBUG" OR status == "DEBUG"),
    none = countIf(loglevel == "NONE" OR status == "NONE")
}
| fieldsAdd debug_pct = round((toDouble(debug) / toDouble(total)) * 100,
decimals: 1)
| fieldsAdd none_pct = round((toDouble(none) / toDouble(total)) * 100,
decimals: 1)
```

```
// Assessment Query 4: Identify logs that need parsing (loglevel = NONE)
fetch logs, from: now() - 1h
| filter loglevel == "NONE" OR status == "NONE"
| fieldsAdd content_preview = substring(content, from: 0, to: 80)
| summarize {count = count()}, by: {content_preview, dt.openpipeline.source}
| sort count desc
| limit 15
```

```
// Assessment Query 5: Potential drop candidates (health checks, heartbeats)
fetch logs, from: now() - 24h
| summarize {
    total = count(),
    health_checks = countIf(matchesPhrase(content, "health") OR
matchesPhrase(content, "healthcheck")),
    heartbeats = countIf(matchesPhrase(content, "heartbeat")),
    debug_logs = countIf(loglevel == "DEBUG" OR status == "DEBUG")
}
| fieldsAdd droppable = health_checks + heartbeats + debug_logs
| fieldsAdd savings_pct = round((toDouble(droppable) / toDouble(total)) *
100, decimals: 1)
```

2. Migration Paths

Path A: Automatic Migration (Recommended)

If you're using OneAgent for log collection, most data automatically flows through OpenPipeline.

Steps:

1. Verify `dt.openpipeline.source = oneagent` in your logs
2. Configure custom pipelines for specific processing needs
3. Update queries to use OpenPipeline fields

Path B: API Migration

If using the Log Ingest API:

Steps:

1. Continue using `/api/v2/logs/ingest` endpoint
2. Logs automatically route through OpenPipeline
3. Configure processing rules as needed

Path C: OTLP Migration

For OpenTelemetry-based logging:

Steps:

1. Point OTLP exporters to `/api/v2/otlp/v1/logs`
2. Logs flow through OpenPipeline automatically
3. Configure pipelines for OTLP-specific processing

```
// Check which data sources are already using OpenPipeline
fetch logs, from: now() - 1h
| filter isNotNull(dt.openpipeline.pipelines)
| summarize {count = count()}, by: {dt.openpipeline.source}
| sort count desc
```

3. Planning Your OpenPipeline Configuration

Pipeline Strategy

Use Case	Pipeline Configuration
Default Processing	Use built-in Default Pipeline
Custom Parsing	Create pipeline with DPL parse rules
PII Masking	Add masking stage before processing

Use Case	Pipeline Configuration
Cost Reduction	Add drop rules for noise
Custom Routing	Route to specific buckets

Bucket Strategy

Log Type	Bucket	Retention
Production Errors	prod_errors	90 days
Application Logs	default_logs	35 days
Debug/Verbose	Drop or 7 days	Minimal
Audit/Compliance	audit_logs	365+ days

```
// Analyze current bucket usage
fetch logs, from: now() - 24h
| summarize {
    log_count = count(),
    error_count = countIf(loglevel == "ERROR" OR status == "ERROR")
}, by: {dt.system.bucket}
| fieldsAdd error_rate = round((toDouble(error_count) / toDouble(log_count))
* 100, decimals: 2)
| sort log_count desc
```

4. Migration Validation

After configuring OpenPipeline, validate that:


- 1.  All expected log sources are present
- 2.  Log counts are consistent
- 3.  Parsing rules extract expected fields
- 4.  Masking rules are applied correctly
- 5.  Routing sends logs to correct buckets

```
// Validation Query 1: Confirm all sources are flowing
fetch logs, from: now() - 1h
| summarize {
    total = count(),
    has_pipeline = countIf(isNotNull(dt.openpipeline.pipelines)),
    has_source = countIf(isNotNull(dt.openpipeline.source)),
    has_bucket = countIf(isNotNull(dt.system.bucket))
}
| fieldsAdd pipeline_coverage = round((toDouble(has_pipeline) /
toDouble(total)) * 100, decimals: 1)
```

```
// Validation Query 2: Check log volume consistency (compare hours)
fetch logs, from: now() - 4h
| makeTimeseries {log_count = count()}, by: {dt.openpipeline.source},
interval: 1h
```

```
// Validation Query 3: Verify entity context is preserved
fetch logs, from: now() - 1h
| summarize {
    total = count(),
    with_host = countIf(isNotNull(dt.entity.host)),
    with_process = countIf(isNotNull(dt.entity.process_group)),
    with_k8s_cluster = countIf(isNotNull(dt.entity.kubernetes_cluster)),
    with_k8s_namespace = countIf(isNotNull(k8s.namespace.name))
}
```

5. Common Migration Patterns



Pattern 1: Parse Log Levels from Content

When logs have `loglevel = NONE`, configure a DQL processor to extract the level from content patterns like `[INFO]` or `[ERROR]`.

Processor Configuration:

- **Type:** DQL
- **Matcher:** `loglevel == "NONE"`
- **Statement:** `parse content, '[' LD:parsed_level ']' | fieldsAdd loglevel = parsed_level`

Pattern 2: Drop Debug Logs

Reduce storage costs by dropping DEBUG-level logs before storage.

Processor Configuration:

- **Type:** Drop
- **Matcher:** `loglevel == "DEBUG"`

Pattern 3: Route Errors to Dedicated Bucket

Send error logs to a dedicated bucket with longer retention for compliance.

Bucket Routing Configuration:

- **Matcher:** `loglevel == "ERROR"`
- **Bucket:** `error_logs`
- **Retention:** 90 days

```
// Simulate parsing log levels from content
fetch logs, from: now() - 1h
| filter loglevel == "NONE" OR status == "NONE"
| parse content, '[' LD:parsed_level ']'
| filter isNotNull(parsed_level)
| fields timestamp, content, parsed_level
| limit 10
```

```
// Identify logs that would be dropped (DEBUG + health checks)
fetch logs, from: now() - 24h
| filter loglevel == "DEBUG"
    OR matchesPhrase(content, "healthcheck")
    OR matchesPhrase(content, "health check")
| summarize {would_drop = count()}, by: {dt.openpipeline.source}
| sort would_drop desc
```

6. Migration Checklist

Pre-Migration

- ☐ Document current log sources and volumes
- ☐ Identify logs requiring custom parsing
- ☐ List sensitive data patterns to mask
- ☐ Define drop rules for noise reduction
- ☐ Plan bucket strategy and retention

During Migration

- ☐ Configure OpenPipeline settings
- ☐ Create custom pipelines as needed
- ☐ Set up masking rules
- ☐ Configure bucket routing
- ☐ Test with sample data

Post-Migration

- ☐ Validate all log sources flowing
 - ☐ Verify volume consistency
 - ☐ Confirm parsing rules working
 - ☐ Test masking is applied
 - ☐ Update dashboards and alerts
 - ☐ Update saved queries to use new fields
-



Summary

In this notebook, you learned:

- ✓ **Assessment queries** to understand your current log environment
 - ✓ **Migration paths** for OneAgent, API, and OTLP sources
 - ✓ **Pipeline planning** strategies for processing and routing
 - ✓ **Validation queries** to confirm successful migration
 - ✓ **Common patterns** for parsing, dropping, and routing
-



Next Steps

Continue to **OPLOGS-03: Querying & Parsing** to learn DQL syntax and DPL parsing patterns.



References

- [OpenPipeline Documentation](#)
 - [Log Ingest API](#)
 - [OTLP Log Ingestion](#)
-

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.