



# OpenPipeline Processing

---

**Series:** OPLOGS | **Notebook:** 3 of 8 | **Created:** December 2025 | **Last Updated:** 01/28/2026

## Configuring Pipeline Stages for Log Transformation

---

This notebook covers OpenPipeline processing stages: parsing, enrichment, metric extraction, event generation, bucket routing, and filtering.

---

## Table of Contents

---

1. [Parsing & Field Extraction](#)
  2. [Metric Extraction from Logs](#)
  3. [Attribute Creation & Enrichment](#)
  4. [Event Generation from Logs](#)
  5. [Bucket Routing](#)
  6. [Filtering & Sampling](#)
  7. [Complete Pipeline Example](#)
  8.  [Summary](#)
  9.  [Next Steps](#)
  10.  [References](#)
- 

## Prerequisites

---

-  Access to a Dynatrace environment with log data
-  OpenPipeline configuration permissions

-  Completed OPLOGS-01 and OPLOGS-02

## 2. Parsing & Field Extraction

Parsing extracts structured fields from unstructured log content **at ingestion time**.

### DPL (Dynatrace Pattern Language) Matchers

#### DPL Pattern Matchers Quick Reference

Dynatrace Pattern Language for Log Parsing

##### Text Matchers

|               |                       |                   |
|---------------|-----------------------|-------------------|
| <b>LD</b>     | Line data (greedy)    | "error in module" |
| <b>WORD</b>   | Single word           | "ERROR"           |
| <b>DATA</b>   | Any text (non-greedy) | "any content"     |
| <b>NSPACE</b> | Non-whitespace        | "user@host.com"   |
| <b>SPACE</b>  | Whitespace characters | " " (tabs/spaces) |

##### Number Matchers

|                |                |                     |
|----------------|----------------|---------------------|
| <b>INT</b>     | Integer        | -1, 42, 1000        |
| <b>LONG</b>    | Long integer   | 9223372036854775807 |
| <b>DOUBLE</b>  | Decimal number | 3.14, -0.5          |
| <b>BOOLEAN</b> | true/false     | true, false         |
| <b>HEX</b>     | Hexadecimal    | 0x1F, DEADBEEF      |

##### Network / Identity

|                |                   |               |
|----------------|-------------------|---------------|
| <b>IPADDR</b>  | IPv4/IPv6 address | 10.0.0.1, ::1 |
| <b>IPV4</b>    | IPv4 only         | 192.168.1.100 |
| <b>IPV6</b>    | IPv6 only         | 2001:db8::1   |
| <b>HTTPURI</b> | HTTP URL path     | /api/v1/users |
| <b>UUID</b>    | UUID format       | alb2c3d4-...  |

##### Time / Date

|                  |                   |                |
|------------------|-------------------|----------------|
| <b>TIMESTAMP</b> | ISO timestamp     | 2024-01-15T... |
| <b>DATE</b>      | Date (YYYY-MM-DD) | 2024-01-15     |
| <b>TIME</b>      | Time (HH:MM:SS)   | 14:30:00       |
| <b>EPOCH</b>     | Unix timestamp    | 1705329600     |
| <b>DURATION</b>  | Time duration     | 2h30m, 1d      |

Syntax: MATCHER:fieldname (e.g., INT:status\_code) or MATCHER (no capture)

| Matcher   | Description              | Example Match    |
|-----------|--------------------------|------------------|
| LD        | Line data (to delimiter) | Any text         |
| INT       | Integer                  | 42 , -17         |
| DOUBLE    | Decimal number           | 3.14 , -0.5      |
| IPADDR    | IP address               | 192.168.1.1      |
| WORD      | Word characters          | hello123         |
| TIMESTAMP | Date/time                | Various formats  |
| JSON      | JSON object/array        | {"key": "value"} |

## OpenPipeline Parse Processor

```
# Parse HTTP access logs
processors:
  - name: parse-http-logs
    type: dql
    source: content
    dql: |
      parse content, "IPADDR:client_ip SPACE LD SPACE '['
TIMESTAMP:request_time ']'
                        SPACE '\"' LD:method SPACE LD:path SPACE LD '\"'
                        SPACE INT:status SPACE INT:bytes"
```

```
// Discover log patterns for parsing design
fetch logs, from: now() - 1h
| fieldsAdd content_preview = substring(content, from: 0, to: 100)
| summarize {count = count()}, by: {content_preview}
| sort count desc
| limit 20
```

```
// Test parse pattern before configuring in OpenPipeline
fetch logs, from: now() - 1h
| filter contains(content, "GET") OR contains(content, "POST")
| parse content, "LD:method SPACE '/' LD:path SPACE INT:status_code"
| filter isNotNull(status_code)
| summarize {count = count()}, by: {method, status_code}
| sort count desc
| limit 15
```

```
// Verify fields already extracted by OpenPipeline
fetch logs, from: now() - 1h
| limit 100
| summarize {
  has_loglevel = countIf(isNotNull(loglevel)),
  has_status = countIf(isNotNull(status)),
  has_trace_id = countIf(isNotNull(trace_id)),
  has_span_id = countIf(isNotNull(span_id))
}
```

### 3. Metric Extraction from Logs

**Extract metrics with dimensions** from log data at ingestion time. This is powerful for:

- Creating SLIs from log patterns
- Building dashboards without log queries
- Enabling metric-based alerting

#### OpenPipeline Metric Extraction

```
# Extract request duration metric from logs
processors:
  - name: extract-request-metric
    type: metric
    enabled: true
    condition: contains(content, "duration=")
    metricKey: log.request.duration
    dimensions:
      - service: k8s.namespace.name
      - method: extracted_method
      - status: extracted_status
    value: extracted_duration_ms
```

#### Common Metric Extraction Patterns

| Log Pattern       | Metric                | Dimensions              |
|-------------------|-----------------------|-------------------------|
| Request completed | log.request.count     | service, method, status |
| Response time     | log.response.duration | endpoint, status_code   |
| Error occurred    | log.error.count       | error_type, service     |
| Queue depth       | log.queue.size        | queue_name              |

```
// Identify logs with numeric values for metric extraction
fetch logs, from: now() - 1h
| filter contains(content, "duration")
      OR contains(content, "latency")
      OR contains(content, "time=")
      OR contains(content, "ms")
| fieldsAdd content_preview = substring(content, from: 0, to: 120)
| summarize {count = count()}, by: {content_preview}
| sort count desc
| limit 15
```

```
// Simulate metric extraction: count by dimensions
fetch logs, from: now() - 1h
| filter isNotNull(k8s.namespace.name)
| summarize {
    request_count = count(),
    error_count = countIf(loglevel == "ERROR")
}, by: {k8s.namespace.name, k8s.workload.name}
| fieldsAdd error_rate = round((error_count * 100.0) / request_count,
decimals: 2)
| sort error_count desc
| limit 15
```

```
// Preview: What dimensions would be valuable?
fetch logs, from: now() - 1h
| summarize {
    unique_namespaces = countDistinct(k8s.namespace.name),
    unique_workloads = countDistinct(k8s.workload.name),
    unique_pods = countDistinct(k8s.pod.name),
    unique_hosts = countDistinct(dt.entity.host),
    unique_sources = countDistinct(dt.openpipeline.source)
}
```

## 4. Attribute Creation & Enrichment

---

Add **computed attributes** to logs for enhanced analysis and filtering.

# OpenPipeline Field Processor

```
# Add computed attributes
processors:
  - name: enrich-environment
    type: fieldsAdd
    fields:
      - name: environment
        value: |
          if(contains(k8s.namespace.name, "prod"), "production",
            else: if(contains(k8s.namespace.name, "staging"), "staging",
              else: "development"))

      - name: severity_score
        value: |
          if(loglevel == "ERROR", 3,
            else: if(loglevel == "WARN", 2,
              else: 1))

      - name: team_owner
        value: |
          if(contains(k8s.namespace.name, "payment"), "platform-team",
            else: if(contains(k8s.namespace.name, "frontend"), "web-team",
              else: "unknown"))
```

## Common Attribute Patterns

| Attribute      | Source           | Use Case           |
|----------------|------------------|--------------------|
| environment    | namespace name   | Filter prod vs dev |
| team_owner     | namespace/labels | Route alerts       |
| severity_score | loglevel         | Prioritization     |
| log_category   | content patterns | Classification     |

```
// Preview attribute creation logic
fetch logs, from: now() - 1h
| filter isNotNull(k8s.namespace.name)
| fieldsAdd environment = if(contains(k8s.namespace.name, "prod"),
"production",
                        else: if(contains(k8s.namespace.name, "staging"),
"staging",
                        else: "development"))
| fieldsAdd severity_score = if(loglevel == "ERROR", 3,
                        else: if(loglevel == "WARN", 2,
                        else: 1))
| summarize {count = count()}, by: {environment, severity_score, loglevel}
| sort environment asc, severity_score desc
```

```
// Categorize logs by content patterns
fetch logs, from: now() - 1h
| fieldsAdd log_category = if(contains(content, "Exception") OR
contains(content, "Error"), "exception",
                        else: if(contains(content, "request") OR
contains(content, "response"), "http",
                        else: if(contains(content, "database") OR
contains(content, "query"), "database",
                        else: if(contains(content, "auth") OR
contains(content, "login"), "security",
                        else: "general"))))
| summarize {count = count()}, by: {log_category, loglevel}
| sort count desc
```

```
// Identify namespaces for team ownership mapping
fetch logs, from: now() - 1h
| filter isNotNull(k8s.namespace.name)
| summarize {log_count = count()}, by: {k8s.namespace.name}
| sort log_count desc
| limit 20
```

## 5. Event Generation from Logs

---

Create **business events** from specific log patterns. Events flow to Grail and can trigger workflows.

# OpenPipeline Event Processor

```
# Generate events from critical log patterns
processors:
  - name: generate-payment-events
    type: bizevents
    enabled: true
    condition: contains(content, "payment") AND contains(content,
"completed")
    eventType: com.example.payment.completed
    attributes:
      - payment_id: extracted_payment_id
      - amount: extracted_amount
      - currency: extracted_currency
      - customer_id: extracted_customer

  - name: generate-error-events
    type: bizevents
    enabled: true
    condition: loglevel == "ERROR" AND contains(content, "critical")
    eventType: com.example.critical.error
    attributes:
      - error_type: extracted_error_type
      - service: k8s.namespace.name
```

## Event Use Cases

| Log Pattern     | Event Type           | Purpose            |
|-----------------|----------------------|--------------------|
| Order completed | order.completed      | Business analytics |
| User signup     | user.registered      | Funnel tracking    |
| Deployment      | deployment.completed | Change tracking    |
| Critical error  | error.critical       | Incident trigger   |



```
// Discover patterns suitable for event generation
fetch logs, from: now() - 1h
| filter contains(content, "completed")
      OR contains(content, "success")
      OR contains(content, "failed")
      OR contains(content, "created")
| fieldsAdd content_preview = substring(content, from: 0, to: 100)
| summarize {count = count()}, by: {content_preview}
| sort count desc
| limit 20
```

```
// Check existing bizevents (if any generated)
fetch bizevents, from: now() - 24h
| summarize {count = count()}, by: {event.type}
| sort count desc
| limit 15
```

```
// Preview: Logs that would become critical error events
fetch logs, from: now() - 1h
| filter loglevel == "ERROR"
| filter contains(content, "critical")
      OR contains(content, "fatal")
      OR contains(content, "severe")
| fieldsAdd content_preview = substring(content, from: 0, to: 100)
| fields timestamp, k8s.namespace.name, content_preview
| sort timestamp desc
| limit 20
```

## 6. Bucket Routing

---

Route logs to **appropriate buckets** based on content, source, or computed attributes.

## Why Bucket Routing Matters

| Bucket       | Logs                | Retention | Cost Impact   |
|--------------|---------------------|-----------|---------------|
| default_logs | Standard            | 35 days   | Baseline      |
| debug_logs   | DEBUG/TRACE         | 7 days    | 80% savings   |
| audit_logs   | Security/compliance | 365 days  | Compliance    |
| error_logs   | Errors only         | 90 days   | Investigation |

## OpenPipeline Route Processor

```
# Route logs to appropriate buckets
processors:
  - name: route-debug-logs
    type: route
    condition: loglevel == "DEBUG" OR loglevel == "TRACE"
    bucket: debug_logs

  - name: route-audit-logs
    type: route
    condition: contains(content, "audit") OR contains(content, "security")
    bucket: audit_logs

  - name: route-error-logs
    type: route
    condition: loglevel == "ERROR" OR loglevel == "FATAL"
    bucket: error_logs
```

```
// Current bucket distribution
fetch logs, from: now() - 1h
| summarize {count = count()}, by: {dt.system.bucket}
| sort count desc
```

```
// Preview routing decisions
fetch logs, from: now() - 1h
| fieldsAdd target_bucket = if(loglevel == "DEBUG" OR loglevel == "TRACE",
"debug_logs",
                                else: if(loglevel == "ERROR" OR loglevel ==
"FATAL", "error_logs",
                                else: if(contains(content, "audit") OR
contains(content, "security"), "audit_logs",
                                else: "default_logs"))
| summarize {count = count()}, by: {target_bucket, loglevel}
| sort count desc
```

```
// Estimate storage savings from routing
fetch logs, from: now() - 24h
| fieldsAdd content_bytes = stringLength(content)
| summarize {
    total_logs = count(),
    total_mb = sum(content_bytes) / 1048576.0,
    debug_logs = countIf(loglevel == "DEBUG" OR loglevel == "TRACE"),
    debug_mb = sum(if(loglevel == "DEBUG" OR loglevel == "TRACE",
content_bytes, else: 0)) / 1048576.0
}
| fieldsAdd debug_pct = round((debug_logs * 100.0) / total_logs, decimals: 1)
| fieldsAdd savings_7d_vs_35d = round(debug_mb * 0.8, decimals: 2)
```

## 7. Filtering & Sampling

---

Reduce log volume by **filtering** or **sampling** at ingestion.

## OpenPipeline Filter Processor

```
# Drop noisy, low-value logs
processors:
  - name: drop-health-checks
    type: filter
    condition: contains(content, "health") AND loglevel == "INFO"
    action: drop

  - name: drop-heartbeats
    type: filter
    condition: contains(content, "heartbeat") OR contains(content,
"keepalive")
    action: drop

  - name: sample-debug-logs
    type: sample
    condition: loglevel == "DEBUG"
    rate: 0.1 # Keep only 10% of DEBUG logs
```

```
// Identify candidates for filtering (high volume, low value)
fetch logs, from: now() - 1h
| filter contains(content, "health")
      OR contains(content, "heartbeat")
      OR contains(content, "alive")
      OR contains(content, "ready")
| summarize {count = count()}, by: {k8s.namespace.name, loglevel}
| sort count desc
| limit 15
```

```
// Calculate potential savings from filtering health checks
fetch logs, from: now() - 24h
| summarize {
  total_logs = count(),
  health_logs = countIf(contains(content, "health") OR contains(content,
"heartbeat")),
  debug_logs = countIf(loglevel == "DEBUG")
}
| fieldsAdd health_pct = round((health_logs * 100.0) / total_logs, decimals:
1)
| fieldsAdd debug_pct = round((debug_logs * 100.0) / total_logs, decimals: 1)
| fieldsAdd potential_reduction = round(((health_logs + debug_logs * 0.9) *
100.0) / total_logs, decimals: 1)
```

## 8. Complete Pipeline Example

---

Here's a comprehensive OpenPipeline configuration:

```
# Complete OpenPipeline configuration for logs
name: production-log-pipeline
enabled: true

processors:
  # 1. FILTER - Remove unwanted logs first
  - name: drop-health-checks
    type: filter
    condition: contains(content, "health") AND loglevel == "INFO"
    action: drop

  # 2. PARSE - Extract structured fields
  - name: parse-http-logs
    type: dql
    condition: contains(content, "HTTP")
    dql: parse content, "LD:method SPACE '/' LD:path SPACE INT:status SPACE
INT:duration_ms"

  # 3. ENRICH - Add computed attributes
  - name: add-environment
    type: fieldsAdd
    fields:
      - name: environment
        value: if(contains(k8s.namespace.name, "prod"), "production", else:
"non-prod")

  # 4. MASK - Protect sensitive data
  - name: mask-emails
    type: mask
    field: content
    pattern: "[A-Za-z0-9._%+-]+@[A-Za-z0-9.-]+\.\.[A-Za-z]{2,}"
    replacement: "[EMAIL-MASKED]"

  # 5. EXTRACT METRICS - Create dimensional metrics
  - name: extract-request-count
    type: metric
    metricKey: log.http.requests
    dimensions:
      - namespace: k8s.namespace.name
      - status: status

  # 6. GENERATE EVENTS - Create business events
  - name: generate-error-events
    type: bizevents
    condition: loglevel == "ERROR"
    eventType: com.app.error
```

```
# 7. ROUTE - Send to appropriate bucket
- name: route-debug
  type: route
  condition: loglevel == "DEBUG"
  bucket: debug_logs
```

```
// Verify current pipeline processing
fetch logs, from: now() - 1h
| summarize {
  total_logs = count(),
  with_pipeline = countIf(isNotNull(dt.openpipeline.pipelines)),
  unique_pipelines = countDistinct(dt.openpipeline.pipelines)
}, by: {dt.openpipeline.source}
| sort total_logs desc
```

```
// Pipeline processing summary
fetch logs, from: now() - 1h
| summarize {count = count()}, by: {dt.openpipeline.pipelines,
dt.system.bucket}
| sort count desc
```



## Summary

In this notebook, you learned:

- ✅ **Parsing** - Extract structured fields at ingestion
- ✅ **Metrics** - Create dimensional metrics from logs
- ✅ **Attributes** - Add computed fields for analysis
- ✅ **Events** - Generate business events from patterns
- ✅ **Routing** - Direct logs to appropriate buckets
- ✅ **Filtering** - Drop/sample low-value logs

## Key Takeaway

**Process at ingestion, not at query time.** OpenPipeline processing is fundamental to cost optimization, data quality, and operational efficiency.

## Next Steps

---

Continue to **OPLOGS-04: Buckets & Data Governance** to learn about storage management and retention policies.

---

## References

---

- [OpenPipeline Overview](#)
  - [OpenPipeline Processors](#)
  - [Metric Extraction](#)
  - [Event Generation](#)
  - [DPL Pattern Language](#)
- 

*This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.*