



Security & Data Protection

Series: OPLOGS | Notebook: 8 of 8 | Created: December 2025

Sensitive Data Discovery, Masking, and Compliance

This notebook covers sensitive data discovery, OpenPipeline masking configuration, security event monitoring, and compliance reporting.

Table of Contents


1. Sensitive Data Discovery
2. OpenPipeline Masking Configuration
3. IP Address Analysis
4. Security Event Monitoring
5. Audit Log Queries
6. Compliance Reporting
7. Masking Best Practices

Prerequisites

- ☒ Access to a Dynatrace environment with log data
- ☒ Completed OPLOGS-01 through OPLOGS-07
- ☒ Understanding of data privacy requirements

1. Sensitive Data Discovery

Before implementing masking, discover what sensitive data exists in your logs.



Common Sensitive Data Patterns

Data Type	Pattern Example	Risk Level
Email addresses	user@domain.com	Medium
Credit cards	4111-1111-1111-1111	Critical
Social Security	123-45-6789	Critical
API keys	sk_live_abc123...	Critical
IP addresses	192.168.1.100	Low-Medium
Phone numbers	+1-555-123-4567	Medium
JWT tokens	eyJhbG...	High

```
// Search for potential email addresses in logs
fetch logs, from: now() - 1h
| filter contains(content, "@")
| filter NOT contains(content, "@dynatrace")
| filter NOT contains(content, "@example")
| fieldsAdd content_preview = substring(content, from: 0, to: 150)
| summarize {count = count()}, by: {content_preview, k8s.namespace.name}
| sort count desc
| limit 20
```

```
// Search for potential API keys or tokens
fetch logs, from: now() - 1h
| filter contains(content, "key=")
      OR contains(content, "token=")
      OR contains(content, "api_key")
      OR contains(content, "apikey")
      OR contains(content, "secret")
| fieldsAdd content_preview = substring(content, from: 0, to: 120)
| summarize {count = count()}, by: {content_preview, k8s.namespace.name}
| sort count desc
| limit 20
```


```
// Search for potential JWT tokens
fetch logs, from: now() - 1h
| filter contains(content, "eyJ") // JWT typically starts with eyJ
| fieldsAdd content_preview = substring(content, from: 0, to: 100)
| summarize {count = count()}, by: {content_preview, k8s.namespace.name}
| sort count desc
| limit 15
```

```
// Search for potential credit card patterns (16 digits)
fetch logs, from: now() - 1h
| filter contains(content, "card")
      OR contains(content, "payment")
      OR contains(content, "visa")
      OR contains(content, "mastercard")
| fieldsAdd content_preview = substring(content, from: 0, to: 120)
| summarize {count = count()}, by: {content_preview}
| sort count desc
| limit 15
```

```
// Search for password-related entries (should never be logged!)
fetch logs, from: now() - 1h
| filter contains(content, "password")
      OR contains(content, "passwd")
      OR contains(content, "pwd=")
| fieldsAdd content_preview = substring(content, from: 0, to: 100)
| summarize {count = count()}, by: {content_preview, k8s.namespace.name}
| sort count desc
| limit 20
```

2. OpenPipeline Masking Configuration

OpenPipeline provides built-in masking processors to protect sensitive data at ingestion time.



Masking Processor Types

Processor	Use Case
Mask value	Replace with fixed pattern (e.g., <code>***MASKED***</code>)
Hash value	One-way hash for correlation without exposure
Remove field	Completely remove the field
Pattern mask	Partial masking (e.g., <code>****1234</code>)

Configuration Path

1. Settings → Log processing → OpenPipeline
2. Select or create pipeline
3. Add "Mask value" or "Hash value" processor
4. Configure field and pattern

DPL Patterns for Masking

```
# Email pattern
[A-Za-z0-9._%+-]+@[A-Za-z0-9.-]+\.[A-Za-z]{2,}
```

```
# Credit card pattern
\b(?:\d[ -]*?){13,16}\b
```

```
# SSN pattern (US)
\b\d{3}-\d{2}-\d{4}\b
```

```
# IP address pattern
\b(?:\d{1,3}\.){3}\d{1,3}\b
```

```
// Verify masking is working (look for masked patterns)
fetch logs, from: now() - 1h
| filter contains(content, "MASKED")
      OR contains(content, "***")
      OR contains(content, "[REDACTED]")
| fieldsAdd content_preview = substring(content, from: 0, to: 120)
| summarize {count = count()}, by: {content_preview}
| sort count desc
| limit 15
```

```
// Check which pipelines are processing logs
fetch logs, from: now() - 1h
| summarize {count = count()}, by: {dt.openpipeline.pipelines}
| sort count desc
```

3. IP Address Analysis

IP addresses may require masking depending on your compliance requirements.

```
// Find logs containing IP addresses
fetch logs, from: now() - 1h
| parse content, "IPADDR:ip_found"
| filter isNotNull(ip_found)
| summarize {count = count()}, by: {ip_found, k8s.namespace.name}
| sort count desc
| limit 20
```

```
// Classify IP addresses (internal vs external)
fetch logs, from: now() - 1h
| parse content, "IPADDR:ip_found"
| filter isNotNull(ip_found)
| fieldsAdd ip_type = if(startsWith(ip_found, "10."), "RFC1918-10",
                        else: if(startsWith(ip_found, "192.168."), "RFC1918-192",
                                else: if(startsWith(ip_found, "172."), "RFC1918-172",
                                        else: if(startsWith(ip_found, "127."), "LOOPBACK",
                                                else: "EXTERNAL"))))
| summarize {count = count()}, by: {ip_type}
| sort count desc
```

4. Security Event Monitoring

Monitor logs for security-relevant events.

```
// Authentication-related logs
fetch logs, from: now() - 1h
| filter contains(content, "login")
      OR contains(content, "auth")
      OR contains(content, "signin")
      OR contains(content, "logout")
| fieldsAdd content_preview = substring(content, from: 0, to: 100)
| summarize {count = count()}, by: {content_preview, k8s.namespace.name}
| sort count desc
| limit 20
```

```
// Failed authentication attempts
fetch logs, from: now() - 1h
| filter (contains(content, "failed") OR contains(content, "denied") OR
contains(content, "unauthorized"))
      AND (contains(content, "login") OR contains(content, "auth") OR
contains(content, "access"))
| fieldsAdd content_preview = substring(content, from: 0, to: 120)
| summarize {count = count()}, by: {content_preview, k8s.namespace.name}
| sort count desc
| limit 20
```

```
// Security-related errors over time
fetch logs, from: now() - 24h
| filter loglevel == "ERROR"
| filter contains(content, "security")
      OR contains(content, "unauthorized")
      OR contains(content, "forbidden")
      OR contains(content, "denied")
| makeTimeseries {security_errors = count()}, interval: 30m
```

```
// Access pattern anomalies - hourly access distribution
fetch logs, from: now() - 24h
| filter contains(content, "access") OR contains(content, "request")
| fieldsAdd hour_bucket = bin(timestamp, 1h)
| summarize {access_count = count()}, by: {hour_bucket, k8s.namespace.name}
| sort hour_bucket asc
```

5. Audit Log Queries

Track administrative and configuration changes.

```
// Administrative action logs
fetch logs, from: now() - 24h
| filter contains(content, "created")
      OR contains(content, "deleted")
      OR contains(content, "updated")
      OR contains(content, "modified")
| filter contains(content, "user")
      OR contains(content, "admin")
      OR contains(content, "config")
| fieldsAdd content_preview = substring(content, from: 0, to: 100)
| summarize {
  count = count(),
  first_seen = min(timestamp),
  last_seen = max(timestamp)
}, by: {content_preview, k8s.namespace.name}
| sort count desc
| limit 20
```

```
// Configuration change events
fetch logs, from: now() - 24h
| filter contains(content, "config")
      OR contains(content, "setting")
      OR contains(content, "environment")
| filter contains(content, "change")
      OR contains(content, "update")
      OR contains(content, "reload")
| fieldsAdd content_preview = substring(content, from: 0, to: 120)
| summarize {count = count()}, by: {content_preview, k8s.namespace.name}
| sort count desc
| limit 15
```

```
// Permission-related events
fetch logs, from: now() - 24h
| filter contains(content, "permission")
      OR contains(content, "role")
      OR contains(content, "privilege")
      OR contains(content, "rbac")
| fieldsAdd content_preview = substring(content, from: 0, to: 100)
| summarize {count = count()}, by: {content_preview, k8s.namespace.name}
| sort count desc
| limit 15
```

6. Compliance Reporting

```
// Data access summary
fetch logs, from: now() - 24h
| filter contains(content, "access") OR contains(content, "read") OR
contains(content, "view")
| summarize {
    total_accesses = count(),
    unique_sources = countDistinct(dt.openpipeline.source)
}, by: {k8s.namespace.name}
| sort total_accesses desc
| limit 15
```



```
// Log retention verification
fetch logs, from: now() - 7d
| summarize {
    earliest = min(timestamp),
    latest = max(timestamp),
    total_logs = count()
}, by: {dt.system.bucket}
| fieldsAdd retention_days = (latest - earliest) / 86400000000000
```

```
// Sensitive field exposure audit
fetch logs, from: now() - 1h
| summarize {
    total = count(),
    with_email_pattern = countIf(contains(content, "@")),
    with_key_pattern = countIf(contains(content, "key=") OR contains(content,
"token=")),
    with_password_ref = countIf(contains(content, "password") OR
contains(content, "passwd"))
}
| fieldsAdd email_exposure_pct = round((with_email_pattern * 100.0) / total,
decimals: 2)
| fieldsAdd key_exposure_pct = round((with_key_pattern * 100.0) / total,
decimals: 2)
| fieldsAdd password_exposure_pct = round((with_password_ref * 100.0) /
total, decimals: 2)
```

7. Masking Best Practices

Do's

1. **Mask at ingestion** - Use OpenPipeline processors to mask before storage
2. **Use hashing** - For fields that need correlation but not visibility
3. **Audit regularly** - Run discovery queries to find new sensitive data patterns
4. **Document patterns** - Keep a registry of masked data types and patterns
5. **Test in staging** - Verify masking works before production deployment

Don'ts

1. **Don't log passwords** - Ever, even "for debugging"
2. **Don't log full credit cards** - Maximum last 4 digits

3. **Don't rely on application masking alone** - OpenPipeline is your safety net
4. **Don't skip internal data** - Internal IP/usernames may still be sensitive
5. **Don't forget API keys** - Rotate if accidentally logged

OpenPipeline Masking Configuration Example

```
# OpenPipeline processor configuration
processors:
  - name: mask-emails
    type: mask
    field: content
    pattern: "[A-Za-z0-9._%+-]+@[A-Za-z0-9.-]+\.[A-Za-z]{2,}"
    replacement: "[EMAIL-MASKED]"

  - name: mask-api-keys
    type: mask
    field: content
    pattern: "(api_key|apikey|key)=[A-Za-z0-9]{20,}"
    replacement: "$1=[KEY-MASKED]"
```

```
// Security health summary
fetch logs, from: now() - 1h
| summarize {
    total_logs = count(),
    security_events = countIf(contains(content, "security") OR
contains(content, "auth")),
    failed_events = countIf(contains(content, "failed") OR contains(content,
"denied")),
    potential_pii = countIf(contains(content, "@") OR contains(content,
"password"))
}
| fieldsAdd security_event_rate = round((security_events * 100.0) /
total_logs, decimals: 2)
| fieldsAdd failure_rate = round((failed_events * 100.0) / total_logs,
decimals: 2)
| fieldsAdd pii_risk_rate = round((potential_pii * 100.0) / total_logs,
decimals: 2)
```



Summary

In this notebook, you learned:

- ✓ **Sensitive data discovery** - Finding PII, tokens, and credentials
 - ✓ **OpenPipeline masking** - Configuration and patterns
 - ✓ **IP address analysis** - Classification and masking needs
 - ✓ **Security monitoring** - Auth events, failures, anomalies
 - ✓ **Audit logging** - Admin actions, config changes
 - ✓ **Compliance reporting** - Data access and exposure audits
-



Next Steps

Continue to **OPLOGS-07: Buckets & Cost Optimization** for storage management.



References

- [OpenPipeline Data Masking](#)
 - [Log Security Best Practices](#)
 - [GDPR and Compliance](#)
-

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.