ORGNZ-02: Understanding Grail Buckets

> **Series:** ORGNZ | **Notebook:** 2 of 10 | **Created:** January 2026 | **Last Updated:** 04/03/2026

Overview

Buckets are logical storage containers in Dynatrace Grail where records are stored. Think of a bucket as a folder in a filesystem. Use buckets to group telemetry data that naturally belongs together, such as data from the same region, environment, or with the same sensitivity classification.

Prerequisites

Requirement	Details
Dynatrace Environment	SaaS environment with Grail enabled
Permissions	`storage:buckets:read`, `storage:logs:read`
Knowledge	Completed ORGNZ-01 (Introduction)
Data	At least 1 hour of log data across default buckets

Table of Contents

-------------|----------|
Dynatrace Environment	SaaS environment with Grail enabled
Permissions	`storage:buckets:read` for viewing, `storage:bucket-definitions:write` for creating
Knowledge	Completed ORGNZ-01

Bucket Administration Permissions

Permission	Purpose
`storage:bucket-definitions:read`	View bucket configurations
`storage:bucket-definitions:write`	Create and modify buckets
`storage:bucket-definitions:delete`	Remove buckets
`storage:bucket-definitions:truncate`	Truncate bucket data

Learning Objectives

By the end of this notebook, you will:
- Understand bucket fundamentals and supported data types
- Know bucket limits and query constraints
- Query bucket information using DQL
- Recognize when custom buckets are needed

What Are Buckets?

Buckets organize data to meet performance, compliance, and retention requirements:

Function	Description
Data Grouping	Organize large datasets into meaningful segments
Streamlined Organization	Records that should be handled together are stored together
Query Performance	Reduced queried dataset improves query speed
Retention Management	Each bucket has its own retention policy
Access Control	Bucket-level IAM policies for team isolation
Cost Attribution	Enable data cost tracking by organizational unit

Supported Data Types (Tables)

Each bucket is associated with exactly **one data type**:

Data Type	Default Bucket	Description
logs	`default_logs`	Log records from all sources
metrics	`default_metrics`	Metric data points
spans	`default_spans`	Distributed trace spans
events	`default_events`	Platform events
bizevents	`default_bizevents`	Business events
security_events	`default_securityevents`	Security-related events

> **Important**: You cannot mix data types in the same bucket. A logs bucket only stores logs.

Bucket Limits

Environment Limits

Limit	Value	Notes
Maximum buckets per environment	80	Default limit; can request increase
Typical capacity	Up to 5 TB/day per table	With default bucket limit

Bucket Size Guidelines

Metric	Value	Impact
Optimal ingest per bucket	~1 TB/day	Best query performance
Acceptable ingest	1-3 TB/day	Limited query window
Maximum ingest	3 TB/day	Performance degrades above this

Retention Limits

Limit	Value
Minimum retention	1 day
Maximum retention	3,657 days (~10 years + 1 week)

Query Constraints

Understanding query limits is critical for bucket planning:

Query Limits

Limit	Value	Impact
Maximum data scanned	500 GB	Limits queryable time window

Limit	Value	Impact
Maximum records returned	1,000	Use aggregations for larger datasets
Maximum response payload	1 MB	Large result sets may be truncated

Queryable Window by Ingest Volume

The 500 GB scan limit determines how far back you can query:

Daily Ingest	Queryable Window
500 GB/day	~24 hours
1 TB/day	~12 hours
3 TB/day	~4 hours
10 TB/day	~1 hour

> **Implication**: High-volume buckets may only allow querying recent data. Use aggregations, time filters, and strategic bucket partitioning to work within these limits.

Default Buckets

Dynatrace provides default buckets with varying retention:

Default Bucket	Data Type	Default Retention	Notes
`default_logs`	logs	35 days	Standard log retention
`default_metrics`	metrics	15 months	Extended for trend analysis
`default_spans`	spans	10 days	Short-term APM data

Default Bucket	Data Type	Default Retention	Notes
`default_events`	events	35 days	Platform events
`default_bizevents`	bizevents	35 days	Business events
`default_securityevents`	security_events	1 year	Security events
`default_securityevents_builtin`	security_events	3 years	Built-in security events

> **Note**: Query your environment's `dt.system.buckets` to verify current retention settings.

Querying Bucket Information

> **Lab Exercise:** Complete Exercise 1 in **ORGNZ-02 LAB** for hands-on practice with this concept.

Querying Data from Buckets

> **Lab Exercise:** Complete Exercises 2-3 in **ORGNZ-02 LAB** for hands-on practice with these concepts.

Bucket Characteristics

Immutability Rules

Characteristic	Mutable?	Notes
Bucket name	No	Cannot be changed after creation
Display name	Yes	Can be updated anytime
Retention period	Yes	**Caution**: Lowering deletes data
Data table type	No	Fixed at creation

Data Handling

Operation	Supported?	Alternative
Move data between buckets	No	Route new data via OpenPipeline
Selective record deletion	No	Wait for retention to expire
Bucket deletion	Yes	Deletes ALL data permanently

Bucket Naming Rules

Rule	Requirement
First character	Must be a letter
Allowed characters	Lowercase alphanumeric, underscores, hyphens
Case	Lowercase only
Reserved names	Cannot use `default_*` prefix

Valid Examples

```
team_logs_30d
finance-audit-logs
prod_metrics_90d
app123_spans
```

Invalid Examples

```
123_logs          # Starts with number
Team_Logs         # Contains uppercase
default_custom    # Uses reserved prefix
my logs           # Contains space
```

When to Create Custom Buckets

Need	Solution
Different retention periods	Create bucket with specific retention
Cost attribution by team/LOB	Separate buckets per cost center
Simple team isolation	Bucket-level IAM policies
Regulatory compliance	Dedicated buckets for compliance data
Query performance	Isolate high-volume data sources
Debug/verbose logs	Short-retention bucket to reduce costs

Next Steps

Continue with the ORGNZ series:
- **ORGNZ-03**: Bucket Strategy and Design

References

- Grail Buckets
- Data Retention
- DQL Reference

This notebook was AI-generated from Dynatrace documentation and enterprise best practices. It is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.