

ORGNZ-05 LAB: Bucket-Level Access Control - Hands-on Exercises

Series: ORGNZ — Organize Data: Buckets, Segments, Security | **Notebook:** 5 of 10 | **Type:** LAB | **Created:** February 2026 | **Last Updated:** 02/19/2026

Overview

This lab notebook contains 3 hands-on exercises extracted from **ORGNZ-05: Bucket-Level Access Control**. Complete the lecture notebook first, then work through these exercises to reinforce the concepts with real DQL queries against your Dynatrace environment.

Table of Contents

- [Exercise 1: Required Permissions](#)
 - [Exercise 2: Required Permissions](#)
 - [Exercise 3: Required Permissions](#)
 - [Lab Summary](#)
-

Prerequisites

Requirement	Details
Completed	ORGNZ-05: Bucket-Level Access Control (lecture notebook)
Dynatrace Environment	SaaS tenant with Grail enabled

Requirement	Details
Permissions	logs.read , metrics.read , entities.read , spans.read

Exercise 1: Required Permissions

ORGNZ-05: Bucket-Level Access Control

Series: ORGNZ — Organize Data: Buckets, Segments, Security | **Notebook:** 5 of 10 | **Created:** January 2026 | **Last Updated:** 01/28/2026

Bucket-level access control provides a straightforward way to isolate data by team, application, or business unit. By granting permissions to specific buckets, you can ensure teams only access data relevant to their responsibilities.

| Requirement | Details |
|---

- 1. Bucket Permission Fundamentals
- 2. Policy Examples
- 3. Team Isolation Pattern
- 4. Default Buckets Policy
- 5. Verifying Bucket Access
- 6. Best Practices
- 7. When Bucket-Level Isn't Enough

-----|-----|
| **Dynatrace Account** | Account-level administrative access |

| **Permissions** | IAM policy management permissions |
| **Knowledge** | Completed ORGNZ-04 |

By the end of this notebook, you will:

- Create IAM policies for bucket-level access
- Use bucket naming patterns in policies
- Implement team isolation using buckets
- Understand bucket permission best practices

All bucket access policies must start with `storage:buckets:read` :

```
ALLOW storage:buckets:read WHERE <condition>;  
ALLOW storage:logs:read; // Then allow table access
```

Step	Purpose	Example
1. Bucket access	Define which buckets	<code>storage:buckets:read WHERE bucket-name = 'x'</code>
2. Table access	Define which data types	<code>storage:logs:read , storage:metrics:read</code>

Grant a team access to their specific bucket:

```
{  
  "name": "platform-team-logs-access",  
  "description": "Platform team can access platform logs bucket",  
  "statementQuery": "ALLOW storage:buckets:read WHERE storage:bucket-name =  
'team_platform_logs'; ALLOW storage:logs:read;",  
  "tags": ["team:platform"]  
}
```

Grant access to a list of specific buckets:


```
{  
  "name": "finance-multi-bucket-access",  
  "description": "Finance team can access multiple buckets",  
  "statementQuery": "ALLOW storage:buckets:read WHERE storage:bucket-name IN  
( 'finance_logs', 'finance_metrics', 'finance_audit' ); ALLOW  
storage:logs:read, storage:metrics:read;",  
  "tags": ["team:finance"]  
}
```

Grant access to all buckets matching a prefix:

```
{
  "name": "prod-infrastructure-access",
  "description": "Access to all production infrastructure buckets",
  "statementQuery": "ALLOW storage:buckets:read WHERE storage:bucket-name STARTSWITH 'prod_infra_'; ALLOW storage:logs:read, storage:metrics:read, storage:spans:read;",
  "tags": ["env:production"]
}
```

Use pattern matching for complex bucket names:

```
{
  "name": "database-team-access",
  "description": "Database team can access any database-related bucket",
  "statementQuery": "ALLOW storage:buckets:read WHERE storage:bucket-name MATCH ('*-database-*'); ALLOW storage:logs:read;",
  "tags": ["team:database"]
}
```



Step 1: Create team buckets

```
team_platform_logs
team_checkout_logs
team_payments_logs
```

Step 2: Route data via OpenPipeline

```
processors:
  - type: route
    rules:
      - condition: "host.group starts-with 'platform-'"
        destination: "team_platform_logs"
      - condition: "service.name contains 'checkout'"
        destination: "team_checkout_logs"
```

Step 3: Create IAM policies per team

```
// Platform team policy
ALLOW storage:buckets:read WHERE storage:bucket-name STARTSWITH
"team_platform_";
ALLOW storage:logs:read;

// Checkout team policy
ALLOW storage:buckets:read WHERE storage:bucket-name STARTSWITH
"team_checkout_";
ALLOW storage:logs:read;
```

For users who need access to all default buckets:

```
ALLOW storage:buckets:read WHERE storage:bucket-name STARTSWITH "default_";
ALLOW storage:events:read, storage:logs:read, storage:metrics:read,
storage:entities:read, storage:bizevents:read, storage:spans:read;
```

```
ALLOW storage:buckets:read WHERE storage:bucket-name STARTSWITH "default_";
ALLOW storage:buckets:read WHERE storage:bucket-name = "audit_logs_365d";
ALLOW storage:logs:read, storage:metrics:read, storage:spans:read;
```

```
// List buckets you have access to
fetch dt.system.buckets
| fields name, display_name, dt.system.table, retention_days
| sort name asc
```

Exercise 2: Required Permissions

```
// Query data from specific bucket to verify access
fetch logs, bucket: "team_platform_logs"
| limit 10
```

Exercise 3: Required Permissions

```
// Check data distribution across accessible buckets
fetch logs, from:-1h
| summarize count = count(), by:{dt.system.bucket}
| sort count desc
```

Lab Summary

You have completed 3 hands-on exercises for **ORGNZ-05: Bucket-Level Access Control**.

Exercises Completed

- ☐ Exercise 1: Required Permissions
- ☐ Exercise 2: Required Permissions
- ☐ Exercise 3: Required Permissions

Next Steps

Continue with **ORGNZ-06** for the next notebook.

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.