

ORGNZ-09 LAB: Enterprise Data Organization Patterns - Hands-on Exercises

Series: ORGNZ — Organize Data: Buckets, Segments, Security | **Notebook:** 9 of 10 | **Type:** LAB | **Created:** February 2026 | **Last Updated:** 04/25/2026

Overview

This lab notebook contains 3 hands-on exercises extracted from **ORGNZ-09: Enterprise Data Organization Patterns**. Complete the lecture notebook first, then work through these exercises to reinforce the concepts with real DQL queries against your Dynatrace environment.

Table of Contents

1. [Exercise 1: Complete Architecture Example](#)
 2. [Exercise 2: Complete Architecture Example](#)
 3. [Exercise 3: Complete Architecture Example](#)
 4. [Lab Summary](#)
-

Prerequisites

Requirement	Details
Completed	ORGNZ-09: Enterprise Data Organization Patterns (lecture notebook)
Dynatrace Environment	SaaS tenant with Grail enabled
Permissions	logs.read , metrics.read , entities.read , spans.read

Exercise 1: Complete Architecture Example

ORGNZ-09: Enterprise Data Organization Patterns

Series: ORGNZ — Organize Data: Buckets, Segments, Security | **Notebook:** 9 of 10 | **Created:** January 2026 | **Last Updated:** 04/25/2026

This notebook brings together buckets, segments, and security context into comprehensive enterprise patterns. Learn how to combine these mechanisms for effective data governance at scale.

| Requirement | Details |
|---

1. Combining All Three Mechanisms
2. Enterprise Pattern 1: Line of Business Isolation
3. Enterprise Pattern 2: Environment-Based Tiering
4. Enterprise Pattern 3: Multi-Cloud Organization
5. Enterprise Pattern 4: Kubernetes-Native Organization

6. Decision Framework
7. Implementation Checklist
8. Auditing Your Organization
9. Best Practices Summary
10. Series Summary

-----|-----|

Dynatrace Account	Account-level administrative access
Permissions	Full IAM and bucket management permissions
Knowledge	Completed ORGNZ-01 through ORGNZ-08

By the end of this notebook, you will:

- Combine buckets, segments, and security context effectively
- Design enterprise-scale data organization architectures
- Implement common organizational patterns
- Plan data governance strategies

For comprehensive data governance, use all three mechanisms:

Mechanism	Purpose	When to Use
Buckets	Physical storage, retention, cost	Different retention needs, cost attribution
Segments	Logical filtering, views	Team-focused data views, cross-signal filtering
Security Context	Access enforcement	Fine-grained security requirements

Finance Team Requirements:

Bucket: `finance_logs_365d`

- 1 year retention for compliance
- Cost attributed to Finance


Segment: `finance-team-view`

- Filtered view of finance data
- Includes related services and hosts

Security Context: `lob:finance`

- IAM policy: `ALLOW WHERE security_context = 'lob:finance'`
- Only finance team members can access

Complete isolation by business unit:



Step 1: Create buckets

```
finance_logs_365d (logs, 365 days)
finance_metrics_90d (metrics, 90 days)
retail_logs_90d (logs, 90 days)
```

Step 2: Configure OpenPipeline routing

```
processors:
  - type: route
    rules:
      - condition: "host.group starts-with 'finance-'"
        destination: "finance_logs_365d"
      - condition: "host.group starts-with 'retail-'"
        destination: "retail_logs_90d"
```

Step 3: Set security context

```
processors:
  - type: security-context
    rules:
      - condition: "host.group starts-with 'finance-'"
        context: "lob:finance"
      - condition: "host.group starts-with 'retail-'"
        context: "lob:retail"
```

Step 4: Create IAM policies

```
// Finance policy
ALLOW storage:buckets:read WHERE storage:bucket-name STARTSWITH "finance_";
ALLOW storage:logs:read, storage:metrics:read WHERE
storage:dt.security_context = "lob:finance";
```

Different access levels for production vs non-production:

```
Production:
├─ Buckets: prod_logs, prod_metrics, prod_spans
├─ Security Context: env:production
├─ Policy: ALLOW WHERE storage:dt.security_context = 'env:production'
├─ Access: Restricted to production support team
└─ Audit: Full audit logging enabled

Non-Production:
├─ Buckets: default_logs, default_metrics, default_spans
├─ Security Context: env:dev, env:staging, env:qa
├─ Policy: ALLOW WHERE storage:dt.security_context MATCH ('env:dev*',
'env:staging*')
├─ Access: Broader developer access
└─ Audit: Reduced audit requirements
```

Group	Production Access	Non-Prod Access
SRE Team	Full	Full
Production Support	Read-only	None
Developers	None	Full
QA	None	Read-only

Organize by cloud provider and account:

AWS:

```
|— Bucket: aws_logs_35d
|— Security Context: cloud:aws/<account-id>;
|— Policy: ALLOW WHERE storage:aws.account.id = '123456789012'
```

Azure:

```
|— Bucket: azure_logs_35d
|— Security Context: cloud:azure/<subscription>;
|— Policy: ALLOW WHERE storage:azure.subscription = 'sub-id'
```

GCP:

```
|— Bucket: gcp_logs_35d
|— Security Context: cloud:gcp/<project>;
|— Policy: ALLOW WHERE storage:gcp.project.id = 'my-project'
```

Leverage Kubernetes constructs for organization:

Cluster: main-cluster

```
|— Namespace: production
|   |— Security Context: k8s:main-cluster/production
|   |— Policy: ALLOW WHERE storage:k8s.namespace.name = 'production'
|
|— Namespace: staging
|   |— Security Context: k8s:main-cluster/staging
|   |— Policy: ALLOW WHERE storage:k8s.namespace.name = 'staging'
|
|— Namespace: development
|   |— Security Context: k8s:main-cluster/development
|   |— Policy: ALLOW WHERE storage:k8s.namespace.name = 'development'
```

```


name: checkout-application
displayName: "Checkout App (All Environments)"

includes:
  - type: logs
    filter: "k8s.deployment.name == 'checkout'"

  - type: spans
    filter: "service.name == 'checkout'"

```

Use this framework to choose your organization strategy:



- ☐ Identified retention requirements
- ☐ Mapped organizational structure (teams, LOBs, cost centers)
- ☐ Defined access control requirements
- ☐ Planned bucket naming convention
- ☐ Designed security context schema
- ☐ Created custom buckets
- ☐ Configured OpenPipeline routing rules
- ☐ Configured OpenPipeline security context processors

- [] Verified data flow to correct buckets
- [] Created IAM policies
- [] Assigned policies to groups
- [] Tested access with sample users
- [] Documented all policies
- [] Created segments for team views
- [] Tested segment filters
- [] Documented segment purposes
- [] Planned access review process
- [] Configured audit logging
- [] Created runbook for access management

```
// Audit bucket distribution
fetch logs, from:-1h
| summarize
    recordCount = count(),
    by:{dt.system.bucket}
| sort recordCount desc
```

Exercise 2: Complete Architecture Example

```
// Audit security context coverage
fetch logs, from:-1h
| summarize
    total = count(),
    withContext = countIf(isNotNull(dt.security_context)),
    withoutContext = countIf(isNull(dt.security_context))
| fieldsAdd coverage = round(toDouble(withContext) / toDouble(total) * 100,
decimals: 2)
```


Exercise 3: Complete Architecture Example

```
// Audit bucket retention settings
fetch dt.system.buckets
| fields name, dt.system.table, retention_days
| sort dt.system.table asc, retention_days desc
```

Lab Summary

You have completed 3 hands-on exercises for **ORGNZ-09: Enterprise Data Organization Patterns**.

Exercises Completed

- [] Exercise 1: Complete Architecture Example
- [] Exercise 2: Complete Architecture Example
- [] Exercise 3: Complete Architecture Example

Next Steps

Continue with **ORGNZ-10** for the next notebook.

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.