ORGNZ-09: Enterprise Data Organization Patterns

> **Series:** ORGNZ | **Notebook:** 9 of 9 | **Created:** January 2026 | **Last Updated:** 01/28/2026

Overview

This notebook brings together buckets, segments, and security context into comprehensive enterprise patterns. Learn how to combine these mechanisms for effective data governance at scale.

Prerequisites

| Requirement | Details |
|---

Table of Contents

10. Series Summary

----------|----------|
Dynatrace Account	Account-level administrative access
Permissions	Full IAM and bucket management permissions
Knowledge	Completed ORGNZ-01 through ORGNZ-08

Learning Objectives

By the end of this notebook, you will:
- Combine buckets, segments, and security context effectively
- Design enterprise-scale data organization architectures
- Implement common organizational patterns
- Plan data governance strategies

Combining All Three Mechanisms

For comprehensive data governance, use all three mechanisms:

Mechanism	Purpose	When to Use
Buckets	Physical storage, retention, cost	Different retention needs, cost attribution
Segments	Logical filtering, views	Team-focused data views, cross-signal filtering
Security Context	Access enforcement	Fine-grained security requirements

Complete Architecture Example

```
Finance Team Requirements:

Bucket: finance_logs_365d
  - 1 year retention for compliance
  - Cost attributed to Finance

Segment: finance-team-view
  - Filtered view of finance data
  - Includes related services and hosts

Security Context: lob:finance
  - IAM policy: ALLOW WHERE security_context = 'lob:finance'
  - Only finance team members can access
```

Enterprise Pattern 1: Line of Business Isolation

Complete isolation by business unit:

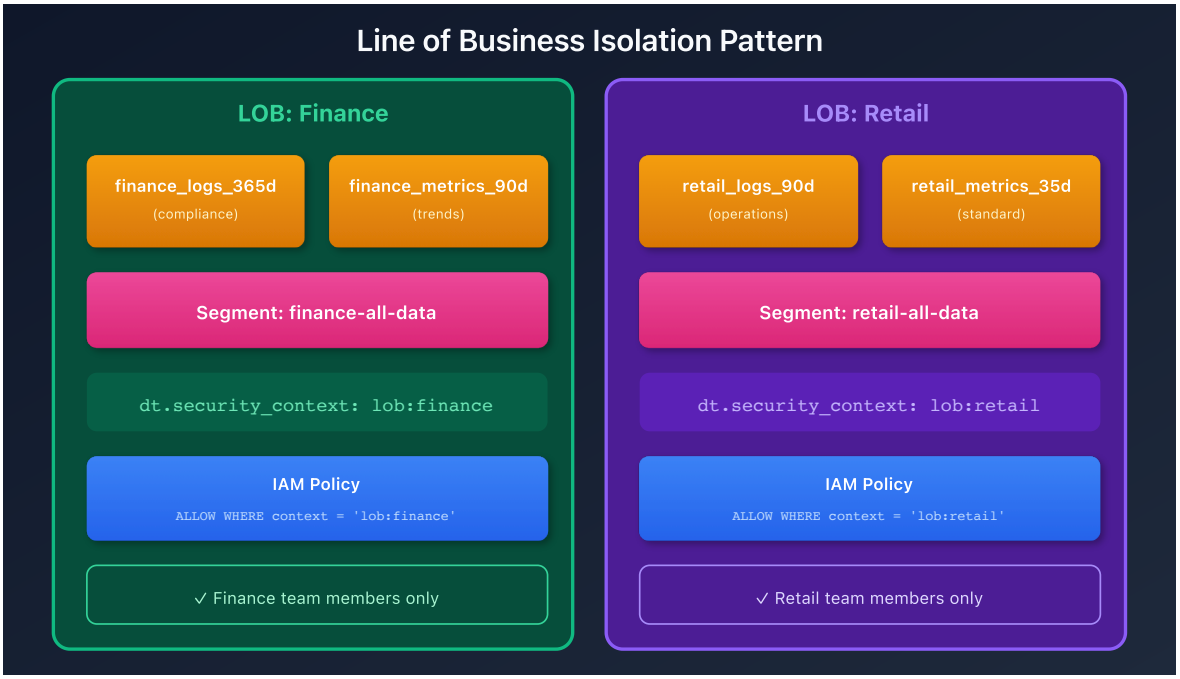

Implementation

Step 1: Create buckets

```
finance_logs_365d (logs, 365 days)
finance_metrics_90d (metrics, 90 days)
retail_logs_90d (logs, 90 days)
```

Step 2: Configure OpenPipeline routing

```
processors:
  - type: route
    rules:
      - condition: "host.group starts-with 'finance-'"
        destination: "finance_logs_365d"
      - condition: "host.group starts-with 'retail-'"
        destination: "retail_logs_90d"
```

Step 3: Set security context

```
processors:
  - type: security-context
    rules:
      - condition: "host.group starts-with 'finance-'"
        context: "lob:finance"
      - condition: "host.group starts-with 'retail-'"
        context: "lob:retail"
```

Step 4: Create IAM policies

```
// Finance policy
ALLOW storage:buckets:read WHERE storage:bucket-name STARTSWITH "finance_";
ALLOW storage:logs:read, storage:metrics:read WHERE
storage:dt.security_context = "lob:finance";
```

Enterprise Pattern 2: Environment-Based Tiering

Different access levels for production vs non-production:

```
Production:
├── Buckets: prod_logs, prod_metrics, prod_spans
├── Security Context: env:production
├── Policy: ALLOW WHERE storage:dt.security_context = 'env:production'
├── Access: Restricted to production support team
└── Audit: Full audit logging enabled

Non-Production:
├── Buckets: default_logs, default_metrics, default_spans
├── Security Context: env:dev, env:staging, env:qa
├── Policy: ALLOW WHERE storage:dt.security_context MATCH ('env:dev*',
'env:staging*')
├── Access: Broader developer access
└── Audit: Reduced audit requirements
```

Tiered Access Groups

Group	Production Access	Non-Prod Access
SRE Team	Full	Full
Production Support	Read-only	None
Developers	None	Full
QA	None	Read-only

Enterprise Pattern 3: Multi-Cloud Organization

Organize by cloud provider and account:

```
AWS:
├── Bucket: aws_logs_35d
├── Security Context: cloud:aws/&lt;account-id&gt;
└── Policy: ALLOW WHERE storage:aws.account.id = '123456789012'

Azure:
├── Bucket: azure_logs_35d
├── Security Context: cloud:azure/&lt;subscription&gt;
└── Policy: ALLOW WHERE storage:azure.subscription = 'sub-id'

GCP:
├── Bucket: gcp_logs_35d
├── Security Context: cloud:gcp/&lt;project&gt;
└── Policy: ALLOW WHERE storage:gcp.project.id = 'my-project'
```

Enterprise Pattern 4: Kubernetes-Native Organization

Leverage Kubernetes constructs for organization:

```
Cluster: main-cluster
├── Namespace: production
|   ├── Security Context: k8s:main-cluster/production
|   └── Policy: ALLOW WHERE storage:k8s.namespace.name = 'production'
|
├── Namespace: staging
|   ├── Security Context: k8s:main-cluster/staging
|   └── Policy: ALLOW WHERE storage:k8s.namespace.name = 'staging'
|
└── Namespace: development
    ├── Security Context: k8s:main-cluster/development
    └── Policy: ALLOW WHERE storage:k8s.namespace.name = 'development'
```

Cross-Namespace Segment

```
name: checkout-application
displayName: "Checkout App (All Environments)"

includes:
  - type: logs
    filter: "k8s.deployment.name == 'checkout'"

  - type: spans
    filter: "service.name == 'checkout'"
```

Decision Framework

Use this framework to choose your organization strategy:

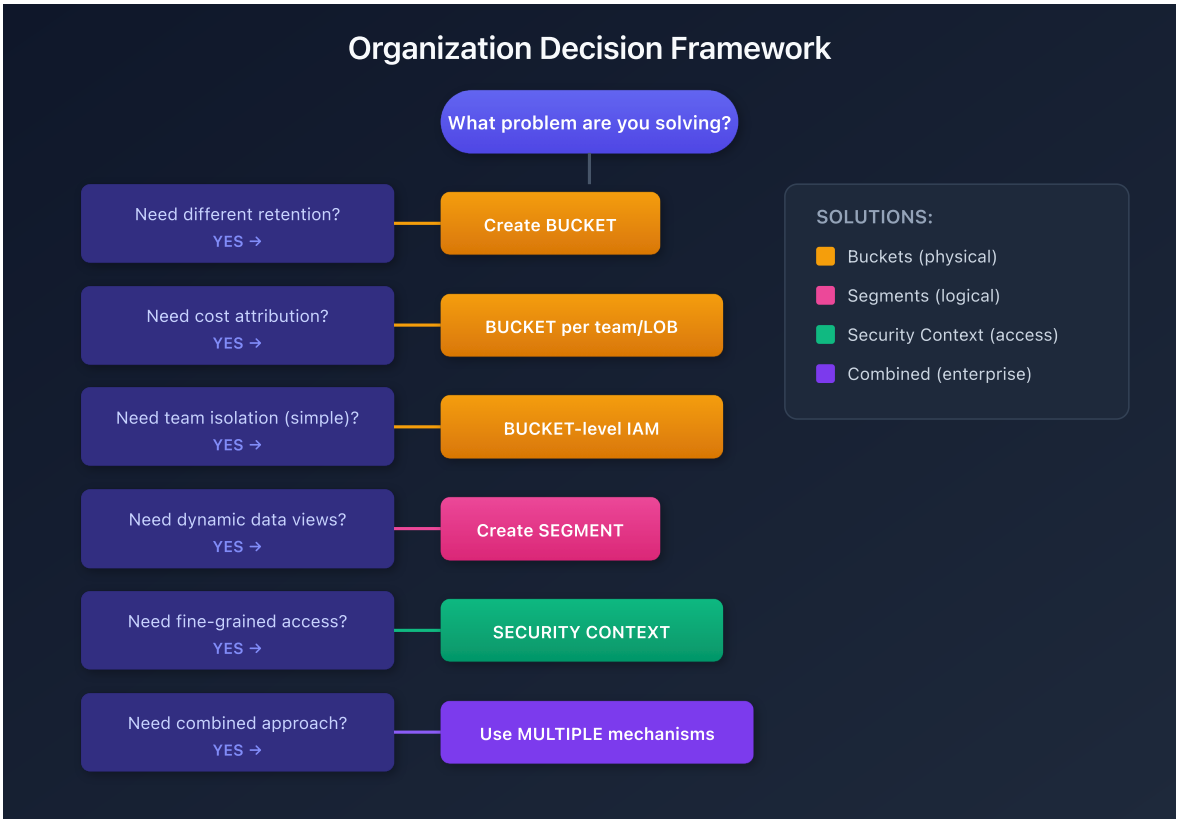

Implementation Checklist

Phase 1: Planning

- [ ] Identified retention requirements
- [ ] Mapped organizational structure (teams, LOBs, cost centers)
- [ ] Defined access control requirements
- [ ] Planned bucket naming convention
- [ ] Designed security context schema

Phase 2: Infrastructure

- [ ] Created custom buckets
- [ ] Configured OpenPipeline routing rules
- [ ] Configured OpenPipeline security context processors
- [ ] Verified data flow to correct buckets

Phase 3: Access Control

- [ ] Created IAM policies
- [ ] Assigned policies to groups
- [ ] Tested access with sample users
- [ ] Documented all policies

Phase 4: Views

- [ ] Created segments for team views
- [ ] Tested segment filters
- [ ] Documented segment purposes

Phase 5: Governance

- [ ] Planned access review process
- [ ] Configured audit logging
- [ ] Created runbook for access management

Auditing Your Organization

```
// Audit bucket distribution
fetch logs
| summarize
    recordCount = count(),
    by:{dt.system.bucket}
| sort recordCount desc
```

```
// Audit security context coverage
fetch logs
| summarize
    total = count(),
    withContext = countIf(isNotNull(dt.security_context)),
    withoutContext = countIf(isNull(dt.security_context))
| fieldsAdd coverage = round(toDouble(withContext) / toDouble(total) * 100,
decimals: 2)
```

```
// Audit bucket retention settings
fetch dt.system.buckets
| fields bucket.name, bucket.table, bucket.retentionDays
| sort bucket.table asc, bucket.retentionDays desc
```

Best Practices Summary

Area	Best Practice
Buckets	Keep ingest <1 TB/day; use consistent naming
Segments	Test rules with DQL before deploying
Security Context	Use hierarchical encoding for flexibility
IAM Policies	Assign to groups, not individuals
Documentation	Document all policies and contexts
Reviews	Schedule regular access reviews

Series Summary

The ORGNZ series covered:

Notebook	Topic	Key Takeaway
ORGNZ-01	Introduction	Why organize data; three pillars
ORGNZ-02	Grail Buckets	Bucket fundamentals and limits
ORGNZ-03	Bucket Strategy	Naming, retention, design patterns
ORGNZ-04	Permissions Overview	Permission levels and policy structure
ORGNZ-05	Bucket-Level Access	IAM policies for bucket isolation
ORGNZ-06	Security Context	Fine-grained access with dt.security_context
ORGNZ-07	Advanced Permissions	Record and field-level patterns
ORGNZ-08	Grail Segments	Dynamic data organization
ORGNZ-09	Enterprise Patterns	Combined approaches at scale

References

- Organize data
- Permissions in Grail
- Advanced permission setup
- Enhance data management with Grail

This notebook was AI-generated from Dynatrace documentation and enterprise best practices. It is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.