

Metrics Instrumentation

Series: OTEL | **Notebook:** 5 of 8 | **Created:** January 2026 | **Last Updated:** 02/09/2026

Creating Custom Metrics with OpenTelemetry

OpenTelemetry metrics provide quantitative measurements of your application's behavior over time. This notebook covers metric types, instrumentation patterns, and integration with Dynatrace.

Table of Contents

- 1. [Metric Types](#)
- 2. [Instrument Selection](#)
- 3. [Creating Metrics](#)
- 4. [Metric Attributes](#)
- 5. [Aggregation and Views](#)
- 6. [Prometheus Integration](#)
- 7. [Best Practices](#)

Prerequisites

Requirement	Details
Knowledge	OTEL-01 and OTEL-03
Environment	Collector with metrics pipeline

1. Metric Types

OTel Metric Instruments

Instrument	Aggregation	Use Case	Example
Counter	Sum (monotonic)	Cumulative events	Total requests
UpDownCounter	Sum (non-monotonic)	Current values	Active connections
Histogram	Distribution	Value ranges	Response time
Gauge	Last value	Current measurement	Temperature

Sync vs Async

Type	When Recorded	Use Case
Synchronous	On each operation	Request handling
Asynchronous (Observable)	On collection	System stats, gauges

2. Instrument Selection

Decision Guide

Question	If Yes	If No
Value only goes up?	Counter	UpDownCounter or Histogram
Measuring duration/size?	Histogram	Counter or UpDownCounter
Current state?	Gauge	Counter or Histogram
Recording in request path?	Sync	Async (Observable)

Common Patterns

Metric	Instrument	Why
Request count	Counter	Monotonically increasing
Request duration	Histogram	Distribution needed
Active requests	UpDownCounter	Goes up and down
Queue size	ObservableGauge	Current state
CPU usage	ObservableGauge	Periodic measurement

3. Creating Metrics

Python Examples

```
from opentelemetry import metrics
from opentelemetry.sdk.metrics import MeterProvider
from opentelemetry.sdk.metrics.export import PeriodicExportingMetricReader
from opentelemetry.exporter.otlp.proto.grpc.metric_exporter import
OTLPMetricExporter

# Setup
exporter = OTLPMetricExporter(endpoint="http://collector:4317")
reader = PeriodicExportingMetricReader(exporter,
export_interval_millis=60000)
provider = MeterProvider(metric_readers=[reader])
metrics.set_meter_provider(provider)

# Get meter
meter = metrics.get_meter(__name__)

# Counter
request_counter = meter.create_counter(
    name="http.server.request.count",
    description="Total HTTP requests",
    unit="1"
)

# Histogram
request_duration = meter.create_histogram(
    name="http.server.request.duration",
    description="HTTP request duration",
    unit="ms"
)

# UpDownCounter
active_requests = meter.create_up_down_counter(
    name="http.server.active_requests",
    description="Active HTTP requests",
    unit="1"
)
```

Using Metrics

```
import time

def handle_request(request):
    # Increment active requests
    active_requests.add(1, {"method": request.method})
    start = time.time()

    try:
        response = process_request(request)
        status = response.status_code
    except Exception:
        status = 500
        raise
    finally:
        # Record duration
        duration_ms = (time.time() - start) * 1000
        request_duration.record(duration_ms, {
            "method": request.method,
            "status": str(status)
        })

        # Increment counter
        request_counter.add(1, {
            "method": request.method,
            "status": str(status)
        })

        # Decrement active requests
        active_requests.add(-1, {"method": request.method})

    return response
```

Observable (Async) Metrics

```
import psutil

# Observable gauge for CPU
def get_cpu_usage(options):
    yield metrics.Observation(psutil.cpu_percent(), {})

cpu_gauge = meter.create_observable_gauge(
    name="system.cpu.usage",
    callbacks=[get_cpu_usage],
    description="CPU usage percentage",
    unit="%"
)

# Observable counter for network bytes
last_bytes = {"sent": 0, "recv": 0}

def get_network_io(options):
    net = psutil.net_io_counters()
    yield metrics.Observation(net.bytes_sent, {"direction": "sent"})
    yield metrics.Observation(net.bytes_recv, {"direction": "recv"})

network_counter = meter.create_observable_counter(
    name="system.network.io",
    callbacks=[get_network_io],
    description="Network I/O bytes",
    unit="By"
)
```

4. Metric Attributes

Common Attributes

Attribute	Example	Use
<code>http.method</code>	<code>GET</code> , <code>POST</code>	HTTP method breakdown
<code>http.status_code</code>	<code>200</code> , <code>500</code>	Response status
<code>service.name</code>	<code>checkout-api</code>	Service identification
<code>environment</code>	<code>production</code>	Environment

Cardinality Considerations

High Cardinality (Avoid)	Low Cardinality (Good)
User ID	Region
Request ID	HTTP method
Timestamp	Status code class
Full URL	Route template

Attribute Best Practices

```
# Good - Low cardinality
request_counter.add(1, {
    "method": "GET",
    "route": "/api/users/{id}", # Template, not actual ID
    "status": "2xx" # Status class
})

# Bad - High cardinality
request_counter.add(1, {
    "user_id": "user-12345", # Millions of values
    "url": "/api/users/12345", # Every request different
    "timestamp": "2026-01-26T10:00:00" # Every second different
})
```

5. Aggregation and Views

Histogram Buckets

Configure explicit bucket boundaries:

```

from opentelemetry.sdk.metrics import MeterProvider
from opentelemetry.sdk.metrics.view import View,
ExplicitBucketHistogramAggregation

# Custom histogram buckets for response times
duration_view = View(
    instrument_name="http.server.request.duration",
    aggregation=ExplicitBucketHistogramAggregation(
        boundaries=[5, 10, 25, 50, 100, 250, 500, 1000, 2500, 5000, 10000]
    )
)

provider = MeterProvider(
    metric_readers=[reader],
    views=[duration_view]
)

```

Drop Metrics via Views

```

from opentelemetry.sdk.metrics.view import DropAggregation

# Drop noisy metrics
drop_view = View(
    instrument_name="internal.debug.*",
    aggregation=DropAggregation()
)

```

6. Prometheus Integration

Scraping Prometheus Metrics

Collector configuration:


```

receivers:
  prometheus:
    config:
      scrape_configs:
        - job_name: 'my-app'
          scrape_interval: 15s
          static_configs:
            - targets: ['app:8080']
          metrics_path: /metrics

service:
  pipelines:
    metrics:
      receivers: [prometheus, otlp]
      processors: [batch]
      exporters: [otlphttp]

```

Exposing OTel Metrics as Prometheus

```

from opentelemetry.exporter.prometheus import PrometheusMetricReader
from prometheus_client import start_http_server

# Start Prometheus endpoint
start_http_server(8080)

# Use Prometheus reader
reader = PrometheusMetricReader()
provider = MeterProvider(metric_readers=[reader])

```

```

// Query OTel metrics in Dynatrace
timeseries avg(http.server.request.duration), from:-1h, by:{http.method}
| limit 10

```

```

// Request rate by status
timeseries rate = sum(http.server.request.count), from:-1h, by:
{http.status_code}
| limit 10

```

7. Best Practices

Naming Conventions

Pattern	Example	Description
<code><domain>.<component>.<metric></code>	<code>http.server.request.count</code>	Hierarchical naming
Use lowercase	<code>request_count</code> not <code>RequestCount</code>	Consistency
Include units	<code>duration_ms</code> , <code>size_bytes</code>	Clarity

Performance Tips

Tip	Reason
Limit attribute cardinality	Memory, query performance
Use appropriate export interval	Balance freshness vs. load
Batch metrics	Reduce network calls
Use observable for system stats	Avoid polling in hot path

What to Measure

Category	Metrics
RED	Rate, Errors, Duration (for services)
USE	Utilization, Saturation, Errors (for resources)
Business	Revenue, conversions, user actions

Summary

In this notebook, you learned:

- Metric types: Counter, UpDownCounter, Histogram, Gauge
 - Sync vs async instruments
 - Creating and recording metrics
 - Attribute best practices and cardinality
 - Views and aggregation configuration
 - Prometheus integration patterns
-

References

- [OTel Metrics Specification](#)
 - [OTel Python Metrics](#)
 - [Dynatrace OTLP Metrics Ingest](#)
 - [Configure OTLP Metrics](#)
-

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.