

Logs and Events

Series: OTEL | **Notebook:** 6 of 8 | **Created:** January 2026 | **Last Updated:** 01/28/2026

Integrating Logs with OpenTelemetry

OpenTelemetry logs bridge traditional logging with distributed tracing, enabling correlation between log events and traces. This notebook covers log instrumentation, correlation, and best practices.

Table of Contents

- 1. [Log Bridge Pattern](#)
- 2. [Log-Trace Correlation](#)
- 3. [Collector Log Processing](#)
- 4. [File-Based Log Collection](#)
- 5. [Structured Logging](#)
- 6. [Best Practices](#)

Prerequisites

Requirement	Details
Knowledge	OTEL-01 and OTEL-04
Environment	Collector with logs pipeline

1. OTel Logs Overview

Log Data Model

Field	Description	Example
timestamp	When event occurred	2026-01-26T10:00:00Z
observed_timestamp	When log was collected	2026-01-26T10:00:01Z
severity_number	Numeric severity (1-24)	17 (ERROR)
severity_text	Severity string	ERROR
body	Log message	Failed to connect
attributes	Structured fields	{"user.id": "123"}
resource	Resource attributes	{"service.name": "api"}
trace_id	Associated trace	abc123...
span_id	Associated span	def456...

Severity Levels

Number	Level	Use Case
1-4	TRACE	Detailed debugging
5-8	DEBUG	Development info
9-12	INFO	Normal operations
13-16	WARN	Potential issues
17-20	ERROR	Errors occurred
21-24	FATAL	Critical failures

2. Log Bridge Pattern

OTel doesn't replace your logging library—it bridges logs to OTel format.

Python Logging Bridge

```
import logging
from opentelemetry import _logs
from opentelemetry.sdk._logs import LoggerProvider, LoggingHandler
from opentelemetry.sdk._logs.export import BatchLogRecordProcessor
from opentelemetry.exporter.otlp.proto.grpc._log_exporter import
OTLPLogExporter

# Setup OTel logs
logger_provider = LoggerProvider()
exporter = OTLPLogExporter(endpoint="http://collector:4317")
logger_provider.add_log_record_processor(BatchLogRecordProcessor(exporter))
_logs.set_logger_provider(logger_provider)

# Bridge to Python logging
handler = LoggingHandler(
    level=logging.INFO,
    logger_provider=logger_provider
)

# Attach to standard logger
logging.getLogger().addHandler(handler)
logging.getLogger().setLevel(logging.INFO)

# Use standard logging - logs go to OTel
logger = logging.getLogger(__name__)
logger.info("Application started")
logger.error("Connection failed", extra={"host": "db.example.com"})
```

Java Log4j Bridge

```
<!-- pom.xml -->
<dependency>
  <groupId>io.opentelemetry.instrumentation</groupId>
  <artifactId>opentelemetry-log4j-appender-2.17</artifactId>
  <version>1.26.0-alpha</version>
</dependency>
```

```
<!-- log4j2.xml -->
<Appenders>
    <OpenTelemetry name="OpenTelemetry" />
</Appenders>
<Loggers>
    <Root level="info">
        <AppenderRef ref="OpenTelemetry" />
    </Root>
</Loggers>
```

3. Log-Trace Correlation

Automatic Correlation

When using OTel tracing with the log bridge, trace context is automatically added:

```
from opentelemetry import trace
import logging

logger = logging.getLogger(__name__)
tracer = trace.get_tracer(__name__)

with tracer.start_as_current_span("process_order"):
    # Log automatically gets trace_id and span_id
    logger.info("Processing order", extra={"order_id": "12345"})
```

Manual Correlation

If not using OTel log bridge:

```

from opentelemetry import trace

def get_trace_context():
    span = trace.get_current_span()
    ctx = span.get_span_context()
    return {
        "trace_id": format(ctx.trace_id, '032x'),
        "span_id": format(ctx.span_id, '016x')
    }

# Add to log format
ctx = get_trace_context()
logger.info(f"[trace_id={ctx['trace_id']}] Processing order")

```

Correlation in Dynatrace

When logs include trace_id:

- View logs from trace details
- Jump to trace from log entry
- Filter logs by trace context

```

// Find logs with trace correlation
fetch logs, from:-1h
| filter isNotNull(trace_id)
| fields timestamp, trace_id, span_id, content, loglevel
| sort timestamp desc
| limit 20

```

```

// Find logs for a specific trace
fetch logs, from:-1h
| filter trace_id == "REPLACE_WITH_TRACE_ID"
| fields timestamp, span_id, content, loglevel
| sort timestamp asc

```

4. Collector Log Processing

OTLP Log Receiver

```
receivers:
  otlp:
    protocols:
      grpc:
        endpoint: 0.0.0.0:4317
      http:
        endpoint: 0.0.0.0:4318
```

Log Processing

```
processors:
  # Add resource attributes
  resource:
    attributes:
      - key: environment
        value: production
        action: upsert

  # Parse JSON logs
  transform:
    log_statements:
      - context: log
        statements:
          - merge_maps(attributes, ParseJSON(body), "insert")

  # Filter debug logs
  filter:
    logs:
      log_record:
        - 'severity_number < 9' # Drop TRACE and DEBUG

service:
  pipelines:
    logs:
      receivers: [otlp]
      processors: [resource, transform, filter, batch]
      exporters: [otlphttp]
```

5. File-Based Log Collection

Filelog Receiver

```
receivers:
  filelog:
    include:
      - /var/log/app/*.log
    exclude:
      - /var/log/app/debug.log
    start_at: beginning
    include_file_path: true
    include_file_name: true
    operators:
      # Parse JSON logs
      - type: json_parser
        if: 'body matches "^\\{'

      # Parse timestamp
      - type: time_parser
        parse_from: attributes.timestamp
        layout: '%Y-%m-%dT%H:%M:%S.%LZ'

      # Set severity
      - type: severity_parser
        parse_from: attributes.level
        mapping:
          error: error
          warn: warn
          info: info
          debug: debug
```

Container Logs

```
receivers:
  filelog:
    include:
      - /var/log/containers/*.log
    operators:
      # Parse container log format
      - type: regex_parser
        regex: '^(?P<time>[^\s]+) (?P<stream>stdout|stderr) (?P<logtag>[^\s]*) (?P<log>.*)$'

      # Extract k8s metadata from filename
      - type: regex_parser
        parse_from: attributes["log.file.name"]
        regex: '^(?P<pod_name>[^\s]+)_(?P<namespace>[^\s]+)_(?P<container>.+)\.log$'
```


6. Structured Logging

JSON Logging

```
import json
import logging

class JSONFormatter(logging.Formatter):
    def format(self, record):
        log_obj = {
            "timestamp": self.formatTime(record),
            "level": record.levelname,
            "message": record.getMessage(),
            "logger": record.name,
        }
        if hasattr(record, 'extra_fields'):
            log_obj.update(record.extra_fields)
        return json.dumps(log_obj)

# Use JSON formatter
handler = logging.StreamHandler()
handler.setFormatter(JSONFormatter())
logger.addHandler(handler)

# Log with extra fields
logger.info("Order processed", extra={"extra_fields": {
    "order_id": "12345",
    "customer_id": "c-789",
    "total": 99.99
}})
```

structlog

```
import structlog

structlog.configure(
    processors=[
        structlog.processors.add_log_level,
        structlog.processors.TimeStamper(fmt="iso"),
        structlog.processors.JSONRenderer()
    ]
)

logger = structlog.get_logger()
logger.info("order_processed", order_id="12345", total=99.99)
```

7. Best Practices

Log Content

Do	Don't
Use structured fields	Log PII or secrets
Include relevant context	Log at DEBUG in prod
Add trace correlation	Use inconsistent levels
Log actionable events	Log every detail

Performance

Tip	Why
Batch log exports	Reduce network calls
Filter at source	Reduce volume
Use async logging	Don't block requests
Set appropriate levels	Control verbosity

Log Levels Guide

Level	When to Use	Example
ERROR	Unexpected failures	DB connection failed
WARN	Recoverable issues	Retry succeeded
INFO	Business events	Order completed
DEBUG	Development details	Query parameters

Summary

In this notebook, you learned:

- OTel logs data model and severity levels
- Log bridge pattern for existing loggers
- Automatic and manual log-trace correlation
- Collector log processing and transformation
- File-based log collection
- Structured logging best practices

References

- [OTel Logs Specification](#)
- [Filelog Receiver](#)
- [Dynatrace Log Ingest](#)

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.