

S2D-04: Alert Migration - Davis

Anomaly Detectors

Series: S2D | Notebook: 4 of 9 | Created: January 2026 | Last Updated: 01/30/2026

Overview

This notebook explains how to translate Splunk alert conditions into Dynatrace Davis Anomaly Detectors. Understanding the fundamental differences between scheduled queries and continuous monitoring is essential for successful alert migration.




Table of Contents

- [1. Splunk vs Dynatrace Alerting](#)
- [2. The Sliding Window Concept](#)

3. [Threshold Translation Formula](#)
 4. [Migration Strategies](#)
 5. [Concept Mapping](#)
 6. [Creating a Davis Anomaly Detector](#)
 7. [Migration Decision Tree](#)
 8. [Validation Query](#)
-

Prerequisites

Requirement	Details
Dynatrace Environment	SaaS with Grail enabled
Permissions	<code>logs.read</code> , <code>settings.write</code>
Knowledge	Understanding of source Splunk alert logic

Learning Objectives

By the end of this notebook, you will be able to:

1. Explain the difference between Splunk scheduled alerts and Davis Anomaly Detectors
2. Apply the threshold translation formula
3. Choose the appropriate migration strategy for different alert types
4. Configure Davis Anomaly Detector parameters correctly

Splunk vs Dynatrace Alerting

Splunk Alerting Model

In Splunk, alerts are **scheduled queries** that:

- Run at defined intervals (e.g., every 5 minutes)
- Evaluate a query over a timeframe (e.g., last 15 minutes)

- Trigger when results exceed a threshold (e.g., count > 10)
- Apply suppression to prevent repeated alerts

Dynatrace Alerting Model

Davis Anomaly Detectors provide **continuous monitoring** that:

- Evaluate every minute (1-minute granularity)
- Use a sliding window approach
- Require multiple violating samples to trigger
- Require dealerting samples to close

Key Parameters

Parameter	Description	Maximum
Threshold	Value triggering a violation	Unlimited
Sliding Window	Time range evaluated	60 minutes
Violating Samples	Minutes above threshold to alert	60
Dealerting Samples	Minutes below threshold to close	60

The Sliding Window Concept

Davis Anomaly Detectors continuously evaluate metrics using a sliding window.

Example Configuration:

- Metric: Error Log Count
- Threshold: 5 errors/minute
- Sliding Window: 15 minutes
- Violating Samples: 3
- Dealerting Samples: 5

Behavior:

1. Every minute, the detector checks if error count > 5
2. If 3 of the last 15 minutes exceed the threshold → Alert opens
3. Alert stays open until 5 consecutive minutes are below threshold

Threshold Translation Formula

The key insight for translation is that both platforms should require roughly the same number of events to trigger:

$$[\text{Splunk Threshold}] = [\text{Dynatrace Threshold}] \times [\text{Dynatrace Violating Samples}]$$

Example Translation

Splunk Alert:

- Runs every 5 minutes
- Timeframe: last 15 minutes
- Threshold: 60 error logs
- Suppress: 15 minutes

Dynatrace Options (all satisfy 60 total errors):

Option	DT Threshold	Violating Samples	Product
A	15	4	60
B	10	6	60
C	30	2	60
D	20	3	60

Migration Strategies

Option 1: Alert Reimagined (Recommended)

Strategy: Calculate new threshold and violating samples to most closely replicate Splunk behavior.

Parameters:

- Threshold: 3
- Sliding Window: 15
- Violating Samples: 2 (triggers at 6 total errors spread across 2 minutes)
- Dealerting Samples: 15

Pros:

- Best balance between sensitivity and stability
- Closest to expected alert behavior

Cons:

- Requires threshold recalculation
- Mindset shift for teams used to Splunk

```
// Option 1: Alert Reimagined – DQL for Davis Anomaly Detector
// Threshold: 3, Violating Samples: 2, Sliding Window: 15
fetch logs, from:-24h
| filter loglevel == "ERROR"
| filter matchesPhrase(k8s.cluster.name, "production")
| filter matchesPhrase(k8s.deployment.name, "checkout-service")
| makeTimeseries count = count(), by:{dt.entity.cloud_application},
interval:1m
```

Option 2: Conservative Threshold

Strategy: Use a 1-minute threshold equal to Splunk threshold, with high violating samples.

Parameters:

- Threshold: 6 (same as Splunk total)
- Sliding Window: 15
- Violating Samples: 1
- Dealerting Samples: 15

Pros:

- Easy to understand and explain
- Preserves original threshold number

Cons:

- Very sensitive to brief spikes
- May alert more frequently than Splunk

```
// Option 2: Conservative - Same threshold, minimum samples
// Threshold: 6, Violating Samples: 1, Sliding Window: 15
fetch logs, from:-24h
| filter loglevel == "ERROR"
| filter matchesPhrase(k8s.cluster.name, "production")
| filter matchesPhrase(k8s.deployment.name, "checkout-service")
| makeTimeseries count = count(), by:{dt.entity.cloud_application},
interval:1m
```

Option 3: Aggressive Threshold

Strategy: Divide Splunk threshold by sliding window for per-minute threshold.

Parameters:

- Threshold: 0.4 ($6 \div 15 = 0.4$ per minute)
- Sliding Window: 15
- Violating Samples: 15 (all samples must violate)
- Alerting Samples: 15

Pros:

- Very stable, requires sustained issues
- Less sensitive to spikes

Cons:

- May miss short-duration issues
- Delayed alerting compared to Splunk

Option 4: Data-Driven

Strategy: Analyze actual data patterns to determine optimal thresholds.

Steps:

1. Query historical data to understand baseline patterns
2. Identify normal vs. problematic periods
3. Set thresholds that would have caught past incidents
4. Tune based on observed alert frequency

Pros:

- Most accurate for the specific environment
- Data-backed decisions

Cons:

- Requires more analysis effort

- Needs historical data availability

```
// Data-Driven Analysis: View error patterns over time
fetch logs, from:now()-7d
| filter loglevel == "ERROR"
| filter matchesPhrase(k8s.deployment.name, "checkout-service")
| makeTimeseries count = count(), interval:1h
| fieldsAdd
    daily_avg = arrayAvg(count),
    daily_max = arrayMax(count),
    daily_min = arrayMin(count)
```

Concept Mapping

Splunk Concept	Dynatrace Equivalent	Notes
Query Timeframe	Sliding Window	Both define evaluation period
Suppress Duration	Dealerting Samples	Time before alert can re-trigger
Threshold	Threshold × Violating Samples	Must calculate product
Schedule Frequency	Always 1 minute	Cannot be changed

Creating a Davis Anomaly Detector

Step 1: Create the DQL Query

Your query must return timeseries data with a 1-minute interval:

```
// Template for Davis Anomaly Detector DQL
fetch logs, from:-24h
| filter loglevel == "ERROR"
| filter matchesPhrase(k8s.cluster.name, "your-cluster")
| filter matchesPhrase(k8s.namespace.name, "your-namespace")
| makeTimeseries count = count(), by:{dt.entity.cloud_application},
interval:1m
```

Step 2: Configure Detector Settings

Setting	Value	Rationale
Name	[AppName] High Error Count	Clear identification
Threshold	Calculated per formula	See translation formula
Sliding Window	Match Splunk timeframe	Usually 15-60 minutes
Violating Samples	Depends on strategy	See migration options
Dealerting Samples	Match suppress duration	Prevents alert flapping

Step 3: Configure Event Properties

Property	Example	Purpose
Event Name	[AppName] P2 - {host.name} - High Errors	Alert identification
Event Type	ERROR_EVENT	Severity classification
Description	Dynamic with affected entity	Context for responders

Migration Decision Tree

Use this decision tree to choose your migration approach:

1. Is the Splunk alert timeframe > 1 hour?

- Yes → Use Workflows (see S2D-05) or ArrayMovingSum (see S2D-06)
- No → Continue

2. Is the alert scheduled infrequently (< 6 times/day)?

- Yes → Consider Workflows (see S2D-05)
- No → Continue

3. Do you have historical data for analysis?

- Yes → Use Option 4 (Data-Driven)
- No → Continue

4. Is minimizing false positives critical?

- Yes → Use Option 1 (Alert Reimagined) or Option 3 (Aggressive)
- No → Use Option 2 (Conservative)

Validation Query

Before creating a Davis Anomaly Detector, validate your query returns expected results:

```
// Validate alert query returns timeseries data
fetch logs, from:now()-24h
| filter loglevel == "ERROR"
| filter matchesPhrase(k8s.deployment.name, "checkout-service")
| makeTimeseries count = count(), by:{dt.entity.cloud_application},
interval:1m
| limit 100
```

Next Steps

- **S2D-05: Alert Migration - Workflows** - For alerts that don't fit the Davis model
- **S2D-06: ArrayMovingSum** - For timeframes exceeding 1 hour

References

- [Davis Anomaly Detectors](#)
 - [Custom Events for Alerting](#)
 - [Alerting Profiles](#)
-

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.