

S2S-08: Step 8 — Enable: Parallel Operation and Stakeholder Handover

Series: S2S | **Notebook:** 8 of 9 | **Phase:** Run | **Step:** Enable | **Created:** March 2026 | **Last Updated:** 04/04/2026

Overview

With data pipelines, SLOs, and alerting configured in the target tenant (Step 7), this step focuses on the human side of migration: managing parallel operation, communicating changes to stakeholders, establishing Dynatrace Intelligence baselines, and preparing users for the cutover. Parallel operation is the most expensive phase of the migration — you are paying for two tenants — so the goal is to keep it as short as possible while ensuring confidence in the target tenant.

S2S Migration Journey — 3 Phases / 9 Steps

Plan: 1. Discover | 2. Strategize | 3. Design

Upgrade: 4. Prepare | 5. Execute | 6. Integrate

Run: 7. Expand | **8. Enable** | 9. Optimize

Table of Contents

- [1. Parallel Operation Strategy](#)
- [2. Dynatrace Intelligence Baseline Establishment](#)
- [3. SLO Continuity](#)
- [4. Stakeholder Communication](#)
- [5. Cost Management During Parallel Period](#)

6. User Training and Documentation

7. Step Completion Checklist

Prerequisites

Requirement	Details
Step 7 Complete	OpenPipeline rules, SLOs, and alerting deployed in target tenant
Agents Reporting	OneAgents and DynaKube operators sending data to target tenant
Both Tenants Active	Source and target tenants both receiving data for validation
Stakeholder List	Identified audiences for migration communication
Training Environment	Target tenant accessible for user training and exploration

1. Parallel Operation Strategy

During parallel operation, both the source and target tenants receive monitoring data. This period exists to validate the target tenant, build Dynatrace Intelligence baselines, and give users confidence before cutover.

Parallel Operation Timeline

Week	Activity	Key Milestones
Week 1	Data validation	Confirm hosts, services, logs, and spans in target match source counts
Week 2	Dynatrace Intelligence learning	Baselines begin forming; initial anomaly detection active
Week 3	SLO evaluation	7-day rolling SLOs have full data; compare source vs target values

Week	Activity	Key Milestones
Week 4	User acceptance	Key users validate dashboards, workflows, and alerts in target
Weeks 5–6	Cutover preparation	Go/no-go decision, communication, maintenance window scheduled
Weeks 7–8	Post-cutover buffer	Source tenant in read-only for historical queries
Week 9+	Decommission	Source tenant deactivated and deprovisioned

Operation Models

Model	Description	Best For	Duration
Full Parallel	Both tenants receive all data from all agents	High-risk environments, regulatory requirements	4–6 weeks
Phased	Migrate groups of hosts/services incrementally	Large environments (1000+ hosts)	6–10 weeks
Reference	Source receives data, target receives from a representative subset only	Cost-constrained, low-risk migrations	2–4 weeks

Cost impact: Full parallel doubles your DPS consumption. Work with your Dynatrace account team to negotiate temporary parallel licensing.

Validate Host Count Parity

Run this query in both source and target tenants. The host counts should match (allowing for normal churn):

```
// Host count validation – run in both source and target tenants
fetch dt.entity.host
| summarize total_hosts = count()
| append [
    fetch dt.entity.host
    | fieldsAdd provider = if(isNotNull(awsNameTag), then: "AWS",
        else: if(isNotNull(azureResourceGroupName), then: "Azure", else: "On-
Premises"))
    | summarize count = count(), by:{provider}
]

// Alternative: Smartscape on Grail (entity.name → name)
// smartscapeNodes HOST
// | summarize total_hosts = count()
// | append [
// smartscapeNodes HOST
// | fieldsAdd provider = if(isNotNull(awsNameTag), then: "AWS",
// else: if(isNotNull(azureResourceGroupName), then: "Azure", else: "On-
Premises"))
// | summarize count = count(), by:{provider}
// ]
```

Validate Log Continuity

Compare log volumes between source and target to ensure data pipelines are operating at parity:

```
// Log volume over the last 24 hours by loglevel
// Run in both source and target – compare totals
fetch logs, from:-24h
| summarize record_count = count(), by:{loglevel}
| sort record_count desc
```

2. Dynatrace Intelligence Baseline Establishment

Dynatrace Intelligence requires historical data to establish baselines for anomaly detection. In the target tenant, Dynatrace Intelligence starts from zero — there is no way to transfer learned baselines.

Baseline Types and Learning Times

Baseline Type	Learning Period	What It Detects
Response time	1–2 weeks	Service slowdowns
Error rate	1–2 weeks	Error rate increases
Throughput	2–4 weeks	Traffic anomalies (requires weekly patterns)
Infrastructure	1 week	CPU, memory, disk anomalies
Custom events	2–4 weeks	Depends on event frequency

Accelerating Baseline Learning

You cannot directly speed up Dynatrace Intelligence learning, but you can reduce noise:

Action	Impact
Suppress maintenance alerts during migration	Prevents false positives from polluting baselines
Ensure consistent traffic patterns	Avoid load testing or unusual deployments during baseline period
Configure sensitivity thresholds	Set anomaly detection sensitivity to match source tenant settings
Migrate anomaly detection settings first	Ensures Dynatrace Intelligence uses the same thresholds from day one

Expected Behavior During Baseline Period

Week	Dynatrace Intelligence Behavior	Action Required
Week 1	Many false positives (no baseline context)	Suppress or route to staging alert channel
Week 2	Infrastructure baselines forming, fewer false positives	Monitor and triage manually

Week	Dynatrace Intelligence Behavior	Action Required
Week 3	Service baselines forming, anomaly detection improving	Compare detected problems with source tenant
Week 4+	Baselines stable for most metrics	Ready for production alerting

Monitor Detected Problem Trends

Track detected problem volume in the target tenant to see baselines stabilize over time:

```
// detected problem trend over the last 7 days (target tenant)
// Expect high volume in week 1, decreasing as baselines stabilize
fetch dt.davis.problems, from:-7d
| summarize problem_count = count(), by:{day = bin(timestamp, 1d)}
| sort day asc
```

```
// detected problems by category – identify which areas are generating noise
fetch dt.davis.problems, from:-7d
| filter event.status == "ACTIVE" or event.status == "CLOSED"
| summarize problem_count = count(), by:{event.category}
| sort problem_count desc
```

3. SLO Continuity

SLOs with rolling evaluation windows will show incomplete or misleading values during the transition period. This is expected and must be communicated to stakeholders.

Impact During Transition

SLO Window	Impact	Duration of Impact
7-day rolling	Partial data for first 7 days	1 week
30-day rolling	Severely incomplete for first 30 days	4 weeks
Calendar month	No data until month boundary	Up to 30 days

Mitigation Strategies

Strategy	Implementation	When to Use
Temporary short window	Change 30-day SLOs to 7-day during transition	When stakeholders need early visibility
Dual reporting	Report SLOs from both tenants until target window fills	When SLO values are reported externally
SLO holiday	Suspend SLO reporting during transition with documented justification	When stakeholders accept a gap
Timestamp shift	Start SLO tracking at beginning of evaluation window	For calendar-based SLOs

4. Stakeholder Communication

Migration success depends as much on communication as on technical execution. Different audiences need different messages.

Communication Plan by Audience

Audience	What They Need to Know	When	Channel
Executive leadership	Timeline, risk summary, cost impact, go/no-go decision	Week 1 + weekly status	Email + steering committee
SRE / Platform team	Technical details, runbooks, rollback procedures	Throughout	Slack + wiki
Application developers	New tenant URL, updated API tokens, dashboard locations	2 weeks before cutover	Email + team meeting
Service desk	Updated escalation paths, known issues during transition	1 week before cutover	Runbook update

Audience	What They Need to Know	When	Channel
External stakeholders	Expected alert changes, SLO reporting gaps	2 weeks before cutover	Email + SLA documentation

Expected Impact Summary

Include this table in your stakeholder communication:

Area	Impact	Duration	Mitigation
Dynatrace Intelligence	Increased false positives in target tenant	2–4 weeks	Alerts routed to staging channel
SLOs	Incomplete evaluation in target tenant	1–4 weeks (depends on window)	Dual reporting from both tenants
Historical data	Not available in target tenant	Permanent	Source tenant available as read-only during buffer period
Dashboards	New URLs, possible layout differences	One-time	User training and URL redirect documentation
API tokens	All tokens regenerated	One-time	Token distribution before cutover

5. Cost Management During Parallel Period

Parallel operation is expensive. Both tenants consume DPS for every monitored host, log record, metric, and span.

Dual Licensing Impact

Cost Component	Impact During Parallel	Optimization
Host units	2x (agents report to both)	Phased migration reduces overlap

Cost Component	Impact During Parallel	Optimization
Log ingest	2x (same logs to both tenants)	Route verbose logs to source only
Span ingest	2x	Reduce span sampling ratio in source
Metric ingest	2x for custom metrics	Minimal for built-in metrics
DEM units	2x for RUM/synthetic	Disable synthetic in source once target is validated

Cost Optimization Strategies

Strategy	Savings	Risk
Shorten parallel period (2 weeks vs 4 weeks)	50% reduction	Less baseline time for Dynatrace Intelligence
Phased migration (migrate by group)	Only migrated hosts are dual-cost	Longer total timeline
Reduce source tenant data collection	20–40% on source during parallel	Reduced source visibility
Negotiate temporary parallel licensing	Cost-neutral	Requires account team engagement

Recommendation: Negotiate parallel licensing with your Dynatrace account team **before** starting the parallel phase. Most account teams can provide temporary licensing for migration periods.

6. User Training and Documentation

Users need to know what changes, what stays the same, and where to find things in the target tenant.

Persona-Based Training

Persona	Training Focus	Format	Duration
Platform admin	IAM, OpenPipeline, bucket management, API token generation	Hands-on workshop	2 hours
SRE	Dashboards, SLOs, alerting, Dynatrace Intelligence differences	Live demo + Q&A	1 hour
Developer	New tenant URL, notebook access, DQL querying	Self-service guide + office hours	30 min
Service desk	Escalation paths, known issues, FAQ	Runbook update	30 min

Documentation Updates Required

Document	Update Required
Runbooks	New tenant URLs, API endpoints, token references
Wiki / knowledge base	Dashboard links, notebook links, SLO references
CI/CD pipelines	API token environment variables, deployment scripts
Incident response playbooks	Alert routing, escalation contacts
Onboarding documentation	New user access procedures, IAM group requests

FAQ Template

Question	Answer
Why are my dashboards showing less data?	The target tenant does not have historical data. Data accumulates from the migration date forward.
Why am I getting more alerts than usual?	Dynatrace Intelligence is establishing baselines. False positives are expected for 2–4 weeks.
Where is my old dashboard?	Dashboards have been migrated. Find them at <code><target-tenant-url>/ui/dashboards</code> .

Question	Answer
Do I need a new API token?	Yes. All API tokens must be regenerated in the target tenant.
When will SLOs be accurate?	SLOs need one full evaluation window of data. 7-day SLOs: 1 week. 30-day SLOs: 4 weeks.

7. Step Completion Checklist

Before proceeding to **Step 9 — Optimize**, confirm that you have completed each item:

Checkpoint	Status
Parallel operation model selected (full, phased, or reference)	[]
Host count parity validated between source and target	[]
Log and span volumes validated between source and target	[]
Dynatrace Intelligence problem trend monitored — false positives decreasing	[]
SLO continuity strategy documented (temporary window, dual reporting, or holiday)	[]
Stakeholder communication sent to all audiences	[]
Cost optimization negotiated with account team	[]
User training completed for all personas	[]
Documentation updated (runbooks, wiki, CI/CD, playbooks)	[]
Go/no-go criteria defined for cutover decision	[]

Next Step

S2S-09: Step 9 — Optimize — Execute the final cutover, validate migration completeness, decommission the source tenant, and capture lessons learned.

Completed	Next	Remaining
~~1. Discover~~ → ~~2. Strategize~~ → ~~3. Design~~ → ~~4. Prepare~~ → ~~5. Execute~~ → ~~6. Integrate~~ → ~~7. Expand~~ → ~~8. Enable~~	9. Optimize	—

Summary

In Step 8, you:

- Established a parallel operation strategy with a clear timeline (weeks 1–9+)
- Monitored Dynatrace Intelligence baseline establishment and tracked false positive trends
- Documented SLO continuity strategy for the transition period
- Communicated migration impact to all stakeholder audiences
- Optimized cost during the dual-tenant parallel period
- Trained users by persona and updated all supporting documentation

Additional Resources

- [Dynatrace Intelligence Anomaly Detection](#)
- [SLO Configuration](#)
- [Dynatrace Notifications and Alerting](#)
- [DPS Licensing](#)

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.