

S2S-09: Data Continuity and Parallel Operation

Series: S2S | **Notebook:** 9 of 12 | **Created:** March 2026 | **Last Updated:** 03/23/2026

Overview

Historical data cannot be migrated between SaaS tenants. The only way to maintain continuity is to run both tenants in parallel during the transition. This notebook covers parallel operation strategy, SLO baseline establishment, and stakeholder communication.

Table of Contents

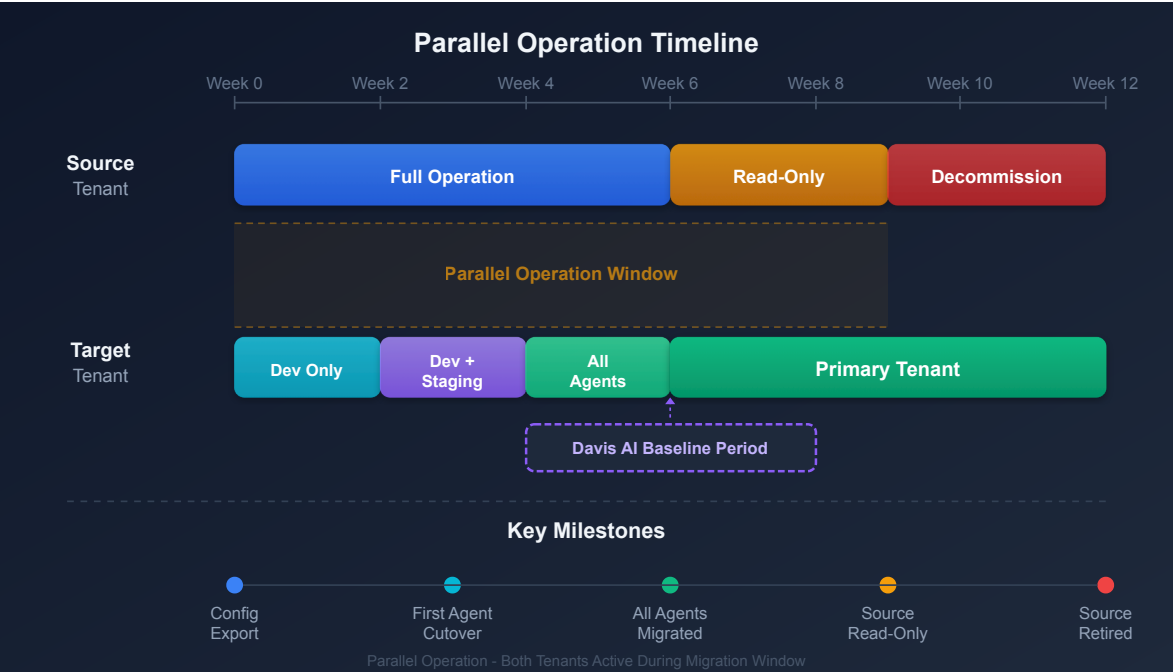
- [1. Parallel Operation Strategy](#)
 - [2. Davis AI Baseline Establishment](#)
 - [3. SLO Continuity](#)
 - [4. Stakeholder Communication](#)
 - [5. Cost Management During Parallel Period](#)
-

Prerequisites

Requirement	Details
Both Tenants	Source and target both active and licensed
Licensing	Budget for parallel operation (typically 30-90 days)
Monitoring	At least some agents reporting to target tenant

1. Parallel Operation Strategy

Timeline



Parallel Operation Models

Model	Description	Duration	Cost
Full Parallel	All agents report to both tenants	4-8 weeks	2x host units
Phased Parallel	Agents migrate in waves (dev → staging → prod)	6-12 weeks	1.1-1.5x host units
Reference Parallel	Source is read-only, target is primary	2-4 weeks	1x + read-only cost

Recommendation: Phased parallel is the best balance of cost and risk. You never have all agents on both tenants simultaneously.

2. Davis AI Baseline Establishment

Davis AI needs time to learn normal behavior patterns in the target tenant:

Baseline Type	Time to Establish	Factors
Availability	2-3 days	Binary metric, fast learning
Response time	1-2 weeks	Requires weekday/weekend patterns
Error rate	1-2 weeks	Needs regular traffic patterns
Resource utilization	2-4 weeks	Needs full business cycle

Accelerating Baseline Learning

- Migrate dev/staging first to give Davis AI a head start
- Ensure representative traffic reaches services in target
- Avoid migration during atypical periods (holidays, sales events)

3. SLO Continuity

SLOs in the target tenant start with no historical evaluation:

SLO Type	Impact	Mitigation
Rolling week	7 days of no data	SLO appears at 0% until data accumulates
Rolling month	30 days of no data	Create with reduced window initially
Calendar-based	Varies	Align cutover with calendar boundary

Strategy

1. Create SLOs in target with same definitions
2. Set initial evaluation window to match available data (e.g., `ROLLING_3_DAYS`)
3. Expand to full window after sufficient data (e.g., `ROLLING_WEEK`)
4. Document the baseline gap for stakeholders

4. Stakeholder Communication

Communication Plan

Audience	When	Message
Engineering teams	4 weeks before	"Migration timeline, agent cutover schedule, expected impacts"
SRE/On-call	2 weeks before	"Alert routing changes, dual-tenant monitoring during transition"
Management	1 week before	"SLO baseline gap, Davis AI relearning period, risk mitigation"
All users	Day of cutover	"New tenant URL, SSO changes, bookmark updates"

Expected Impact Summary

Item	Impact	Duration
Davis AI alerts	May fire false positives during baseline learning	2-4 weeks
SLO reporting	Gap in historical evaluation	7-30 days
Dashboards	Some tiles may show "No data" initially	Hours to days
Problem history	Not available in target	Permanent (document key problems)

5. Cost Management During Parallel Period

Strategy	Savings	Trade-off
Migrate dev first, decommission source dev	Early savings	Less parallel validation
Negotiate parallel license with Dynatrace	Significant	Requires contract discussion
Use DPS licensing model	Pay-per-use	Cost scales with actual usage
Reduce source to monitoring-only	Reduced DPS consumption	No alerting from source

Next Steps

Continue to **S2S-10: SLO and Alerting Migration** to configure SLOs and notification rules in the target tenant.

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.