

# S2S-99: Best Practice Summary

Series: S2S | Notebook: 99 | Created: March 2026 | Last Updated: 04/16/2026

## Overview

This notebook consolidates **112 actionable best practices** from the S2S series into a single reference organized around the 9-step migration framework. Each practice is definitive: it tells you exactly what to set, when, and why. Use this as a pre-flight checklist and ongoing reference throughout your SaaS-to-SaaS migration.

The S2S series follows a **9-step framework** (Discover through Optimize) supported by an **11-step Order of Operations** that governs execution sequence.

## The 9-Step Migration Framework

| Step | Name       | Focus                                       |
|------|------------|---------------------------------------------|
| 1    | Discover   | Migration scenarios and inventory           |
| 2    | Strategize | Define your migration approach              |
| 3    | Design     | Target tenant architecture                  |
| 4    | Prepare    | Export and pre-stage                        |
| 5    | Execute    | Configuration import and agent cutover      |
| 6    | Integrate  | Cloud, dashboards, and workflows            |
| 7    | Expand     | OpenPipeline, SLOs, and alerting            |
| 8    | Enable     | Parallel operation and stakeholder handover |
| 9    | Optimize   | Cutover validation and decommission         |

## The 11-Step Order of Operations

| # | Operation | Description             |
|---|-----------|-------------------------|
| 1 | Assess    | Inventory source tenant |

| #  | Operation           | Description                             |
|----|---------------------|-----------------------------------------|
| 2  | <b>Provision</b>    | Target SaaS tenant and access (SSO/IAM) |
| 3  | <b>Export</b>       | Configuration from source tenant        |
| 4  | <b>Migrate</b>      | Configuration to target tenant          |
| 5  | <b>Rebuild</b>      | Dashboards, SLOs, alerts                |
| 6  | <b>Redirect</b>     | OneAgents and operators to target       |
| 7  | <b>Reconnect</b>    | Cloud integrations and extensions       |
| 8  | <b>Migrate</b>      | Any remaining configuration             |
| 9  | <b>Validate</b>     | Data flow and performance               |
| 10 | <b>Cutover</b>      | Full switch to target tenant            |
| 11 | <b>Decommission</b> | Source tenant                           |

**OneAgent Attribute Enrichment (1.331+):** OneAgent can enrich all telemetry (metrics, spans, logs, events) with primary fields ( `dt.security_context` , `dt.cost.costcenter` ) and primary tags ( `primary_tags.environment` , `primary_tags.team` ) at the source. More efficient than auto-tags — feeds directly into OpenPipeline routing, bucket assignment, and Grail permissions. Configure via `oneagentctl --set-host-tag` or `--set-host-tag` at install time. See [docs](#).

## Table of Contents

1. [Step 1: Discover — Migration Scenarios and Inventory](#)
2. [Step 2: Strategize — Define Your Migration Approach](#)
3. [Step 3: Design — Target Tenant Architecture](#)
4. [Step 4: Prepare — Export and Pre-Stage](#)
5. [Step 5: Execute — Configuration Import and Agent Cutover](#)
6. [Step 6: Integrate — Cloud, Dashboards, and Workflows](#)
7. [Step 7: Expand — OpenPipeline, SLOs, and Alerting](#)
8. [Step 8: Enable — Parallel Operation and Stakeholder Handover](#)
9. [Step 9: Optimize — Cutover Validation and Decommission](#)
10. [The Five Principles of S2S Migration](#)

## Prerequisites

| Requirement          | Details                                                |
|----------------------|--------------------------------------------------------|
| <b>S2S Series</b>    | Read S2S-01 through S2S-09 for full context            |
| <b>Source Tenant</b> | Active Dynatrace SaaS with admin access                |
| <b>Target Tenant</b> | Provisioned Dynatrace SaaS with admin access           |
| <b>CLI Tools</b>     | Monaco CLI v2.x, Terraform v1.5+ with provider ~> 1.91 |

## 1. Step 1: Discover — Migration Scenarios and Inventory

Source: S2S-01

| # | Best Practice                                      | Recommended Setting/Value                                                                                                                                                                          | Priority        |
|---|----------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------|
| 1 | Complete entity discovery before anything else     | Run <code>fetch dt.entity.host \   summarize count()</code> for hosts, services, process groups, and applications to establish baseline numbers                                                    | <b>Critical</b> |
| 2 | Identify your migration scenario                   | Classify as consolidation (many-to-one), split (one-to-many), regional relocation, or cloud provider change — each has distinct tooling and planning requirements                                  | <b>Critical</b> |
| 3 | Select Monaco for configuration, Terraform for IAM | Monaco v2 covers all 8 config types (settings, document, automation, bucket, segment, slo-v2, openpipeline, classic api); Terraform is required exclusively for IAM policies, groups, and bindings | <b>Critical</b> |
| 4 | Plan for entity ID changes                         | Entity IDs (HOST-xxx, SERVICE-xxx) are tenant-specific and will change; identify every dashboard, SLO, and alert that hardcodes an entity ID                                                       | <b>Critical</b> |
| 5 | Document the 90/10 manual items                    | 90% of config migrates automatically; budget 90% of your effort for the remaining 10% — entity ID remapping, integration repointing, IAM redesign, credential recreation                           | <b>Critical</b> |
| 6 | Build a comprehensive endpoint inventory           | Document every agent URL, API call, webhook, and integration endpoint that references the source tenant URL                                                                                        | <b>Critical</b> |

| #  | Best Practice                                 | Recommended Setting/Value                                                                                                                                         | Priority           |
|----|-----------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------|
| 7  | Triage detected problems before migration     | Active problems (especially frequent/duplicate events) carry noise to the target tenant; suppress or tune anomaly detection for sources with >500 active problems | <b>Critical</b>    |
| 8  | Identify configuration debt to leave behind   | Catalog stale maintenance windows, disabled notification rules, inactive synthetic monitors, unused management zones — do not migrate these                       | <b>Critical</b>    |
| 9  | Inventory Extensions 2.0 separately           | Extensions cannot be exported via Monaco or Terraform; list all extensions and plan for manual reinstallation from the Dynatrace Hub                              | <b>Recommended</b> |
| 10 | Audit configuration changes from last 30 days | Query audit logs to identify recently changed settings that may not be in your last export                                                                        | <b>Recommended</b> |

## 2. Step 2: Strategize — Define Your Migration Approach

Source: S2S-02

| #  | Best Practice                                           | Recommended Setting/Value                                                                                                                        | Priority        |
|----|---------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------|-----------------|
| 11 | Choose phased approach for environments with >500 hosts | Migrate dev first (weeks 1-2), staging (weeks 3-4), production (weeks 5-6); gives Dynatrace Intelligence time to build baselines                 | <b>Critical</b> |
| 12 | Define success criteria before starting                 | Document measurable targets: host count parity, service count parity, log volume parity, SLO evaluation accuracy, alerting coverage              | <b>Critical</b> |
| 13 | Engage your account team for licensing early            | Contact Dynatrace to discuss temporary parallel licensing; negotiate before the migration window opens                                           | <b>Critical</b> |
| 14 | Plan a 2-4 week parallel operation period               | Historical data does not migrate; parallel is the only path to data continuity and Dynatrace Intelligence baseline stability                     | <b>Critical</b> |
| 15 | Establish a configuration freeze window                 | No settings modifications in the source tenant during cutover; communicate freeze dates to all teams                                             | <b>Critical</b> |
| 16 | Migrate multi-source consolidations sequentially        | Complete Source 1 → validate → Source 2 → validate; never migrate two sources in parallel; migrate the simpler source first (non-K8s before K8s) | <b>Critical</b> |

| #  | Best Practice                                      | Recommended Setting/Value                                                                                            | Priority           |
|----|----------------------------------------------------|----------------------------------------------------------------------------------------------------------------------|--------------------|
| 17 | Align migration timing with stable traffic periods | Avoid holidays, sales events, or other atypical traffic patterns that would distort Dynatrace Intelligence baselines | <b>Recommended</b> |
| 18 | Expect the 90/10 rule                              | 90% of config migrates automatically; budget 90% of effort for the remaining 10%                                     | <b>Critical</b>    |
| 19 | Prepare a rollback plan                            | Document how to revert agents to the source tenant if issues arise during cutover                                    | <b>Recommended</b> |

### 3. Step 3: Design — Target Tenant Architecture

Source: S2S-03

| #  | Best Practice                                    | Recommended Setting/Value                                                                                                                                                             | Priority           |
|----|--------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------|
| 20 | Redesign IAM during consolidation                | Use the migration as the opportunity to fix group sprawl and overprivileged policies; redesign from scratch for the target tenant                                                     | <b>Critical</b>    |
| 21 | Deploy data foundation before governance         | Buckets and enrichment rules must exist before OpenPipeline routing, segments, SLOs, and IAM policies that reference them                                                             | <b>Critical</b>    |
| 22 | Replace entity IDs with tag-based filters        | Convert <code>entityId("HOST-xxxx")</code> to <code>tag(app:checkout), tag(env:production)</code> in dashboards, SLOs, and notification rules before migration                        | <b>Critical</b>    |
| 23 | Design bucket naming for consolidation           | Use prefix naming (e.g., <code>us-east_app_logs</code> , <code>eu-west_app_logs</code> ) to prevent collisions when consolidating multiple tenants                                    | <b>Recommended</b> |
| 24 | Plan the 24-item deployment order                | Buckets (1st) → Enrichment (2nd) → OpenPipeline (3rd) → Segments (4th) → Gen2 config (5th-13th) → Alerting/SLOs (14th-18th) → Dashboards/Workflows (19th-21st) → IAM last (22nd-24th) | <b>Critical</b>    |
| 25 | Replace management zone ID references with names | Convert <code>mzId(123)</code> to <code>mzName("Production")</code> because zone IDs change between tenants                                                                           | <b>Critical</b>    |

| #  | Best Practice                                                  | Recommended Setting/Value                                                                                                       | Priority    |
|----|----------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------|-------------|
| 26 | Prefix resources with source identifier during consolidation   | Set <code>name = "\${source_prefix} - \${resource_name}"</code> on all dashboards, rules, and SLOs to prevent naming collisions | Recommended |
| 27 | Tag everything by business unit before splitting a tenant      | Apply auto-tags identifying ownership before splitting one tenant into many                                                     | Recommended |
| 28 | Map cloud provider identity constructs when changing providers | AWS IAM roles to Azure Service Principals, Azure AD Groups to GCP Google Groups, etc.                                           | Recommended |

## 4. Step 4: Prepare — Export and Pre-Stage

Source: S2S-04, S2S-10

| #  | Best Practice                                                             | Recommended Setting/Value                                                                                                                 | Priority    |
|----|---------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------|-------------|
| 29 | Export via <code>monaco download</code>                                   | Run <code>monaco download manifest.yaml --environment source-tenant</code> to export all configuration in one operation                   | Critical    |
| 30 | Use the S2S-10 migration scripts for SUA packaging                        | Automates Monaco download + SUA-compatible <code>.tar.gz</code> packaging; Bash and PowerShell versions provided                          | Recommended |
| 31 | Package exports as <code>.tar.gz</code> for SUA — never <code>.zip</code> | The SaaS Upgrade Assistant rejects <code>.zip</code> archives; must be <code>.tar.gz</code> with <code>exportMetadata.json</code> at root | Critical    |
| 32 | Store exported configuration in Git                                       | Commit the Monaco export directory to a Git repository immediately after download                                                         | Critical    |
| 33 | Validate export completeness                                              | Compare <code>ls projects/full-export/settings/ \</code>   <code>wc -l</code> against Settings API <code>totalCount</code> per schema     | Critical    |
| 34 | Validate export contains no secrets                                       | Run <code>grep -r "dt0c01" projects/</code> after export and confirm 0 matches                                                            | Critical    |

| #  | Best Practice                                                    | Recommended Setting/Value                                                                                                                                      | Priority           |
|----|------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------|
| 35 | Run <code>monaco validate manifest.yaml</code> before any deploy | Catches JSON validity errors, missing references, and dependency issues before they reach the target                                                           | <b>Critical</b>    |
| 36 | Provision target tenant and verify admin access                  | Confirm SSO, API token creation, and environment admin role before proceeding                                                                                  | <b>Critical</b>    |
| 37 | Configure SSO with full SAML message signing                     | Create a new SAML application in your IdP for the target tenant — new Entity ID, new ACS URL; never reuse the source SAML app                                  | <b>Critical</b>    |
| 38 | Test SSO with a pilot user before cutover                        | SAML configuration issues are the #1 day-of-cutover blocker                                                                                                    | <b>Critical</b>    |
| 39 | Deploy ActiveGates in the target tenant                          | ActiveGates cannot be reconfigured like OneAgents; install fresh from the target tenant UI                                                                     | <b>Critical</b>    |
| 40 | Prepare K8s operator manifests for the target tenant             | Update DynaKube API to <code>v1beta5</code> or <code>v1beta6</code> ; use Operator v1.8.1+ from <code>oci://public.ecr.aws/dynatrace/dynatrace-operator</code> | <b>Critical</b>    |
| 41 | Recreate network zones in the target                             | Export with <code>monaco download --specific-settings builtin:networkzones</code> and deploy to the target before agent migration                              | <b>Critical</b>    |
| 42 | Create new OAuth clients and API tokens in the target            | Source credentials cannot be exported; create fresh clients with matching scopes                                                                               | <b>Critical</b>    |
| 43 | Limit Azure AD group claims to 150 groups per SAML assertion     | Azure AD has a hard limit; use group filtering if your org exceeds this                                                                                        | <b>Recommended</b> |

## 5. Step 5: Execute — Configuration Import and Agent Cutover

Source: S2S-05

| #  | Best Practice                                                                                                                                  | Recommended Setting/Value                                                                                                                         | Priority           |
|----|------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------|--------------------|
| 44 | Follow the 24-item deployment order                                                                                                            | Buckets → Enrichment → OpenPipeline → Segments → Gen2 config → Alerting/SLOs → Dashboards/Workflows → IAM last                                    | <b>Critical</b>    |
| 45 | Deploy IAM groups, policies, and bindings last                                                                                                 | IAM <code>WHERE</code> clauses reference schemas, buckets, and security contexts that must already exist                                          | <b>Critical</b>    |
| 46 | Use <code>monaco deploy manifest.yaml --environment target</code> for import                                                                   | Monaco resolves dependencies automatically within a single run (except IAM)                                                                       | <b>Critical</b>    |
| 47 | Run <code>monaco deploy --dry-run</code> before every real deploy                                                                              | Preview what will change before committing to the target tenant                                                                                   | <b>Critical</b>    |
| 48 | Deploy IAM via Terraform with <code>account-idm-read</code> , <code>account-idm-write</code> , and <code>iam-policies-management</code> scopes | These three OAuth scopes are required to manage groups, policies, and bindings                                                                    | <b>Critical</b>    |
| 49 | Reconfigure OneAgent with <code>oneagentctl --do not reinstall</code>                                                                          | Run <code>oneagentctl --set-server=&lt;target-url&gt;/communication --set-tenant=&lt;target-id&gt; --set-tenant-token=&lt;target-token&gt;</code> | <b>Critical</b>    |
| 50 | Route firewall-restricted hosts through ActiveGates                                                                                            | Use <code>--set-server="https://&lt;ag-host&gt;:9999/communication"</code> for on-prem or DMZ hosts; semicolon-separate multiple AGs for failover | <b>Critical</b>    |
| 51 | Restart application processes after agent reconfiguration                                                                                      | Application processes must restart for the agent to report to the new tenant                                                                      | <b>Critical</b>    |
| 52 | Migrate non-production first                                                                                                                   | Dev agents first (weeks 1-2), staging (weeks 3-4), production last (weeks 5-6)                                                                    | <b>Critical</b>    |
| 53 | Validate after each migration wave                                                                                                             | Verify host count, service count, and data flow in the target after each environment wave                                                         | <b>Critical</b>    |
| 54 | Automate fleet-wide <code>oneagentctl</code> with config management                                                                            | Use Ansible, Chef, Puppet, AWS SSM, or Azure Automation for large fleets                                                                          | <b>Recommended</b> |



| #  | Best Practice                                                                               | Recommended Setting/Value                                                                           | Priority           |
|----|---------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------|--------------------|
| 55 | Budget 540m CPU and 872Mi memory per node for Dynatrace overhead                            | OneAgent (100m/512Mi) + CSI Driver (440m/360Mi) per node; control plane adds 650m/320Mi per cluster | <b>Recommended</b> |
| 56 | Update IRSA (AWS), Managed Identity (Azure), or Workload Identity (GCP) to reference target | Cloud-specific K8s identity must point to the new tenant                                            | <b>Critical</b>    |

## 6. Step 6: Integrate — Cloud, Dashboards, and Workflows

Source: S2S-06

| #  | Best Practice                                                                  | Recommended Setting/Value                                                                                                    | Priority        |
|----|--------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------|-----------------|
| 57 | Recreate cloud provider credentials in the target tenant                       | monaco download cannot export AWS, Azure, or K8s credentials; create new IAM roles, service principals, and service accounts | <b>Critical</b> |
| 58 | Create new AWS IAM role with trust policy for the target tenant                | Reference the target tenant's Dynatrace account ID and external ID from the AWS integration wizard                           | <b>Critical</b> |
| 59 | Create new Azure App Registration with Monitoring Reader role                  | New Tenant ID, Client ID, Client Secret for the target tenant                                                                | <b>Critical</b> |
| 60 | Create new GCP Service Account with Monitoring Viewer and Compute Viewer roles | Generate new JSON key for the target tenant integration                                                                      | <b>Critical</b> |
| 61 | Deploy new AWS Log Forwarder Lambda pointing to the target                     | Update S3 event notifications to trigger the new Lambda function                                                             | <b>Critical</b> |
| 62 | Update webhook URLs that reference source-tenant-specific endpoints            | Verify custom webhooks accept traffic from the target tenant's IP range                                                      | <b>Critical</b> |
| 63 | Migrate dashboards via Monaco and audit for hardcoded entity IDs               | Search dashboard JSON for entityId( , HOST- , SERVICE- , PROCESS_GROUP- patterns and replace with tag-based filters          | <b>Critical</b> |

| #  | Best Practice                                                                      | Recommended Setting/Value                                                                                                     | Priority           |
|----|------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------|--------------------|
| 64 | Reassign dashboard ownership to a target tenant user or service account            | Source user accounts may not exist in the target; set <code>owner</code> field to a valid target identity                     | <b>Critical</b>    |
| 65 | Create new workflow service user (actor) in the target tenant                      | Source service users cannot be migrated; create new service user and set as workflow actor                                    | <b>Critical</b>    |
| 66 | Recreate Synthetic private locations in the target tenant                          | Private Synthetic locations are tenant-specific; update location references in all monitor configs                            | <b>Critical</b>    |
| 67 | Use a classic API token (not OAuth) for Synthetic monitor migration                | Synthetic monitors require classic token as of provider v1.88.0+                                                              | <b>Critical</b>    |
| 68 | Reinstall Extensions 2.0 from the Dynatrace Hub                                    | Neither Monaco nor Terraform supports Extensions 2.0 export/import; extensions require a host-based ActiveGate, not K8s-based | <b>Critical</b>    |
| 69 | Standardize tag naming across cloud providers                                      | Use consistent keys ( <code>app</code> , <code>env</code> , <code>team</code> ) regardless of AWS/Azure/GCP                   | <b>Recommended</b> |
| 70 | Webhook integrations (Teams, Slack, PagerDuty, ServiceNow, Jira) use the same URLs | No URL changes needed for standard integrations; just recreate the notification rule in the target                            | <b>Recommended</b> |

## 7. Step 7: Expand — OpenPipeline, SLOs, and Alerting

Source: S2S-07

| #  | Best Practice                                          | Recommended Setting/Value                                     | Priority        |
|----|--------------------------------------------------------|---------------------------------------------------------------|-----------------|
| 71 | Deploy Grail buckets before OpenPipeline routing rules | Routing rules that target nonexistent buckets cause data loss | <b>Critical</b> |

| #  | Best Practice                                                             | Recommended Setting/Value                                                                                                                                                                         | Priority           |
|----|---------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------|
| 72 | Deploy in order:<br>buckets →<br>enrichment rules →<br>OpenPipeline rules | This three-step sequence ensures data flows correctly from the first byte                                                                                                                         | <b>Critical</b>    |
| 73 | Wait for K8s enrichment propagation before deploying routing rules        | Enrichment rules typically take ~45 min to propagate on EKS; routing rules that depend on enriched attributes fail silently until propagation completes — data lands in <code>default_logs</code> | <b>Critical</b>    |
| 74 | Add a CI/CD validation gate between enrichment and routing deploys        | Query for enriched fields ( <code>isNotNull(app.namespace)</code> ) before applying routing rules; do not rely on <code>monaco deploy</code> alone (it does not wait for propagation)             | <b>Recommended</b> |
| 75 | Use prefix naming for bucket consolidation                                | E.g., <code>us-east_app_logs</code> , <code>eu-west_app_logs</code> to prevent collisions when consolidating tenants                                                                              | <b>Recommended</b> |
| 76 | Match retention policies exactly for regulated environments               | Set to same days as source (e.g., 360 days for PCI); verify compliance requirements before cutover                                                                                                | <b>Critical</b>    |
| 77 | Update bucket references in OpenPipeline routing rules                    | Routing rules from export contain source bucket names; replace with target bucket names                                                                                                           | <b>Critical</b>    |
| 78 | Convert SLO entity IDs to tag-based filters                               | Replace <code>type(SERVICE),entityId(SERVICE-xxx)</code> with <code>type(SERVICE),tag(app:checkout),tag(env:production)</code>                                                                    | <b>Critical</b>    |

| #  | Best Practice                                        | Recommended Setting/Value                                                                                                    | Priority           |
|----|------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------|--------------------|
| 79 | Set SLO evaluation windows to 30 days initially      | Start with <code>ROLLING_3_DAYS</code> , expand to <code>ROLLING_WEEK</code> after 7 days, then to full window after 30 days | <b>Critical</b>    |
| 80 | Align SLO cutover with calendar boundaries           | If using calendar-based SLOs, switch at the start of a month or quarter to avoid partial evaluation                          | <b>Recommended</b> |
| 81 | Port alerting profile severity and tag filters as-is | Ensure the referenced tags exist in the target tenant before deploying alerting profiles                                     | <b>Critical</b>    |
| 82 | Verify maintenance window timezone settings          | Cron expressions are timezone-sensitive; confirm schedules are correct after import                                          | <b>Recommended</b> |
| 83 | Update maintenance window entity scopes              | Source entity references in scope definitions must be remapped to target tenant entities                                     | <b>Critical</b>    |

## 8. Step 8: Enable — Parallel Operation and Stakeholder Handover

Source: S2S-08

| #  | Best Practice                                                     | Recommended Setting/Value                                                                                              | Priority        |
|----|-------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------|-----------------|
| 84 | Run parallel operation for 2-4 weeks minimum                      | Historical data does not migrate; parallel is the only path to continuity                                              | <b>Critical</b> |
| 85 | Allow 7-14 days for Dynatrace Intelligence baselines to stabilize | Availability: 2-3 days; response time and error rate: 1-2 weeks; resource utilization: 2-4 weeks (full business cycle) | <b>Critical</b> |

| #  | Best Practice                                     | Recommended Setting/Value                                                                            | Priority           |
|----|---------------------------------------------------|------------------------------------------------------------------------------------------------------|--------------------|
| 86 | Use phased parallel model to control cost         | Migrate agents in waves so you never run 2x host units simultaneously; expect 1.1–1.5x cost          | <b>Critical</b>    |
| 87 | Communicate to all personas                       | Engineering teams: 4 weeks before cutover; SRE/on-call: 2 weeks before; executives: weekly status    | <b>Critical</b>    |
| 88 | Communicate the SLO baseline gap to stakeholders  | SLOs start at 0% in the target and accumulate data over 7–30 days depending on evaluation window     | <b>Critical</b>    |
| 89 | Warn about Dynatrace Intelligence false positives | Baselines are relearning during first 2–4 weeks; expect noisy alerting until stabilized              | <b>Recommended</b> |
| 90 | Negotiate dual licensing with Dynatrace           | Contact your account team to discuss temporary parallel licensing to reduce cost                     | <b>Recommended</b> |
| 91 | Configure SCIM provisioning for the target tenant | Set new SCIM URL and token; users sync automatically on first login                                  | <b>Recommended</b> |
| 92 | Export problem history before decommission        | Problem history cannot migrate; export critical incidents as documentation for post-mortem reference | <b>Recommended</b> |
| 93 | Keep a fallback local admin account               | If SAML is misconfigured on day of cutover, local admin bypasses SSO                                 | <b>Critical</b>    |

## 9. Step 9: Optimize — Cutover Validation and Decommission

Source: S2S-09

| #  | Best Practice                              | Recommended Setting/Value                                                                                                                                           | Priority        |
|----|--------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------|
| 94 | Complete go/no-go checklist before cutover | Verify host count, service count, log volume, span volume, SLO evaluation, alerting coverage — all must match success criteria                                      | <b>Critical</b> |
| 95 | Validate all data flows via DQL            | Run <code>fetch logs, from:-1h \   summarize count()</code> , <code>fetch spans, from:-1h \   summarize count()</code> , and entity count queries to confirm parity | <b>Critical</b> |

| #   | Best Practice                                                                           | Recommended Setting/Value                                                                                                      | Priority           |
|-----|-----------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------|--------------------|
| 96  | Verify enrichment tags are populating                                                   | Run <code>fetch logs, from:-1h \   filter isNotNull(dt.security_context) \   summarize count()</code> and confirm count > 0    | <b>Critical</b>    |
| 97  | Run final <code>monaco download</code> → <code>monaco deploy sync</code> before cutover | Captures any configuration changes made since initial export                                                                   | <b>Critical</b>    |
| 98  | Get stakeholder sign-off                                                                | SRE, dev, security, management, and compliance each validate their domain: alerting, dashboards, IAM, SLO reporting, retention | <b>Critical</b>    |
| 99  | Disable alerting in source tenant after cutover                                         | Prevents duplicate alerts during the transition tail                                                                           | <b>Critical</b>    |
| 100 | Keep source tenant running 30 days minimum after cutover                                | Read-only mode for reference; do not deprovision until all stakeholders confirm                                                | <b>Critical</b>    |
| 101 | Rotate all credentials within 30 days of cutover                                        | API tokens, OAuth clients, cloud provider keys — revoke any temporary migration credentials                                    | <b>Critical</b>    |
| 102 | Revoke all source tenant API tokens and deactivate OAuth clients                        | Prevents unauthorized access to the deprecated tenant                                                                          | <b>Critical</b>    |
| 103 | Remove SAML/SSO application from IdP for source tenant                                  | Eliminates stale IdP configuration and prevents confusion                                                                      | <b>Recommended</b> |
| 104 | Export final audit logs from source before decommission                                 | Compliance requires audit trail retention; export before access is lost                                                        | <b>Critical</b>    |
| 105 | Update all documentation                                                                | Replace source tenant URLs, API endpoints, and dashboard links in runbooks, CI/CD pipelines, and architecture diagrams         | <b>Critical</b>    |
| 106 | Tune Dynatrace Intelligence anomaly detection during weeks 1-2 post-cutover             | Suppress known false positives; adjust sensitivity for noisy services                                                          | <b>Critical</b>    |

| #   | Best Practice                                                | Recommended Setting/Value                                                                                            | Priority           |
|-----|--------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------|--------------------|
| 107 | Expand SLO evaluation windows to full duration by week 4     | Move from <code>ROLLING_3_DAYS</code> to <code>ROLLING_WEEK</code> to <code>ROLLING_MONTH</code> as data accumulates | <b>Recommended</b> |
| 108 | Consolidate redundant Synthetic monitors                     | If consolidating tenants, deduplicate monitors that tested the same endpoints from different tenants                 | <b>Optional</b>    |
| 109 | Replace all remaining entity ID references                   | Post-migration cleanup ensures configuration is fully portable for future migrations                                 | <b>Recommended</b> |
| 110 | Schedule credential rotation for cloud provider integrations | Rotate within 30 days of cutover for all AWS, Azure, and GCP integrations                                            | <b>Recommended</b> |
| 111 | Conduct a lessons-learned review within 2 weeks of cutover   | Capture planning gaps, tool issues, communication effectiveness, timeline accuracy                                   | <b>Recommended</b> |
| 112 | Set source tenant to read-only retention after cutover       | Preserves historical data for reference without active monitoring cost                                               | <b>Critical</b>    |

## The Five Principles of S2S Migration

1. **Export early, validate often** — do not wait for cutover to discover missing config
2. **Tags over entity IDs** — make all configuration portable before migration
3. **Parallel is not optional** — historical data does not migrate; 30 days minimum
4. **IAM last** — deploy the resources first, then lock them down with policies
5. **Communicate relentlessly** — migration is as much about people as technology

## Summary

This notebook contains **112 best practices** across 9 migration steps. The breakdown by priority:

| Priority           | Count | Guidance                                                                                       |
|--------------------|-------|------------------------------------------------------------------------------------------------|
| <b>Critical</b>    | 81    | Must implement. Skipping any of these risks data loss, access failures, or migration rollback. |
| <b>Recommended</b> | 27    | Implement unless there is a documented reason not to. These reduce risk and cost.              |

| Priority        | Count | Guidance                                                                     |
|-----------------|-------|------------------------------------------------------------------------------|
| <b>Optional</b> | 4     | Implement if resources allow. These improve efficiency but are not blockers. |

Use the 9-step framework table and the 11-step Order of Operations table at the top of this notebook to track progress and govern execution sequence. Each step's best practices align to the corresponding S2S notebook (S2S-01 through S2S-10) for full context and worked examples.

---

*This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.*