S2S-99: Best Practice Summary

> **Series:** S2S | **Notebook:** 99 | **Created:** March 2026 | **Last Updated:** 03/30/2026

Overview

This notebook consolidates **104 actionable best practices** from the S2S series into a single reference organized around the 9-step migration framework. Each practice is definitive: it tells you exactly what to set, when, and why. Use this as a pre-flight checklist and ongoing reference throughout your SaaS-to-SaaS migration.

The S2S series follows a **9-step framework** (Discover through Optimize) supported by an **11-step Order of Operations** that governs execution sequence.

The 9-Step Migration Framework

Step	Name	Focus
1	**Discover**	Migration scenarios and inventory
2	**Strategize**	Define your migration approach
3	**Design**	Target tenant architecture
4	**Prepare**	Export and pre-stage
5	**Execute**	Configuration import and agent cutover
6	**Integrate**	Cloud, dashboards, and workflows
7	**Expand**	OpenPipeline, SLOs, and alerting
8	**Enable**	Parallel operation and stakeholder handover
9	**Optimize**	Cutover validation and decommission

The 11-Step Order of Operations

#	Operation	Description
1	**Assess**	Inventory source tenant

#	Operation	Description
2	**Provision**	Target SaaS tenant and access (SSO/IAM)
3	**Export**	Configuration from source tenant
4	**Migrate**	Configuration to target tenant
5	**Rebuild**	Dashboards, SLOs, alerts
6	**Redirect**	OneAgents and operators to target
7	**Reconnect**	Cloud integrations and extensions
8	**Migrate**	Any remaining configuration
9	**Validate**	Data flow and performance
10	**Cutover**	Full switch to target tenant
11	**Decommission**	Source tenant

Table of Contents

Prerequisites

Requirement	Details
S2S Series	Read S2S-01 through S2S-09 for full context
Source Tenant	Active Dynatrace SaaS with admin access

Requirement	Details
Target Tenant	Provisioned Dynatrace SaaS with admin access
CLI Tools	Monaco CLI v2.x, Terraform v1.5+ with provider ~> 1.91

1. Step 1: Discover — Migration Scenarios and Inventory

Source: S2S-01

#	Best Practice	Recommended Setting/Value	Priority
1	Complete entity discovery before anything else	Run `fetch dt.entity.host | summarize count()` for hosts, services, process groups, and applications to establish baseline numbers	**Critical**
2	Identify your migration scenario	Classify as consolidation (many-to-one), split (one-to-many), regional relocation, or cloud provider change — each has distinct tooling and planning requirements	**Critical**
3	Select Monaco for configuration, Terraform for IAM	Monaco v2 covers all 8 config types (settings, document, automation, bucket, segment, slo-v2, openpipeline, classic api); Terraform is required exclusively for IAM policies, groups, and bindings	**Critical**
4	Plan for entity ID changes	Entity IDs (HOST-xxx, SERVICE-xxx) are tenant-specific and will change; identify every dashboard, SLO, and alert that hardcodes an entity ID	**Critical**
5	Document the 90/10 manual items	90% of config migrates automatically; budget 90% of your effort for the remaining 10% — entity ID remapping, integration repointing, IAM redesign, credential recreation	**Critical**
6	Build a comprehensive endpoint inventory	Document every agent URL, API call, webhook, and integration endpoint that references the source tenant URL	**Critical**
7	Inventory Extensions 2.0 separately	Extensions cannot be exported via Monaco or Terraform; list all extensions and plan for manual reinstallation from the Dynatrace Hub	**Recommended**
8	Audit configuration changes from last 30 days	Query audit logs to identify recently changed settings that may not be in your last export	**Recommended**

2. Step 2: Strategize — Define Your Migration Approach

Source: S2S-02

#	Best Practice	Recommended Setting/Value	Priority
9	Choose phased approach for environments with >500 hosts	Migrate dev first (weeks 1-2), staging (weeks 3-4), production (weeks 5-6); gives Davis AI time to build baselines	**Critical**
10	Define success criteria before starting	Document measurable targets: host count parity, service count parity, log volume parity, SLO evaluation accuracy, alerting coverage	**Critical**
11	Engage your account team for licensing early	Contact Dynatrace to discuss temporary parallel licensing; negotiate before the migration window opens	**Critical**
12	Plan a 2-4 week parallel operation period	Historical data does not migrate; parallel is the only path to data continuity and Davis AI baseline stability	**Critical**
13	Establish a configuration freeze window	No settings modifications in the source tenant during cutover; communicate freeze dates to all teams	**Critical**
14	Align migration timing with stable traffic periods	Avoid holidays, sales events, or other atypical traffic patterns that would distort Davis AI baselines	**Recommended**
15	Expect the 90/10 rule	90% of config migrates automatically; budget 90% of effort for the remaining 10%	**Critical**
16	Prepare a rollback plan	Document how to revert agents to the source tenant if issues arise during cutover	**Recommended**

3. Step 3: Design — Target Tenant Architecture

Source: S2S-03

#	Best Practice	Recommended Setting/Value	Priority
17	Redesign IAM during consolidation	Use the migration as the opportunity to fix group sprawl and overprivileged policies; redesign from scratch for the target tenant	**Critical**

#	Best Practice	Recommended Setting/Value	Priority
18	Deploy data foundation before governance	Buckets and enrichment rules must exist before OpenPipeline routing, segments, SLOs, and IAM policies that reference them	**Critical**
19	Replace entity IDs with tag-based filters	Convert `entityId("HOST-xxxx")` to `tag(app:checkout),tag(env:production)` in dashboards, SLOs, and notification rules before migration	**Critical**
20	Design bucket naming for consolidation	Use prefix naming (e.g., `us-east_app_logs`, `eu-west_app_logs`) to prevent collisions when consolidating multiple tenants	**Recommended**
21	Plan the 24-item deployment order	Buckets (1st) → Enrichment (2nd) → OpenPipeline (3rd) → Segments (4th) → Gen2 config (5th-13th) → Alerting/SLOs (14th-18th) → Dashboards/Workflows (19th-21st) → IAM last (22nd-24th)	**Critical**
22	Replace management zone ID references with names	Convert `mzId(123)` to `mzName("Production")` because zone IDs change between tenants	**Critical**
23	Prefix resources with source identifier during consolidation	Set `name = "${source_prefix} - ${resource_name}"` on all dashboards, rules, and SLOs to prevent naming collisions	**Recommended**
24	Tag everything by business unit before splitting a tenant	Apply auto-tags identifying ownership before splitting one tenant into many	**Recommended**
25	Map cloud provider identity constructs when changing providers	AWS IAM roles to Azure Service Principals, Azure AD Groups to GCP Google Groups, etc.	**Recommended**

4. Step 4: Prepare — Export and Pre-Stage

Source: S2S-04

#	Best Practice	Recommended Setting/Value	Priority
26	Export via `monaco download`	Run `monaco download manifest.yaml --environment source-tenant` to export all configuration in one operation	**Critical**
27	Store exported configuration in Git	Commit the Monaco export directory to a Git repository immediately after download	**Critical**
28	Validate export completeness	Compare `ls projects/full-export/settings/ | wc -l` against Settings API `totalCount` per schema	**Critical**
29	Validate export contains no secrets	Run `grep -r "dt0c01" projects/` after export and confirm 0 matches	**Critical**
30	Run `monaco validate manifest.yaml` before any deploy	Catches JSON validity errors, missing references, and dependency issues before they reach the target	**Critical**
31	Provision target tenant and verify admin access	Confirm SSO, API token creation, and environment admin role before proceeding	**Critical**
32	Configure SSO with full SAML message signing	Create a new SAML application in your IdP for the target tenant — new Entity ID, new ACS URL; never reuse the source SAML app	**Critical**
33	Test SSO with a pilot user before cutover	SAML configuration issues are the #1 day-of-cutover blocker	**Critical**
34	Deploy ActiveGates in the target tenant	ActiveGates cannot be reconfigured like OneAgents; install fresh from the target tenant UI	**Critical**
35	Prepare K8s operator manifests for the target tenant	Update DynaKube API to `v1beta5` or `v1beta6`; use Operator v1.8.1+ from `oci://public.ecr.aws/dynatrace/dynatrace-operator`	**Critical**
36	Recreate network zones in the target	Export with `monaco download --specific-settings builtin:networkzones` and deploy to the target before agent migration	**Critical**
37	Create new OAuth clients and API tokens in the target	Source credentials cannot be exported; create fresh clients with matching scopes	**Critical**

#	Best Practice	Recommended Setting/Value	Priority
38	Limit Azure AD group claims to 150 groups per SAML assertion	Azure AD has a hard limit; use group filtering if your org exceeds this	**Recommended**

5. Step 5: Execute — Configuration Import and Agent Cutover

Source: S2S-05

#	Best Practice	Recommended Setting/Value	Priority
39	Follow the 24-item deployment order	Buckets → Enrichment → OpenPipeline → Segments → Gen2 config → Alerting/SLOs → Dashboards/Workflows → IAM last	**Critical**
40	Deploy IAM groups, policies, and bindings last	IAM `WHERE` clauses reference schemas, buckets, and security contexts that must already exist	**Critical**
41	Use `monaco deploy manifest.yaml --environment target` for import	Monaco resolves dependencies automatically within a single run (except IAM)	**Critical**
42	Run `monaco deploy --dry-run` before every real deploy	Preview what will change before committing to the target tenant	**Critical**
43	Deploy IAM via Terraform with `account-idm-read`, `account-idm-write`, and `iam-policies-management` scopes	These three OAuth scopes are required to manage groups, policies, and bindings	**Critical**
44	Reconfigure OneAgent with `oneagentctl` — do not reinstall	Run `oneagentctl --set-server=<target-url>/communication --set-tenant=<target-id> --set-tenant-token=<target-token>`	**Critical**
45	Restart application processes after agent reconfiguration	Application processes must restart for the agent to report to the new tenant	**Critical**
46	Migrate non-production first	Dev agents first (weeks 1-2), staging (weeks 3-4), production last (weeks 5-6)	**Critical**

#	Best Practice	Recommended Setting/Value	Priority
47	Validate after each migration wave	Verify host count, service count, and data flow in the target after each environment wave	**Critical**
48	Automate fleet-wide `oneagentctl` with config management	Use Ansible, Chef, Puppet, AWS SSM, or Azure Automation for large fleets	**Recommended**
49	Budget 540m CPU and 872Mi memory per node for Dynatrace overhead	OneAgent (100m/512Mi) + CSI Driver (440m/360Mi) per node; control plane adds 650m/320Mi per cluster	**Recommended**
50	Update IRSA (AWS), Managed Identity (Azure), or Workload Identity (GCP) to reference target	Cloud-specific K8s identity must point to the new tenant	**Critical**

6. Step 6: Integrate — Cloud, Dashboards, and Workflows

Source: S2S-06

#	Best Practice	Recommended Setting/Value	Priority
51	Recreate cloud provider credentials in the target tenant	`monaco download` cannot export AWS, Azure, or K8s credentials; create new IAM roles, service principals, and service accounts	**Critical**
52	Create new AWS IAM role with trust policy for the target tenant	Reference the target tenant's Dynatrace account ID and external ID from the AWS integration wizard	**Critical**
53	Create new Azure App Registration with `Monitoring Reader` role	New Tenant ID, Client ID, Client Secret for the target tenant	**Critical**
54	Create new GCP Service Account with `Monitoring Viewer` and `Compute Viewer` roles	Generate new JSON key for the target tenant integration	**Critical**
55	Deploy new AWS Log Forwarder Lambda pointing to the target	Update S3 event notifications to trigger the new Lambda function	**Critical**

#	Best Practice	Recommended Setting/Value	Priority
56	Update webhook URLs that reference source-tenant-specific endpoints	Verify custom webhooks accept traffic from the target tenant's IP range	**Critical**
57	Migrate dashboards via Monaco and audit for hardcoded entity IDs	Search dashboard JSON for `entityId(` , `HOST-` , `SERVICE-` , `PROCESS_GROUP-` patterns and replace with tag-based filters	**Critical**
58	Reassign dashboard ownership to a target tenant user or service account	Source user accounts may not exist in the target; set `owner` field to a valid target identity	**Critical**
59	Create new workflow service user (actor) in the target tenant	Source service users cannot be migrated; create new service user and set as workflow actor	**Critical**
60	Recreate Synthetic private locations in the target tenant	Private Synthetic locations are tenant-specific; update location references in all monitor configs	**Critical**
61	Use a classic API token (not OAuth) for Synthetic monitor migration	Synthetic monitors require classic token as of provider v1.88.0+	**Critical**
62	Reinstall Extensions 2.0 from the Dynatrace Hub	Neither Monaco nor Terraform supports Extensions 2.0 export/import; extensions require a host-based ActiveGate, not K8s-based	**Critical**
63	Standardize tag naming across cloud providers	Use consistent keys (`app` , `env` , `team`) regardless of AWS/Azure/GCP	**Recommended**
64	Webhook integrations (Teams, Slack, PagerDuty, ServiceNow, Jira) use the same URLs	No URL changes needed for standard integrations; just recreate the notification rule in the target	**Recommended**

7. Step 7: Expand — OpenPipeline, SLOs, and Alerting

Source: S2S-07

#	Best Practice	Recommended Setting/Value	Priority
65	Deploy Grail buckets before OpenPipeline routing rules	Routing rules that target nonexistent buckets cause data loss	**Critical**
66	Deploy in order: buckets → enrichment rules → OpenPipeline rules	This three-step sequence ensures data flows correctly from the first byte	**Critical**
67	Use prefix naming for bucket consolidation	E.g., `us-east_app_logs`, `eu-west_app_logs` to prevent collisions when consolidating tenants	**Recommended**
68	Match retention policies exactly for regulated environments	Set to same days as source (e.g., 360 days for PCI); verify compliance requirements before cutover	**Critical**
69	Update bucket references in OpenPipeline routing rules	Routing rules from export contain source bucket names; replace with target bucket names	**Critical**
70	Convert SLO entity IDs to tag-based filters	Replace `type(SERVICE),entityId(SERVICE-xxx)` with `type(SERVICE),tag(app:checkout),tag(env:production)`	**Critical**
71	Set SLO evaluation windows to 30 days initially	Start with `ROLLING_3_DAYS`, expand to `ROLLING_WEEK` after 7 days, then to full window after 30 days	**Critical**
72	Align SLO cutover with calendar boundaries	If using calendar-based SLOs, switch at the start of a month or quarter to avoid partial evaluation	**Recommended**

#	Best Practice	Recommended Setting/Value	Priority
73	Port alerting profile severity and tag filters as-is	Ensure the referenced tags exist in the target tenant before deploying alerting profiles	**Critical**
74	Verify maintenance window timezone settings	Cron expressions are timezone-sensitive; confirm schedules are correct after import	**Recommended**
75	Update maintenance window entity scopes	Source entity references in scope definitions must be remapped to target tenant entities	**Critical**

8. Step 8: Enable — Parallel Operation and Stakeholder Handover

Source: S2S-08

#	Best Practice	Recommended Setting/Value	Priority
76	Run parallel operation for 2-4 weeks minimum	Historical data does not migrate; parallel is the only path to continuity	**Critical**
77	Allow 7-14 days for Davis AI baselines to stabilize	Availability: 2-3 days; response time and error rate: 1-2 weeks; resource utilization: 2-4 weeks (full business cycle)	**Critical**
78	Use phased parallel model to control cost	Migrate agents in waves so you never run 2x host units simultaneously; expect 1.1-1.5x cost	**Critical**
79	Communicate to all personas	Engineering teams: 4 weeks before cutover; SRE/on-call: 2 weeks before; executives: weekly status	**Critical**
80	Communicate the SLO baseline gap to stakeholders	SLOs start at 0% in the target and accumulate data over 7-30 days depending on evaluation window	**Critical**
81	Warn about Davis AI false positives	Baselines are relearning during first 2-4 weeks; expect noisy alerting until stabilized	**Recommended**

#	Best Practice	Recommended Setting/Value	Priority
82	Negotiate dual licensing with Dynatrace	Contact your account team to discuss temporary parallel licensing to reduce cost	**Recommended**
83	Configure SCIM provisioning for the target tenant	Set new SCIM URL and token; users sync automatically on first login	**Recommended**
84	Export problem history before decommission	Problem history cannot migrate; export critical incidents as documentation for post-mortem reference	**Recommended**
85	Keep a fallback local admin account	If SAML is misconfigured on day of cutover, local admin bypasses SSO	**Critical**

9. Step 9: Optimize — Cutover Validation and Decommission

Source: S2S-09

#	Best Practice	Recommended Setting/Value	Priority
86	Complete go/no-go checklist before cutover	Verify host count, service count, log volume, span volume, SLO evaluation, alerting coverage — all must match success criteria	**Critical**
87	Validate all data flows via DQL	Run `fetch logs, from:-1h | summarize count()`, `fetch spans, from:-1h | summarize count()`, and entity count queries to confirm parity	**Critical**
88	Verify enrichment tags are populating	Run `fetch logs, from:-1h | filter isNotNull(dt.security_context) | summarize count()` and confirm count > 0	**Critical**
89	Run final `monaco download` → `monaco deploy` sync before cutover	Captures any configuration changes made since initial export	**Critical**
90	Get stakeholder sign-off	SRE, dev, security, management, and compliance each validate their domain: alerting, dashboards, IAM, SLO reporting, retention	**Critical**
91	Disable alerting in source tenant after cutover	Prevents duplicate alerts during the transition tail	**Critical**

#	Best Practice	Recommended Setting/Value	Priority
92	Keep source tenant running 30 days minimum after cutover	Read-only mode for reference; do not deprovision until all stakeholders confirm	**Critical**
93	Rotate all credentials within 30 days of cutover	API tokens, OAuth clients, cloud provider keys — revoke any temporary migration credentials	**Critical**
94	Revoke all source tenant API tokens and deactivate OAuth clients	Prevents unauthorized access to the deprecated tenant	**Critical**
95	Remove SAML/SSO application from IdP for source tenant	Eliminates stale IdP configuration and prevents confusion	**Recommended**
96	Export final audit logs from source before decommission	Compliance requires audit trail retention; export before access is lost	**Critical**
97	Update all documentation	Replace source tenant URLs, API endpoints, and dashboard links in runbooks, CI/CD pipelines, and architecture diagrams	**Critical**
98	Tune Davis AI anomaly detection during weeks 1-2 post-cutover	Suppress known false positives; adjust sensitivity for noisy services	**Critical**
99	Expand SLO evaluation windows to full duration by week 4	Move from `ROLLING_3_DAYS` to `ROLLING_WEEK` to `ROLLING_MONTH` as data accumulates	**Recommended**
100	Consolidate redundant Synthetic monitors	If consolidating tenants, deduplicate monitors that tested the same endpoints from different tenants	**Optional**
101	Replace all remaining entity ID references	Post-migration cleanup ensures configuration is fully portable for future migrations	**Recommended**
102	Schedule credential rotation for cloud provider integrations	Rotate within 30 days of cutover for all AWS, Azure, and GCP integrations	**Recommended**
103	Conduct a lessons-learned review within 2 weeks of cutover	Capture planning gaps, tool issues, communication effectiveness, timeline accuracy	**Recommended**

#	Best Practice	Recommended Setting/Value	Priority
104	Set source tenant to read-only retention after cutover	Preserves historical data for reference without active monitoring cost	**Critical**

The Five Principles of S2S Migration

1. **Export early, validate often** — do not wait for cutover to discover missing config
2. **Tags over entity IDs** — make all configuration portable before migration
3. **Parallel is not optional** — historical data does not migrate; 30 days minimum
4. **IAM last** — deploy the resources first, then lock them down with policies
5. **Communicate relentlessly** — migration is as much about people as technology

Summary

This notebook contains **104 best practices** across 9 migration steps. The breakdown by priority:

Priority	Count	Guidance
Critical	73	Must implement. Skipping any of these risks data loss, access failures, or migration rollback.
Recommended	27	Implement unless there is a documented reason not to. These reduce risk and cost.
Optional	4	Implement if resources allow. These improve efficiency but are not blockers.

Use the 9-step framework table and the 11-step Order of Operations table at the top of this notebook to track progress and govern execution sequence. Each step's best practices align to the corresponding S2S notebook (S2S-01 through S2S-09) for full context and worked examples.

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.