

SL2DT-02: Assessment & Inventory

Series: SL2DT | **Notebook:** 2 of 10 | **Created:** April 2026 | **Last Updated:** 04/21/2026

Overview

Goal of this step: produce a complete inventory of the source Sumo Logic account and a cut-scope decision artifact before any migration work begins. The cut-scope decision is the single biggest cost driver — get it right here and the rest of the migration is proportional to actual business value.

Everything downstream (SL2DT-03 ingest design, SL2DT-04 translation pass, SL2DT-05 monitors, SL2DT-06 dashboards) depends on the inventory and taxonomy maps produced in this notebook.

Table of Contents

1. [What You'll Produce](#)
 2. [Prerequisites & Access](#)
 3. [Pull the Sumo Inventory](#)
 4. [Pull the Audit Log — Usage-Based Cut Scope](#)
 5. [Map `\_sourceCategory` Taxonomy](#)
 6. [Inventory Collectors & Ingest Paths](#)
 7. [Inventory FERs \(Field Extraction Rules\)](#)
 8. [Inventory Roles & Users](#)
 9. [Cut-Scope Decision Workshop](#)
 10. [Step Exit Criteria](#)
-

Prerequisites

Requirement	Details
Audience	Migration lead + Sumo admin
Format	Procedural — produce six artifacts checked into migration repo
Sumo access	Admin role, access key + access ID generated
Dynatrace access	Not required yet — SL2DT-03 onward
Prior reading	SL2DT-01 for strategy and cut-scope framework

1. What You'll Produce

By the end of this step, commit these artifacts:

Artifact	Format	Purpose
<code>inventory/dashboards.json</code>	JSON	Every dashboard, panels, queries
<code>inventory/monitors.json</code>	JSON	Every monitor, threshold config, action
<code>inventory/searches.json</code>	JSON	Saved searches + scheduled searches
<code>inventory/fers.json</code>	JSON	Field Extraction Rules
<code>inventory/collectors.json</code>	JSON	Collector + source inventory
<code>inventory/roles.json</code> + <code>users.json</code>	JSON	RBAC model
<code>inventory/audit-usage.csv</code>	CSV	Usage data — last access per asset
<code>taxonomy-map.md</code>	Markdown	<code>_sourceCategory</code> → bucket/attribute/tag decision
<code>cut-scope.md</code>	Markdown	Per-asset migrate/retire/cut decision + owner signoff
<code>inventory-report.md</code>	Markdown	Executive summary — counts, cut-scope %, risk register

These gate SL2DT-03. Do not proceed without cut-scope signoff.

2. Prerequisites & Access

Sumo Access Keys

In Sumo Admin → Security → Access Keys, create a pair with admin scope (for inventory + audit log). Do **not** reuse user credentials.

```
export SUMO_ACCESS_ID="..."
export SUMO_ACCESS_KEY="..."
export SUMO_REGION="us2" # or us1, eu, au, jp, in, de, ca, fed
```

Region maps to the API base URL:

Region	API base
us1	<code>https://api.sumologic.com</code>
us2	<code>https://api.us2.sumologic.com</code>
eu	<code>https://api.eu.sumologic.com</code>
au	<code>https://api.au.sumologic.com</code>
...	see Sumo docs for full list

Auth: `Authorization: Basic <base64(access_id:access_key)>`

Migration Repo

```
mkdir sumo-migration && cd sumo-migration
git init
mkdir inventory
echo "SUMO_ACCESS_ID=" &gt; .env.example
echo "SUMO_ACCESS_KEY=" &gt;&gt; .env.example
echo "SUMO_REGION=us2" &gt;&gt; .env.example
cp .env.example .env # populate
echo ".env" &gt; .gitignore
```

3. Pull the Sumo Inventory

Pull each asset class via the Content API. Save each to `inventory/<type>.json`.

3.1 Dashboards

```
# List dashboards (paginated)
curl -u "$SUMO_ACCESS_ID:$SUMO_ACCESS_KEY" \
  "https://api.$SUMO_REGION.sumologic.com/api/v2/dashboards?limit=200" \
  &gt; inventory/dashboards-list.json

# For each, fetch full definition (requires dashboard ID)
jq -r '.dashboards[].id' inventory/dashboards-list.json | while read id; do
  curl -u "$SUMO_ACCESS_ID:$SUMO_ACCESS_KEY" \
    "https://api.$SUMO_REGION.sumologic.com/api/v2/dashboards/$id" \
    &gt; "inventory/dashboards/$id.json"
done
```

3.2 Monitors

```
curl -u "$SUMO_ACCESS_ID:$SUMO_ACCESS_KEY" \
  "https://api.$SUMO_REGION.sumologic.com/api/v1/monitors" \
  &gt; inventory/monitors.json
```

3.3 Saved & Scheduled Searches

```
# Saved searches live in user/folder content hierarchy
curl -u "$SUMO_ACCESS_ID:$SUMO_ACCESS_KEY" \
  "https://api.$SUMO_REGION.sumologic.com/api/v2/content/folders/personal" \
  &gt; inventory/content-personal.json

# Walk folder tree, save each search. Scheduled searches have a `searchSchedule`
field.
```

3.4 FERs, Collectors, Partitions, Roles, Users

```
curl -u "$SUMO_ACCESS_ID:$SUMO_ACCESS_KEY" \
  "https://api.$SUMO_REGION.sumologic.com/api/v1/extractionRules" \
  &gt; inventory/fers.json

curl -u "$SUMO_ACCESS_ID:$SUMO_ACCESS_KEY" \
  "https://api.$SUMO_REGION.sumologic.com/api/v1/collectors" \
  &gt; inventory/collectors.json

curl -u "$SUMO_ACCESS_ID:$SUMO_ACCESS_KEY" \
  "https://api.$SUMO_REGION.sumologic.com/api/v1/partitions" \
  &gt; inventory/partitions.json

curl -u "$SUMO_ACCESS_ID:$SUMO_ACCESS_KEY" \
  "https://api.$SUMO_REGION.sumologic.com/api/v1/roles" \
  &gt; inventory/roles.json

curl -u "$SUMO_ACCESS_ID:$SUMO_ACCESS_KEY" \
  "https://api.$SUMO_REGION.sumologic.com/api/v1/users" \
  &gt; inventory/users.json
```

3.5 Summary Counts

Record these numbers — they drive the wave plan in SL2DT-01:

```
echo "Dashboards: $(jq '.dashboards | length' inventory/dashboards-list.json)"
echo "Monitors:   $(jq '.monitors | length' inventory/monitors.json)"
echo "FERs:       $(jq '.data | length' inventory/fers.json)"
echo "Collectors: $(jq '.collectors | length' inventory/collectors.json)"
echo "Partitions: $(jq '.data | length' inventory/partitions.json)"
echo "Roles:      $(jq '.data | length' inventory/roles.json)"
echo "Users:      $(jq '.data | length' inventory/users.json)"
```

4. Pull the Audit Log — Usage-Based Cut Scope

The cut-scope decision depends on this data. Pull the last 90 days of audit events that show dashboard/monitor access.

Sumo stores audit events in a special index. Query via the Search Job API:

```
# Create search job
curl -X POST -u "$SUMO_ACCESS_ID:$SUMO_ACCESS_KEY" \
-H "Content-Type: application/json" \
"https://api.$SUMO_REGION.sumologic.com/api/v1/search/jobs" \
-d '{
  "query": "_sourceCategory=\\\"audit\\\" (eventName=DashboardViewed OR
eventName=MonitorRead)",
  "from": "2026-01-21T00:00:00Z",
  "to": "2026-04-21T00:00:00Z",
  "timeZone": "UTC"
}'

# Poll, then download messages. See Sumo API docs for job polling protocol.
```

Export to CSV:

dashboard_id	dashboard_name	unique_users_90d	last_accessed	access_count_90d
abc123	"Prod API Overview"	12	2026-04-20	847
def456	"Old Migration Dash"	0	—	0
...	...	...	...	...

Classification Rules

Based on last-90-day data:

Classification	Rule
Migrate	Last accessed $\leq$ 30d AND unique_users_90d $\geq$ 2
Migrate-on-request	Last accessed 31–90d OR unique_users_90d == 1
Cut	Not accessed in 90d+
Retire	Team dissolved / owner gone / superseded

Typical outcome: 20–30% migrate, 10–20% migrate-on-request, 40–60% cut, 5–10% retire. Your mileage will vary — a regulated-industry customer may have compliance constraints requiring 100% audit-trail-preservation.

5. Map `_sourceCategory` Taxonomy

Pull distinct `_sourceCategory` values and their volumes:

```
*
| count by _sourceCategory
| sort by _count desc
```

Export to CSV. For each top-100 `_sourceCategory` value, decide:

1. **Bucket mapping** — which Dynatrace Grail bucket does this belong to?
2. **Attribute preservation** — is the `_sourceCategory` value itself queryable post-migration (yes, via OpenPipeline enrichment)?
3. **Retention tier** — what retention period applies?

Output: `taxonomy-map.md`

```
| _sourceCategory Pattern | Target Bucket | Retention | Notes |
|-----|-----|-----|-----|
| prod/api/* | custom_logs_prod | 90d | High-volume, regulated |
| prod/web/* | custom_logs_prod | 90d | |
| prod/db/* | custom_logs_prod_db | 180d | Compliance hold |
| preprod/* | custom_logs_preprod | 30d | Low retention |
| dev/* | custom_logs_dev | 7d | Dev-only |
| audit/* | audit_logs | 365d | Compliance |
```

Decision framework — three strategies:

Strategy	Pick when
A — Bucket-primary	Retention/cost tiers differ per scope; IAM boundaries align with scope
B — Attribute-primary	Retention uniform; query-time grouping more important than isolation
C — Entity-tag	<code>_sourceCategory</code> tracks host/app ownership (not content type)

Most customers end up with **A + B**: bucket for coarse isolation, attribute preservation for query-time grouping. See SL2DT-03 for implementation.

6. Inventory Collectors & Ingest Paths

Every Sumo Collector needs a Dynatrace ingest equivalent. Classify each:


Collector Type	Dynatrace Target
Installed (Linux/Windows host)	OneAgent (auto-instruments logs + metrics + processes)
Hosted (log upload endpoint)	OpenPipeline HTTP endpoint
Hosted (syslog listener)	Dynatrace Syslog ingest (OneAgent or OTel Collector)

Collector Type	Dynatrace Target
Hosted (AWS CloudWatch)	Dynatrace AWS Clouds app (native integration)
Hosted (Azure EventHub)	Dynatrace Azure cloud integration
Hosted (GCP Pub/Sub)	Dynatrace GCP cloud integration
Hosted (Kubernetes)	DynaKube + OneAgent + log module
Hosted (HTTP source)	OpenPipeline HTTP endpoint or OTel HTTP receiver

For each collector, document:

- Source format (JSON, plain text, structured)
- Typical volume (MB/day)
- Parse rules applied (FER references)
- Owning team
- Retention policy

Output: `inventory/collectors-mapping.md` — one row per collector with target.



7. Inventory FERs (Field Extraction Rules)

FERs are parse rules applied at ingest. Every FER becomes an OpenPipeline processor.

For each FER, record:

- Scope (which `_sourceCategory` pattern does it match)
- Parse expression (Sumo regex syntax)

- Output fields (which fields it extracts)
- Target OpenPipeline pipeline (based on `_sourceCategory` → bucket mapping)

Output: `inventory/fers-mapping.md`

```
| FER Name | Scope | Parse Expression | Output Fields | Target DT Pipeline |
|-----|-----|-----|-----|-----|
| api-parse-req | prod/api/* | `(?<method>\w+) (?<path>\S+) (?<status>\d+)` | method, path, status | logs_prod_pipeline |
| audit-parse | audit/* | ... | user, action, target | audit_pipeline |
```

The translation work happens in SL2DT-03 (OpenPipeline implementation) and SL2DT-04 (DPL pattern conversion from Sumo regex).

8. Inventory Roles & Users

Pull Sumo's role inventory. For each role:

- Name + description
- Search filter (`_sourceCategory=prod/*` , etc.) — this maps to bucket filter in DT
- Capabilities (view, manage dashboards, manage monitors, admin) — maps to IAM policy permissions
- Assigned users count

Output: `inventory/rbac-mapping.md`

```
| Sumo Role | Search Filter | Capabilities | User Count | Target DT Group | Target DT Policy |
|-----|-----|-----|-----|-----|-----|
| prod-readers | _sourceCategory=prod/* | Read-only | 82 | g-prod-readers | p-prod-read (bucket: custom_logs_prod) |
| prod-admins | _sourceCategory=prod/* | Full | 6 | g-prod-admins | p-prod-admin |
| app-team-api | _sourceCategory=prod/api/* | Read + dashboard mgmt | 28 | g-app-team-api | p-app-team-api |
```

SSO mapping: if Sumo uses SSO (SAML / OIDC), the same IdP can feed Platform IAM Group assignments. Document SSO config in `inventory/sso-config.md` .

See SL2DT-07 for full IAM translation.

9. Cut-Scope Decision Workshop

Run a structured workshop with each team owner. Present:

1. Their dashboards/monitors with usage data
2. Cut-scope classification (migrate / migrate-on-request / cut / retire)
3. Required signoff per classification

Workshop Format

- 90 minutes per team (larger teams may need 2–3 sessions)
- Pre-read: team-specific `cut-scope-{team}.md` with sorted usage table
- In-session: walk through exception list, capture decisions
- Post-session: owner signs committed list

Decision Capture

```
# Cut Scope – Team: Payments Platform

## Signed off by: Jane Doe (Team Lead) on 2026-04-25

### Migrate (12)
- dashboard_id_1 – "Payments API Overview"
- monitor_id_2 – "Failed Transactions > 100/min"
- ...

### Migrate-on-request (5) – revisit Wave 4
- dashboard_id_7 – "Old QA Dashboard" (owner: "maybe useful, revisit before Wave 4")
- ...

### Cut (34)
- dashboard_id_11 – "Superseded by DataDog POC" (not accessed in 6mo)
- ...

### Retire (2)
- monitor_id_55 – Legacy card network integration (deprecated 2024)
- ...
```

Commit these to the migration repo. This is the audit trail for every decision.

Aggregate Metrics

Roll up across teams:

Team	Total Assets	Migrate	Migrate-on-request	Cut	Retire
Payments	53	12	5	34	2
Identity	41	14	7	18	2

Team	Total Assets	Migrate	Migrate-on-request	Cut	Retire
...	...	...	...	...	...
Total	2,847	612	241	1,732	262

Present aggregate to the executive sponsor. This is the basis for the migration budget.

10. Step Exit Criteria

G2 — Inventory + Cut-Scope Complete

All must be true:

- ☐ All 6 inventory JSON files written + committed
- ☐ `audit-usage.csv` pulled for last 90 days
- ☐ `taxonomy-map.md` populated with top-100 `_sourceCategory` → target bucket mappings
- ☐ `cut-scope.md` signed off by every team owner
- ☐ Aggregate cut-scope metrics presented to sponsor
- ☐ `inventory-report.md` includes: counts, cut %, risks, timeline implications

If any check fails, do not proceed to SL2DT-03. Fix the gap first.

Next step: SL2DT-03 — Log Ingest Architecture (bucket design, OneAgent + OTel + OpenPipeline deployment, FER → processor conversion).

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace or Sumo Logic. Always verify information against the official [Dynatrace documentation](#) and [Sumo Logic documentation](#).