SL2DT-99: Summary & Runbook Index

> **Series:** SL2DT — Sumo Logic to Dynatrace | **Notebook:** 10 of 10 | **Created:** April 2026 | **Last Updated:** 04/21/2026

Overview

Goal of this notebook: single-page index for the entire SL2DT series. Use as a reference card during migration work, a hand-off document for incoming teams, and a quick-start for future Sumo-to-Dynatrace engagements.

This notebook is intentionally terse. Each section points to the authoritative detail in SL2DT-01 through SL2DT-09.

Table of Contents

Prerequisites

Requirement	Details
Audience	Anyone onboarding to a Sumo-to-Dynatrace engagement
Prior reading	None required — this notebook is the entry point

1. Migration Framework Summary

Sumo-to-Dynatrace in 9 steps + summary:

```
01 Overview          → strategy, mental model, anti-patterns
02 Inventory         → audit-log-driven cut scope
03 Ingest            → OneAgent + OTel + OpenPipeline + buckets
04 Translate         → SumoQL → DQL, confidence-scored
05 Monitors          → Dynatrace Intelligence / Workflow / Metric Event
decisions
06 Dashboards        → Notebooks (analytical) + Dashboards (operational)
07 Governance        → Platform IAM, bucket-scoped policies
08 Automation        → Terraform + Monaco + CI/CD
09 Cutover           → 3-tier validation, decommission
99 Summary           → this notebook
```

Typical duration: **3–6 months** for a mid-size deployment (100–500 active users, 2k–8k dashboards/monitors). Large enterprises (10k+ assets, multi-tenant) may need 9–12 months.

Target outcomes:

- Sumo contract terminated by deadline
- ≥ 95% monitor + dashboard fidelity
- ≥ 40% new monitors using anomaly detection (not static thresholds)
- Alert noise ≤ Sumo baseline

2. Series Index — Which Notebook Does What

#	Notebook	When to read	Primary artifacts
01	Overview & Migration Strategy	First, always	Strategy doc, risk register
02	Assessment & Inventory	Before any hands-on work	`inventory/*.json`, `cut-scope.md`, `taxonomy-map.md`
03	Log Ingest Architecture	Wave 1	Bucket configs, OpenPipeline pipelines, FER conversions
04	SumoQL → DQL Translation	Wave 1 / 2	`translations/all-queries.csv`
05	Monitor/Alert Conversion	Wave 2	Dynatrace Intelligence detectors, Workflows, Metric Events
06	Dashboard Conversion	Wave 3 / 4	Notebook + Dashboard JSON
07	User Governance & Access	Wave 1 (do early!)	IAM groups/policies/bindings
08	Automation & GitOps	Wave 1 ongoing	Terraform + Monaco configs, CI pipelines
09	Cutover, Validation & Decommission	Wave 5	Cutover runbook, validation report, retro
99	Summary (this notebook)	Any time — reference card	—

3. Critical Decisions Checklist

Six decisions shape the entire migration. Make them early, document them in `DECISIONS.md`:

Decision	Default	When to deviate
`_sourceCategory` strategy	Mixed: bucket for coarse isolation + attribute for query grouping	Bucket-only if retention tiers dominate; attribute-only if uniform retention

Decision	Default	When to deviate
Ingest path per source	OneAgent for hosts, OTel Collector for cloud/HTTP	Sumo Forwarding if OneAgent can't be deployed
Monitor conversion	anomaly detection for metric-backed alerts	Static threshold only when truly static
Dashboard target	Notebooks for analytical, Dashboards for operational	Hybrid for SRE pages
Automation tool split	Terraform for IAM/buckets/pipelines, Monaco for dashboards/monitors	Single-tool if team mandates
Parallel-run duration	2–4 weeks	Longer only with documented cause (rarely a good sign)

4. Artifacts Produced Across All Steps

Cumulative list. All committed to the migration repo.

Artifact	Step	Purpose
`inventory/*.json`	SL2DT-02	Asset inventory
`inventory/audit-usage.csv`	SL2DT-02	Usage data for cut scope
`taxonomy-map.md`	SL2DT-02	`_sourceCategory` → bucket/attribute
`cut-scope.md`	SL2DT-02	Per-asset decision + owner signoff
`terraform/buckets.tf`	SL2DT-03	Bucket configs
`openpipeline/*.json`	SL2DT-03	Pipeline + processor configs

Artifact	Step	Purpose
`translations/all-queries.csv`	SL2DT-04	Per-query DQL + confidence
`monitors/decision-matrix.csv`	SL2DT-05	Dynatrace Intelligence/Workflow/Metric-event decision
`dashboards/*`	SL2DT-06	Notebook/Dashboard JSON
`iam/groups.tf`, `policies.tf`	SL2DT-07	IAM as code
`.github/workflows/*.yml`	SL2DT-08	CI pipelines
`cutover/runbook.md`	SL2DT-09	Cutover step-by-step
`cutover/validation-report.md`	SL2DT-09	3-tier validation results
`cutover/post-cutover-retro.md`	SL2DT-09	Retro after 30-day stable window

5. Failure-Mode Reference

When things go wrong, consult this table first:

Symptom	Root cause category	Fix (notebook)
App teams recreating static thresholds in DT	Dynatrace Intelligence not explained	SL2DT-05 §3
Bucket writes to wrong bucket	OpenPipeline matcher wrong	SL2DT-03 §4
Queries time out	No `from:` or too wide range	SL2DT-04 §3 Rule 1
Sort/filter on timeseries fails	Array not reduced	SL2DT-04 §3 Rule 3
Parse returns null	DPL syntax wrong	SL2DT-03 §8, sumoql-to-dql skill

Symptom	Root cause category	Fix (notebook)
Users can't see their data	Bucket policy wrong	SL2DT-07 §3
ServiceNow incidents not created	Webhook payload not rebuilt	SL2DT-05 §6
Alert volume exploded	baselines still learning	SL2DT-05 §8
Dashboard blank after migration	Time range or bucket mismatch	SL2DT-06 §3
SSO assigns no permissions	Claim → group mapping wrong	SL2DT-07 §4

6. Quick Reference — Common DQL Patterns

Count logs per scope

```
// 6. Quick Reference — Common DQL Patterns
fetch logs, from:-1h
| summarize c = count(), by:{dt.source_entity}
| sort c desc
```

Top errors by host

```
// Top errors by host
fetch logs, from:-1h
| filter dt.source_entity == "prod/api"
| filter loglevel == "ERROR"
| summarize c = count(), by:{host.name}
| sort c desc
| limit 10
```

Error rate timeseries

```
// Error rate timeseries
fetch logs, from:-1h
| filter dt.source_entity == "prod/api"
| fieldsAdd is_error = if(loglevel == "ERROR", 1, else:0)
| makeTimeseries total = count(), errors = sum(is_error), interval:1m
```

Latency percentiles

```
// Latency percentiles
fetch logs, from:-1h
| filter dt.source_entity == "prod/api"
| parse content, "LD 'latency=' INT:latency"
| summarize p50 = percentile(latency, 50),
            p95 = percentile(latency, 95),
            p99 = percentile(latency, 99),
            by:{http.path}
```

Recent detected problems

```
// Recent detected problems
fetch events, from:-24h
| filter event.kind == "DAVIS_PROBLEM"
| sort timestamp desc
| limit 20
| fields timestamp, event.name, dt.davis.problem.severity,
dt.davis.problem.status
```

Timeseries with array reduction

```
// Pattern when you need to filter/sort after makeTimeseries
timeseries avg_cpu = avg(dt.host.cpu.usage), from:-1h, by:{dt.entity.host}
| fieldsAdd max_cpu = arrayMax(avg_cpu)
| filter max_cpu &gt; 80
| sort max_cpu desc
```

Lookup for enrichment

```
// Lookup for enrichment
fetch logs, from:-1h
| filter dt.source_entity == "prod/api"
| summarize c = count(), by:{host.name}
| lookup [fetch dt.entity.host | fieldsAdd hostname = entity.name | fields
hostname, `dt.tags.team`],
    sourceField:host.name, lookupField:hostname,
    fields:{`dt.tags.team`}
```

7. Engagement Timeline Template

Scaled to a 6-month mid-size migration. Adjust for your scope.

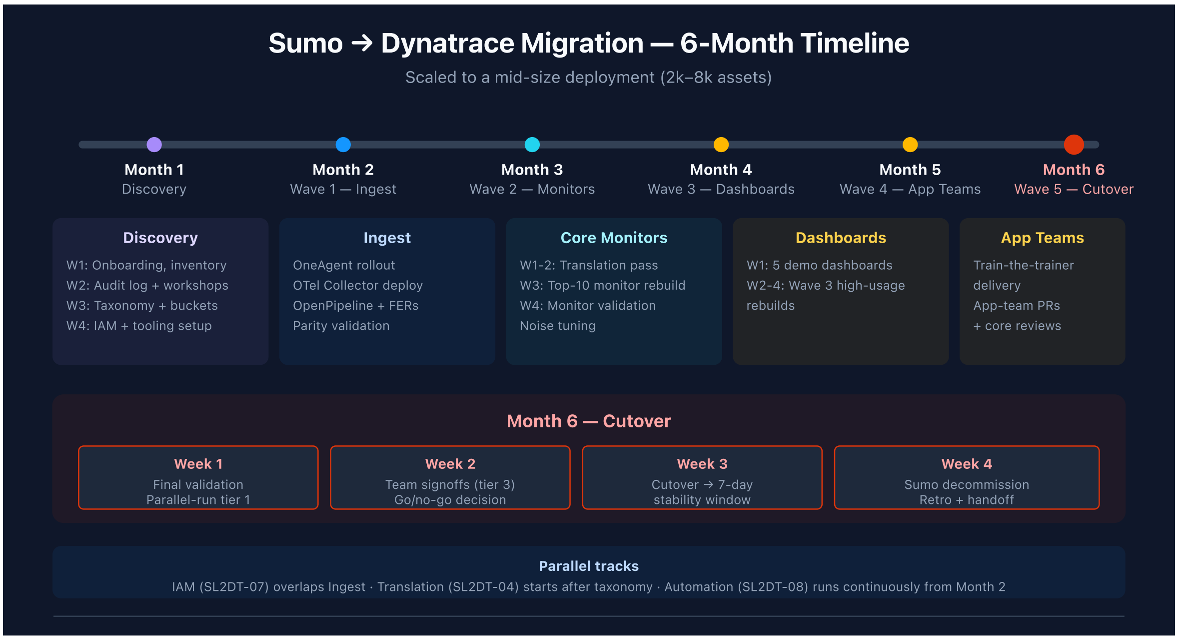

Parallel tracks that can overlap

- IAM (SL2DT-07) runs in parallel with ingest (SL2DT-03)
- Translation (SL2DT-04) can start once taxonomy is mapped, before ingest is complete
- Automation (SL2DT-08) runs continuously from Month 2 onward

8. Companion Assets

Skill

- sumoql-to-dql — translation tables, confidence scoring
- `references/mapping-tables.md` — complete operator/function/field maps
- `references/examples.md` — 15 worked translations
- `references/out-of-scope.md` — features not covered in v1

Docs

- SL2DT HISTORY — development context and decisions
- SL2DT REFERENCE — quick reference, patterns, lessons learned
- SL2DT AGENT-TASKS — brief for building the `Dynatrace-SumoLogic` migration tool

Parallel series

- NR2DT — New Relic → Dynatrace procedural
- NRLC — New Relic component deep dives
- S2D — Splunk → Dynatrace
- M2S — Managed → SaaS
- S2S — SaaS → SaaS

Official documentation

- Sumo Logic Docs
- Dynatrace Docs

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace or Sumo Logic. Always verify information against the official Dynatrace documentation and Sumo Logic documentation.