



Service Dependencies & Flow Analysis

Series: SPANS | **Notebook:** 4 of 8 | **Created:** December 2025 | **Last Updated:** 01/28/2026

Mapping Your Distributed System

This notebook teaches you how to use span data to understand service relationships, analyze request flows, and identify critical dependencies in your system.

Table of Contents

1. [Understanding Service Topology](#)
 2. [Discovering Services](#)
 3. [Mapping Service Dependencies](#)
 4. [Client-Server Call Patterns](#)
 5. [Async Messaging Flows](#)
 6. [Trace Hierarchy Analysis](#)
 7. [Cross-Service Latency Analysis](#)
 8. [Critical Path Analysis](#)
-

Prerequisites


Before starting this notebook, ensure you have:

-  Completed **SPANS-01** through **SPANS-03**

- ☒ Access to a Dynatrace environment with distributed trace data
- ☒ Understanding of span kinds (server, client, producer, consumer)

1. Understanding Service Topology

Service topology shows how your services connect and communicate:



Key Span Types for Topology

Span Kind	Role	What It Tells You
server	Receives requests	Entry points, inbound traffic
client	Makes requests	Outbound calls, dependencies
producer	Sends messages	Async event sources
consumer	Receives messages	Async event handlers

2. Discovering Services

First, discover all services generating spans in your environment:

```
// List all services with span counts
fetch spans, from:-1h
| summarize {
    span_count = count(),
    operations = countDistinct(span.name)
}, by: {service.name}
| sort span_count desc
| limit 30
```

```
// Service role analysis - understand each service's function
fetch spans, from:-1h
| summarize {
    server_spans = countIf(span.kind == "server"),
    client_spans = countIf(span.kind == "client"),
    internal_spans = countIf(span.kind == "internal"),
    producer_spans = countIf(span.kind == "producer"),
    consumer_spans = countIf(span.kind == "consumer")
}, by: {service.name}
| sort server_spans desc
| limit 30
```

```
// Discover all operations/endpoints per service
fetch spans, from:-1h
| filter span.kind == "server"
| summarize {
    request_count = count(),
    avg_duration_ms = avg(duration) / 1000000
}, by: {service.name, span.name}
| sort service.name asc, request_count desc
| limit 50
```

3. Mapping Service Dependencies

Use CLIENT spans to understand which services call which other services:

Key Attributes for Dependencies

Attribute	Description
<code>peer.service</code>	Target service name (if instrumented)

Attribute	Description
<code>server.address</code>	Target host/address
<code>server.port</code>	Target port
<code>http.host</code>	HTTP host header

```
// Map service-to-service calls using CLIENT spans
fetch spans, from:-1h
| filter span.kind == "client"
| summarize {
    call_count = count(),
    avg_latency_ms = avg(duration) / 1000000,
    error_count = countIf(span.status_code == "error")
}, by: {service.name, span.name}
| fieldsAdd error_rate_pct = (error_count * 100.0) / call_count
| sort call_count desc
| limit 50
```

```
// Find services called by other services using peer.service attribute
// peer.service shows the target service name when available
fetch spans, from:-1h
| filter span.kind == "client" and isNotNull(peer.service)
| summarize {
    call_count = count(),
    avg_latency_ms = avg(duration) / 1000000
}, by: {service.name, peer.service}
| sort call_count desc
| limit 30
```

```
// Map dependencies using HTTP host/URL information
fetch spans, from:-1h
| filter span.kind == "client" and isNotNull(server.address)
| summarize {
    call_count = count(),
    avg_latency_ms = avg(duration) / 1000000,
    error_count = countIf(span.status_code == "error")
}, by: {service.name, server.address}
| fieldsAdd error_rate_pct = (error_count * 100.0) / call_count
| sort call_count desc
| limit 30
```

4. Client-Server Call Patterns

Analyze the matching CLIENT and SERVER span pairs to understand inter-service communication:


```
// Analyze outbound calls from each service
fetch spans, from:-1h
| filter span.kind == "client"
| summarize {
    outbound_calls = count(),
    avg_latency_ms = avg(duration) / 1000000,
    p99_latency_ms = percentile(duration, 99) / 1000000,
    error_count = countIf(span.status_code == "error")
}, by: {service.name}
| fieldsAdd error_rate_pct = (error_count * 100.0) / outbound_calls
| sort outbound_calls desc
| limit 20
```

```
// Analyze inbound requests to each service
fetch spans, from:-1h
| filter span.kind == "server"
| summarize {
    inbound_requests = count(),
    avg_latency_ms = avg(duration) / 1000000,
    p99_latency_ms = percentile(duration, 99) / 1000000,
    error_count = countIf(span.status_code == "error")
}, by: {service.name}
| fieldsAdd error_rate_pct = (error_count * 100.0) / inbound_requests
| sort inbound_requests desc
| limit 20
```

```
// Find services with high outbound/inbound ratio (integration heavy)
fetch spans, from:-1h
| summarize {
    inbound = countIf(span.kind == "server"),
    outbound = countIf(span.kind == "client")
}, by: {service.name}
| filter inbound > 0
| fieldsAdd outbound_ratio = outbound / inbound
| sort outbound_ratio desc
| limit 20
```

5. Async Messaging Flows

Analyze asynchronous messaging patterns using PRODUCER and CONSUMER spans:



Key Messaging Attributes

Attribute	Description
messaging.system	Kafka, RabbitMQ, etc.
messaging.destination.name	Topic/queue name
messaging.operation	publish, receive, etc.

```
// Find message producers (services sending async messages)
fetch spans, from:-1h
| filter span.kind == "producer"
| summarize {
  messages_sent = count(),
  avg_duration_ms = avg(duration) / 1000000,
  error_count = countIf(span.status_code == "error")
}, by: {service.name, span.name}
| sort messages_sent desc
| limit 20
```

```
// Find message consumers (services receiving async messages)
fetch spans, from:-1h
| filter span.kind == "consumer"
| summarize {
    messages_received = count(),
    avg_processing_ms = avg(duration) / 1000000,
    error_count = countIf(span.status_code == "error")
}, by: {service.name, span.name}
| fieldsAdd error_rate_pct = (error_count * 100.0) / messages_received
| sort messages_received desc
| limit 20
```

```
// Analyze messaging system usage (Kafka, RabbitMQ, etc.)
fetch spans, from:-1h
| filter span.kind == "producer" or span.kind == "consumer"
| filter isNotNull(messaging.system)
| summarize {
    message_count = count(),
    producers = countIf(span.kind == "producer"),
    consumers = countIf(span.kind == "consumer"),
    error_count = countIf(span.status_code == "error")
}, by: {messaging.system, messaging.destination.name}
| sort message_count desc
| limit 20
```

```
// Map producer to consumer relationships
fetch spans, from:-1h
| filter span.kind == "producer" or span.kind == "consumer"
| summarize {
    span_count = count()
}, by: {service.name, span.kind, messaging.destination.name}
| sort messaging.destination.name, span.kind
| limit 30
```

6. Trace Hierarchy Analysis

Analyze parent-child relationships within traces to understand call depth:

```
// Count spans per trace to understand trace complexity
fetch spans, from:-1h
| summarize {
    span_count = count(),
    services_involved = countDistinct(service.name),
    total_duration_ms = sum(duration) / 1000000
}, by: {trace.id}
| sort span_count desc
| limit 25
```

```
// Examine a complete trace hierarchy
// Replace YOUR_TRACE_ID with an actual trace ID from above
fetch spans, from:-1h
// | filter trace.id == "YOUR_TRACE_ID"
| fieldsAdd duration_ms = duration / 1000000
| fields start_time,
    span.id,
    span.parent_id,
    service.name,
    span.name,
    span.kind,
    duration_ms
| sort start_time asc
| limit 100
```

```
// Find entry points (root spans) and their downstream services
fetch spans, from:-1h
| filter isNull(span.parent_id)
| summarize {
    entry_count = count(),
    avg_duration_ms = avg(duration) / 1000000
}, by: {service.name, span.name}
| sort entry_count desc
| limit 20
```

7. Cross-Service Latency Analysis

Identify latency hot spots between services:


```
// Latency by service-to-service call
fetch spans, from:-1h
| filter span.kind == "client" and isNotNull(server.address)
| summarize {
    call_count = count(),
    avg_ms = avg(duration) / 1000000,
    p95_ms = percentile(duration, 95) / 1000000,
    p99_ms = percentile(duration, 99) / 1000000
}, by: {service.name, server.address}
| sort p95_ms desc
| limit 20
```

```
// Time spent per service in traces
fetch spans, from:-1h
| summarize {
    total_time_ms = sum(duration) / 1000000,
    span_count = count(),
    avg_per_span_ms = avg(duration) / 1000000
}, by: {service.name}
| sort total_time_ms desc
| limit 20
```


```
// Find slowest dependencies (CLIENT spans)
fetch spans, from:-1h
| filter span.kind == "client"
| filter duration > 500ms
| summarize {
    slow_call_count = count(),
    avg_duration_ms = avg(duration) / 1000000,
    max_duration_ms = max(duration) / 1000000
}, by: {service.name, span.name}
| sort avg_duration_ms desc
| limit 20
```

8. Critical Path Analysis

Identify the services and operations that contribute most to end-to-end latency:

Critical Path Analysis

Total Trace Duration: 500ms | Identify the bottleneck to optimize



Optimizing the Payment service (300ms) would reduce trace time by up to 60%. Focus optimization efforts here first.

```
// Find services contributing most to total trace time  
fetch spans, from:-1h  
| summarize {  
  total_self_time_ms = sum(duration) / 1000000,  
  span_count = count(),  
  avg_duration_ms = avg(duration) / 1000000,  
  max_duration_ms = max(duration) / 1000000  
}, by: {service.name}  
| sort total_self_time_ms desc  
| limit 15
```

```
// Find slowest operations across all services  
fetch spans, from:-1h  
| filter span.kind == "server"  
| summarize {  
  call_count = count(),  
  avg_duration_ms = avg(duration) / 1000000,  
  p99_duration_ms = percentile(duration, 99) / 1000000,  
  total_time_ms = sum(duration) / 1000000  
}, by: {service.name, span.name}  
| filter call_count > 10  
| sort p99_duration_ms desc  
| limit 20
```

```
// Identify high-impact optimization candidates
// (high volume + high latency = most benefit from optimization)
fetch spans, from:-1h
| filter span.kind == "server"
| summarize {
    call_count = count(),
    avg_duration_ms = avg(duration) / 1000000,
    total_time_ms = sum(duration) / 1000000
}, by: {service.name, span.name}
| filter call_count > 50
| fieldsAdd impact_score = call_count * avg_duration_ms
| sort impact_score desc
| limit 20
```

Summary

In this notebook, you learned:

- ✓ **Service discovery** - Find all services and their operations
 - ✓ **Dependency mapping** - Use CLIENT spans with `peer.service` and `server.address`
 - ✓ **Client-server patterns** - Analyze inbound/outbound call ratios
 - ✓ **Async messaging** - Track PRODUCER/CONSUMER spans through message queues
 - ✓ **Trace hierarchy** - Understand span parent-child relationships
 - ✓ **Cross-service latency** - Find latency hot spots between services
 - ✓ **Critical path analysis** - Identify bottlenecks for optimization
-

Next Steps

Continue to **SPANS-05: Advanced Span Analytics** to learn:

- Time series analysis and trending
 - Complex aggregations and calculations
 - Building dashboard-ready queries
 - Alerting patterns
-

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.