

SPANS-05: Advanced Span Analytics

Series: SPANS — Distributed Tracing and Spans | **Notebook:** 5 of 8 | **Created:** December 2025 | **Last Updated:** 04/25/2026

Time-Series Analysis and Complex Aggregations

This notebook covers advanced analytical techniques for span data, including time-series analysis, trend detection, and complex aggregations for building dashboards and reports.

Table of Contents

1. [Time-Series with makeTimeseries](#)
 2. [Trend Analysis](#)
 3. [Complex Aggregations](#)
 4. [Comparison Queries](#)
 5. [Dashboard-Ready Queries](#)
-


Prerequisites

Before starting this notebook, ensure you have:

-  Completed previous SPANS notebooks (01-04)
-  Understanding of summarize and aggregation functions
-  Familiarity with time-based filtering

1. Time-Series with makeTimeseries

Use `makeTimeseries` to create time-series data for visualization and trend analysis.



⚠ Important: `makeTimeseries` does NOT support `percentile()` or arithmetic expressions in aggregations. Use time-bucketed `summarize` for those.

```
// Request volume over time by service
fetch spans, from:-24h
| filter span.kind == "server"
| makeTimeseries {
  request_count = count(),
  error_count = countIf(span.status_code == "error")
}, by:{service.name}, interval: 5m
```

```
// Average duration over time (without grouping)
fetch spans, from:-24h
| filter span.kind == "server"
| makeTimeseries {
  request_count = count(),
  avg_duration_ns = avg(duration)
}, interval: 5m
```

```
// For percentile trends, use time-bucketed summarize instead
fetch spans, from:-1h
| filter span.kind == "server"
| fieldsAdd time_bucket = bin(start_time, 10m)
| summarize {
    request_count = count(),
    p95_duration_ms = percentile(duration, 95) / 1ms,
    error_count = countIf(span.status_code == "error")
}, by:{time_bucket, service.name}
| sort time_bucket asc
| limit 200
```

2. Trend Analysis

Identify trends and patterns in your span data over time.

```
// Hourly error rate trend
fetch spans, from:-1h
| filter span.kind == "server"
| fieldsAdd hour_bucket = bin(start_time, 1h)
| summarize {
    total_requests = count(),
    errors = countIf(span.status_code == "error")
}, by:{hour_bucket}
| fieldsAdd error_rate_pct = (errors * 100.0) / total_requests
| sort hour_bucket asc
| limit 48
```

```
// Latency trend by service (10-minute buckets)
fetch spans, from:-1h
| filter span.kind == "server"
| fieldsAdd time_bucket = bin(start_time, 10m)
| summarize {
    request_count = count(),
    avg_duration_ms = avg(duration) / 1ms,
    p90_duration_ms = percentile(duration, 90) / 1ms
}, by:{time_bucket, service.name}
| sort time_bucket asc, service.name
| limit 300
```

```
// Identify services with degrading performance
fetch spans, from:-1h
| filter span.kind == "server"
| fieldsAdd time_bucket = bin(start_time, 30m)
| summarize {
    request_count = count(),
    p95_ms = percentile(duration, 95) / 1ms
}, by:{time_bucket, service.name}
| filter p95_ms > 500
| sort time_bucket desc, p95_ms desc
| limit 50
```

3. Complex Aggregations

Perform multi-dimensional analysis with complex aggregation patterns.

```
// Service health scorecard
fetch spans, from:-1h
| filter span.kind == "server"
| summarize {
    total_requests = count(),
    error_count = countIf(span.status_code == "error"),
    avg_duration_ms = avg(duration) / 1ms,
    p50_duration_ms = percentile(duration, 50) / 1ms,
    p95_duration_ms = percentile(duration, 95) / 1ms,
    p99_duration_ms = percentile(duration, 99) / 1ms,
    max_duration_ms = max(duration) / 1ms
}, by:{service.name}
| fieldsAdd error_rate_pct = (error_count * 100.0) / total_requests
| fieldsAdd health_score = if(error_rate_pct > 5, "Critical",
                             else: if(error_rate_pct > 1, "Warning",
                             else: "Healthy"))
| sort error_rate_pct desc
| limit 30
```

```
// Endpoint-level analysis with multiple metrics
fetch spans, from:-1h
| filter span.kind == "server"
| summarize {
    request_count = count(),
    error_count = countIf(span.status_code == "error"),
    slow_count = countIf(duration > 1000000000), // > 1 second
    avg_duration_ms = avg(duration) / 1ms,
    p95_duration_ms = percentile(duration, 95) / 1ms
}, by:{service.name, span.name}
| fieldsAdd error_rate_pct = (error_count * 100.0) / request_count
| fieldsAdd slow_rate_pct = (slow_count * 100.0) / request_count
| filter request_count > 10
| sort error_rate_pct desc
| limit 50
```

```
// HTTP method distribution with performance metrics
fetch spans, from:-1h
| filter isNotNull(http.request.method)
| summarize {
    request_count = count(),
    error_count = countIf(span.status_code == "error"),
    avg_duration_ms = avg(duration) / 1ms
}, by:{http.request.method, service.name}
| sort request_count desc
| limit 30
```

4. Comparison Queries

Compare metrics across different dimensions to identify outliers and patterns.

```
// Compare span kinds: server vs client performance
fetch spans, from:-1h
| filter in(span.kind, {"server", "client"})
| summarize {
    span_count = count(),
    error_count = countIf(span.status_code == "error"),
    avg_duration_ms = avg(duration) / 1ms,
    p95_duration_ms = percentile(duration, 95) / 1ms
}, by:{span.kind}
| fieldsAdd error_rate_pct = (error_count * 100.0) / span_count
```

```
// Compare success vs error span characteristics
fetch spans, from:-1h
| filter span.kind == "server"
| summarize {
    span_count = count(),
    avg_duration_ms = avg(duration) / 1ms,
    p95_duration_ms = percentile(duration, 95) / 1ms
}, by:{span.status_code, service.name}
| sort service.name, span.status_code
| limit 50
```

```
// HTTP status code distribution by service
fetch spans, from:-1h
| filter isNotNull(http.response.status_code)
| fieldsAdd status_class = if(http.response.status_code >= 500, "5xx",
                             else: if(http.response.status_code >= 400,
                             "4xx",
                             else: if(http.response.status_code >= 300,
                             "3xx",
                             else: "2xx")))
| summarize {count = count()}, by:{status_class, service.name}
| sort service.name, status_class
```

5. Dashboard-Ready Queries

Queries optimized for use in Dynatrace dashboards and reports.

 **Tip:** These queries are designed to produce clean output suitable for dashboard tiles.

```
// Dashboard: Request rate and error rate over time
fetch spans, from:-24h
| filter span.kind == "server"
| makeTimeseries {
    requests = count(),
    errors = countIf(span.status_code == "error")
}, interval: 5m
```

```
// Dashboard: Service health summary (single value tiles)
fetch spans, from:-1h
| filter span.kind == "server"
| summarize {
    total_requests = count(),
    total_errors = countIf(span.status_code == "error"),
    unique_services = countDistinct(service.name),
    avg_latency_ms = avg(duration) / 1ms
}
| fieldsAdd overall_error_rate_pct = (total_errors * 100.0) / total_requests
```

```
// Dashboard: Top 10 services by request volume
fetch spans, from:-1h
| filter span.kind == "server"
| summarize {
    requests = count(),
    errors = countIf(span.status_code == "error"),
    p95_ms = percentile(duration, 95) / 1ms
}, by:{service.name}
| fieldsAdd error_rate = (errors * 100.0) / requests
| sort requests desc
| limit 10
```

```
// Dashboard: Recent errors list
fetch spans, from:-1h
| filter span.status_code == "error"
| fields start_time, service.name, span.name, span.status_message
| sort start_time desc
| limit 20
```

Summary

In this notebook, you learned:

- ✓ **Time-series analysis** with makeTimeseries and its limitations
- ✓ **Time-bucketed summarize** for percentile trends
- ✓ **Trend analysis** to identify patterns over time
- ✓ **Complex aggregations** for multi-dimensional analysis
- ✓ **Comparison queries** to identify outliers
- ✓ **Dashboard-ready queries** for visualization

Next Steps

Continue to **SPANS-06: Security Analysis with Spans** to learn:

- Detecting security-relevant patterns in traces
 - Analyzing authentication and authorization flows
 - Finding anomalous behavior
 - Security audit queries
-

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.